

Deep Learning for Data Science

DS 542

<https://dl4ds.github.io/sp2026/>

Gradient Descent

Announcements

- Re: discussion deadlines are moving to 11:59pm on the day of discussion.
 - Why? The practice is more important than the timing.
 - Still targeting \leq 30 minutes to do, but more time if you need/want it.
- Shared Compute Cluster (SCC) Tutorial next Monday.
 - Please bring your laptop to class.
 - No graded exercise, but will be walking through account setup.

Plan for Today

- Loss functions for multiclass classification (spillover)
- Example of gradient descent
- Basics of gradient descent
- Gradient descent as a statistical process
- Challenges with gradient descent

Loss Function for Regression

If you recast regression as

1. Predicting the mean of a normal distribution with a fixed variance and
2. Optimize output for maximum likelihood,

Then the optimization is equivalent to optimizing with least squared errors (L_2) as your loss function.

Loss Function for Binary Classification

If you are modeling a binary classification problem,

1. The sigmoid function is handy to map arbitrary “scores” into probabilities, so
2. Your loss function is equivalent to

$$L[\phi] = \sum_i -(1 - y_i) \log[1 - \text{sig}[f[x_i|\phi]]] - y_i \log[\text{sig}[f[x_i|\phi]]]$$

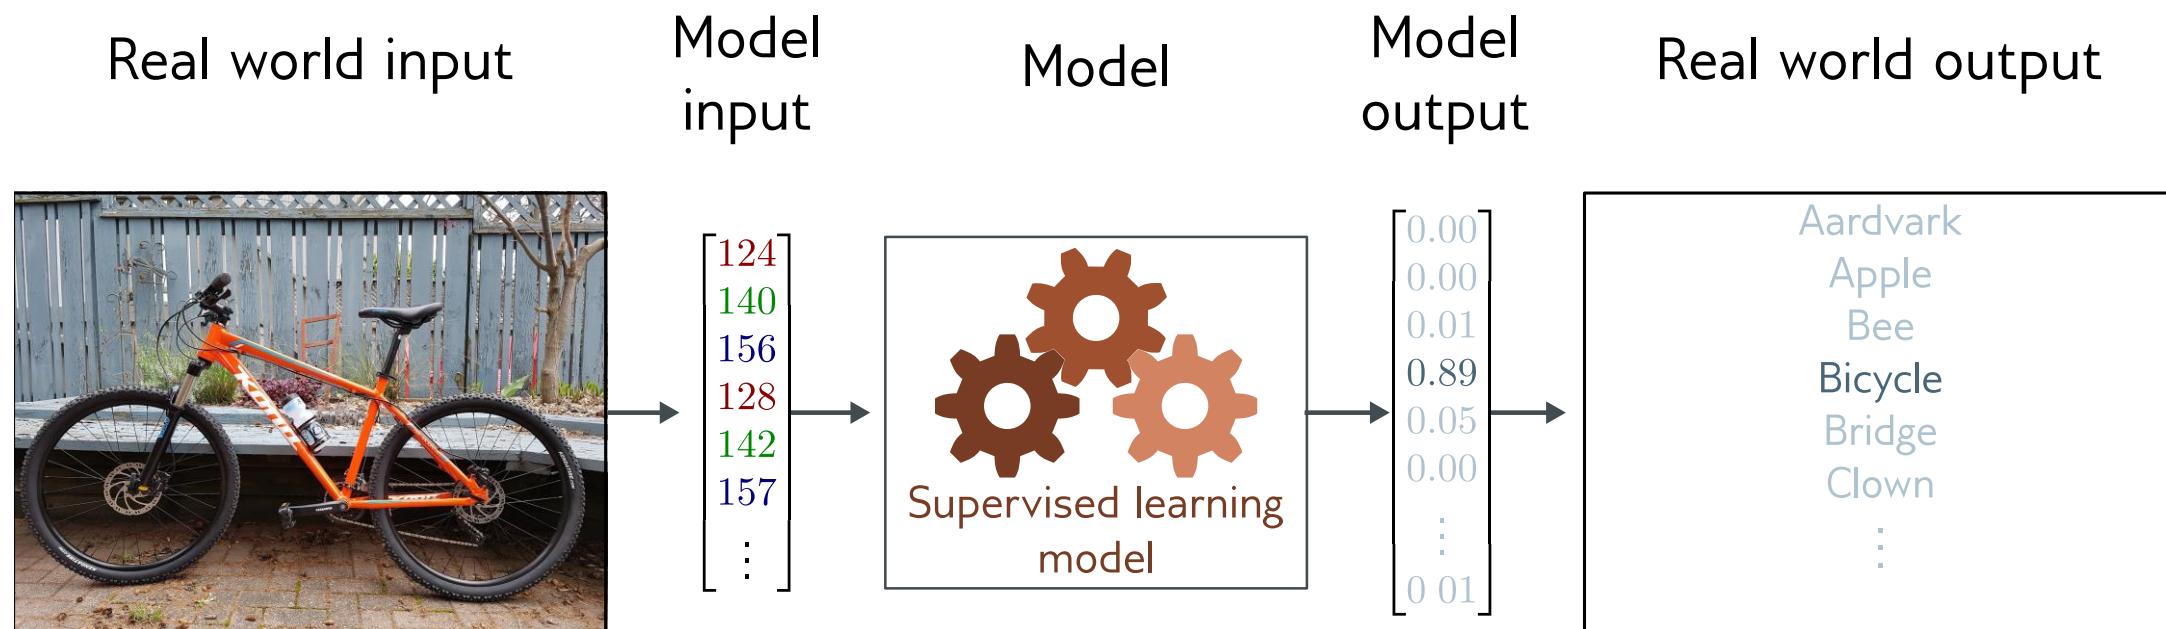
Conceptualizing the Binary Loss Function

$$L[\phi] = \sum_i -(1 - y_i) \log[1 - \text{sig}[f[x_i|\phi]]] - y_i \log[\text{sig}[f[x_i|\phi]]]$$

$$L[\phi] = \sum_i -(1 - y_i) \log \Pr[y_i = 0 | x_i] - y_i \log \Pr[y_i = 1 | x_i]$$

$$L[\phi] = \sum_i -\log \Pr[y_i | x_i]$$

Example 3: multiclass classification



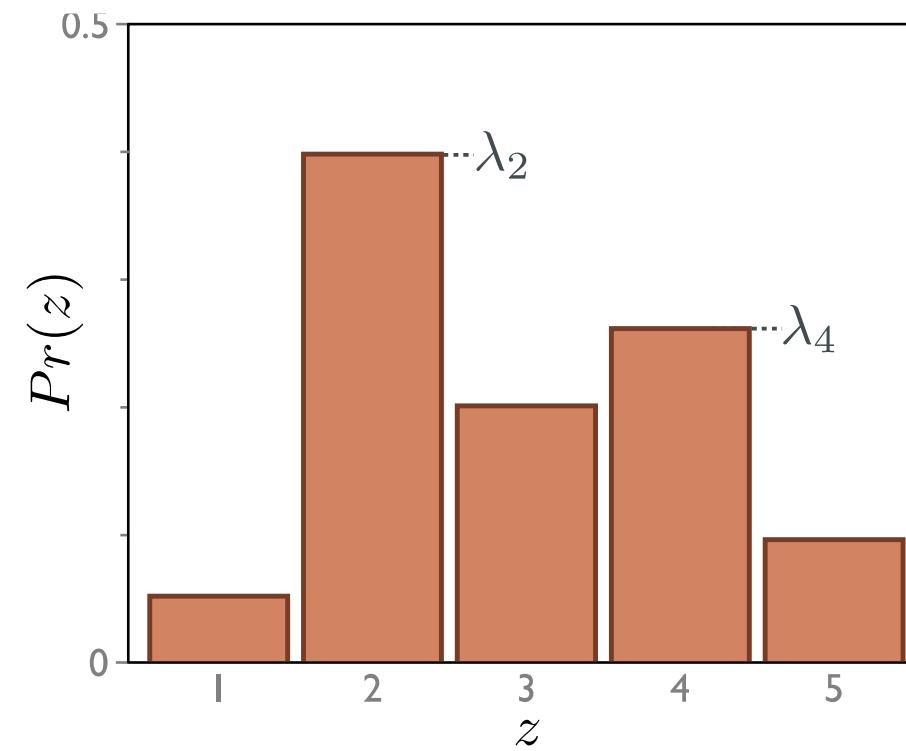
Goal: predict which of K classes $y \in \{1, 2, \dots, K\}$ the input x belongs to.

Example 3: multiclass classification

1. Choose a suitable probability distribution $Pr(\mathbf{y}|\boldsymbol{\theta})$ that is defined over the domain of the predictions \mathbf{y} and has distribution parameters $\boldsymbol{\theta}$.

- Domain: $y \in \{1, 2, \dots, K\}$
- **Categorical distribution**
- K parameters $\lambda_k \in [0, 1]$
- $\sum_k \lambda_k = 1$

$$Pr(y = k) = \lambda_k$$



Example 3: multiclass classification

- Set the machine learning model $\mathbf{f}[\mathbf{x}, \phi]$ to predict one or more of these parameters so $\theta = \mathbf{f}[\mathbf{x}, \phi]$ and $Pr(\mathbf{y}|\theta) = Pr(\mathbf{y}|\mathbf{f}[\mathbf{x}, \phi])$.

Problem:

- Output of neural network can be anything
- Parameters $\lambda_k \in [0,1]$, sum to one

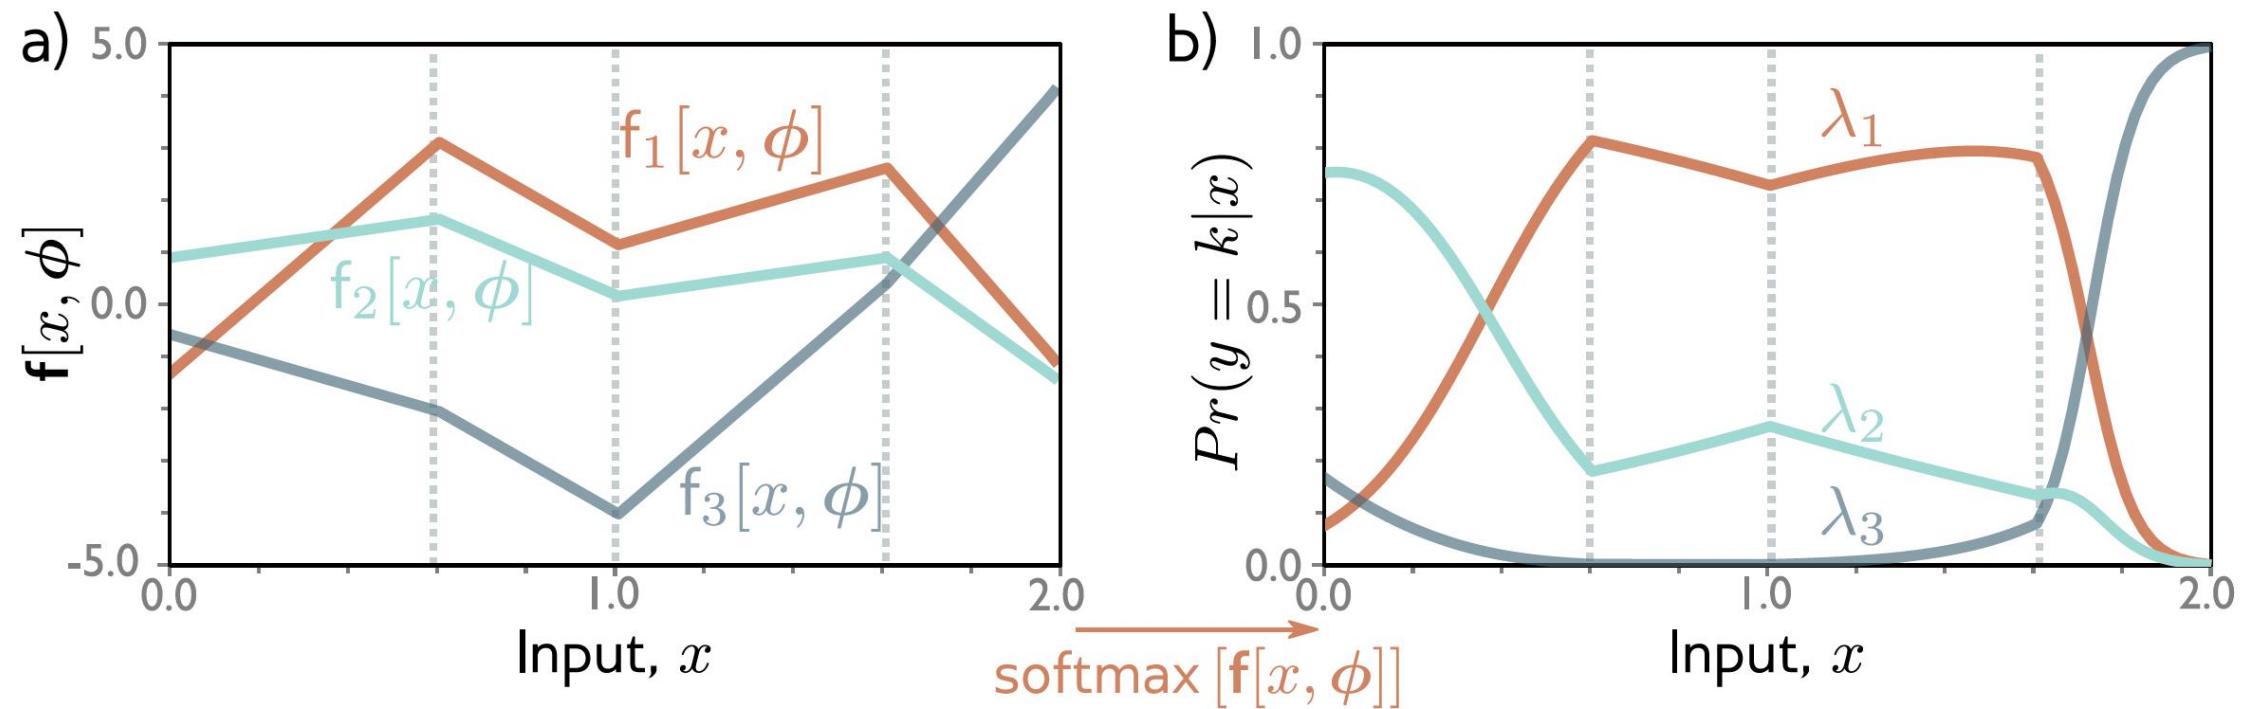
Solution:

- Pass through function that maps “anything” to $[0,1]$ and sums to one

$$\text{softmax}_k[\mathbf{z}] = \frac{\exp[z_k]}{\sum_{k'=1}^K \exp[z_{k'}]}$$

$$Pr(y = k | \mathbf{x}) = \text{softmax}_k[\mathbf{f}[\mathbf{x}, \phi]]$$

Example 3: multiclass classification



$$Pr(y = k|\mathbf{x}) = \text{softmax}_k [\mathbf{f}[\mathbf{x}, \phi]]$$

Example 3: multiclass classification

- To train the model, find the network parameters $\hat{\phi}$ that minimize the negative log-likelihood loss function over the training dataset pairs $\{\mathbf{x}_i, \mathbf{y}_i\}$:

$$\hat{\phi} = \operatorname{argmin}_{\phi} [L[\phi]] = \operatorname{argmin}_{\phi} \left[- \sum_{i=1}^I \log \left[Pr(\mathbf{y}_i | \mathbf{f}[\mathbf{x}_i, \phi]) \right] \right]. \quad (5.7)$$

$$L[\phi] = - \sum_{i=1}^I \log [\text{softmax}_{y_i} [\mathbf{f}[\mathbf{x}_i, \phi]]]$$

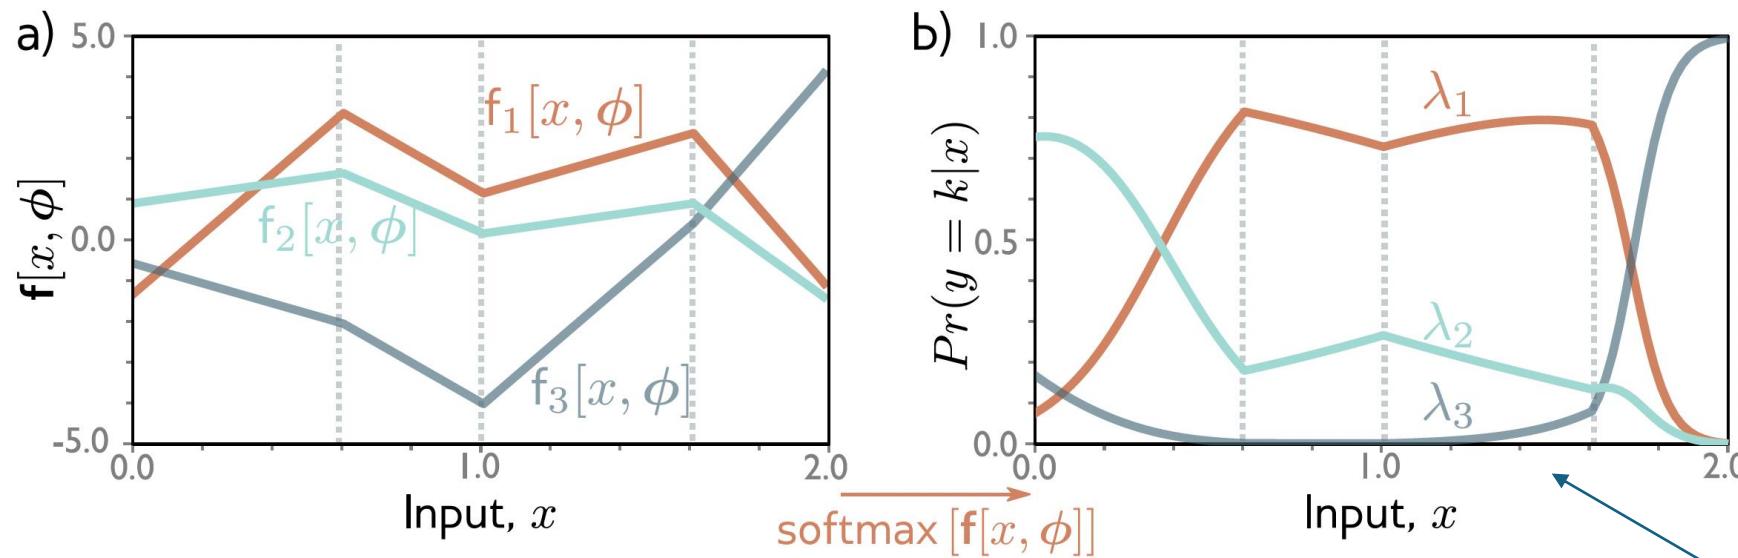
$$\text{softmax}_k[\mathbf{z}] = \frac{\exp[z_k]}{\sum_{k'=1}^K \exp[z_{k'}]}$$

$$= - \sum_{i=1}^I f_{y_i} [\mathbf{x}_i, \phi] - \log \left[\sum_{k=1}^K \exp [f_k [\mathbf{x}_i, \phi]] \right]$$

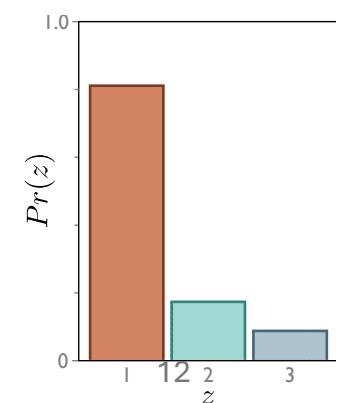
Multiclass cross-entropy loss

Example 3: multiclass classification

4. To perform inference for a new test example \mathbf{x} , return either the full distribution $Pr(\mathbf{y}|\mathbf{f}[\mathbf{x}, \hat{\phi}])$ or the maximum of this distribution.



Choose the class with the largest probability
We also get probability or “confidence”



Any questions?

Multiple outputs

- Treat each output y_d as *independent*:

$$Pr(\mathbf{y}|\mathbf{f}[\mathbf{x}_i, \boldsymbol{\phi}]) = \prod_d Pr(y_d|\mathbf{f}_d[\mathbf{x}_i, \boldsymbol{\phi}])$$

where $\mathbf{f}_d[\mathbf{x}, \boldsymbol{\phi}]$ is the d^{th} set of network outputs

- Negative log likelihood becomes sum of terms:

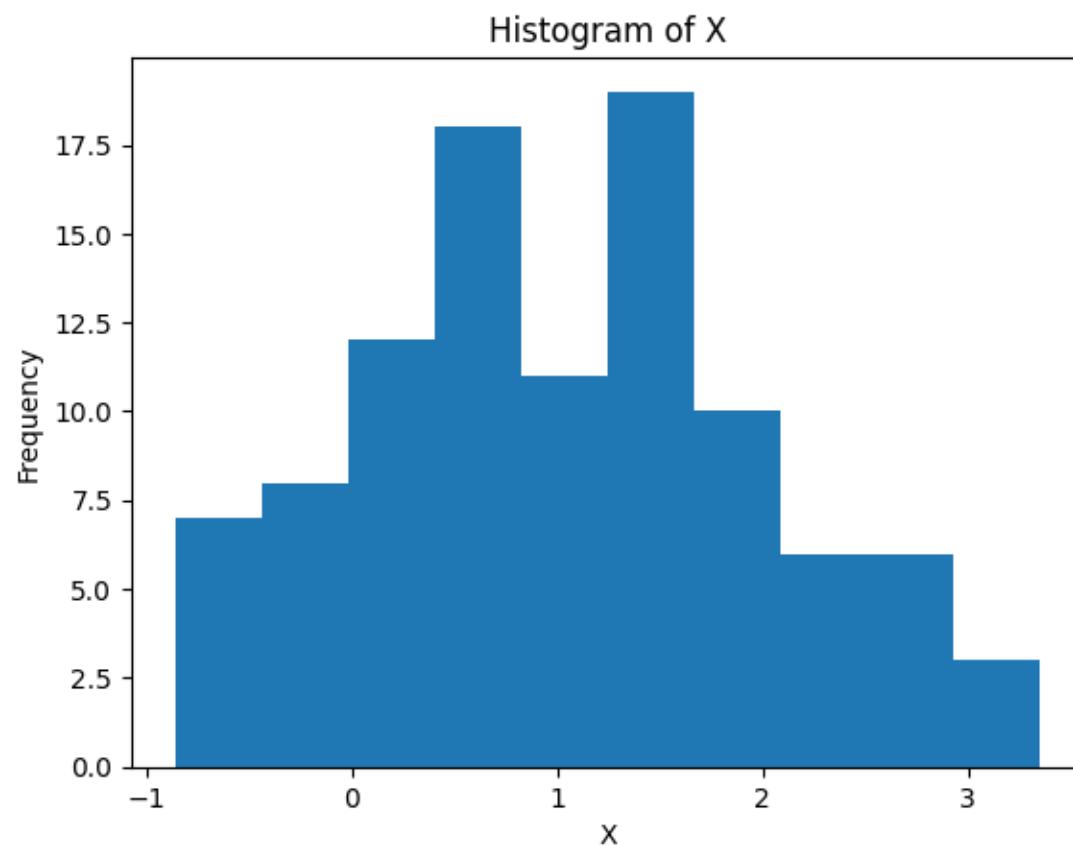
$$L[\boldsymbol{\phi}] = - \sum_{i=1}^I \log [Pr(\mathbf{y}|\mathbf{f}[\mathbf{x}_i, \boldsymbol{\phi}])] = - \sum_{i=1}^I \sum_d \log [Pr(y_{id}|\mathbf{f}_d[\mathbf{x}_i, \boldsymbol{\phi}])]$$

d^{th} output of the i^{th} training example

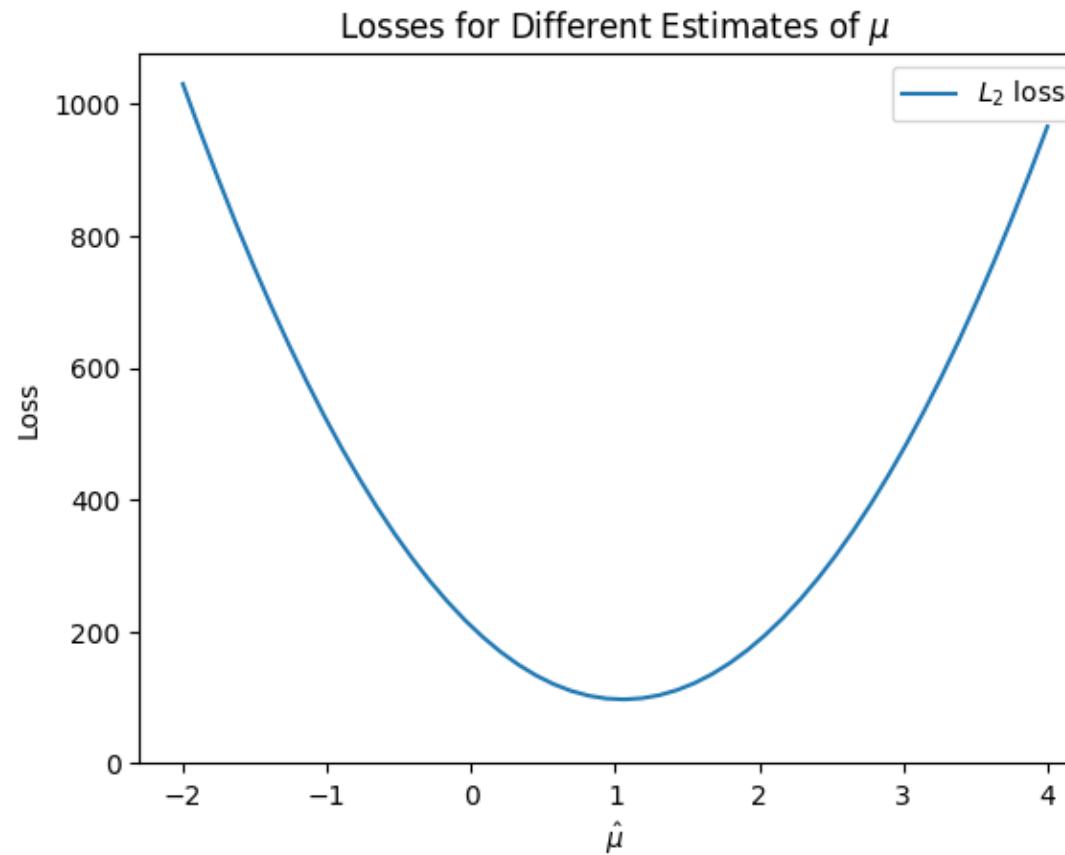
Any questions?

An Example of Gradient Descent

- X is 100 samples from a normal distribution.
 - What were the parameters of that normal distribution?
 - What was the mean of that normal distribution?



Visualizing Gradient Descent

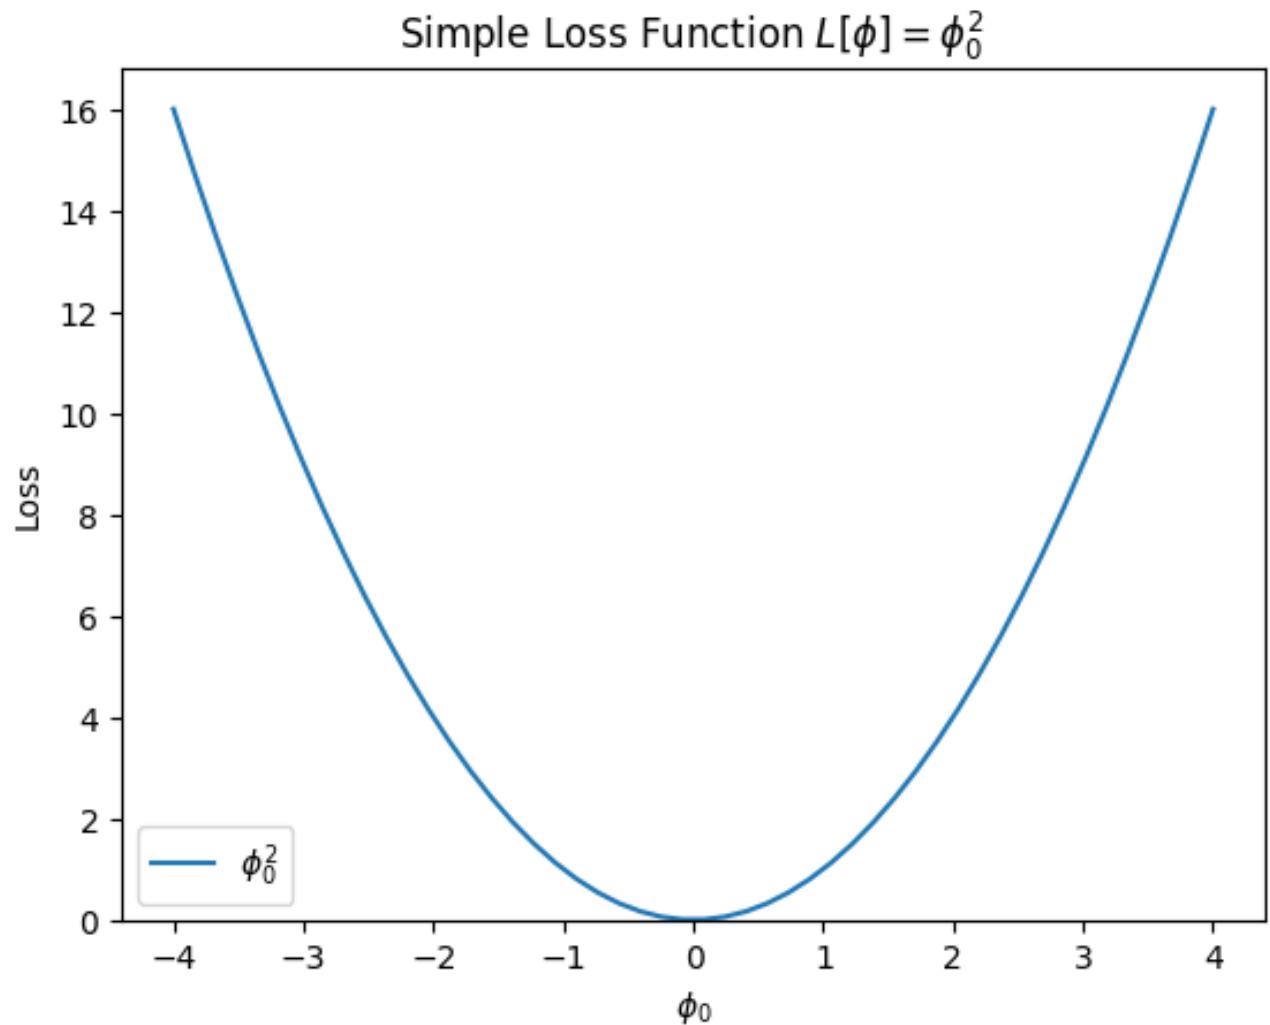


Basics of Gradient Descent

- Given current set of parameters ϕ_t ,
 - Calculate all **partial derivatives** $\frac{\partial L[\phi]}{\partial \phi_i}$ based on current parameters ϕ_t .
 - The vector of these $\frac{\partial L[\phi]}{\partial \phi_i}$ is the **gradient** of the loss function $\nabla L[\phi]$.
 - Update $\phi_{t+1} = \phi_t - \alpha \nabla L[\phi]$
where α is the **learning rate**.

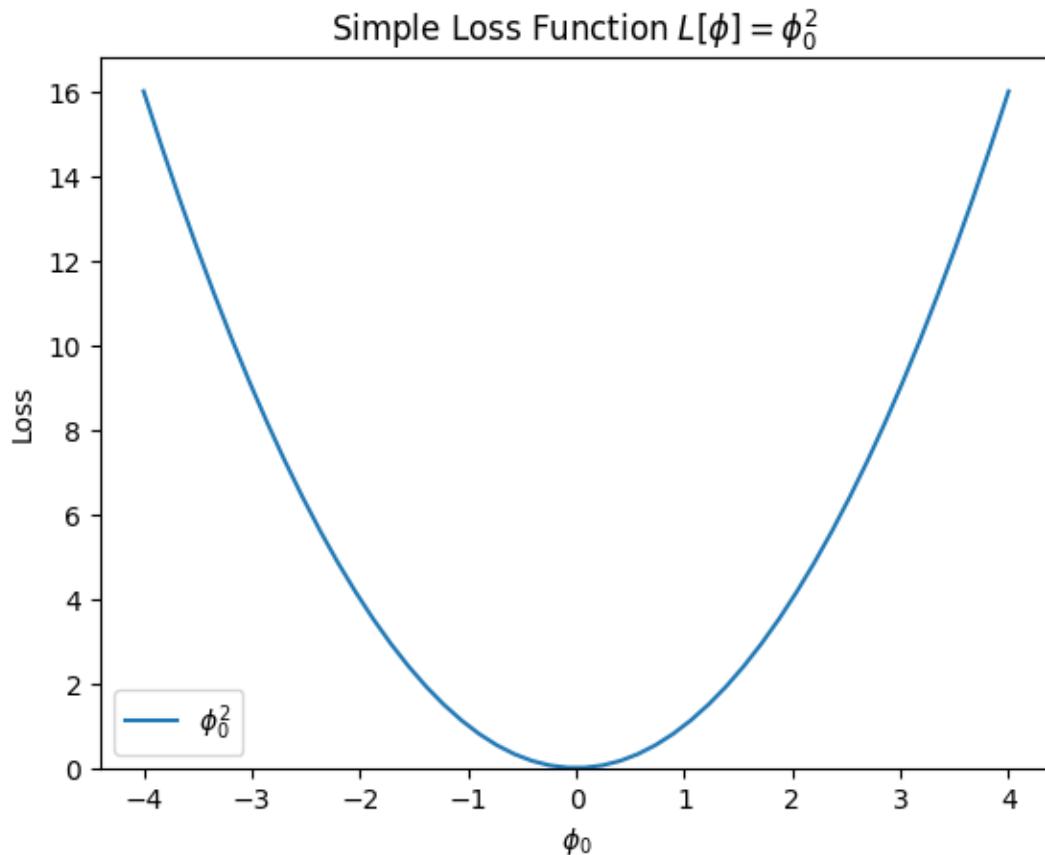
What should the learning rate be?

- Try $\alpha = 1$.
- Too small, and it takes many steps to get close.
- Too big, and it overshoots.



Convex Loss Functions

- Generally, a lot easier to optimize...
- With gradient descent, main issue is not overshooting minimum too much.



Non-Convex Loss Functions

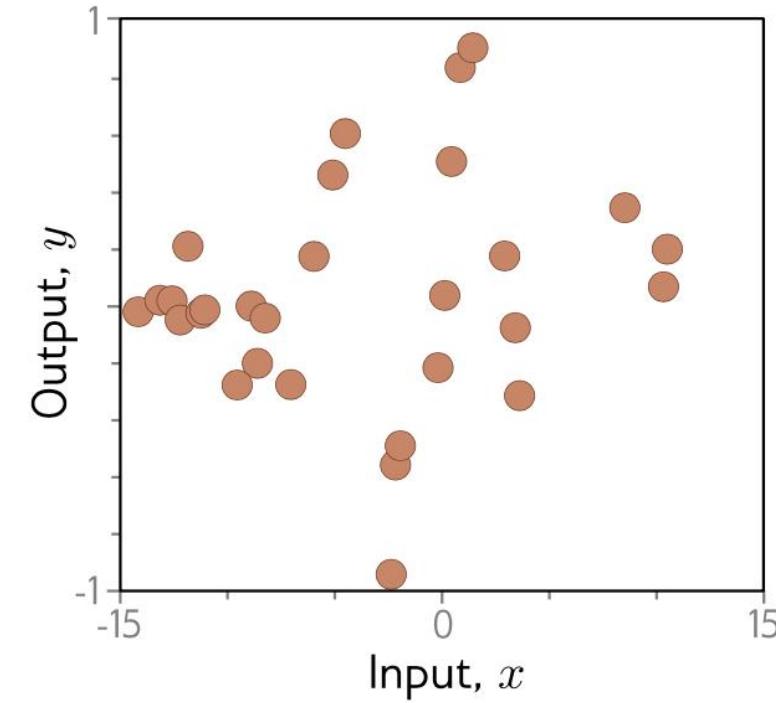
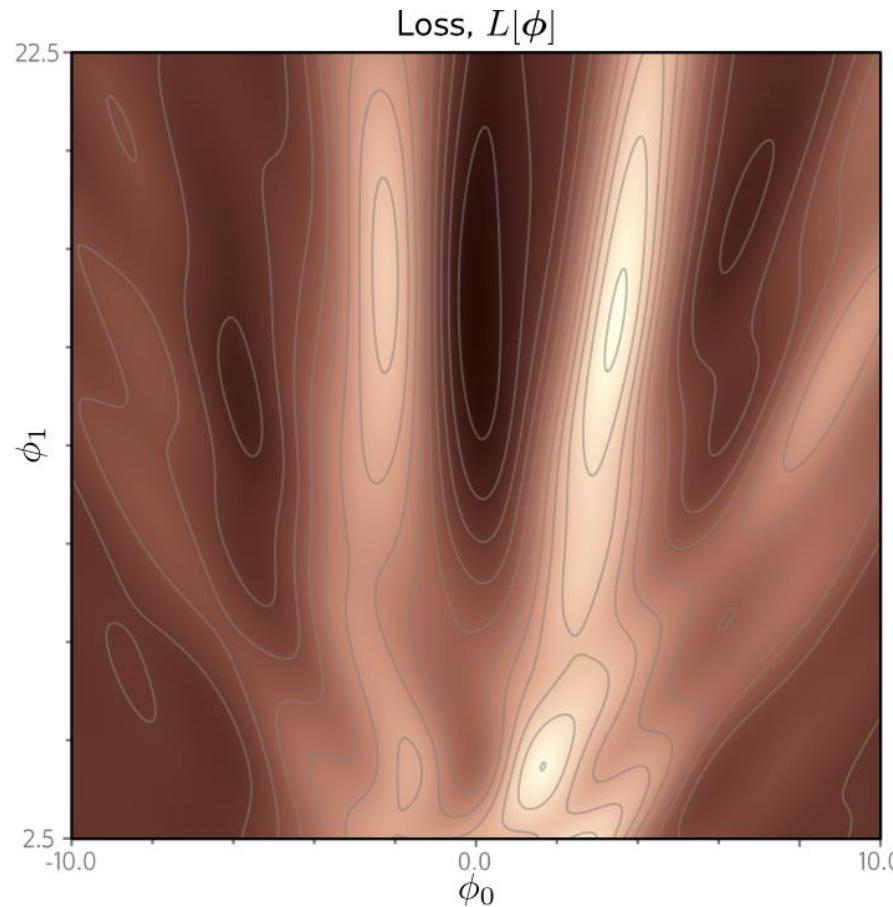
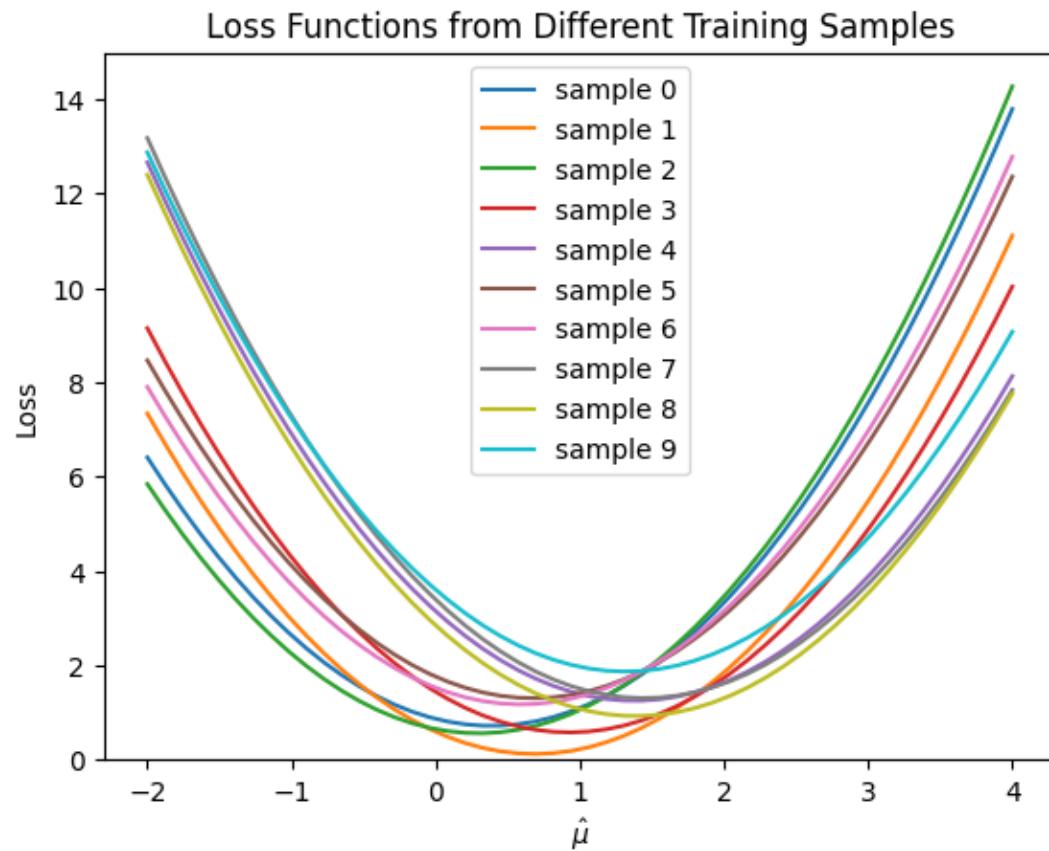


Image Source: Understanding Deep Learning, via <https://udlbook.github.io/udlfigures/>

Any questions?

Gradient Descent as a Statistical Process

- Our training data is a sample of the whole population.
 - Different training samples yield different training loss functions.



Loss Functions for Different Training Samples

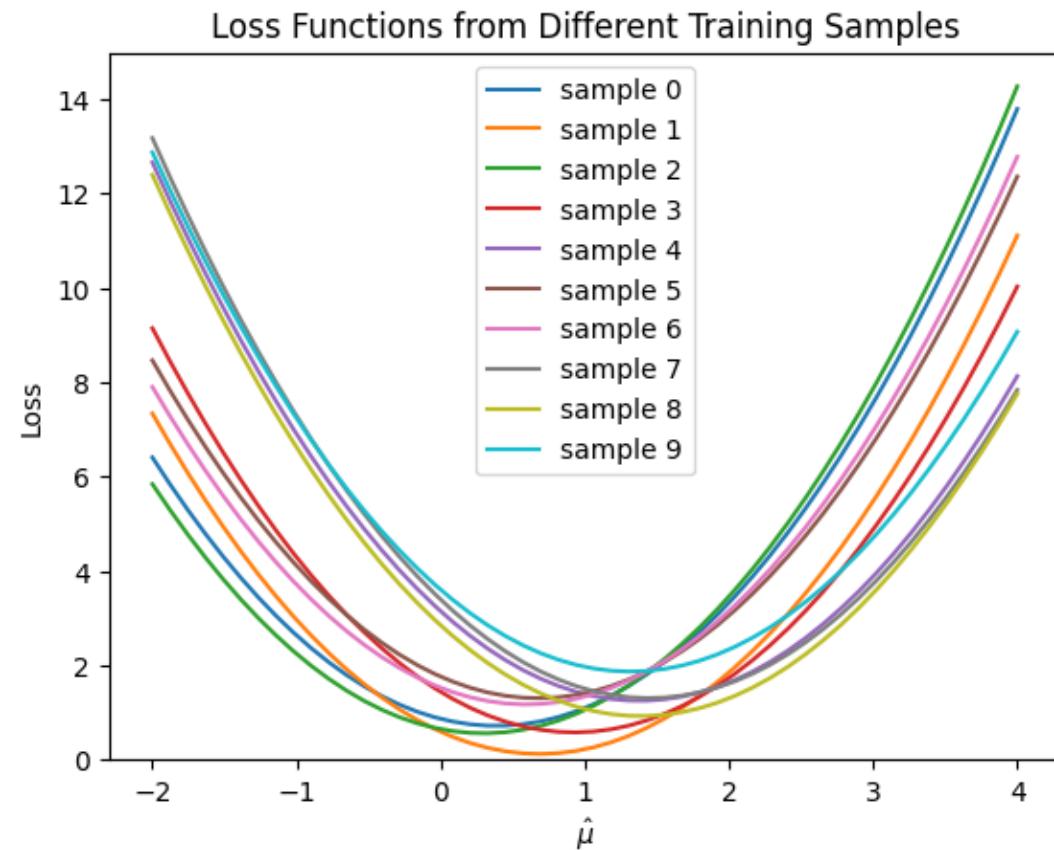
- If we collect different training data sets, will we get different models?

Loss Functions for Samples of the Training Set

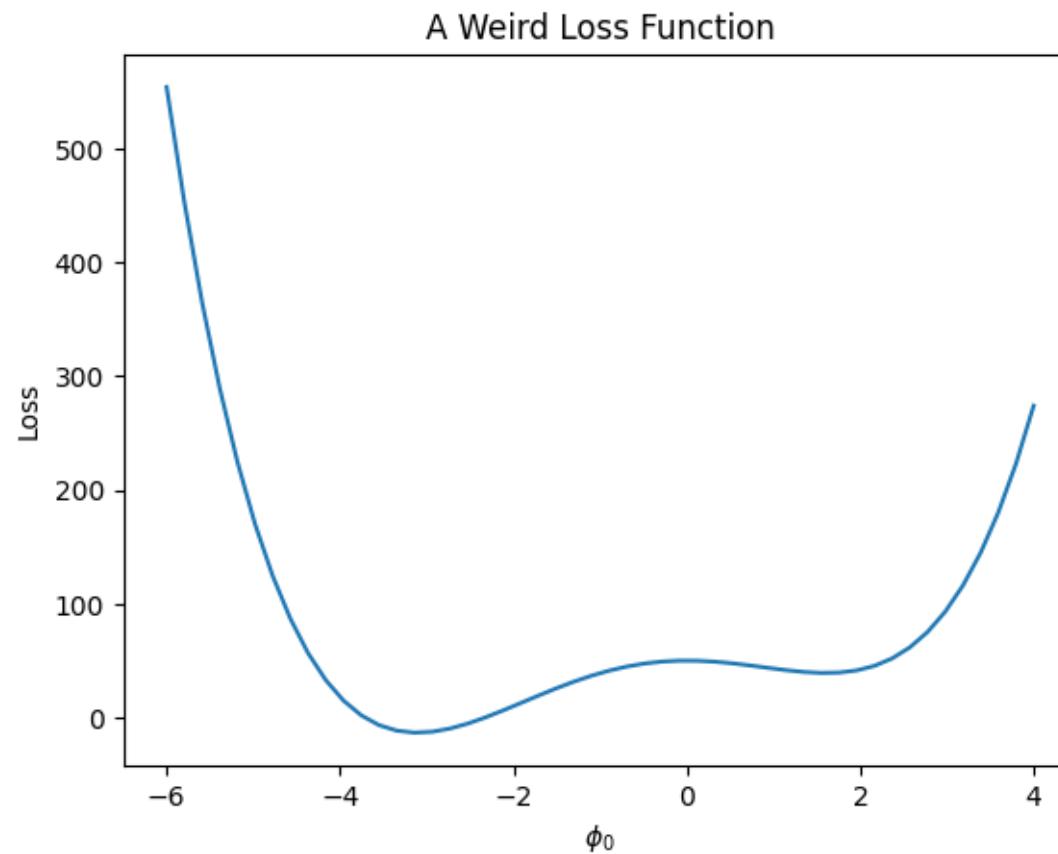
- If we sample the training data, will we get different models?

Comparing Models with Different Training Samples

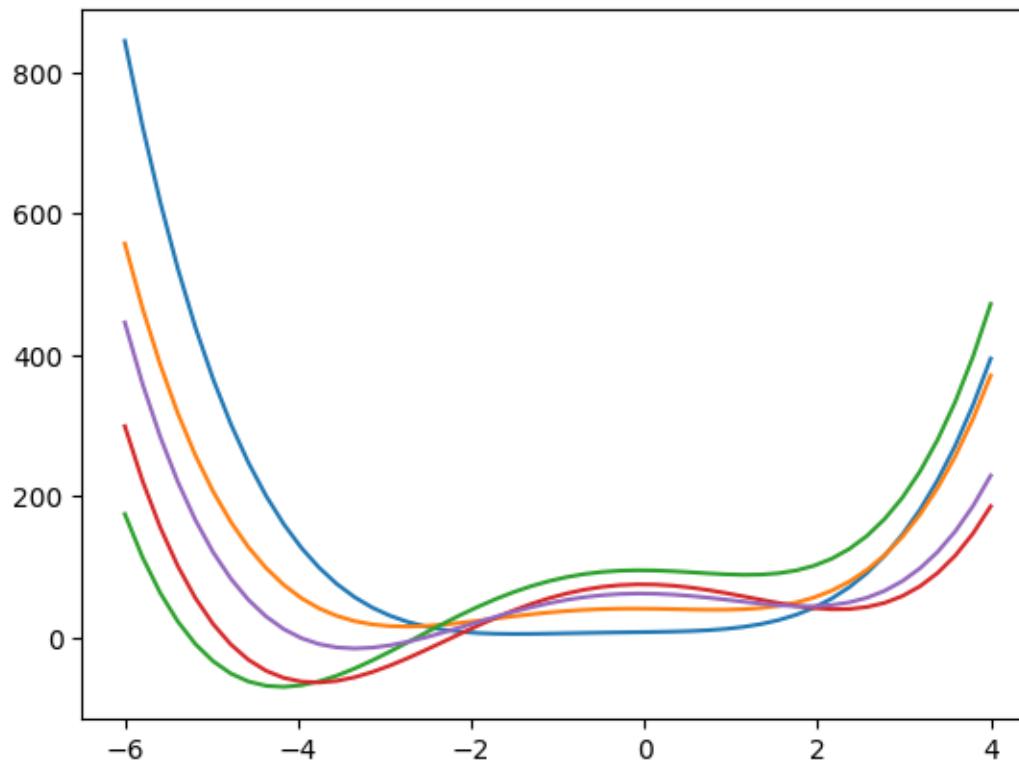
- How far apart are the models of these samples?
- Where do they agree and disagree?



A Weird Loss Function



Local Minima vs Samples of the Training Set



Stochastic Gradient Descent

Idea: Run gradient descent with “mini batches” instead of the full training set.

- E.g. pick a random partition of data into 10 equal-sized batches.
- One epoch = running through all the data once.
 - Vanilla gradient → one parameter update.
 - Stochastic gradient descent → one parameter update per mini batch.

Variation in Sampled Gradients

- Expected mini batch gradient = whole training set gradient.
 - On average, they agree.
 - But with noise from sampling.
- But remember, just taking one step with each mini batch.
 - Not optimizing to mini batch minimum loss.

Local Minima vs Stochastic Gradient Descent

- When far from a local minima, mini batches tend to agree on gradient direction.
- When close to a local minima, mini batches disagree more.
 - Sampling noise.
 - Explore the flat area around the minima.

Speed of Stochastic Gradient Descent

- How fast is this compared to vanilla gradient descent?

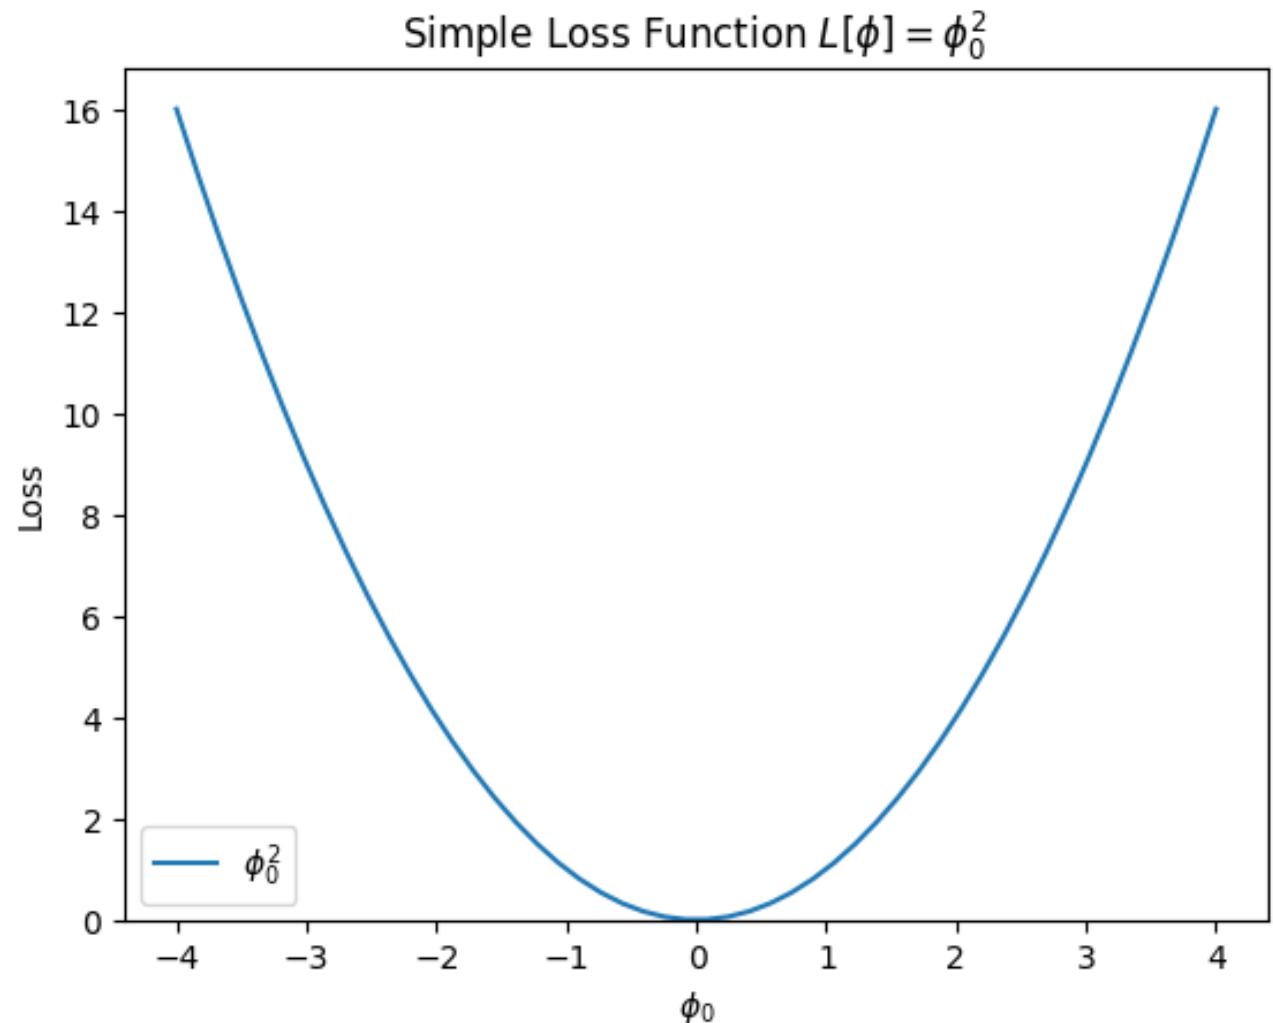
Any questions?

Gradient Descent as a Universal Algorithm

- What's the catch?

How do we pick Learning Rate?

- Remember, $\alpha = 1$ gives an infinite loop.
- Also, be impatient.



Really Bad Linear Regression

- $f(x) = f_1 \left(f_2 \left(f_3 \left(f_4(x) \right) \right) \right)$
- $f_1(x) = a_1x + b_1$
- $f_2(x) = a_2x + b_2$
- $f_3(x) = a_3x + b_3$
- $f_4(x) = a_4x + b_4$
- $f(x)$ is just a linear function?

Really Bad Linear Regression (part 2)

- $f(x) = f_1 \left(f_2 \left(f_3 \left(f_4(x) \right) \right) \right)$
- Initialize all parameters to zero.
- What are the gradients?
- $f_1(x) = a_1x + b_1$
- $f_2(x) = a_2x + b_2$
- $f_3(x) = a_3x + b_3$
- $f_4(x) = a_4x + b_4$
- $f(x)$ is just a linear function?

Really Bad Linear Regression (part 3)

- $f(x) = f_1 \left(f_2 \left(f_3 \left(f_4(x) \right) \right) \right)$
- Initialize all parameters to 100.
- What are the gradients?
- $f_1(x) = a_1x + b_1$
- $f_2(x) = a_2x + b_2$
- $f_3(x) = a_3x + b_3$
- $f_4(x) = a_4x + b_4$
- $f(x)$ is just a linear function?

Really Bad Linear Regression (part 4)

- $f(x) = f_1 \left(f_2 \left(f_3 \left(f_4(x) \right) \right) \right)$
- $f_1(x) = a_1 x + b_1$
- $f_2(x) = a_2 x + b_2$
- $f_3(x) = a_3 x + b_3$
- $f_4(x) = a_4 x + b_4$
- $f(x)$ is just a linear function?
- We will see both these problems with neural networks if we use the wrong initialization.

Any questions?