Deep Learning for Data Science DS 542
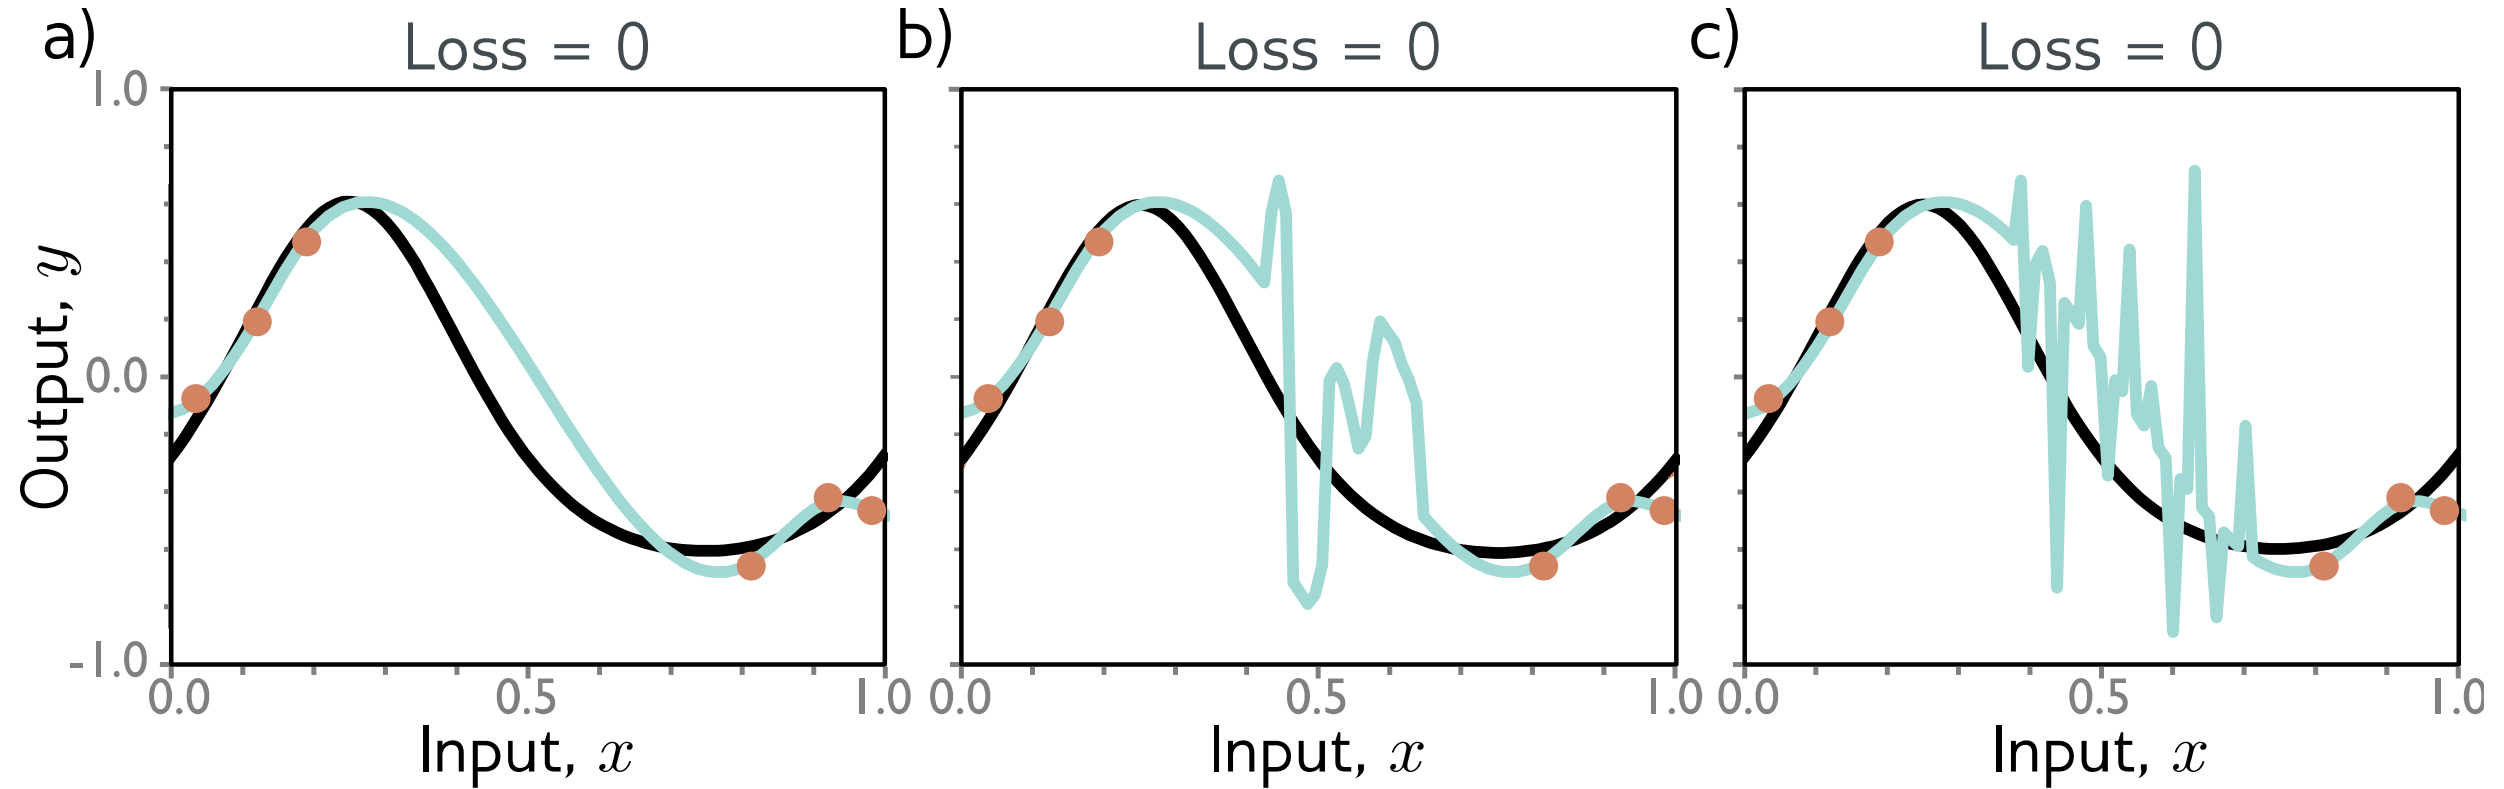
https://dl4ds.github.io/fa2025/

Regularization

Regularization

- Why is there a generalization gap between training and test data?
  - Overfitting (model describes statistical peculiarities)
  - Model unconstrained in areas where there are no training examples
- Regularization = methods to reduce the generalization gap
- Technically means adding terms to loss function
- But colloquially means any method (hack) to reduce gap between training and test data

How to bias for smoothness?

- All of these solutions are equivalent in terms of loss.
- Why should the model choose the smooth solution?
- Tendency of model to choose one solution over another is inductive bias

# Interpolating Polynomials Overfit A Lot

If our "before" model was good, why do we change it so much?

Plan for Today

- Reproducing Double Descent Demystified
- Explicit regularization
- Implicit regularization
- Early stopping
- Ensembling
- Dropout
- Adding noise
- Transfer learning, multi-task learning, self-supervised learning
- Data augmentation

Toy Setup

https://colab.research.google.com/drive/1p6DYrIO6p1zGePiD3NJ9JPfkgsP70ojM

Reproducing Double Descent Demystified...

Based on https://iclr-blogposts.github.io/2024/blog/double-descent-demystified/

Regularizing from the Beginning

Any Questions?

???

Moving on
- Reproducing Double Descent Demystified
- Explicit regularization
- Implicit regularization
- Early stopping
- Ensembling
- Dropout
- Adding noise
- Transfer learning, multi-task learning, self-supervised learning
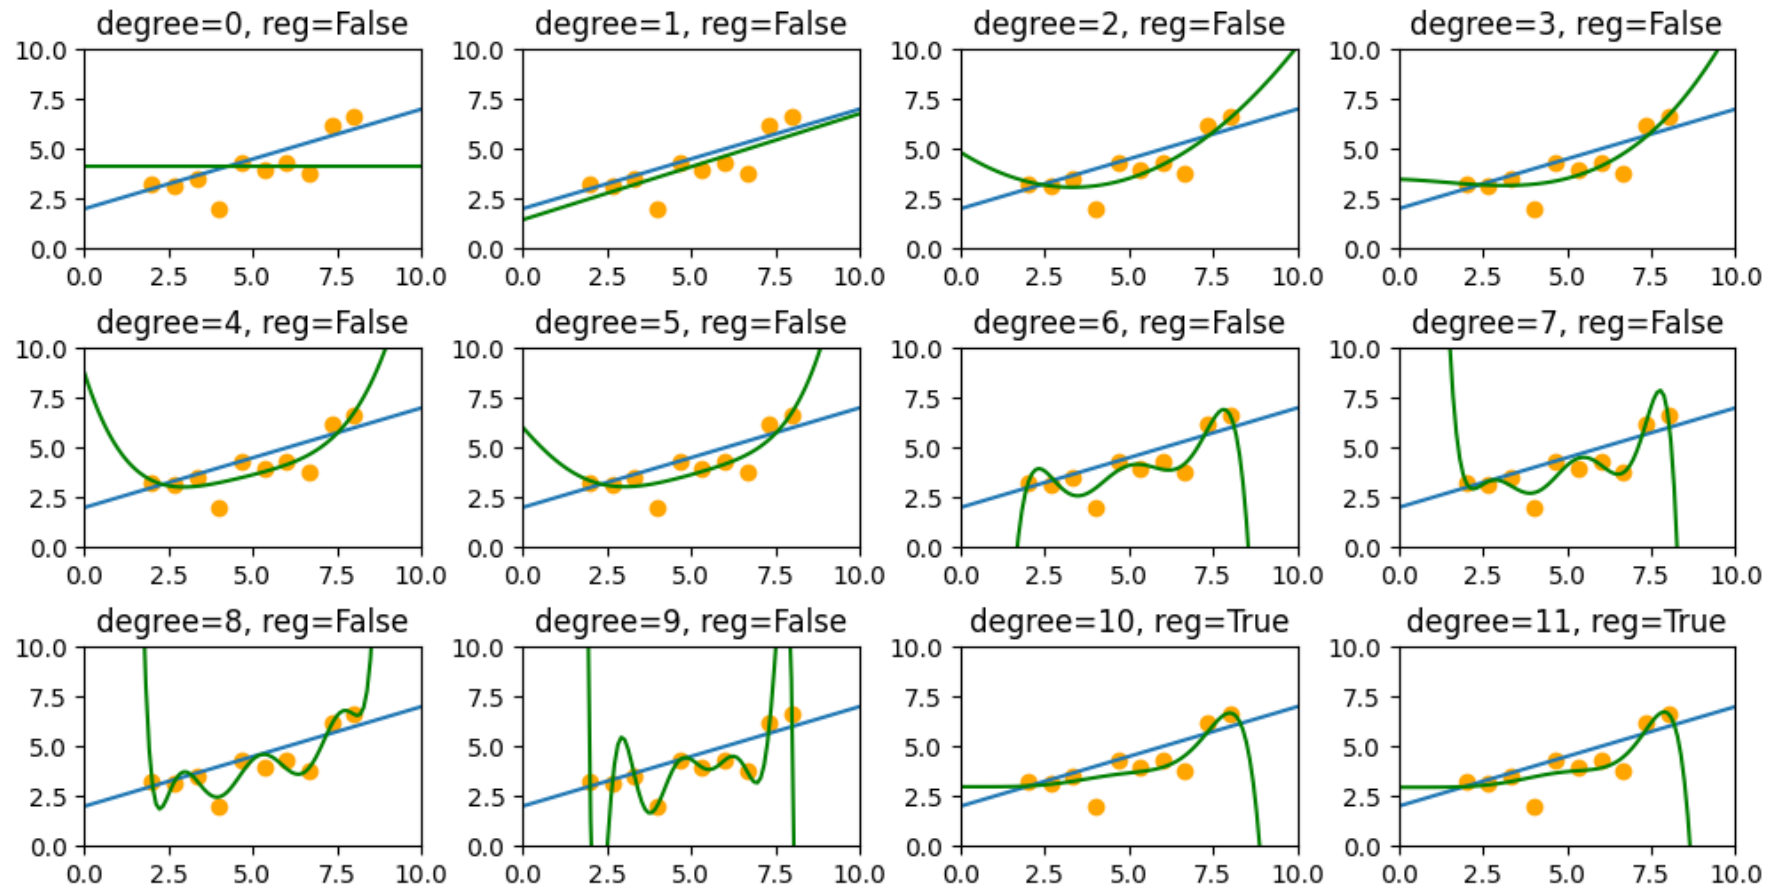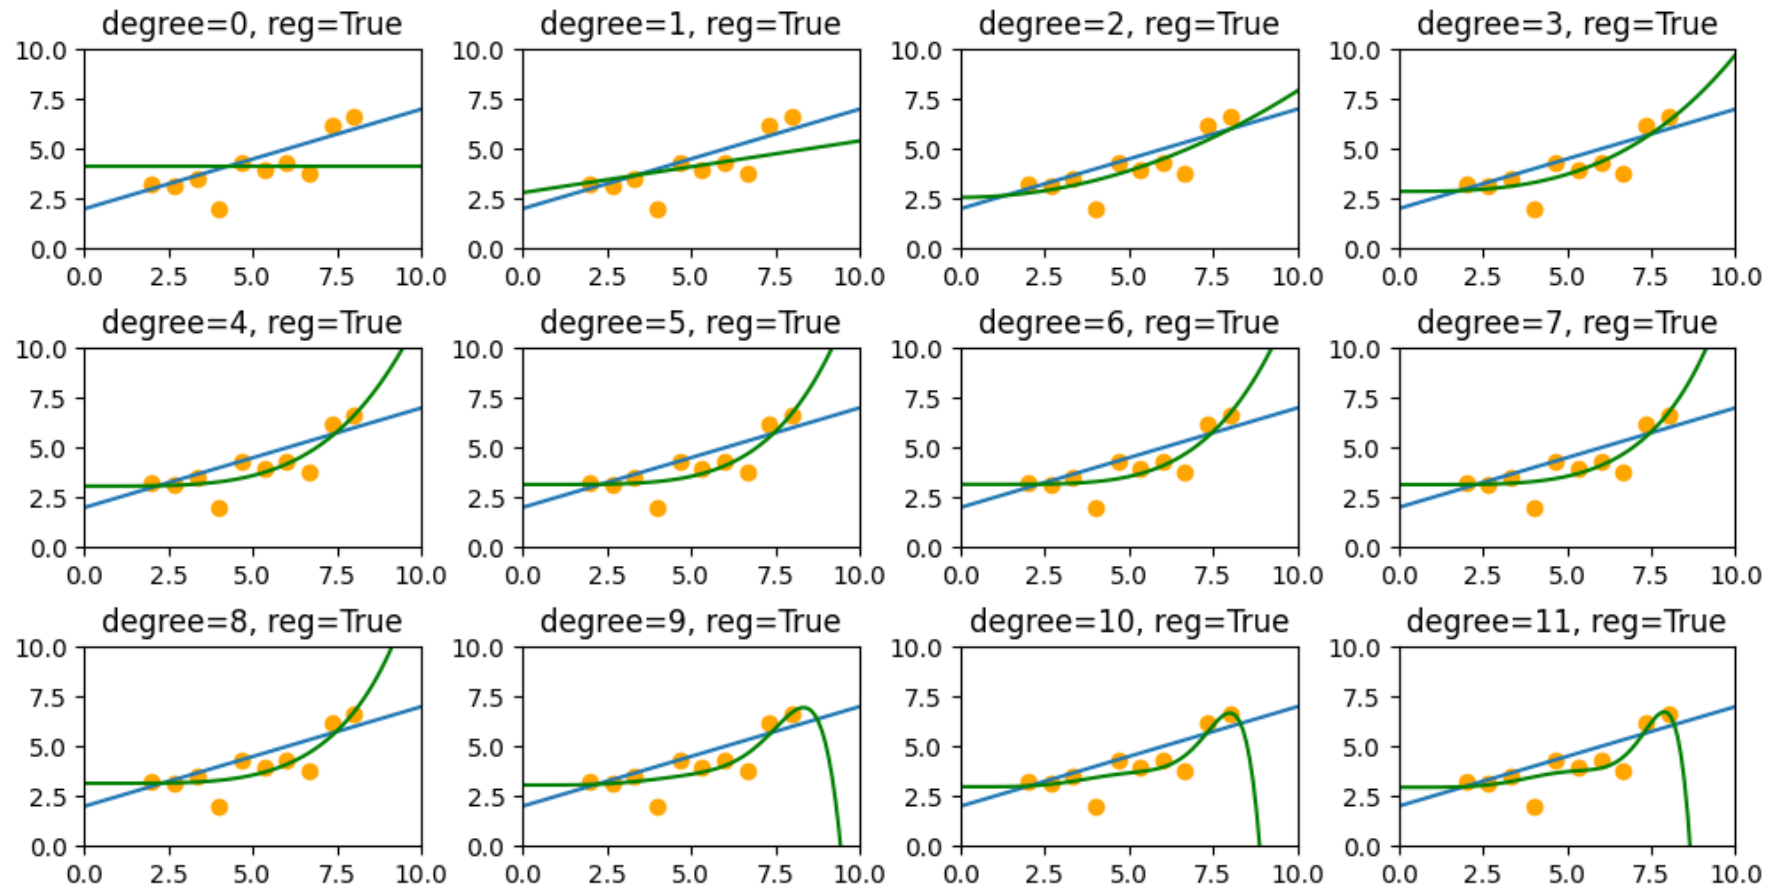- Data augmentation

Explicit regularization

- Standard loss function:

$$\hat{\phi} = \underset{\phi}{\mathrm{argmin}}\big[\mathrm{L}[\phi]\big]$$

$$= \underset{\phi}{\mathrm{argmin}}\left[\sum_{i=1}^{I}\ell_i[\mathbf{x}_i, \mathbf{y}_i]\right]$$

Explicit regularization

- Standard loss function:

$$\hat{\phi} = \underset{\phi}{\operatorname{argmin}}\left[\mathrm{L}[\phi]\right]$$

$$= \underset{\phi}{\operatorname{argmin}}\left[\sum_{i=1}^{I} \ell_i[\mathbf{x}_i, \mathbf{y}_i]\right]$$

- Regularization adds an extra term

$$\hat{\phi} = \underset{\phi}{\operatorname{argmin}}\left[\sum_{i=1}^{I} \ell_i[\mathbf{x}_i, \mathbf{y}_i] + \lambda \cdot \mathrm{g}[\phi]\right]$$

Explicit regularization

- Standard loss function:

$$\hat{\phi} = \underset{\phi}{\text{argmin}}\left[L[\phi]\right]$$

$$= \underset{\phi}{\text{argmin}}\left[\sum_{i=1}^{I} \ell_i[\mathbf{x}_i, \mathbf{y}_i]\right]$$

- Regularization adds an extra term

$$\hat{\phi} = \underset{\phi}{\text{argmin}}\left[\sum_{i=1}^{I} \ell_i[\mathbf{x}_i, \mathbf{y}_i] + \lambda \cdot g[\phi]\right]$$

- Where $g[\phi]$ is smaller for preferred parameters
- $\lambda > 0$ controls the strength of influence

Explicit regularization

a)

Loss function for Gabor model of Lecture 6 and Chapter 6.

⬤ denotes local minima

Explicit regularization

a) Loss

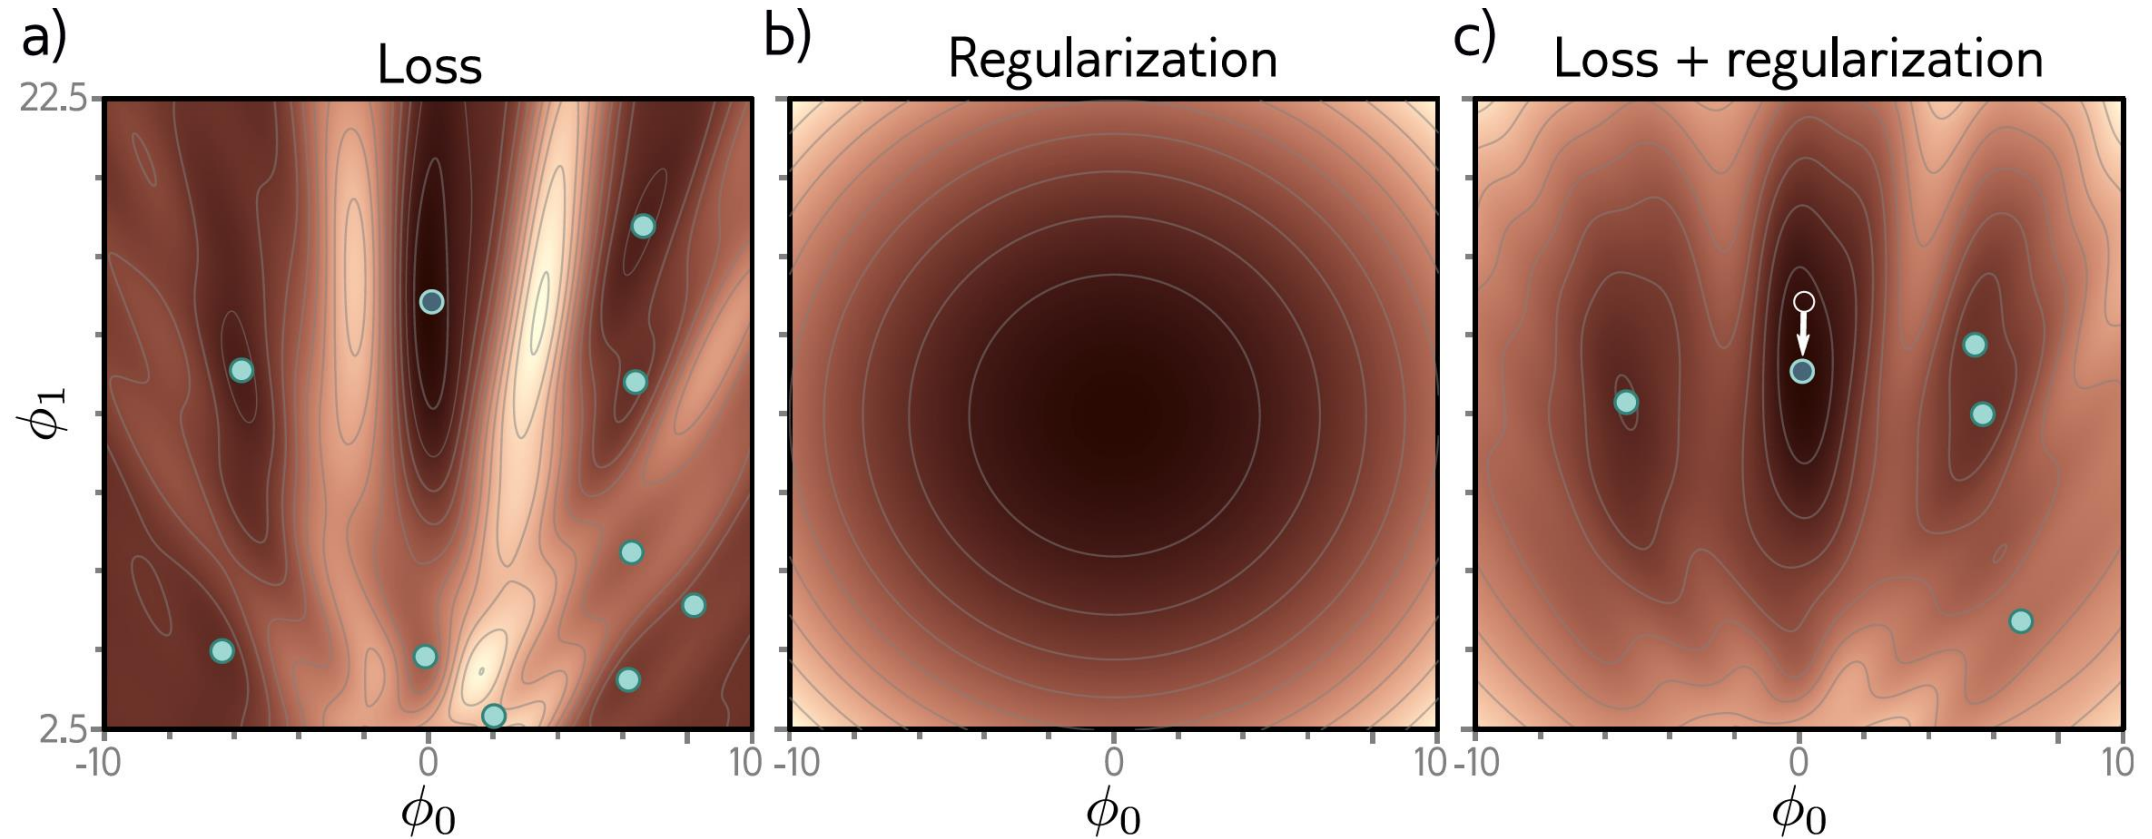b) Regularization

Example of a regularization function that prefers parameters close to 0.

Explicit regularization

Fewer local minima and the absolute minimum has moved.

a) Loss

b) Regularization

c) Loss + regularization

○ denotes local minima

Probabilistic interpretation

- Maximum likelihood:

$$\hat{\boldsymbol{\phi}} = \operatorname*{argmax}_{\boldsymbol{\phi}} \left[\prod_{i=1}^{I} Pr(\mathbf{y}_i | \mathbf{x}_i, \boldsymbol{\phi}) \right]$$

- Regularization is equivalent to adding a prior over parameters

$$\hat{\boldsymbol{\phi}} = \operatorname*{argmax}_{\boldsymbol{\phi}} \left[\prod_{i=1}^{I} Pr(\mathbf{y}_i | \mathbf{x}_i, \boldsymbol{\phi}) Pr(\boldsymbol{\phi}) \right]$$

Maximum a posteriori or MAP criterion

... what you know about parameters *before* seeing the data

Equivalence

- Explicit regularization:

$$\hat{\phi} = \operatorname*{argmin}_{\phi} \left[\sum_{i=1}^{I} \ell_i[\mathbf{x}_i, \mathbf{y}_i] + \lambda \cdot \mathrm{g}[\phi] \right]$$

- Probabilistic interpretation:

$$\hat{\phi} = \operatorname*{argmax}_{\phi} \left[\prod_{i=1}^{I} Pr(\mathbf{y}_i | \mathbf{x}_i, \phi) Pr(\phi) \right]$$

- Converting to Negative Log Likelihood (e.g. $-\log(\cdot)$):
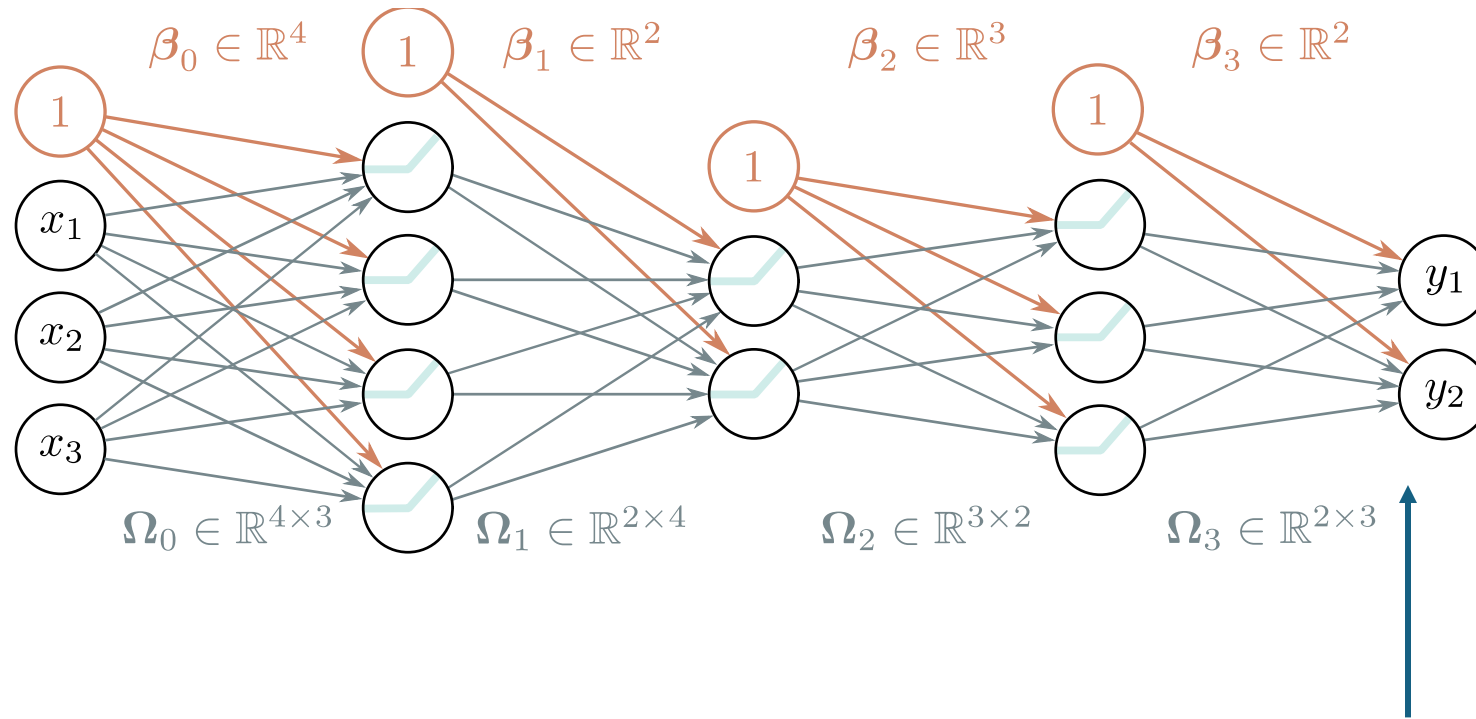
$$\lambda \cdot \mathrm{g}[\phi] = -\log[Pr(\phi)]$$

L2 Regularization

- Most common regularization is L2 regularization
- Favors smaller parameters (like in previous example)

$$\hat{\phi} = \operatorname*{argmin}_{\phi} \left[L[\phi, \{\mathbf{x}_i, \mathbf{y}_i\}] + \lambda \sum_j \phi_j^2 \right]$$

- Also called Tikhonov regularization, ridge regression
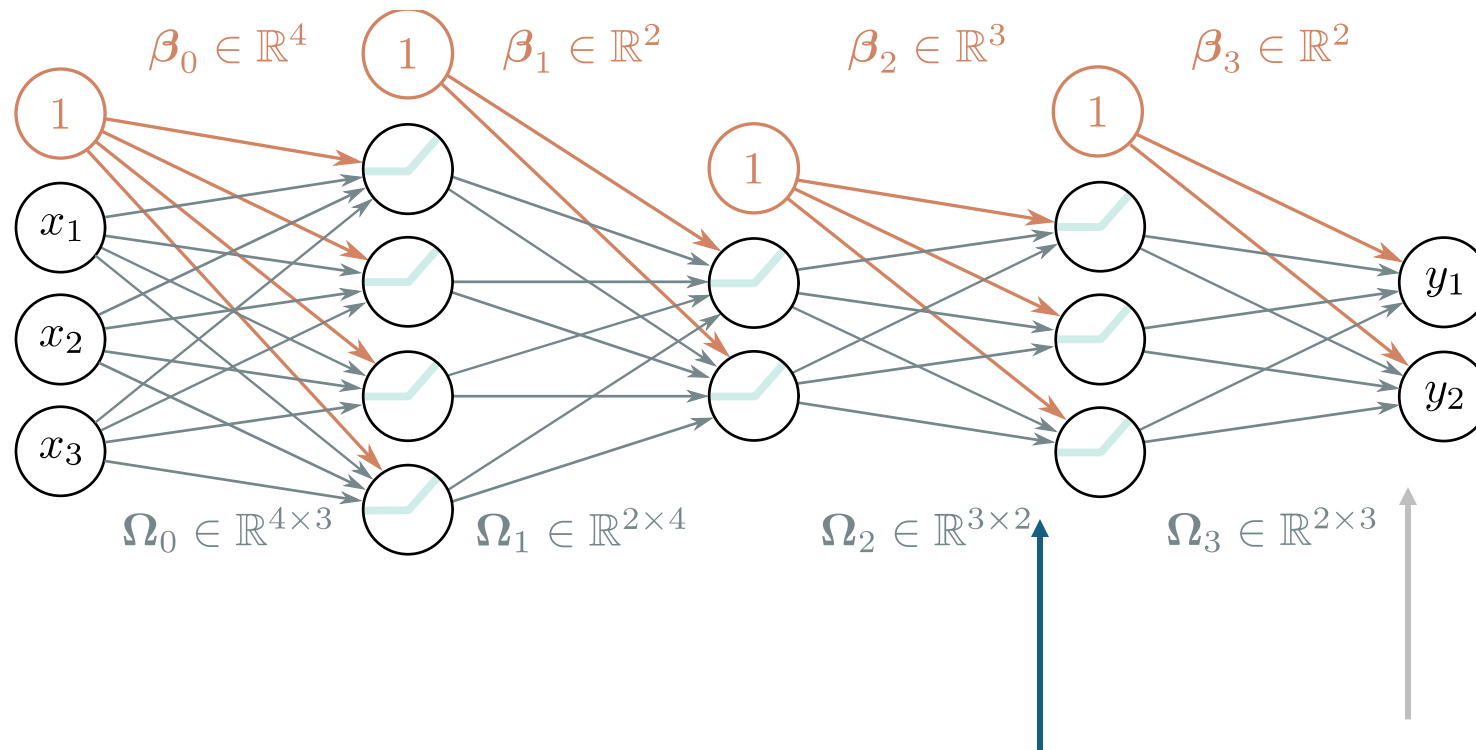- In neural networks, usually just for weights, and called weight decay

Why does L2 regularization help?

$\boldsymbol{\beta}_0 \in \mathbb{R}^4$ $\boldsymbol{\beta}_1 \in \mathbb{R}^2$ $\boldsymbol{\beta}_2 \in \mathbb{R}^3$ $\boldsymbol{\beta}_3 \in \mathbb{R}^2$

$\boldsymbol{\Omega}_0 \in \mathbb{R}^{4\times 3}$ $\boldsymbol{\Omega}_1 \in \mathbb{R}^{2\times 4}$ $\boldsymbol{\Omega}_2 \in \mathbb{R}^{3\times 2}$ $\boldsymbol{\Omega}_3 \in \mathbb{R}^{2\times 3}$

Outputs are weighted linear combination of last layer activations.

Smaller weights attenuate changes.

Why does L2 regularization help?

$\boldsymbol{\beta}_0 \in \mathbb{R}^4$ $\boldsymbol{\beta}_1 \in \mathbb{R}^2$ $\boldsymbol{\beta}_2 \in \mathbb{R}^3$ $\boldsymbol{\beta}_3 \in \mathbb{R}^2$

$\boldsymbol{\Omega}_0 \in \mathbb{R}^{4\times 3}$ $\boldsymbol{\Omega}_1 \in \mathbb{R}^{2\times 4}$ $\boldsymbol{\Omega}_2 \in \mathbb{R}^{3\times 2}$ $\boldsymbol{\Omega}_3 \in \mathbb{R}^{2\times 3}$
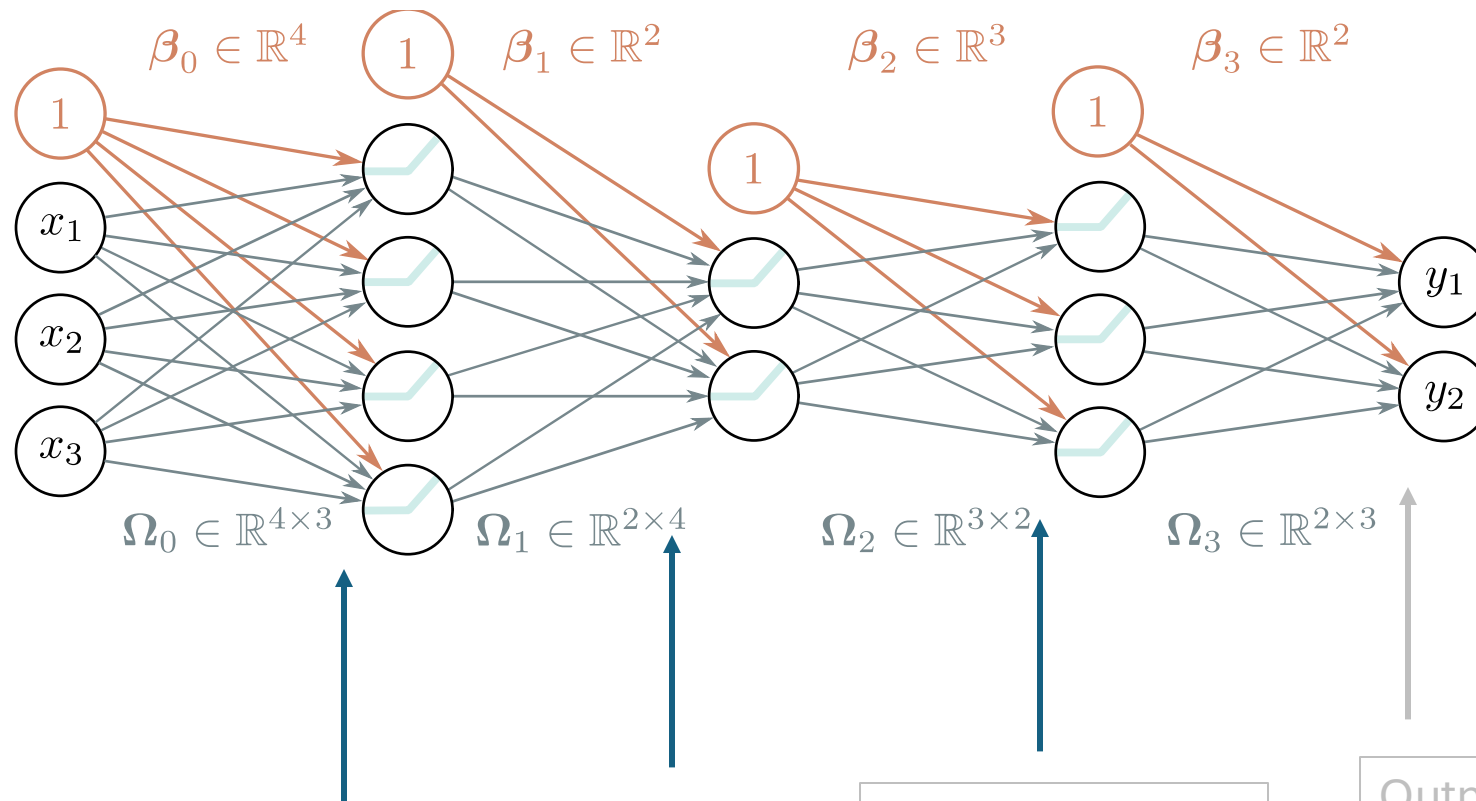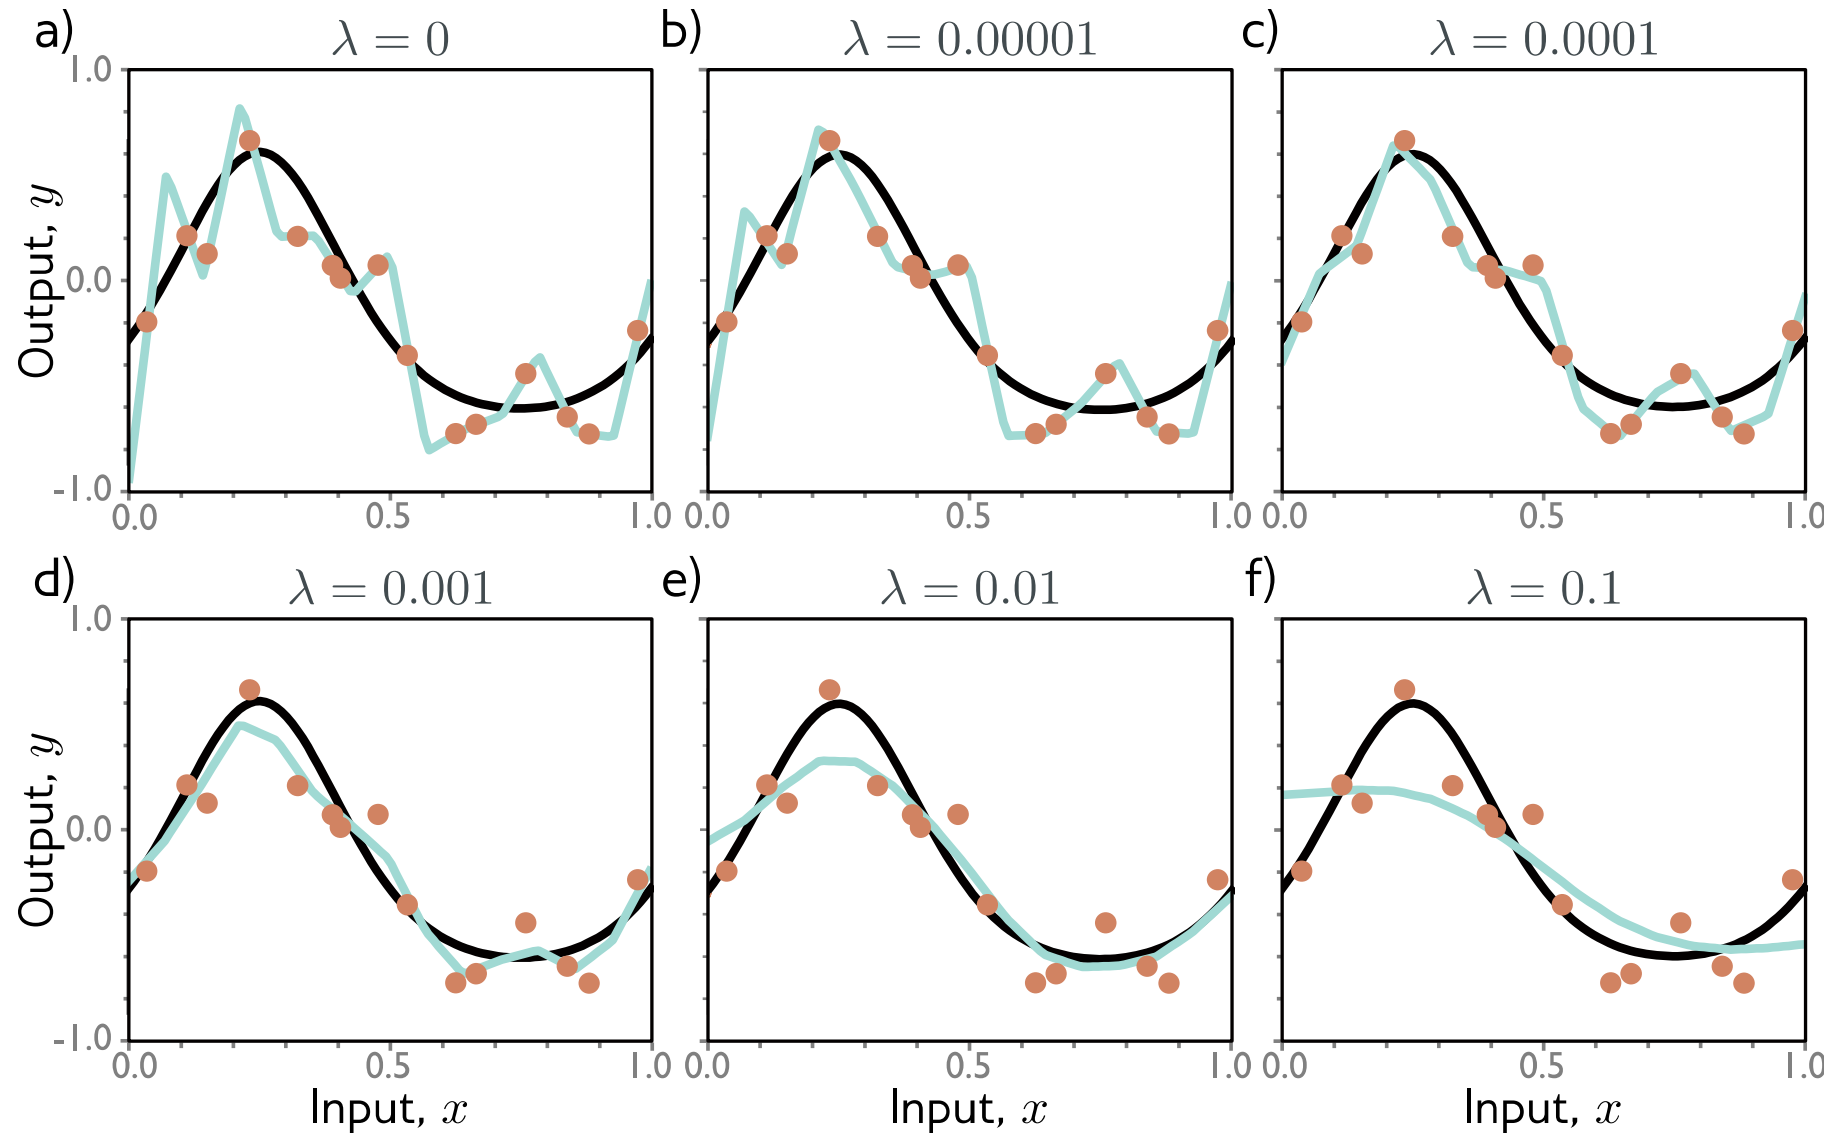
Same for the pre-activations into the last layer

Outputs are weighted linear combination of last layer activations.

Smaller weights attenuate changes.

20

Why does L2 regularization help?

$$\boldsymbol{\beta}_0 \in \mathbb{R}^4 \qquad \boldsymbol{\beta}_1 \in \mathbb{R}^2 \qquad \boldsymbol{\beta}_2 \in \mathbb{R}^3 \qquad \boldsymbol{\beta}_3 \in \mathbb{R}^2$$

$$\boldsymbol{\Omega}_0 \in \mathbb{R}^{4\times3} \qquad \boldsymbol{\Omega}_1 \in \mathbb{R}^{2\times4} \qquad \boldsymbol{\Omega}_2 \in \mathbb{R}^{3\times2} \qquad \boldsymbol{\Omega}_3 \in \mathbb{R}^{2\times3}$$

And so on...

All the way back.

Same for the pre-activations into the last layer

Outputs are weighted linear combination of last layer activations.

Smaller weights attenuate changes.

Why does L2 regularization help?

- Discourages fitting excessively to the training data (overfitting)
- Encourages smoothness between datapoints
  - Specifically by making coefficients smaller, so small input changes have smaller output changes.

L2 regularization (simple net from last lecture)

PyTorch Explicit L2 Regularizer

SGD

CLASS torch.optim.SGD(*params, lr=0.001, momentum=0, dampening=0, weight_decay=0,
         nesterov=False, *, maximize=False, foreach=None, differentiable=False*) [SOURCE]

Implements stochastic gradient descent (optionally with momentum).

Parameters

- **params** (*iterable*) – iterable of parameters to optimize or dicts defining parameter groups
- **lr** (*float, optional*) – learning rate (default: 1e-3)
- **momentum** (*float, optional*) – momentum factor (default: 0)
- **weight_decay** (*float, optional*) – weight decay (L2 penalty) (default: 0)

ADAM

CLASS torch.optim.Adam(*params, lr=0.001, betas=(0.9, 0.999), eps=1e-08,
         weight_decay=0, amsgrad=False, *, foreach=None, maximize=False,
         capturable=False, differentiable=False, fused=None*) [SOURCE]

Implements Adam algorithm.

Parameters

- **params** (*iterable*) – iterable of parameters to optimize or dicts defining parameter groups
- **lr** (*float, Tensor, optional*) – learning rate (default: 1e-3). A tensor LR is not yet supported for all our implementations. Please use a float LR if you are not also specifying fused=True or capturable=True.
- **betas** (*Tuple[float, float], optional*) – coefficients used for computing running averages of gradient and its square (default: (0.9, 0.999))
- **eps** (*float, optional*) – term added to the denominator to improve numerical stability (default: 1e-8)
- **weight_decay** (*float, optional*) – weight decay (L2 penalty) (default: 0)

Any Questions?

???

Moving on

- Reproducing Double Descent Demystified
- Explicit regularization
- Implicit regularization
- Early stopping
- Ensembling
- Dropout
- Adding noise
- Transfer learning, multi-task learning, self-supervised learning
- Data augmentation

Implicit regularization

$$\phi_{t+1} = \phi_t - \alpha \frac{\partial L[\phi_t]}{\partial \phi}$$

$$\lim_{\alpha \to 0}$$

$$\frac{d\phi}{dt} = -\frac{\partial L}{\partial \phi}$$

- In the limit, as $\alpha \to 0$, the gradient descent equation becomes the gradient flow differential equation.
- Doesn't converge to the same place

Implicit regularization

$$\phi_{t+1} = \phi_t - \alpha \frac{\partial L[\phi_t]}{\partial \phi}$$

$$\lim_{\alpha \to 0}$$

$$\frac{d\phi}{dt} = -\frac{\partial L}{\partial \phi}$$
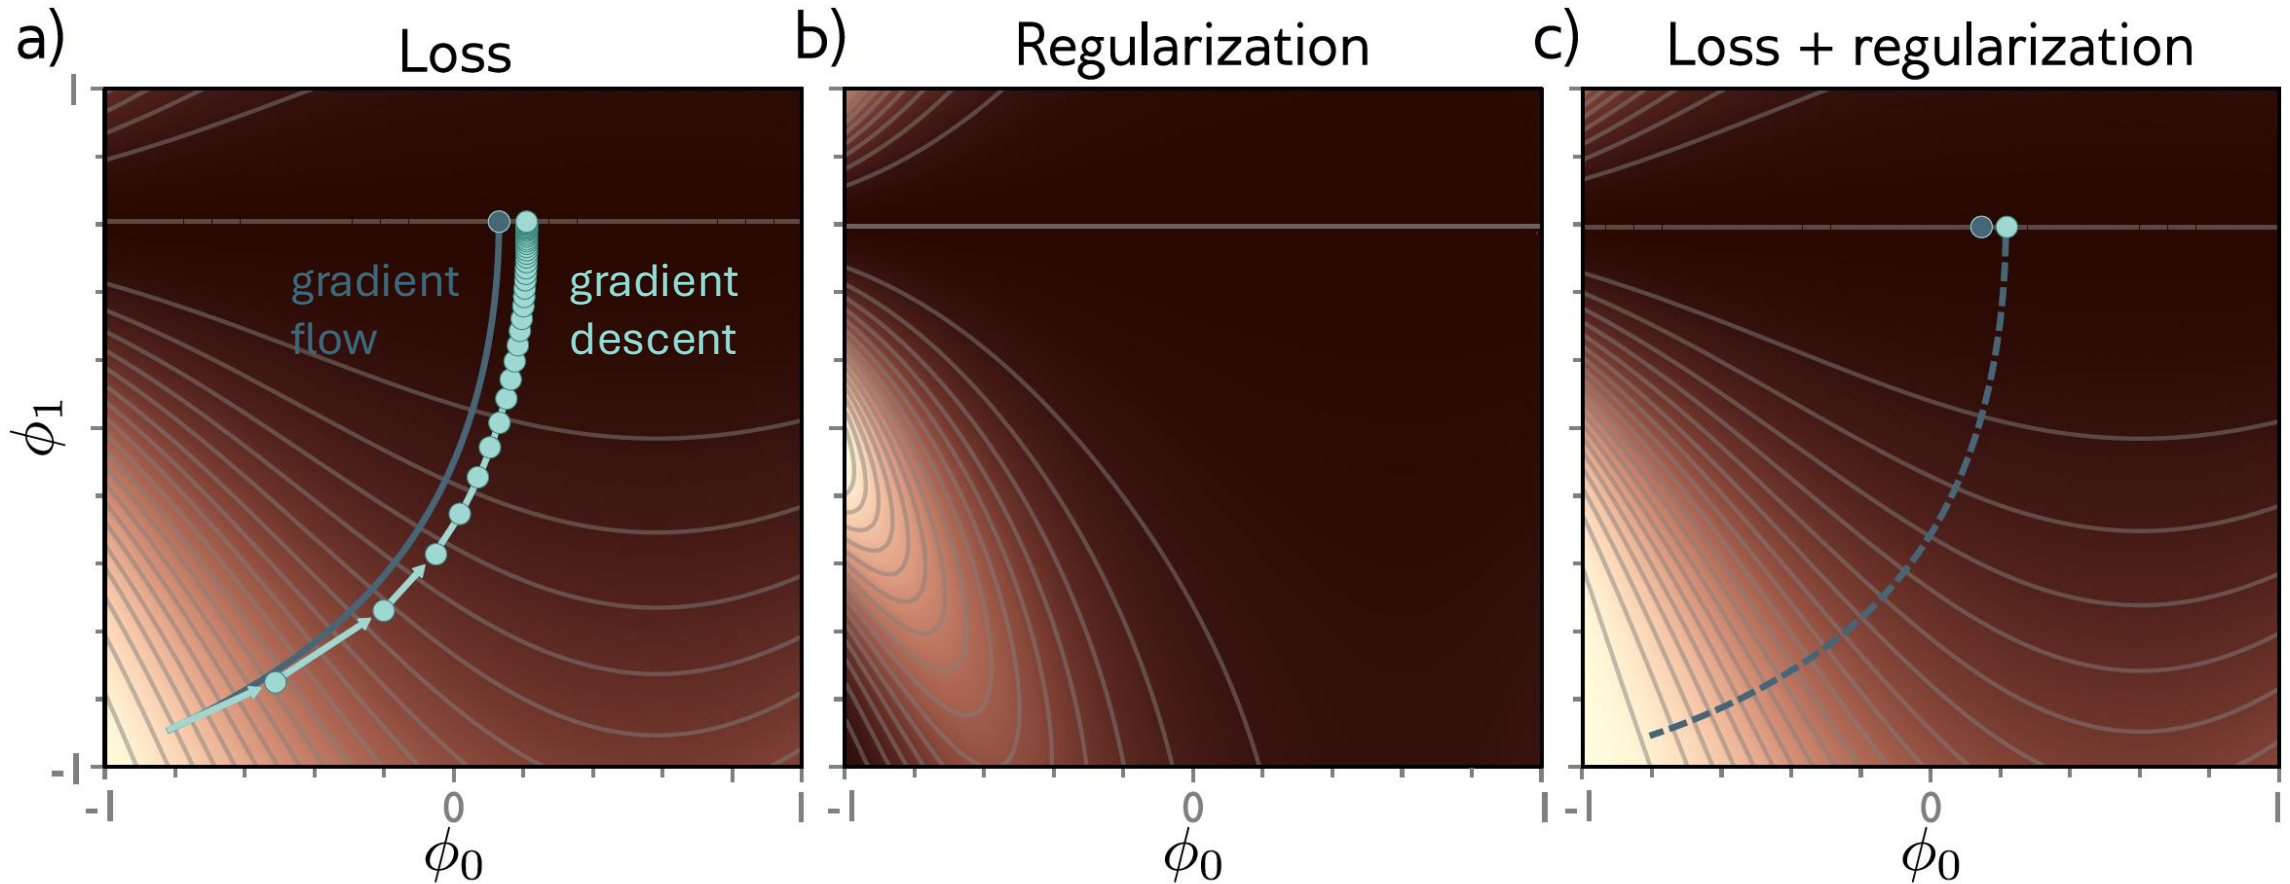
• The implicit regularization can be derived:

$$\tilde{\mathrm{L}}_{GD}[\phi] = \mathrm{L}[\phi] + \underbrace{\frac{\alpha}{4} \left\| \frac{\partial L}{\partial \phi} \right\|^2}$$

See derivation at
end of UDL Ch. 9

Penalty for large
gradients.

Implicit regularization

a) Loss

gradient flow

gradient descent

ϕ_1

ϕ_0

b) Regularization

ϕ_0

c) Loss + regularization

ϕ_0

Gradient descent doesn't converge to same location as (continuous) gradient flow.

Plot of the Implicit regularization $(\sim \|\partial L / \partial \phi\|^2)$ to be added to loss

With regularization, continuous descent converges to same place

28

Implicit regularization of SGD

- Gradient descent disfavors areas where gradients are steep

$$\tilde{\mathrm{L}}_{GD}[\phi] = \mathrm{L}[\phi] + \frac{\alpha}{4} \left\| \frac{\partial L}{\partial \phi} \right\|^2$$

- SGD likes all batches to have similar gradients

$$\tilde{\mathrm{L}}_{SGD}[\phi] = \tilde{\mathrm{L}}_{GD}[\phi] + \frac{\alpha}{4B} \sum_{b=1}^{B} \left\| \frac{\partial L_b}{\partial \phi} - \frac{\partial L}{\partial \phi} \right\|^2$$

Want the batch variance to be small, rather than some batches fitting well and others not well...

Where $L = \frac{1}{I} \sum_{i=1}^{I} \ell_i[\mathbf{x}_i, y_i]$ and $L_b = \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}_b} \ell_i[\mathbf{x}_i, y_i].$

29

Implicit regularization of SGD

- Gradient descent disfavors areas where gradients are steep

$$\tilde{\mathrm{L}}_{GD}[\phi] = \mathrm{L}[\phi] + \frac{\alpha}{4}\left\|\frac{\partial L}{\partial \phi}\right\|^2$$

- SGD likes all batches to have similar gradients

$$\tilde{\mathrm{L}}_{SGD}[\phi] = \tilde{\mathrm{L}}_{GD}[\phi] + \frac{\alpha}{4B}\sum_{b=1}^{B}\left\|\frac{\partial L_b}{\partial \phi} - \frac{\partial L}{\partial \phi}\right\|^2$$

$$= \mathrm{L}[\phi] + \frac{\alpha}{4}\left\|\frac{\partial L}{\partial \phi}\right\|^2 + \frac{\alpha}{4B}\sum_{b=1}^{B}\left\|\frac{\partial L_b}{\partial \phi} - \frac{\partial L}{\partial \phi}\right\|^2$$

- Depends on learning rate – perhaps why larger learning rates generalize better.

Loss and Regularization Surfaces

Original Gabor Model
Loss

$$\frac{\alpha}{4}\left\|\frac{\partial L}{\partial \phi}\right\|^2$$

$$\frac{\alpha}{4B}\sum_{b=1}^{B}\left\|\frac{\partial L_b}{\partial \phi} - \frac{\partial L}{\partial \phi}\right\|^2$$

$$\tilde{\mathrm{L}}_{SGD}[\phi]$$

$$= \mathrm{L}[\phi] + \frac{\alpha}{4}\left\|\frac{\partial L}{\partial \phi}\right\|^2 + \frac{\alpha}{4B}\sum_{b=1}^{B}\left\|\frac{\partial L_b}{\partial \phi} - \frac{\partial L}{\partial \phi}\right\|^2$$

a) MNIST1D no label noise

Full batch, LR = 0.5
Full batch, LR = 0.1
Full batch, LR = 0.05

Test

Train

b) MNIST1D no label noise

Batch size 10, LR = 0.1
Batch Size 100, LR = 0.1
Full batch (4000), LR = 0.1

Test

Train

Generally, performance is
- best for larger learning rates
- best with smaller batches

Recap: Implicit regularization of GD and SGD

- Larger learning rates may lead to better generalization
- SGD seems to favor places where gradients are stable (all batches agreed on slope)
- SGD generalizes better than GD
- Smaller batches in SGD generally perform better than larger ones

Any Questions?

???

Moving on
- Reproducing Double Descent Demystified
- Explicit regularization
- Implicit regularization
- Early stopping
- Ensembling
- Dropout
- Adding noise
- Transfer learning, multi-task learning, self-supervised learning
- Data augmentation

Early stopping

- If we stop training early, weights don't have time to overfit to noise
- Weights start small, don't have time to get large
- Reduces effective model complexity
- Known as early stopping
- Don't have to re-train with different hyper-parameters – just "checkpoint" regularly and pick the model with lowest validation loss

a) Iter = 0, Loss = 32.24
b) Iter = 1000, Loss = 1.64
c) Iter = 5000, Loss = 1.10
d) Iter = 10000, Loss = 0.80
e) Iter = 50000, Loss = 0.36
f) Iter = 200000, Loss = 0.16

Simplified shallow network model with 14 linear regions initialized randomly (cyan curve in (a)) and trained with SGD using a batch size of five and a learning rate of 0.05.

Project 1 Provided Code

Project 1 36 Epochs

Same Seed (1) 500 Epochs

Any Questions?

???

Moving on
- Reproducing Double Descent Demystified
- Explicit regularization
- Implicit regularization
- Early stopping
- Ensembling
- Dropout
- Adding noise
- Transfer learning, multi-task learning, self-supervised learning
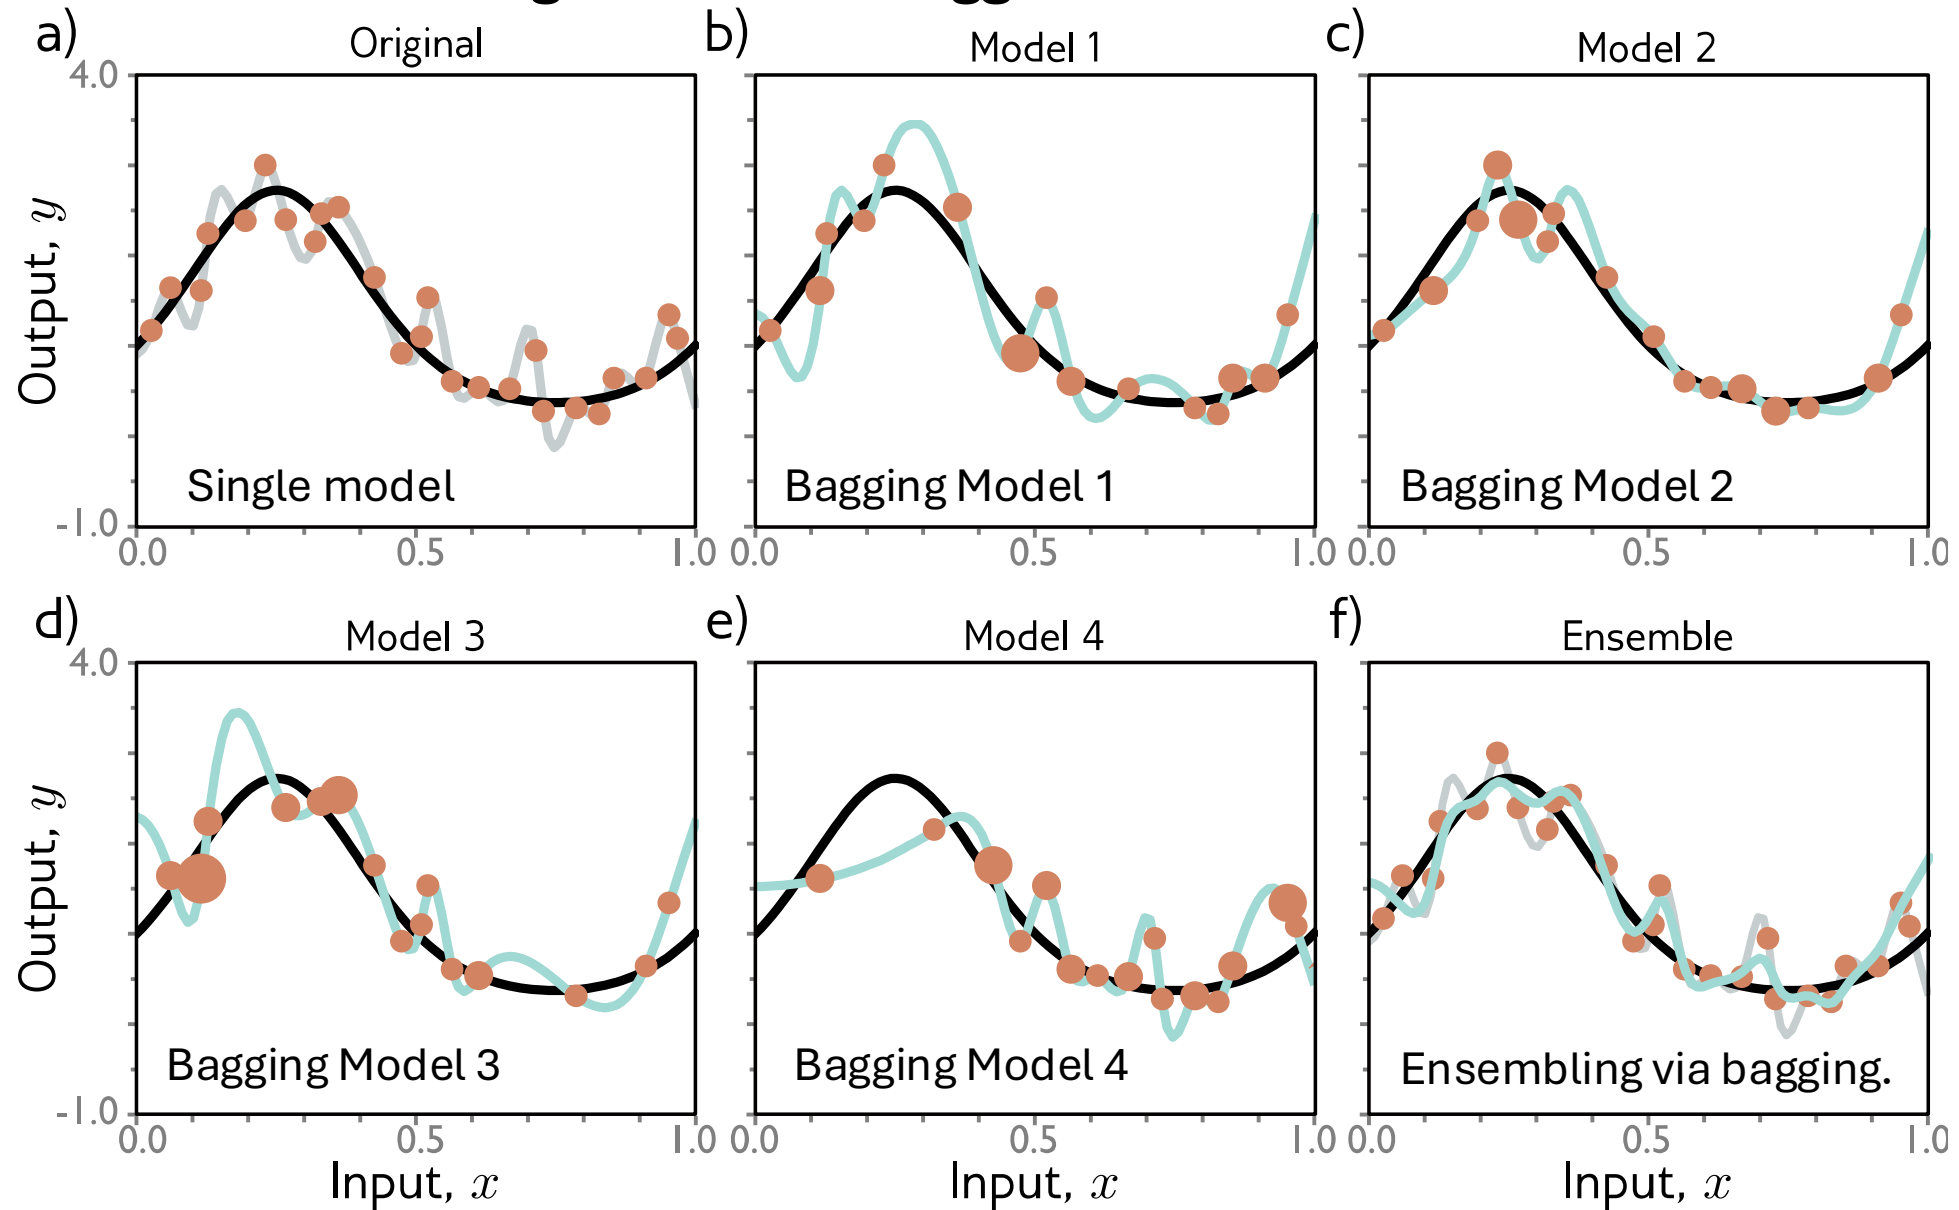- Data augmentation

Ensembling

- Combine several models – an ensemble
- Combining outputs

	Mean	Median/Frequent (Robust)
Regression	Mean of outputs	Median of outputs
Classification	Mean before softmax	Most frequent predicted class

- Can be simply different initializations or even different models
- Or train with different subsets of the data resampled with replacements – bootstrap aggregating (bagging)

Single Model vs Bagged Ensemble

a) Original — Single model

b) Model 1 — Bagging Model 1

c) Model 2 — Bagging Model 2

d) Model 3 — Bagging Model 3

e) Model 4 — Bagging Model 4

f) Ensemble — Ensembling via bagging.

Output, y — Input, x

Size of orange point indicates number of times the data point was re-sampled.

42

Any Questions?

???

Moving on
- Reproducing Double Descent Demystified
- Explicit regularization
- Implicit regularization
- Early stopping
- Ensembling
- Dropout
- Adding noise
- Transfer learning, multi-task learning, self-supervised learning
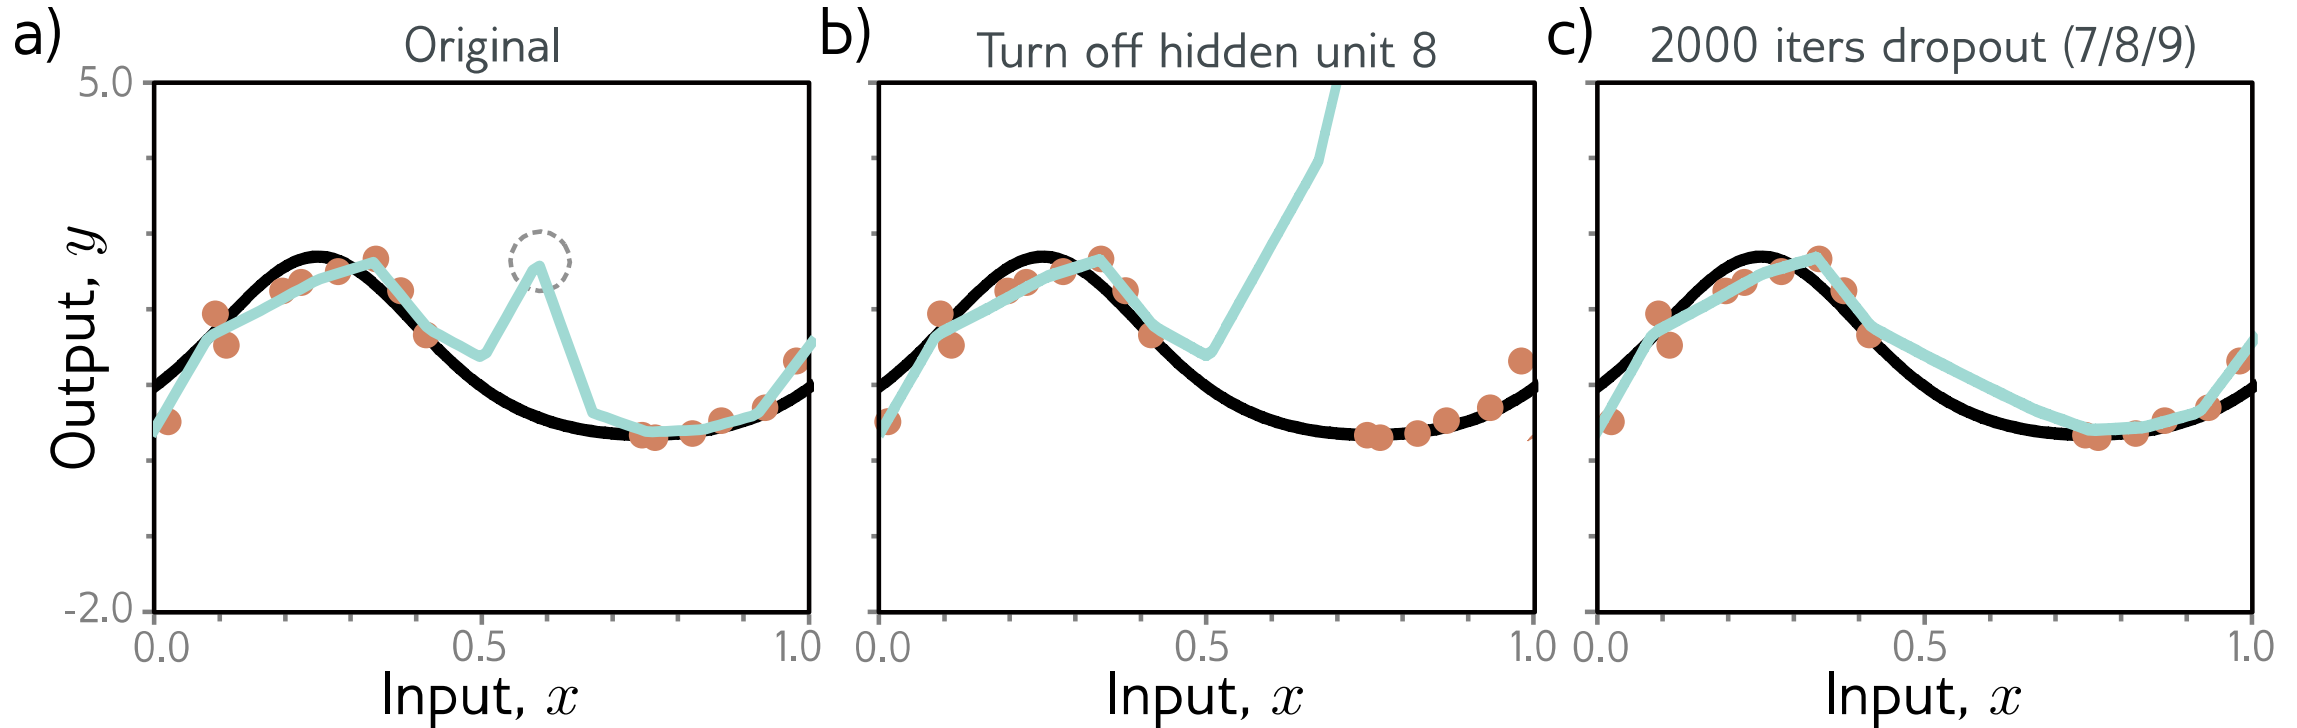- Data augmentation

Dropout

a) Iteration: 1

b) Iteration: 2

c) Iteration: 3

d) Iteration: 4

- Makes the network less dependent on any given hidden unit.
- At test time, all hidden units are active, which was not the case during training
  - Must rescale using *weight scaling inference rule*
    - Multiply weights by (1 – dropout probability) so average contribution is the same.

Dropout

a) Original

b) Turn off hidden unit 8

c) 2000 iters dropout (7/8/9)

- Prevents situations where subsequent hidden units correct for excessive swings from earlier hidden units
- Can eliminate kinks in function that are far from data and don't contribute to training loss

Monte Carlo Dropout for Inference (optional)

- Run the network multiple times with different random subsets of units clamped to zero (as in training).

- Combine the results using an ensembling method,

- This is closely related to ensembling in that every random version of the network is a different model; however, we do not have to train or store multiple networks here.
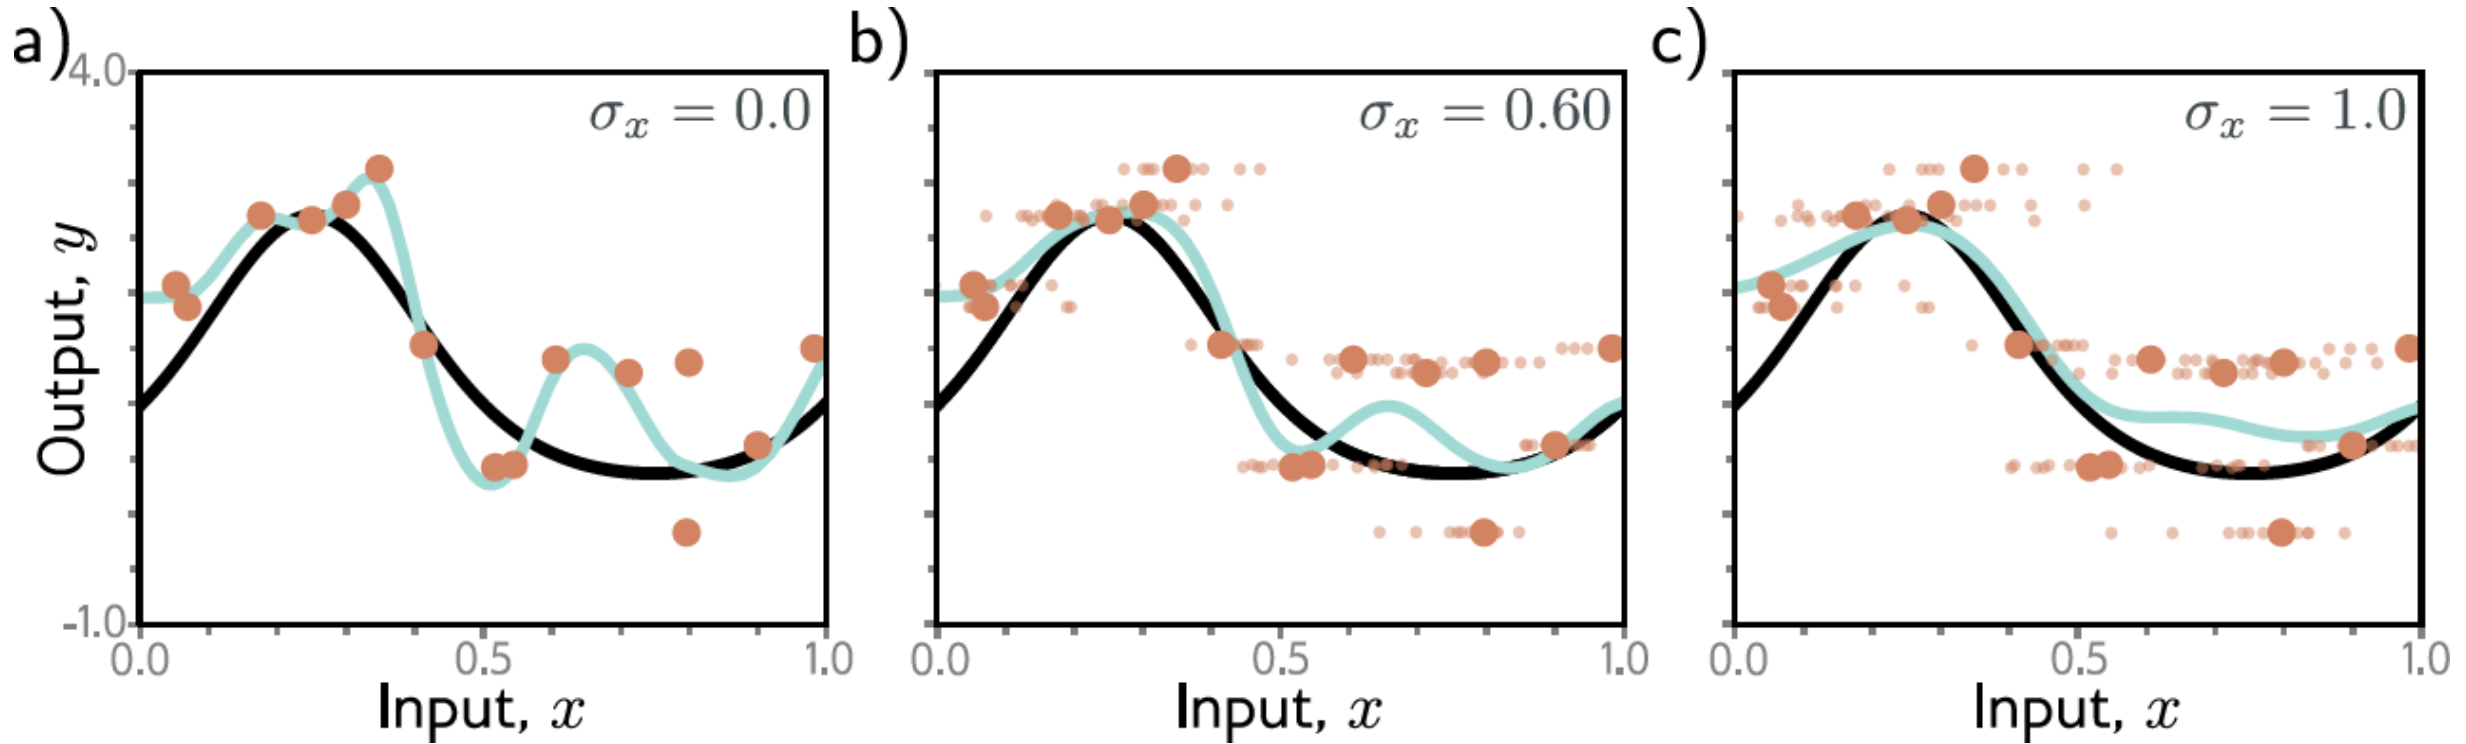
Any Questions?

???

Moving on
- Reproducing Double Descent Demystified
- Explicit regularization
- Implicit regularization
- Early stopping
- Ensembling
- Dropout
- Adding noise
- Transfer learning, multi-task learning, self-supervised learning
- Data augmentation

Adding noise

Adding noise to input with different variances.

- to inputs – induces weight regularization (see Exercise 9.3 in UDL)
- to weights – makes robust to small weight perturbations
- to outputs (labels) – reduces "overconfident" probability for target class

Any Questions?

???

Moving on

- Reproducing Double Descent Demystified
- Explicit regularization
- Implicit regularization
- Early stopping
- Ensembling
- Dropout
- Adding noise
- Transfer learning, multi-task learning, self-supervised learning
- Data augmentation

Transfer & Multitask Learning, Augmentation

- Strictly speaking not regularization, but can help improve generalization when dataset sizes are limited

Transfer Learning

(1) Train the model for segmentation
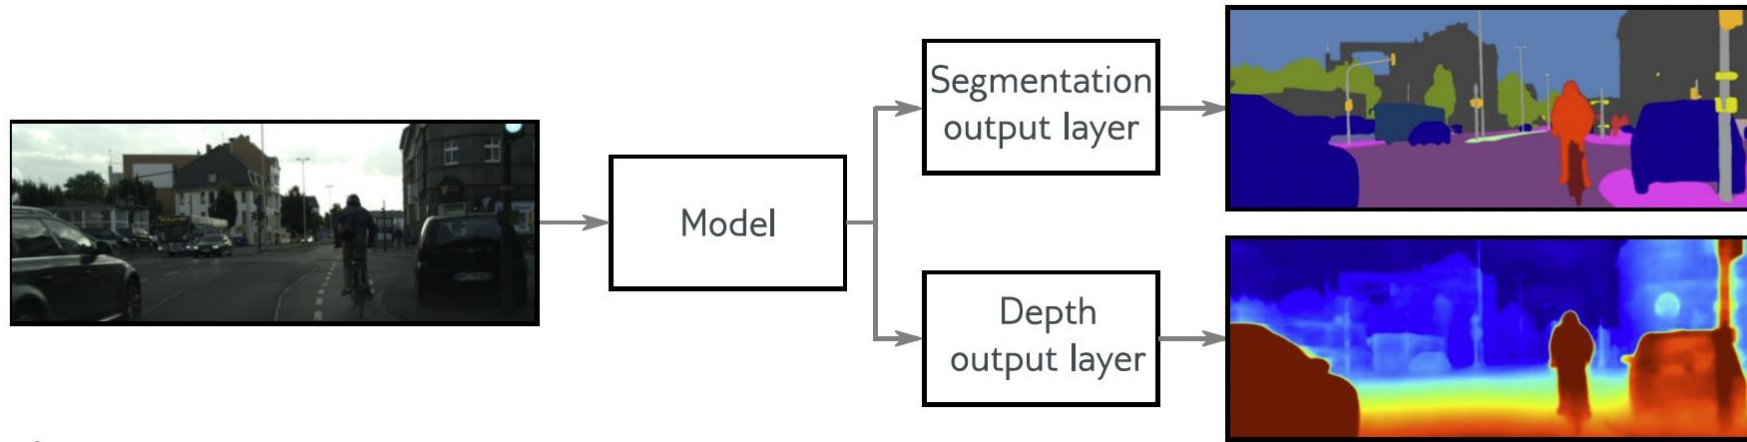
Assume we have lots of segmentation training data

Assume we have limited depth training data

(2) Replace the final layers to match the new task and

(3) Either:
a) Freeze the rest of the layers and train the final layers
b) Fine tune the entire model

Multi-Task Learning

- Train the model for 2 or more tasks simultaneously
  - Weighted combo of loss functions

$$L_{total} = \alpha \cdot L_{segmentation} + \beta \cdot L_{depth}$$

- Less likely to overfit to training data of one task

- Can be harder to get training to converge. Might have to vary the individual task loss weightings, α and β.

Self-Supervised Learning

The animal didn't cross th██████et because it was too tired.

- Mask out part of the training data
- Train model to try to infer missing data
  - masked data is the target
- ➔ Model learns characteristics of the data

- Then apply transfer learning

Any Questions?

???

Moving on
- Reproducing Double Descent Demystified
- Explicit regularization
- Implicit regularization
- Early stopping
- Ensembling
- Dropout
- Adding noise
- Transfer learning, multi-task learning, self-supervised learning
- Data augmentation

Data augmentation



a) Original
b) Flip
c) Rotate and crop
d) Vertical stretch
e) Color balance
f) Blur
g) Vignette
h) Pincushion

Image augmentation in PyTorch

```python
import torch
import torchvision.transforms as transforms


# Define augmentation pipeline
transform = transforms.Compose([
transforms.RandomHorizontalFlip(p=0.5),
transforms.RandomVerticalFlip(p=0.3),
transforms.RandomRotation(degrees=30),
transforms.ColorJitter(brightness=0.5, contrast=0.5, saturation=0.5),
transforms.RandomAffine(degrees=20, translate=(0.2, 0.2), shear=10),
transforms.RandomPerspective(distortion_scale=0.5, p=0.5),
transforms.ToTensor(), # Convert image to tensor
])

# Apply transformations multiple times to visualize augmentation
augmented_image = transform(image)
```

https://pytorch.org/vision/main/transforms.html

Data Augmentation: MNIST1D

Examples in training set: 4000
Examples in test set: 1000
Length of each example: 40

- Randomly circularly rotate
- randomly scale between 0.8 and 1.2

Regularization overview

Make function smoother

Increase data

Explicit L2 regularization

Apply noise to inputs

Data augmentation

Multi-task learning

Early stopping

Apply noise to outputs (label smoothing)

Transfer learning

Ensembling

Dropout

Implicit regularization

Bayesian approach

Apply noise to weights

Combine multiple models

Find wider minima

Any Questions?

???

- Reproducing Double Descent Demystified
- Explicit regularization
- Implicit regularization
- Early stopping
- Ensembling
- Dropout
- Adding noise
- Transfer learning, multi-task learning, self-supervised learning
- Data augmentation