Deep Learning for Data Science DS 542

https://dl4ds.github.io/sp2026/

Residual Networks

Plan for Today

- Finish up CNN examples (from last time)
- Residual networks

- Image classification
- Object detection
- Semantic segmentation

ImageNet 1K database

Fei-Fei Li
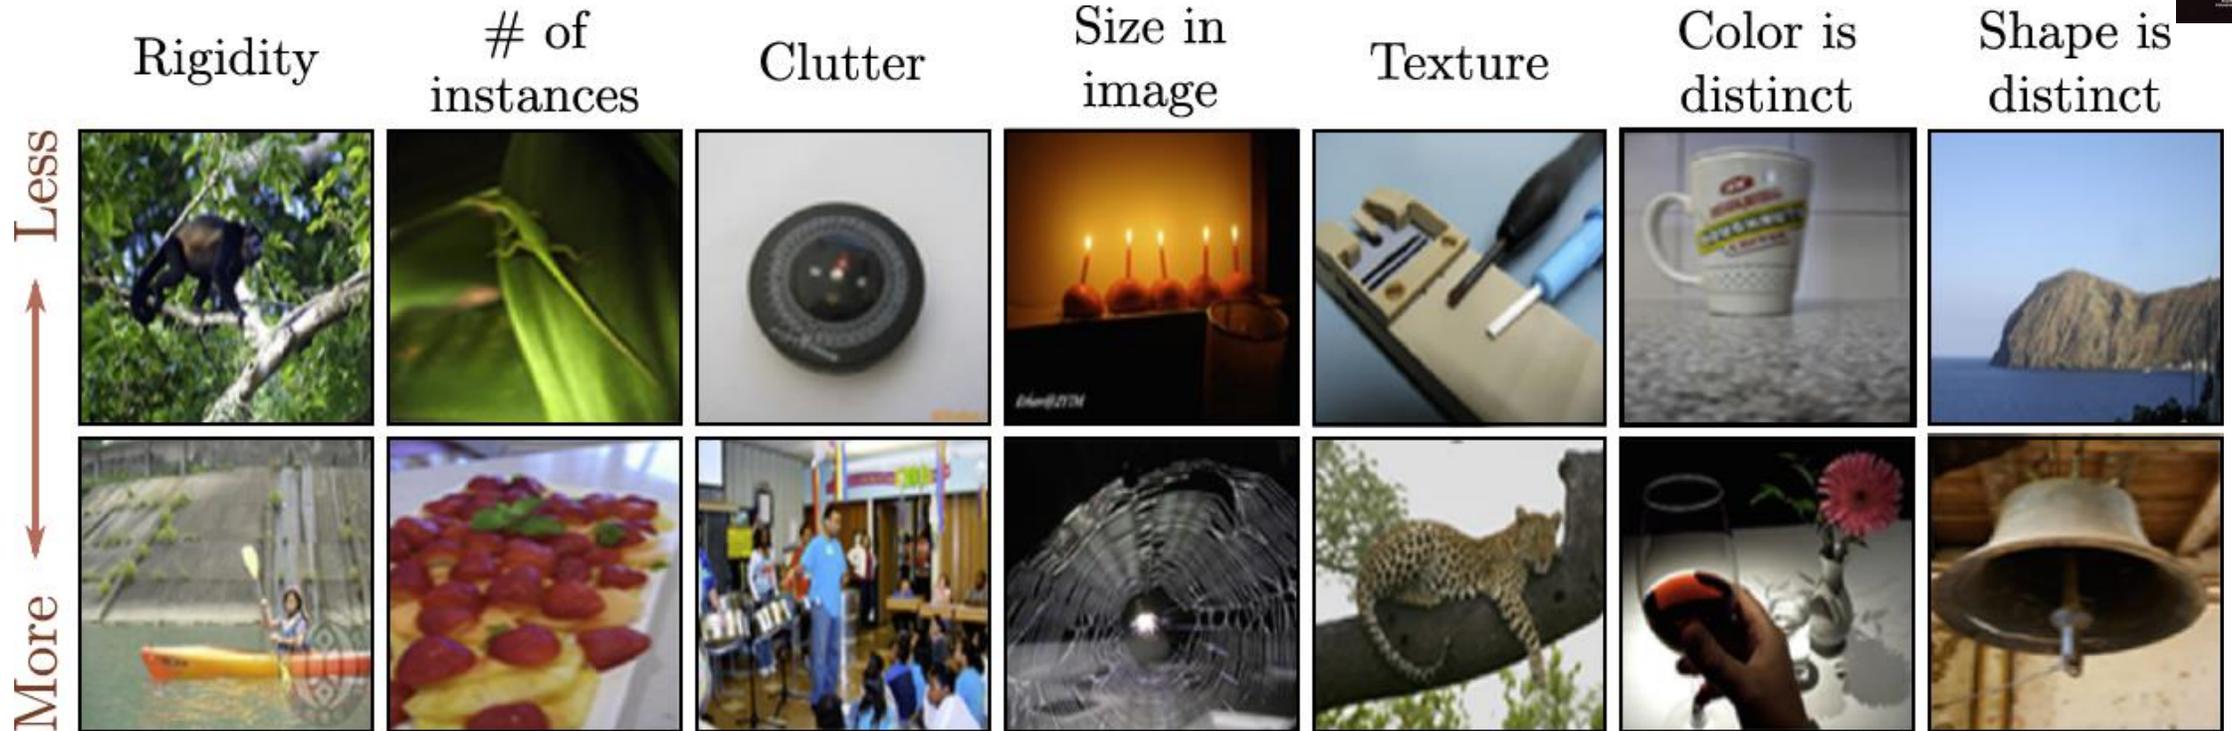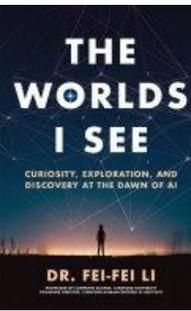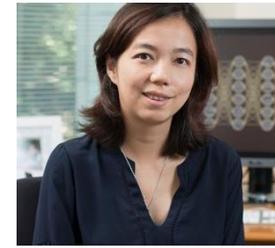
- 224 x 224 images

- 1,281,167 training images, 50,000 validation images, and 100,000 test images

- 1000 classes

AlexNet (2012)

$224 \times 224 \times 3$

$55 \times 55 \times 96$

$27 \times 27 \times 256$

$13 \times 13 \times 384$

$13 \times 13 \times 256$

4096 4096

1000 1000

Conv11×11
Stride 4

MaxPool
Conv 5 × 5

MaxPool

Conv 3 × 3

MaxPool + FC

FC

FC
Softmax
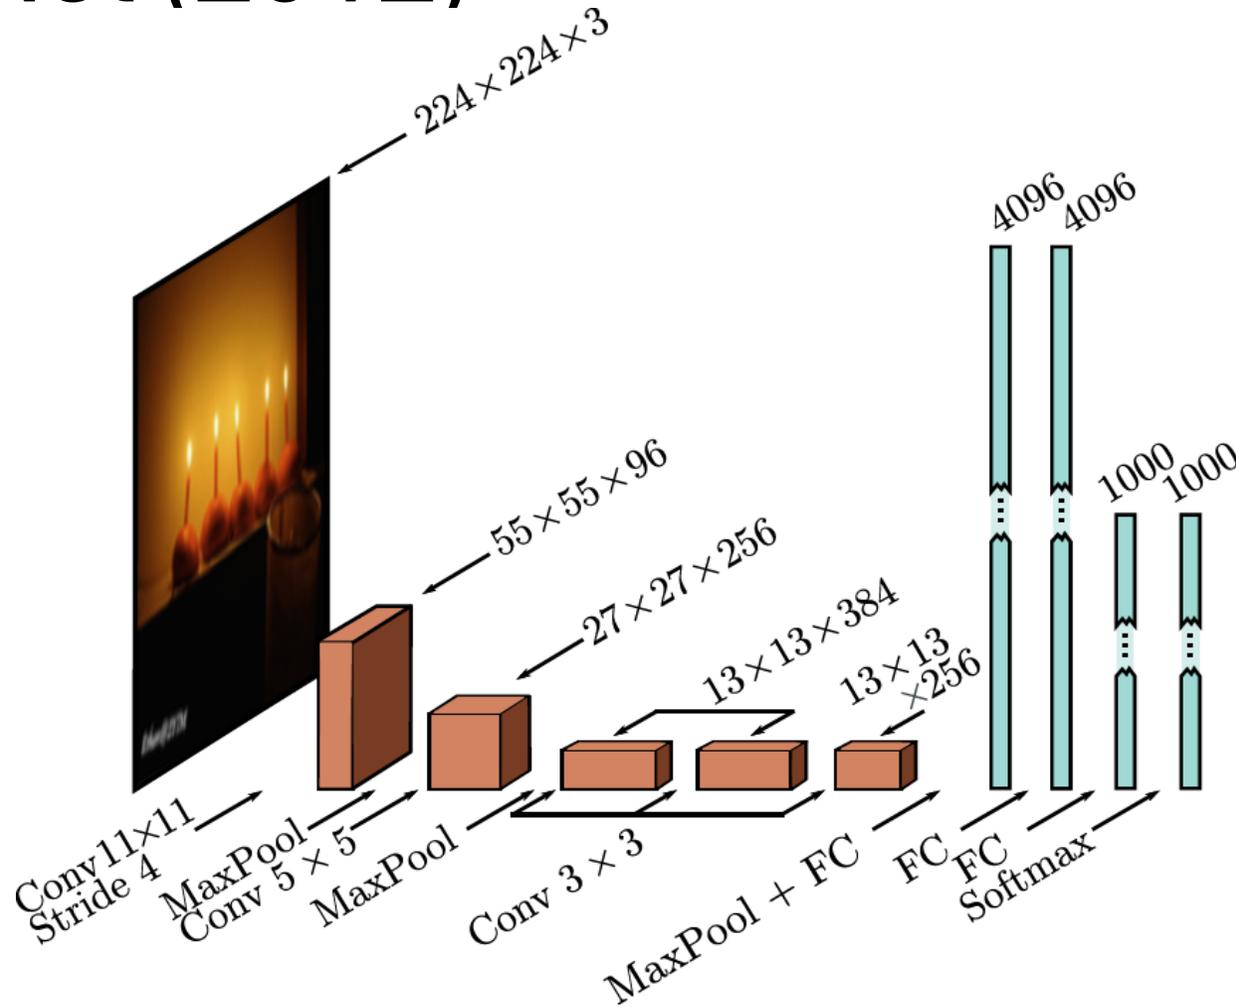
Almost all the 60 million parameters parameters are in fully connected layers

AlexNet (2012)

Figure 2: An illustration of the architecture of our CNN, explicitly showing the delineation of responsibilities between the two GPUs. One GPU runs the layer-parts at the top of the figure while the other runs the layer-parts at the bottom. The GPUs communicate only at certain layers. The network's input is 150,528-dimensional, and the number of neurons in the network's remaining layers is given by 253,440–186,624–64,896–64,896–43,264–4096–4096–1000.

Won the 2012 Large-Scale Vision Recognition Challenge (ILSVRC) by a big margin.

Took between five and six days to train on two GTX 580 3GB GPUs with manually optimized compute kernels.

A. Krizhevsky, I. Sutskever, and G. E. Hinton, "ImageNet classification with deep convolutional neural networks," *Commun. ACM*, vol. 60, no. 6, pp. 84–90, May 2012, doi: 10.1145/3065386.

AlexNet

Cat image input
(not actual image)

Activations (feature maps)
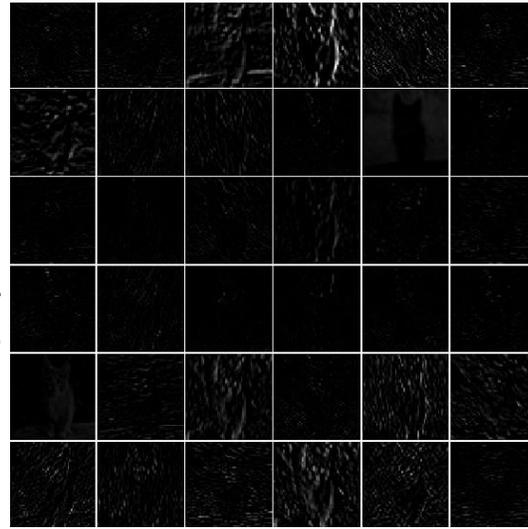
1st Layer

2nd Layer
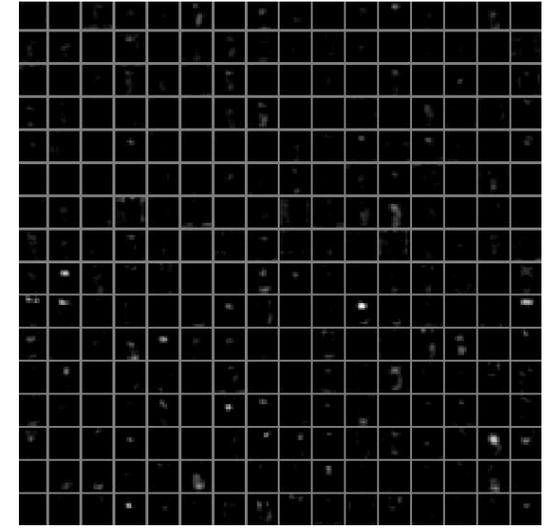
5th Layer

• • •

Filter Kernels

• • •

https://cs231n.github.io/understanding-cnn/

6
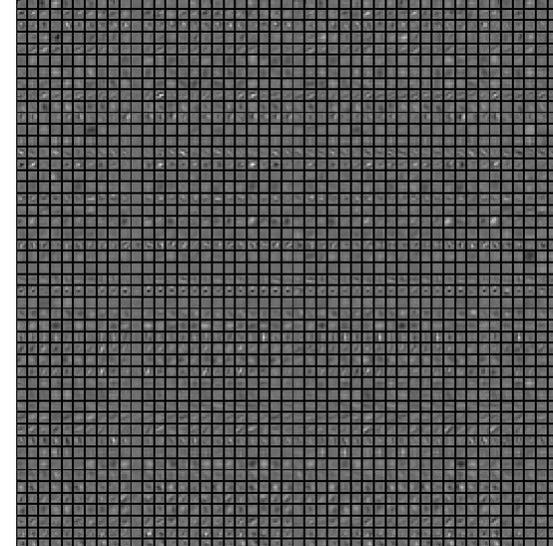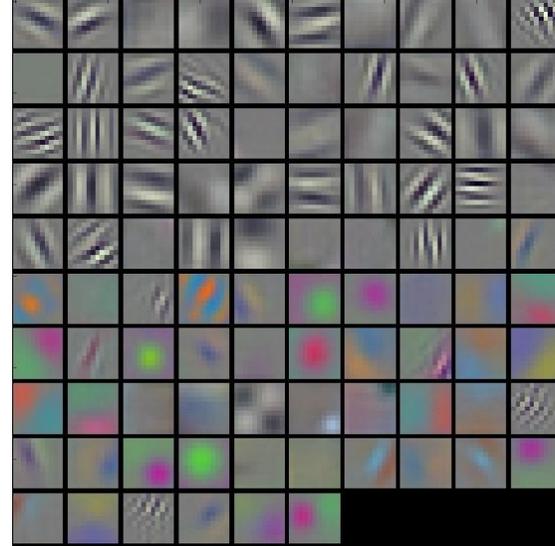
Details

- At test time average results from five different cropped and mirrored versions of the image

- SGD with a momentum coefficient of 0.9 and batch size of 128.

- L2 (weight decay) regularizer used.

- This system achieved a 16.4% top-5 error rate and a 38.1% top-1 error rate.

VGG (2015)
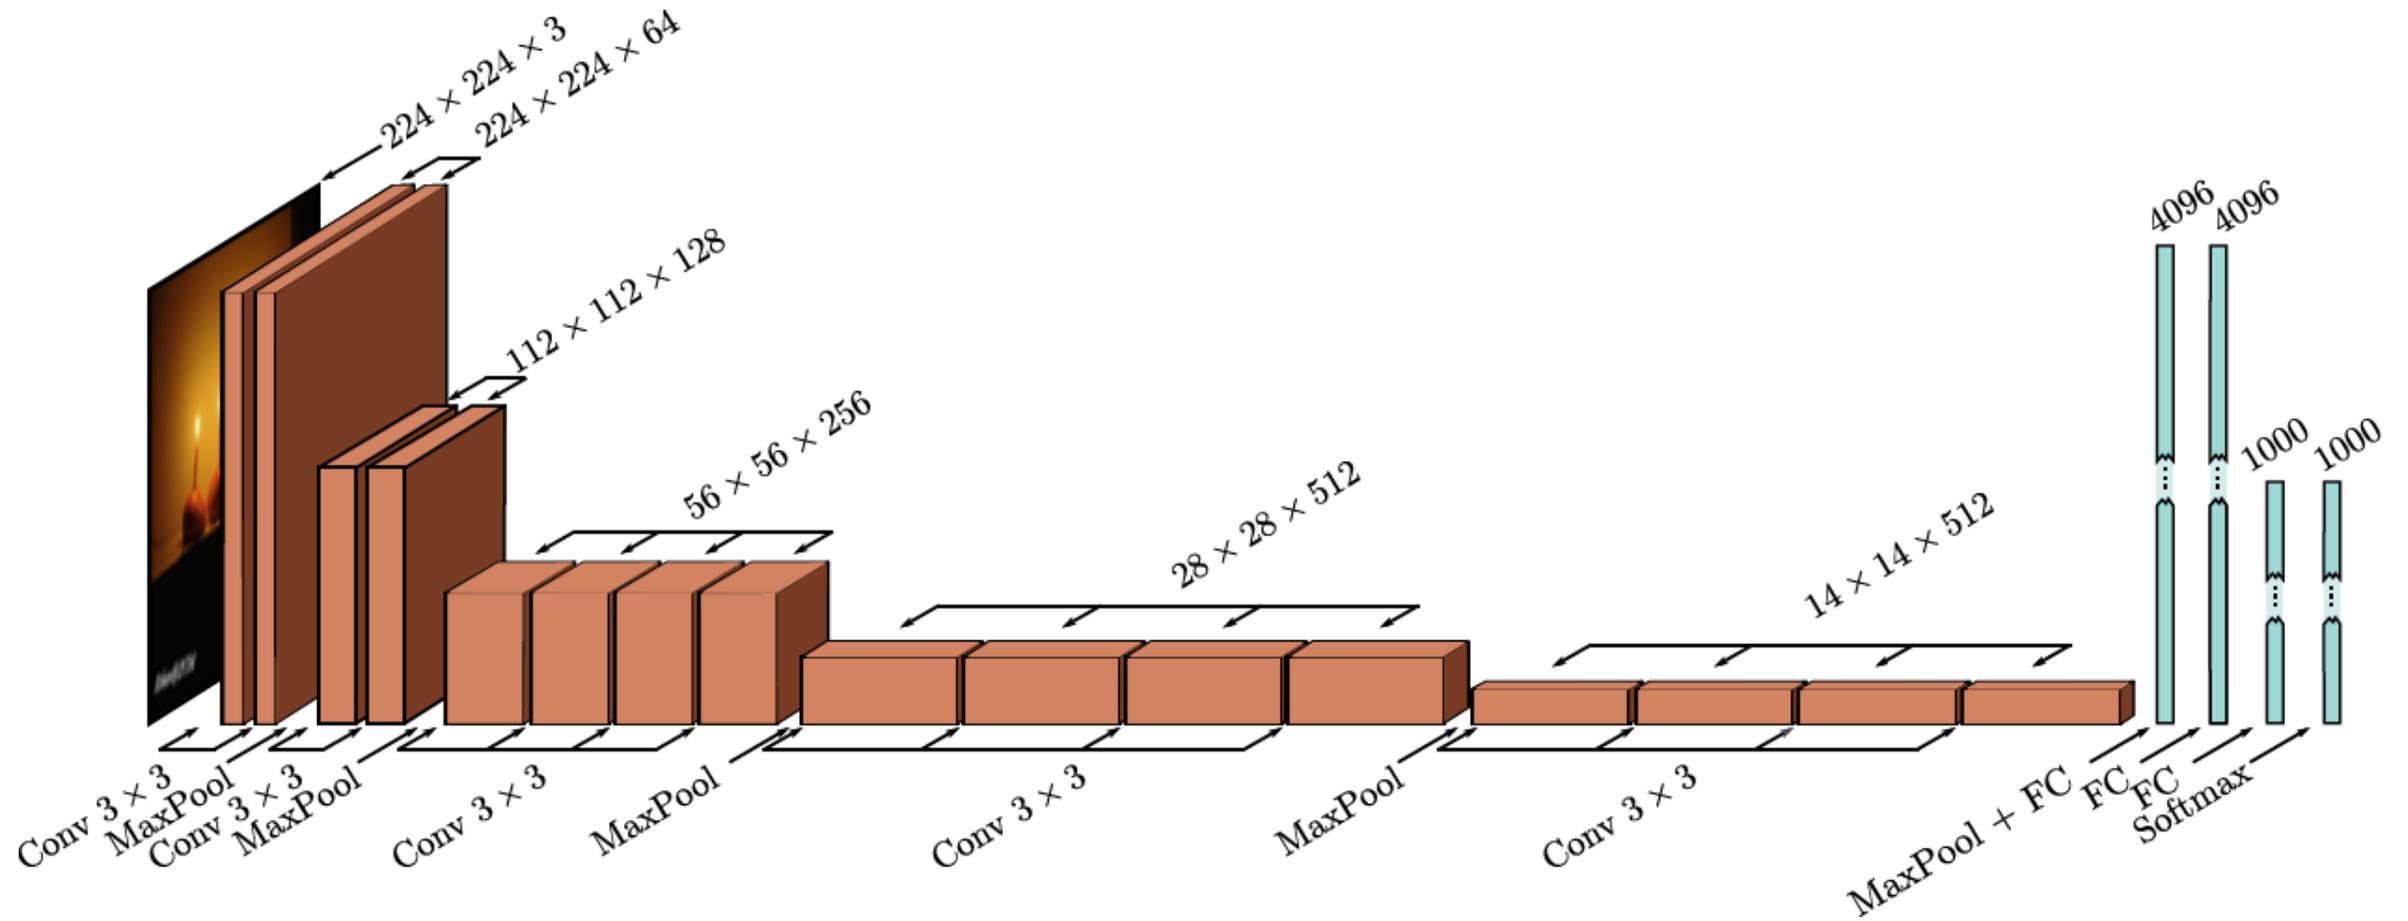
Details

- 19 hidden layers
- 144 million parameters
- 6.8% top-5 error rate, 23.7% top-1 error rate

ImageNet History

Any Questions?

???
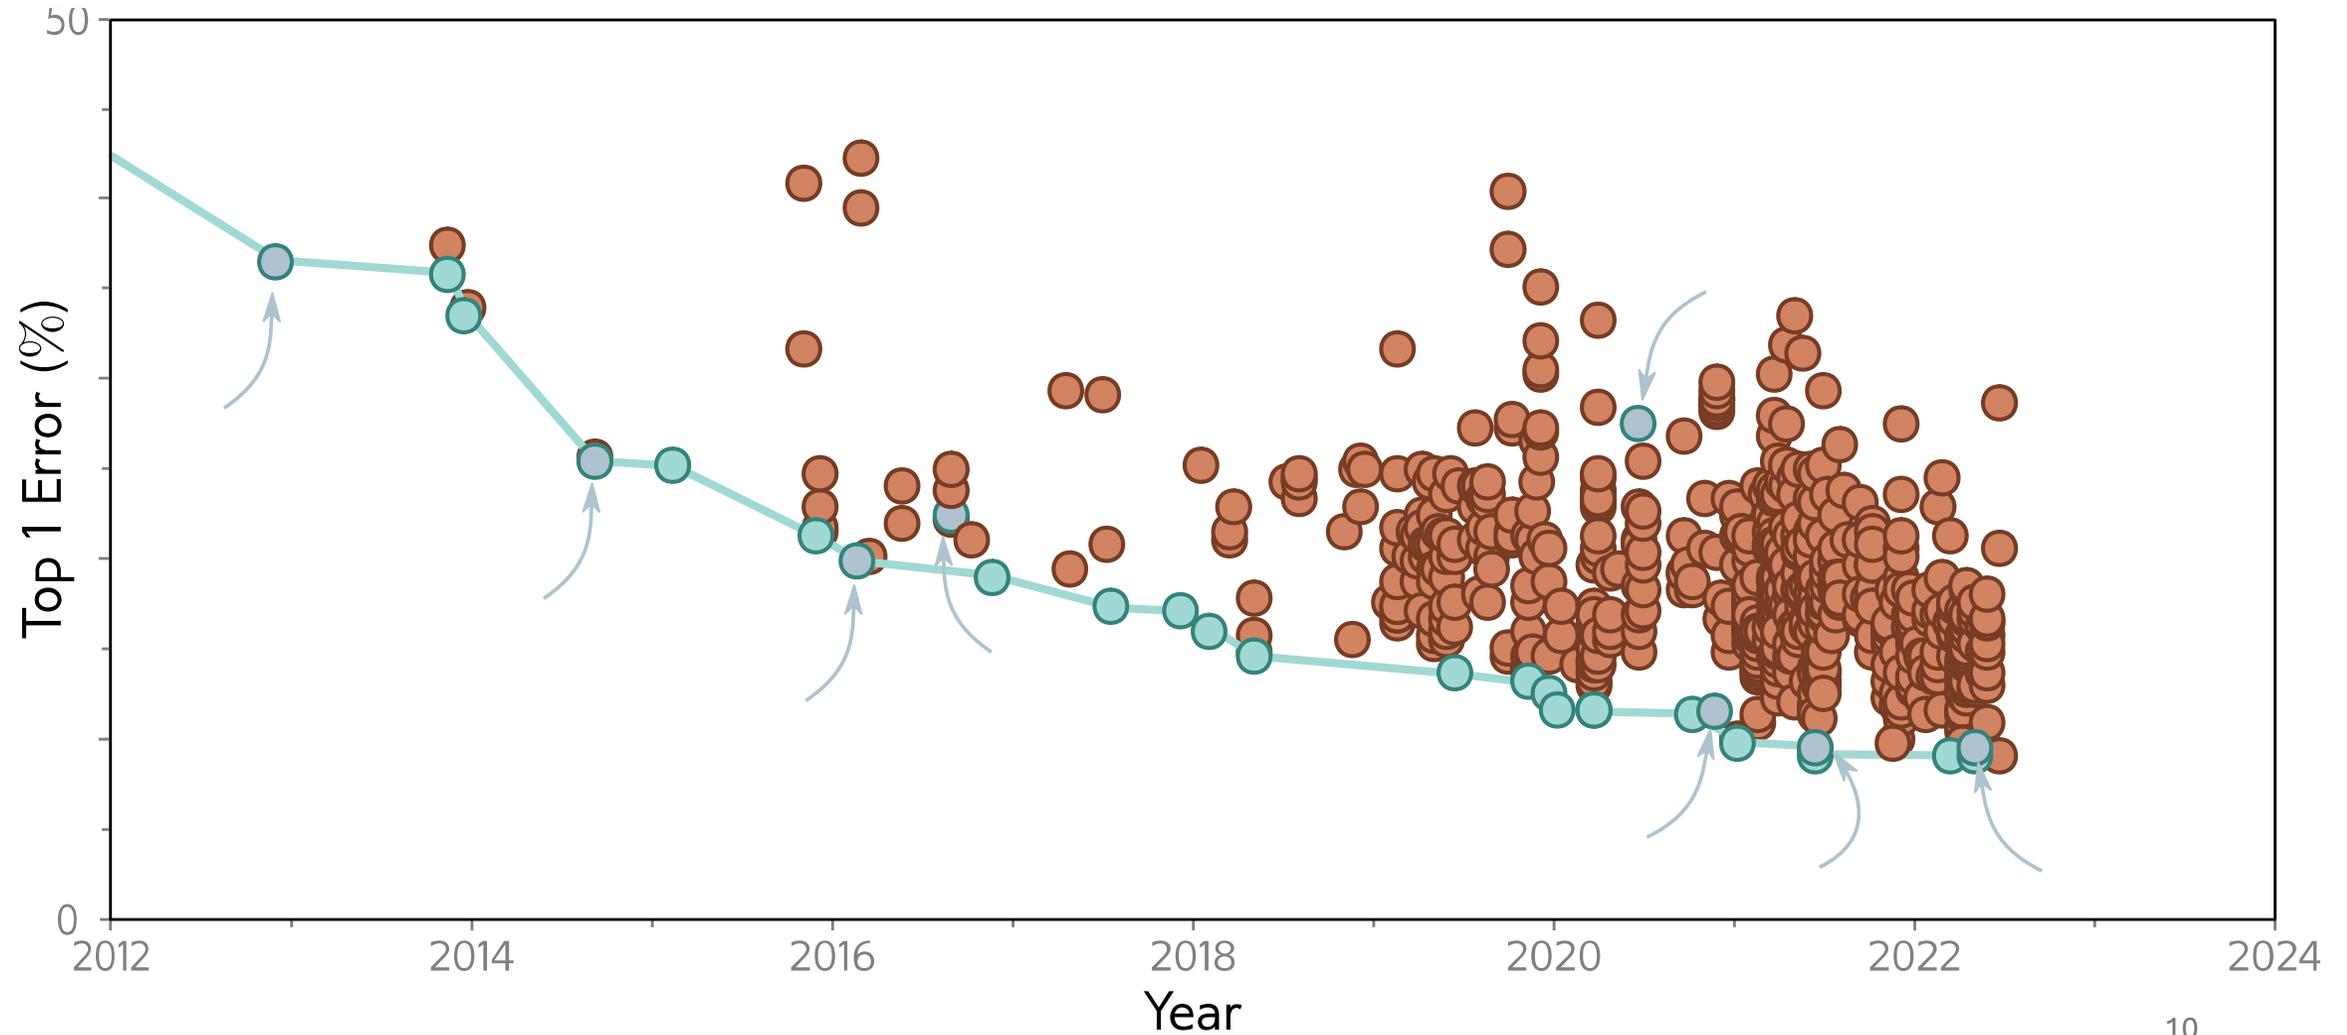
Moving on

- Image classification
- Object detection
- Semantic segmentation

You Only Look Once (YOLO)

- Network similar to VGG (448x448 input)
- 7×7 grid of locations
- Predict class at each location
- Predict 2 bounding boxes at each location
  - Five parameters –x,y, height, width, and confidence
- Momentum, weight decay, dropout, and data augmentation
- Heuristic at the end to threshold and decide final boxes – (non maximum suppression)
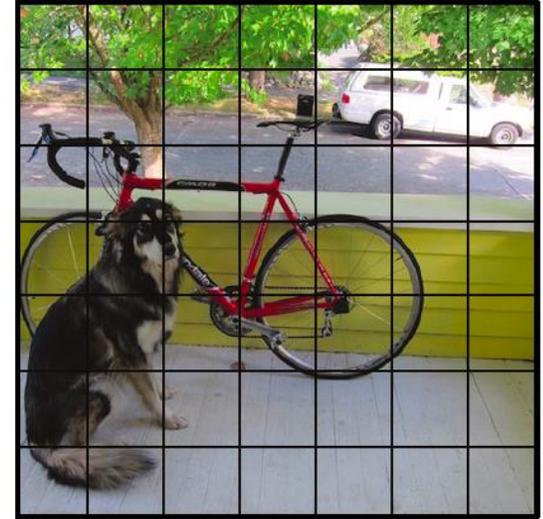
Object detection (YOLO)
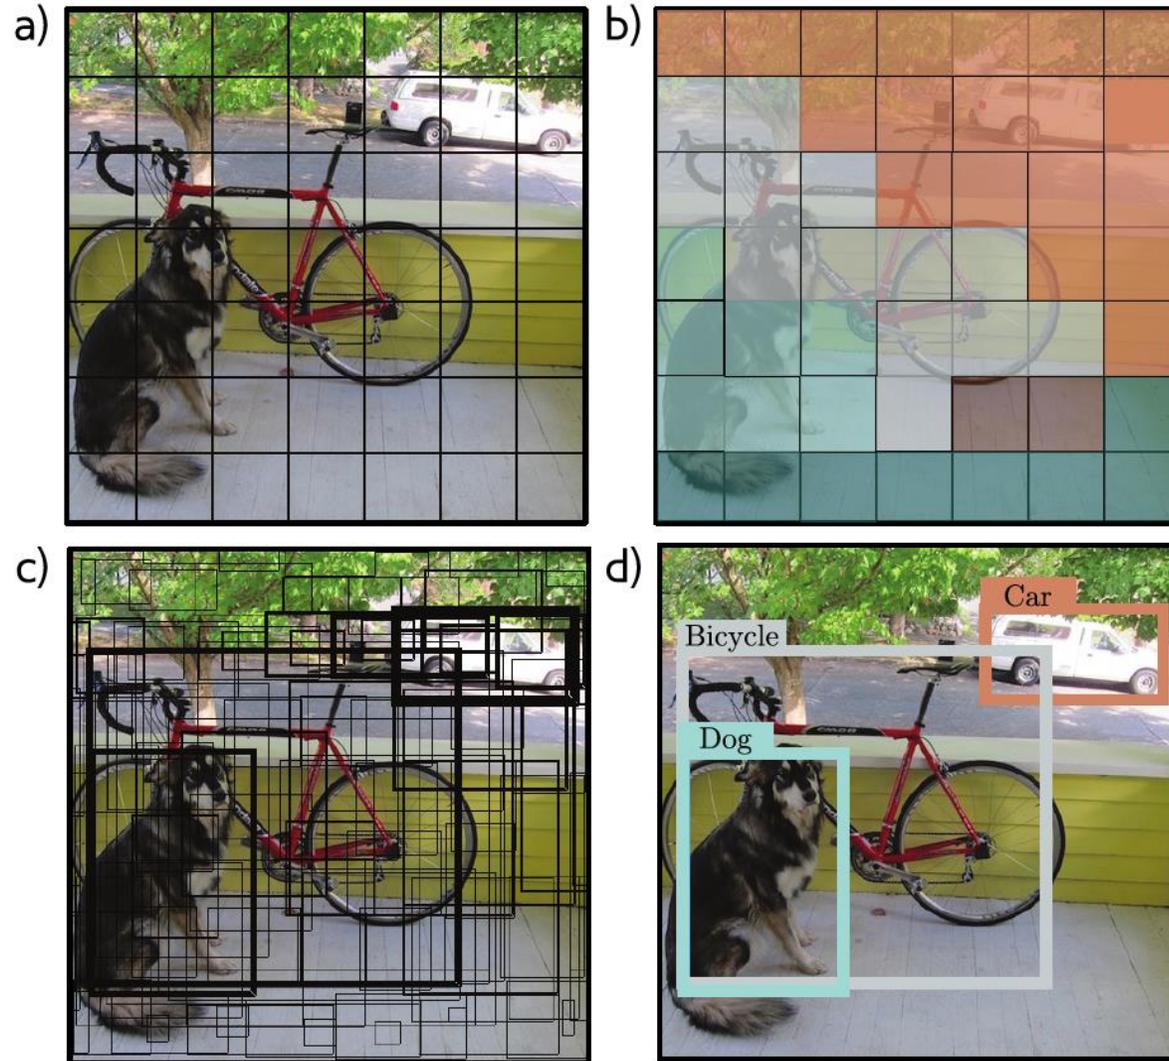
Results

Any Questions?

???
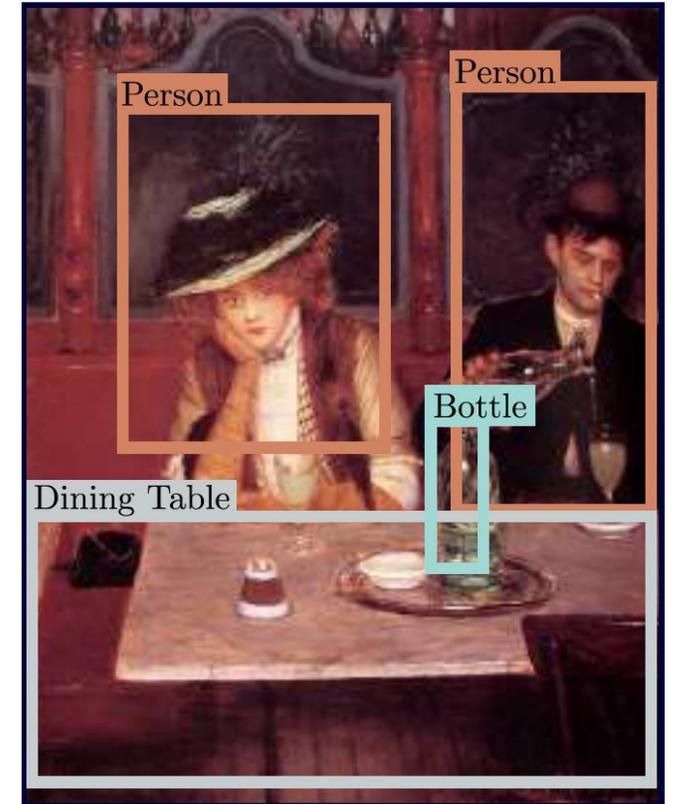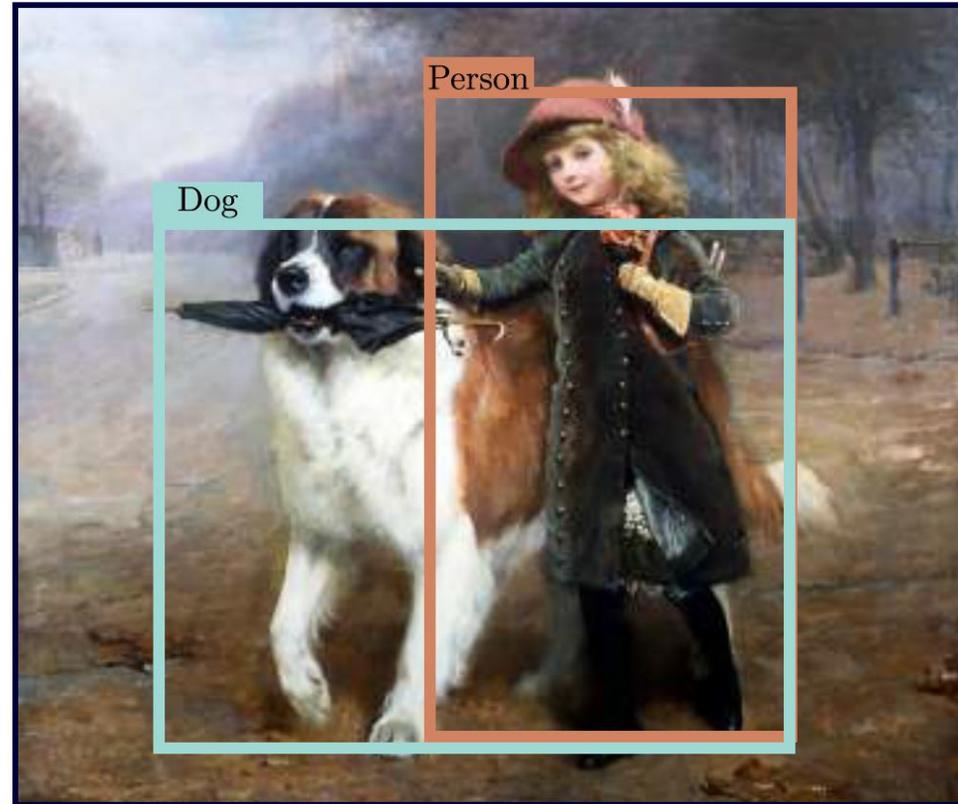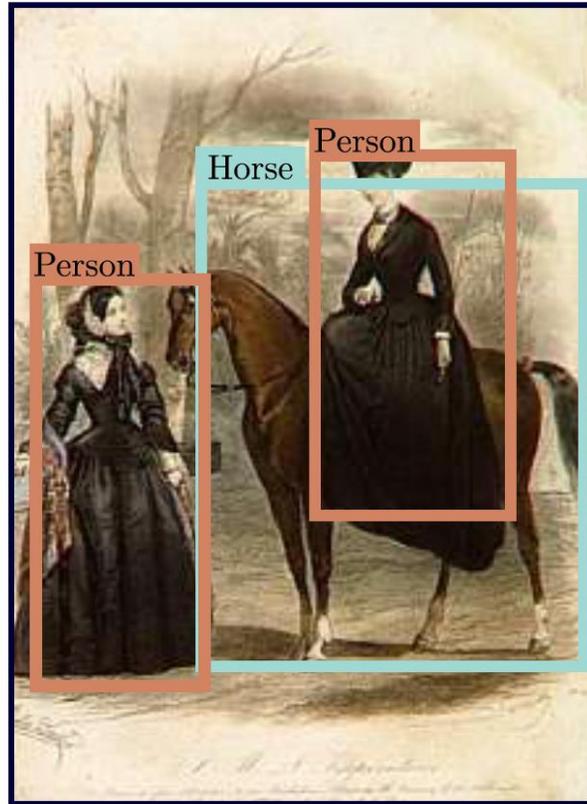
Moving on
- Image classification
- Object detection
- Semantic segmentation

Semantic Segmentation (2015)

Encoder

Decoder
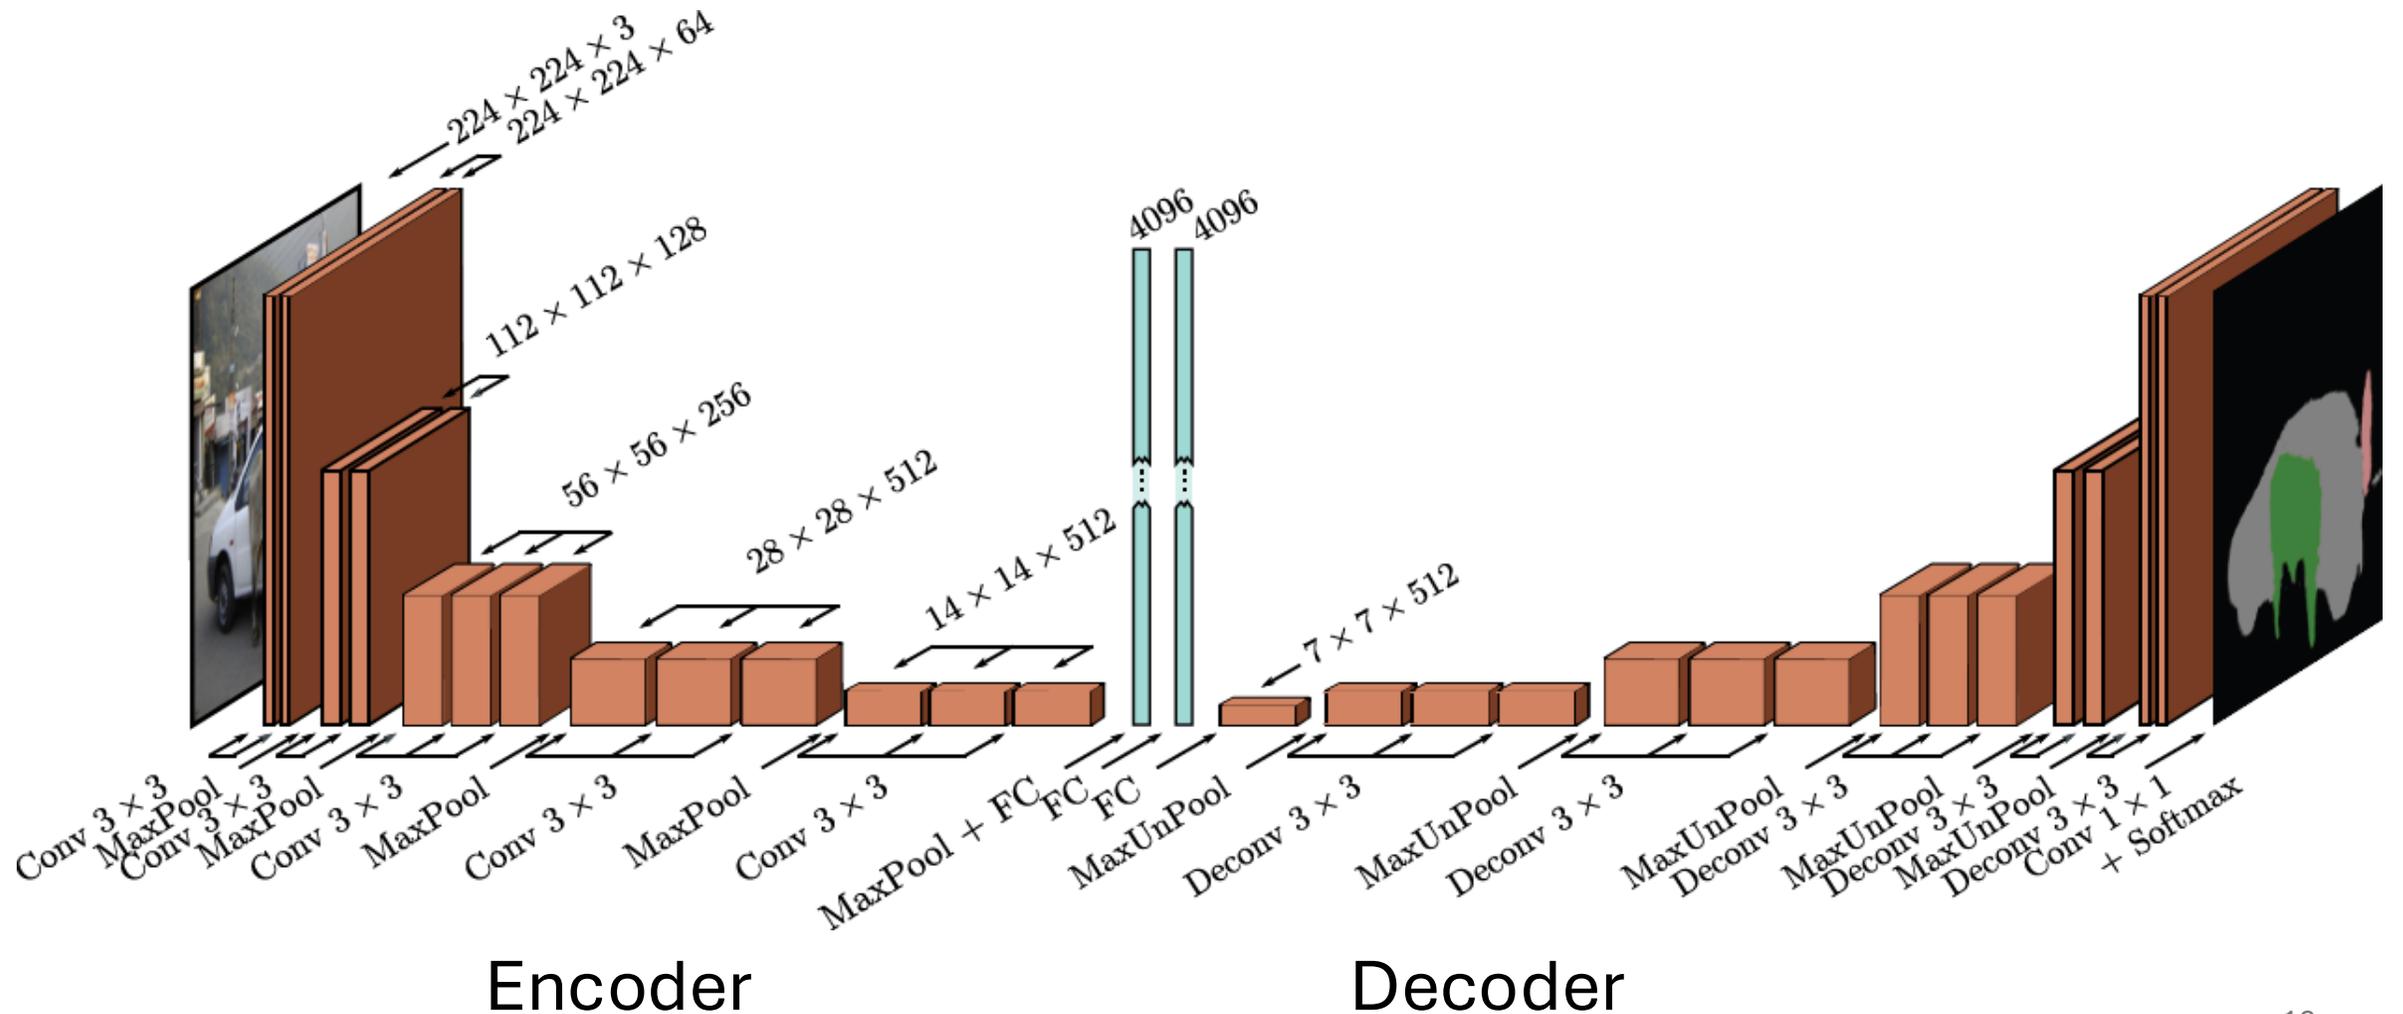
Semantic segmentation results
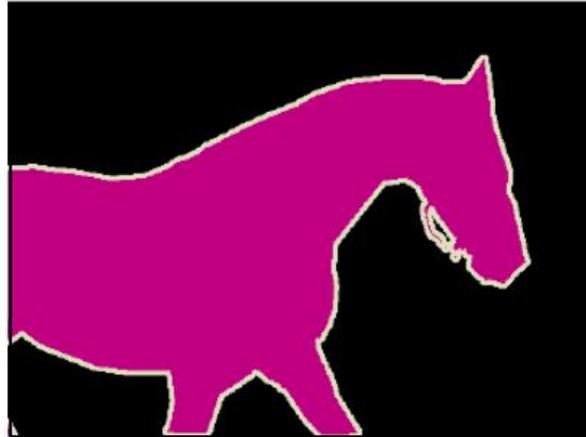
Input · Ground truth · Result
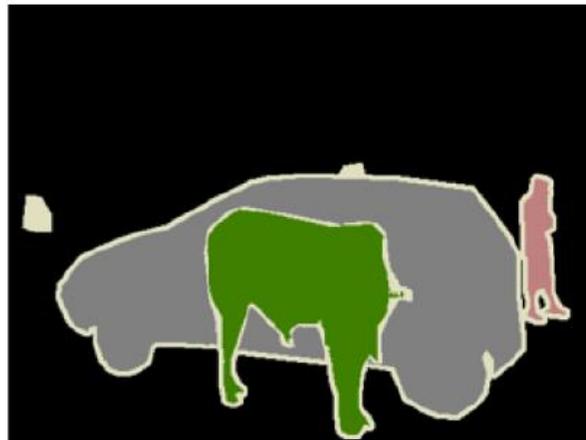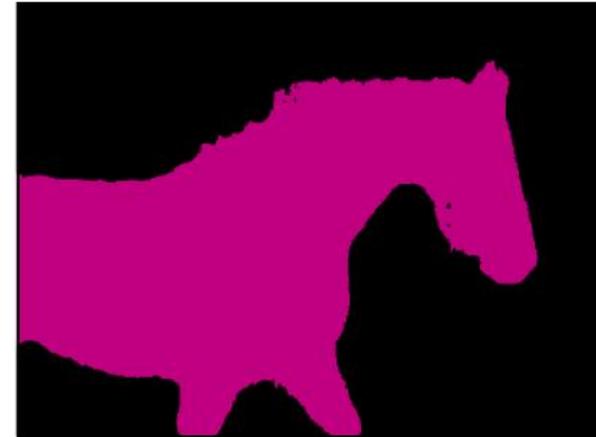
Any Questions?

- Finish up CNN examples (from last time)
- Residual networks

- Challenges of deep networks
- Residual connections and residual blocks
- Exploding gradients in residual networks
- Batch normalization
- Common residual architectures

Previously we saw a sequential network:

$$\mathbf{h}_1 = \mathbf{f}_1[\mathbf{x}, \phi_1]$$
$$\mathbf{h}_2 = \mathbf{f}_2[\mathbf{h}_1, \phi_2]$$
$$\mathbf{h}_3 = \mathbf{f}_3[\mathbf{h}_2, \phi_3]$$
$$\mathbf{y} = \mathbf{f}_4[\mathbf{h}_3, \phi_4]$$
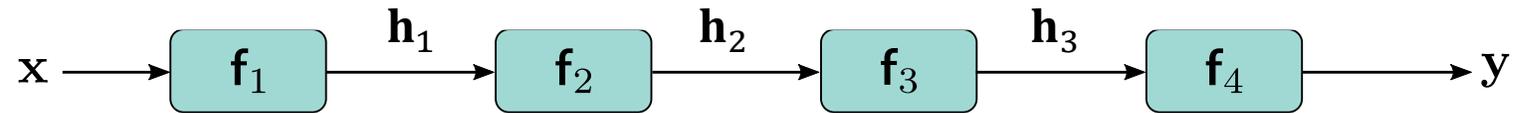
Fully connected network:

$$h_i = \mathrm{a}\left[\beta_i + \sum_{j=1}^{D} \omega_{ij} x_j\right]$$

Convolutional network (e.g. 1 channel → 1 channel):

$$h_i = \mathrm{a}\left[\beta + \omega_1 x_{i-1} + \omega_2 x_i + \omega_3 x_{i+1}\right]$$
$$= \mathrm{a}\left[\beta + \sum_{j=1}^{3} \omega_j x_{i+j-2}\right]$$

Previously we saw a sequential network:

$$\mathbf{h}_1 = \mathbf{f}_1[\mathbf{x}, \boldsymbol{\phi}_1]$$
$$\mathbf{h}_2 = \mathbf{f}_2[\mathbf{h}_1, \boldsymbol{\phi}_2]$$
$$\mathbf{h}_3 = \mathbf{f}_3[\mathbf{h}_2, \boldsymbol{\phi}_3]$$
$$\mathbf{y} = \mathbf{f}_4[\mathbf{h}_3, \boldsymbol{\phi}_4]$$
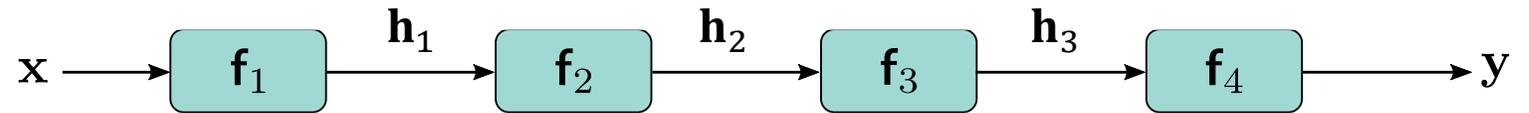
Can think of as a sequence of nested functions:

$$\mathbf{y} = \mathbf{f}_4\left[\mathbf{f}_3\left[\mathbf{f}_2\left[\mathbf{f}_1[\mathbf{x}, \boldsymbol{\phi}_1], \boldsymbol{\phi}_2\right], \boldsymbol{\phi}_3\right], \boldsymbol{\phi}_4\right]$$

More layers are better...
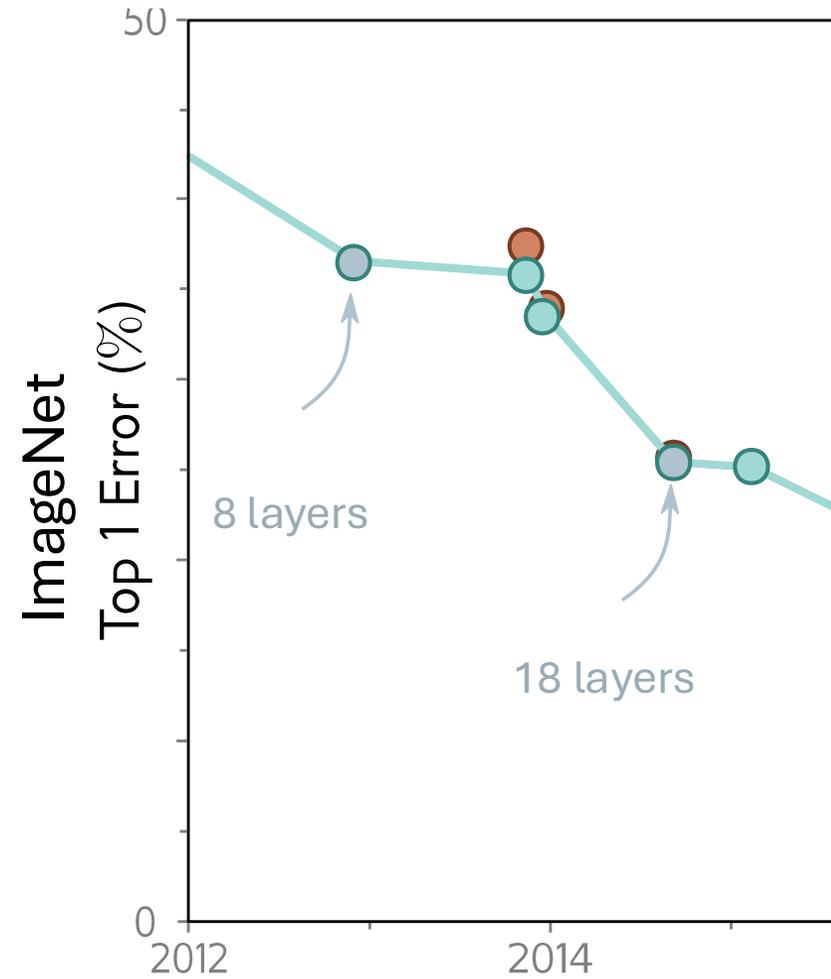
More layers are better... up to a point
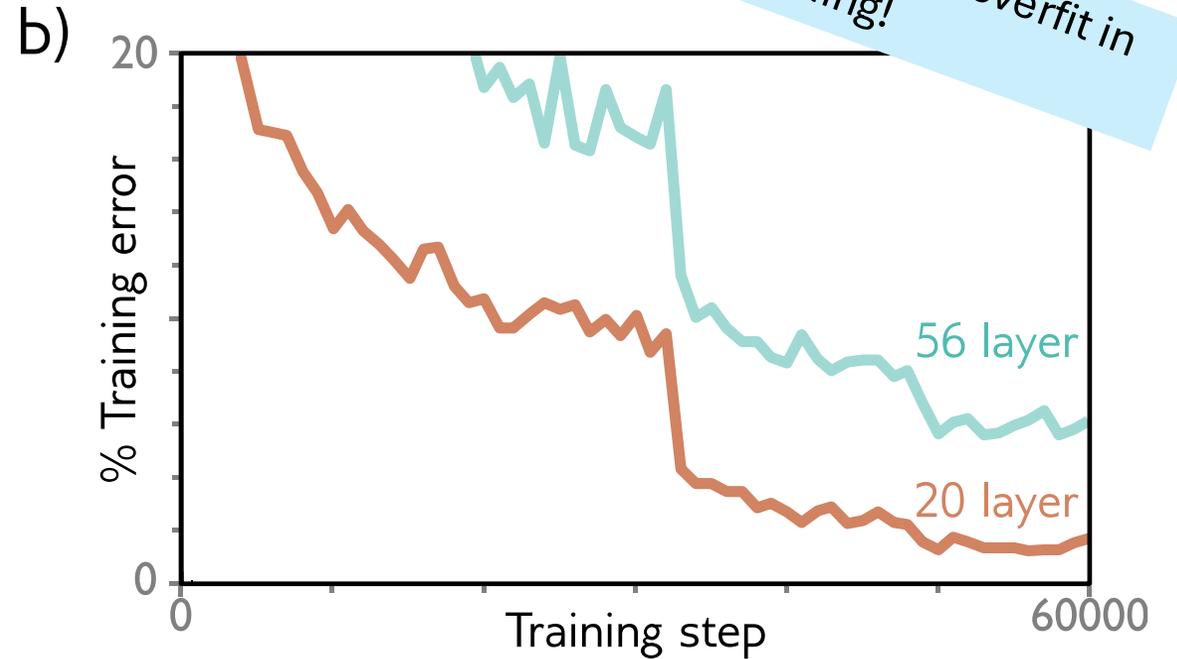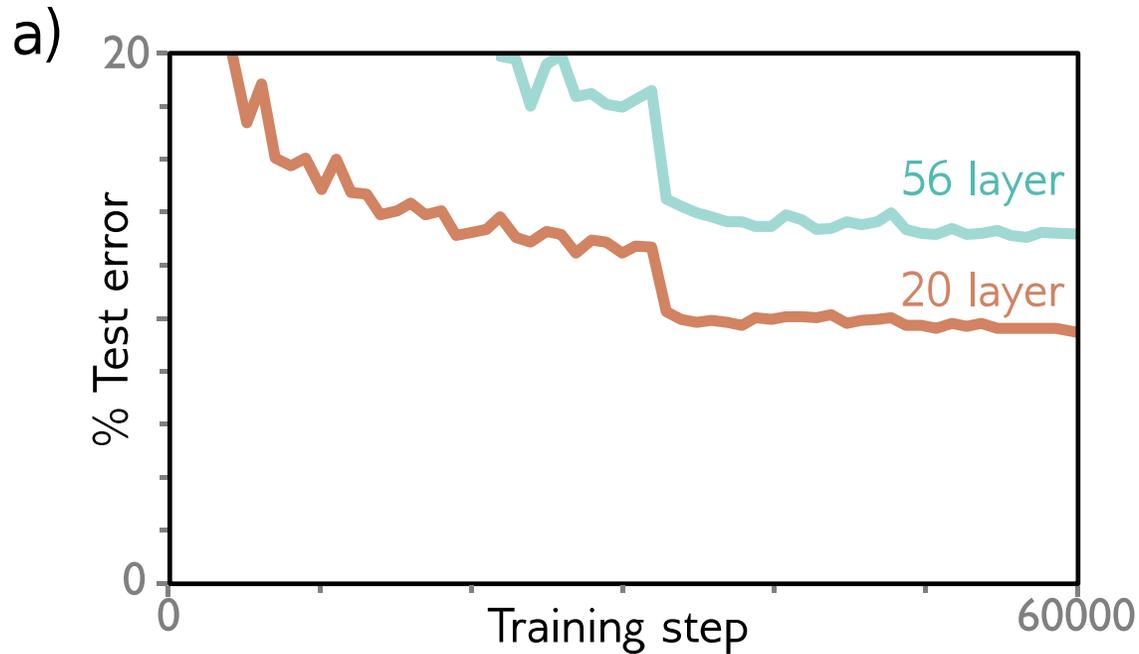
a) Convolutional Network on CIFAR10

b) Can't even overfit in training!

What's going on?

Not completely understood, but...

Take a look at $\partial y/\partial x$ for shallow and deep networks.

a) 1 hidden layer, 200 hidden units — Gradient, dy/dx vs Input, x

b) 24 hidden layers, 200 hidden units each layer — Gradient, dy/dx vs Input, x

c) Autocorrelation vs Δx for 2 hidden layers, 4 hidden layers, 24 hidden layers

Gradients of deeper networks are much less correlated!

Autocorrelation

A small step in gradient descent may jump to wildly different valued gradient!

What's going on?

Not completely understood, but...

Take a look at $\partial y / \partial x$ for shallow and deep networks.

a)

1 hidden layer
200 hidden units

b)

24 hidden layers
200 hidden units each layer

c)

Gradients of deeper networks are much less correlated!

2 hidden layers

4 hidden layers

24 hidden layers

Autocorrelation

A small step in gradient descent may jump to wildly different valued gradient!

What's going on?

$$x \longrightarrow \boxed{f_1} \longrightarrow \boxed{f_2} \longrightarrow \boxed{f_3} \longrightarrow \boxed{f_4} \longrightarrow y$$
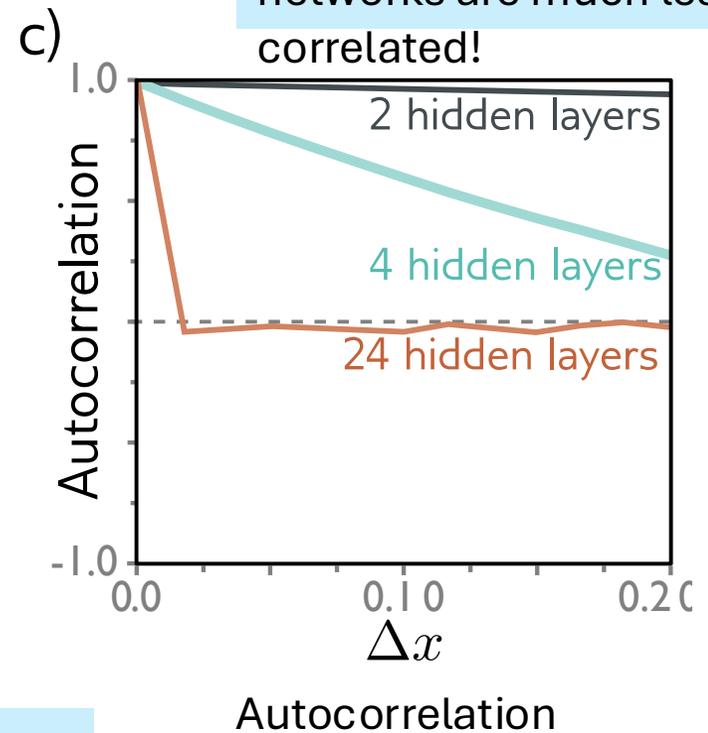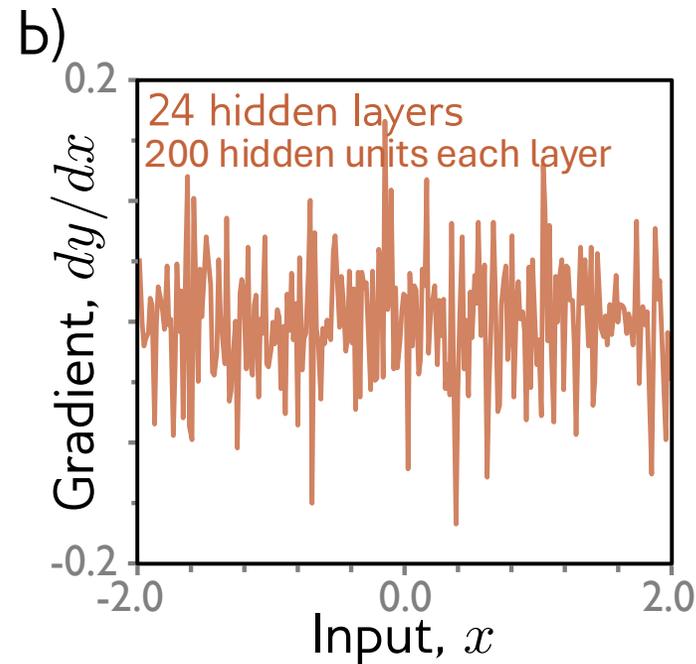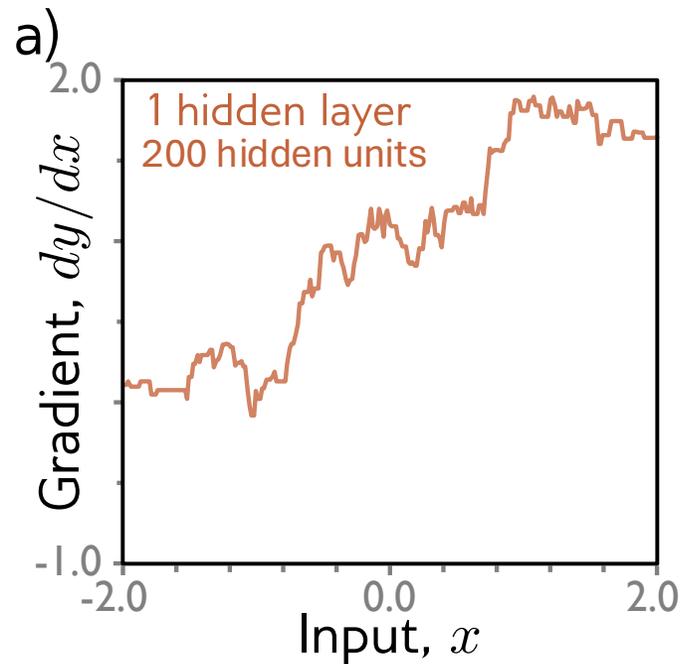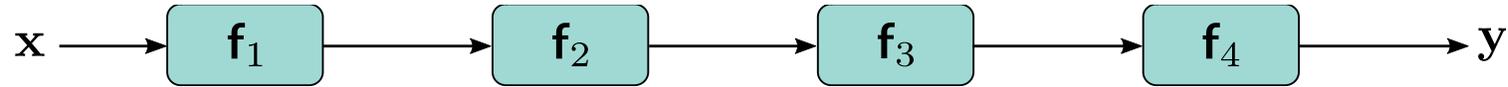
$$\mathbf{y} = \mathbf{f}_4 \left[\mathbf{f}_3 \left[\mathbf{f}_2 \left[\mathbf{f}_1 [\mathbf{x}, \boldsymbol{\phi}_1], \boldsymbol{\phi}_2 \right], \boldsymbol{\phi}_3 \right], \boldsymbol{\phi}_4 \right]$$

The derivative of the output y w.r.t. the first layer f_1 is, by the chain rule:

$$\frac{\partial \mathbf{y}}{\partial \mathbf{f}_1} = \frac{\partial \mathbf{f}_4}{\partial \mathbf{f}_3} \frac{\partial \mathbf{f}_3}{\partial \mathbf{f}_2} \frac{\partial \mathbf{f}_2}{\partial \mathbf{f}_1}$$

f_1 impacts f_2 impacts f_3, etc...

Any Questions?

- Finish up CNN examples (from last time)
- Residual networks

- Challenges of deep networks
- Residual connections and residual blocks
- Exploding gradients in residual networks
- Batch normalization
- Common residual architectures

Solution: Residual connections

$$\mathbf{h}_1 = \mathbf{f}_1[\mathbf{x}, \phi_1]$$
$$\mathbf{h}_2 = \mathbf{f}_2[\mathbf{h}_1, \phi_2]$$
$$\mathbf{h}_3 = \mathbf{f}_3[\mathbf{h}_2, \phi_3]$$
$$\mathbf{y} = \mathbf{f}_4[\mathbf{h}_3, \phi_4]$$

Residual network:

$$\mathbf{h}_1 = \mathbf{x} + \mathbf{f}_1[\mathbf{x}, \phi_1]$$
$$\mathbf{h}_2 = \mathbf{h}_1 + \mathbf{f}_2[\mathbf{h}_1, \phi_2]$$
$$\mathbf{h}_3 = \mathbf{h}_2 + \mathbf{f}_3[\mathbf{h}_2, \phi_3]$$
$$\mathbf{y} = \mathbf{h}_3 + \mathbf{f}_4[\mathbf{h}_3, \phi_4]$$

K. He, X. Zhang, S. Ren, and J. Sun, "Deep Residual Learning for Image Recognition," *arXiv:1512.03385 [cs]*, Dec. 2015, http://arxiv.org/abs/1512.03385
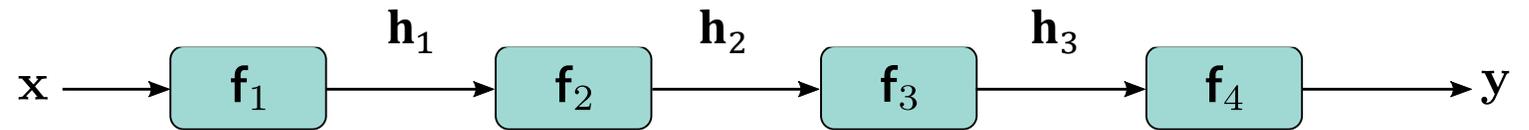
Residual Network

$$\mathbf{h}_1 = \mathbf{x} + \mathbf{f}_1[\mathbf{x}, \phi_1]$$
$$\mathbf{h}_2 = \mathbf{h}_1 + \mathbf{f}_2[\mathbf{h}_1, \phi_2]$$
$$\mathbf{h}_3 = \mathbf{h}_2 + \mathbf{f}_3[\mathbf{h}_2, \phi_3]$$
$$\mathbf{y} = \mathbf{h}_3 + \mathbf{f}_4[\mathbf{h}_3, \phi_4]$$

Residual connections
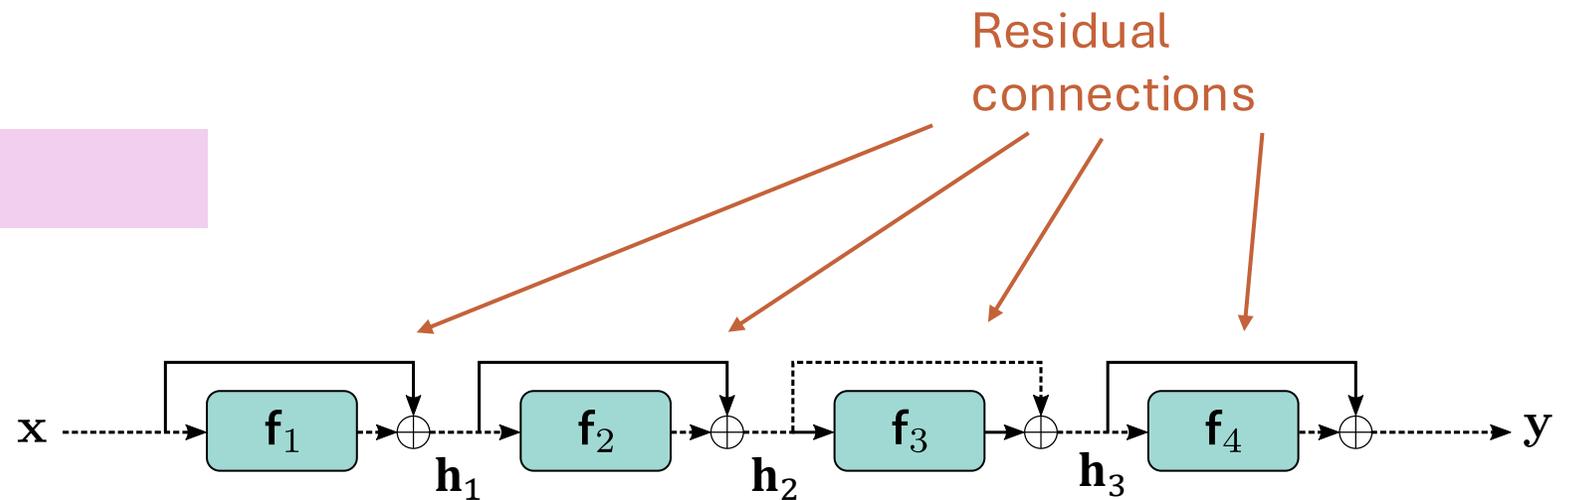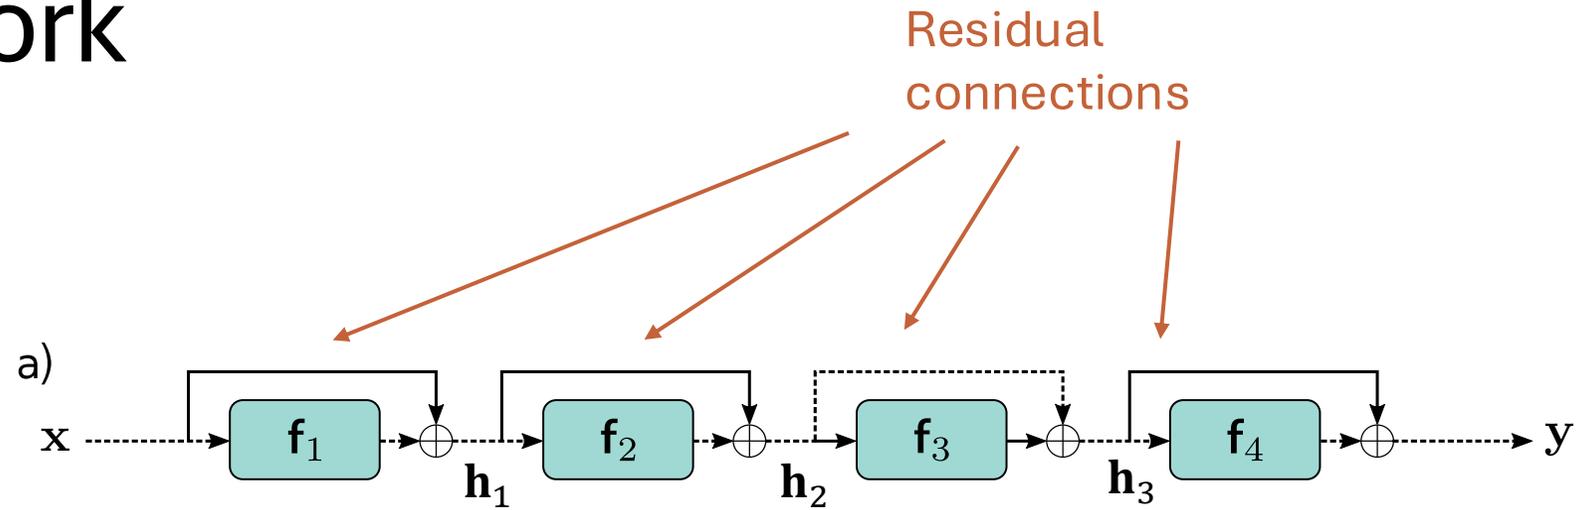
Substituting in:

$$
\begin{aligned}
\mathbf{y} = \mathbf{x} + &\; \mathbf{f}_1[\mathbf{x}] \\
+ &\; \mathbf{f}_2\big[\mathbf{x} + \mathbf{f}_1[\mathbf{x}]\big] \\
+ &\; \mathbf{f}_3\Big[\mathbf{x} + \mathbf{f}_1[\mathbf{x}] + \mathbf{f}_2\big[\mathbf{x} + \mathbf{f}_1[\mathbf{x}]\big]\Big] \\
+ &\; \mathbf{f}_4\Big[\mathbf{x} + \mathbf{f}_1[\mathbf{x}] + \mathbf{f}_2\big[\mathbf{x} + \mathbf{f}_1[\mathbf{x}]\big] + \mathbf{f}_3\big[\mathbf{x} + \mathbf{f}_1[\mathbf{x}] + \mathbf{f}_2[\mathbf{x} + \mathbf{f}_1[\mathbf{x}]]\big]\Big]
\end{aligned}
$$

28

Residual Network

We can unravel all the possible paths

The output is the sum of the input plus 4 partial networks.

$$
\begin{aligned}
\mathbf{y} = \mathbf{x} + \ & \mathbf{f}_1[\mathbf{x}] \\
+ \ & \mathbf{f}_2\big[\mathbf{x} + \mathbf{f}_1[\mathbf{x}]\big] \\
+ \ & \mathbf{f}_3\Big[\mathbf{x} + \mathbf{f}_1[\mathbf{x}] + \mathbf{f}_2\big[\mathbf{x} + \mathbf{f}_1[\mathbf{x}]\big]\Big] \\
+ \ & \mathbf{f}_4\Big[\mathbf{x} + \mathbf{f}_1[\mathbf{x}] + \mathbf{f}_2\big[\mathbf{x} + \mathbf{f}_1[\mathbf{x}]\big] + \mathbf{f}_3\big[\mathbf{x} + \mathbf{f}_1[\mathbf{x}] + \mathbf{f}_2[\mathbf{x} + \mathbf{f}_1[\mathbf{x}]]\big]\Big]
\end{aligned}
$$

29
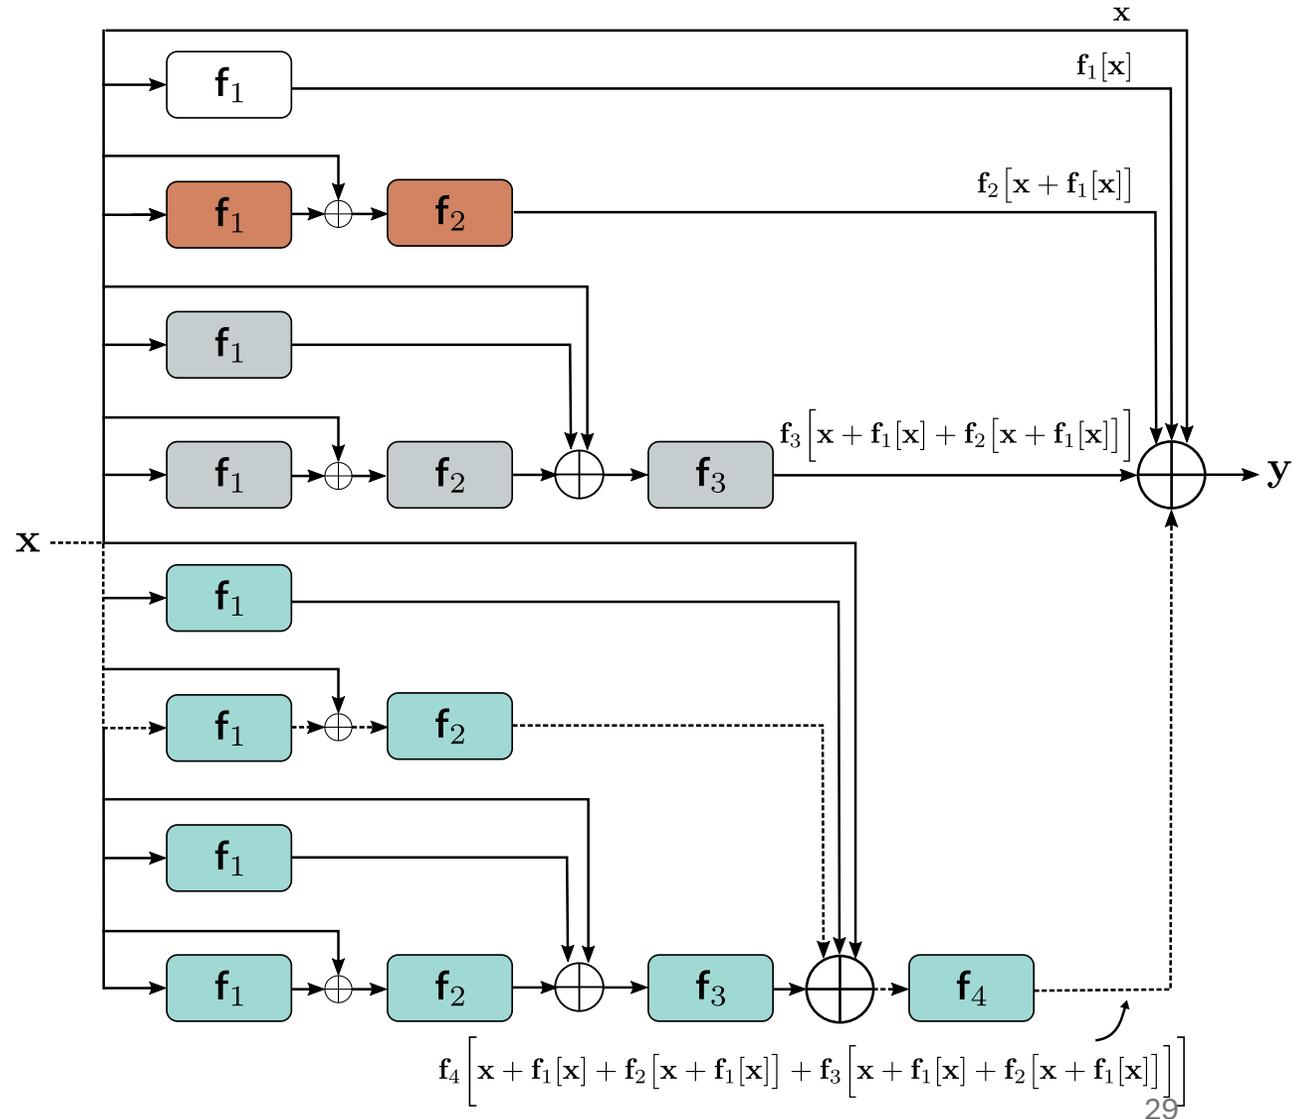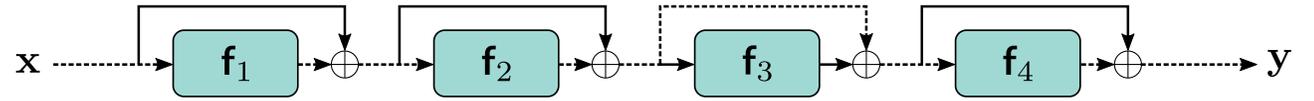
Residual Network as Ensemble of Networks
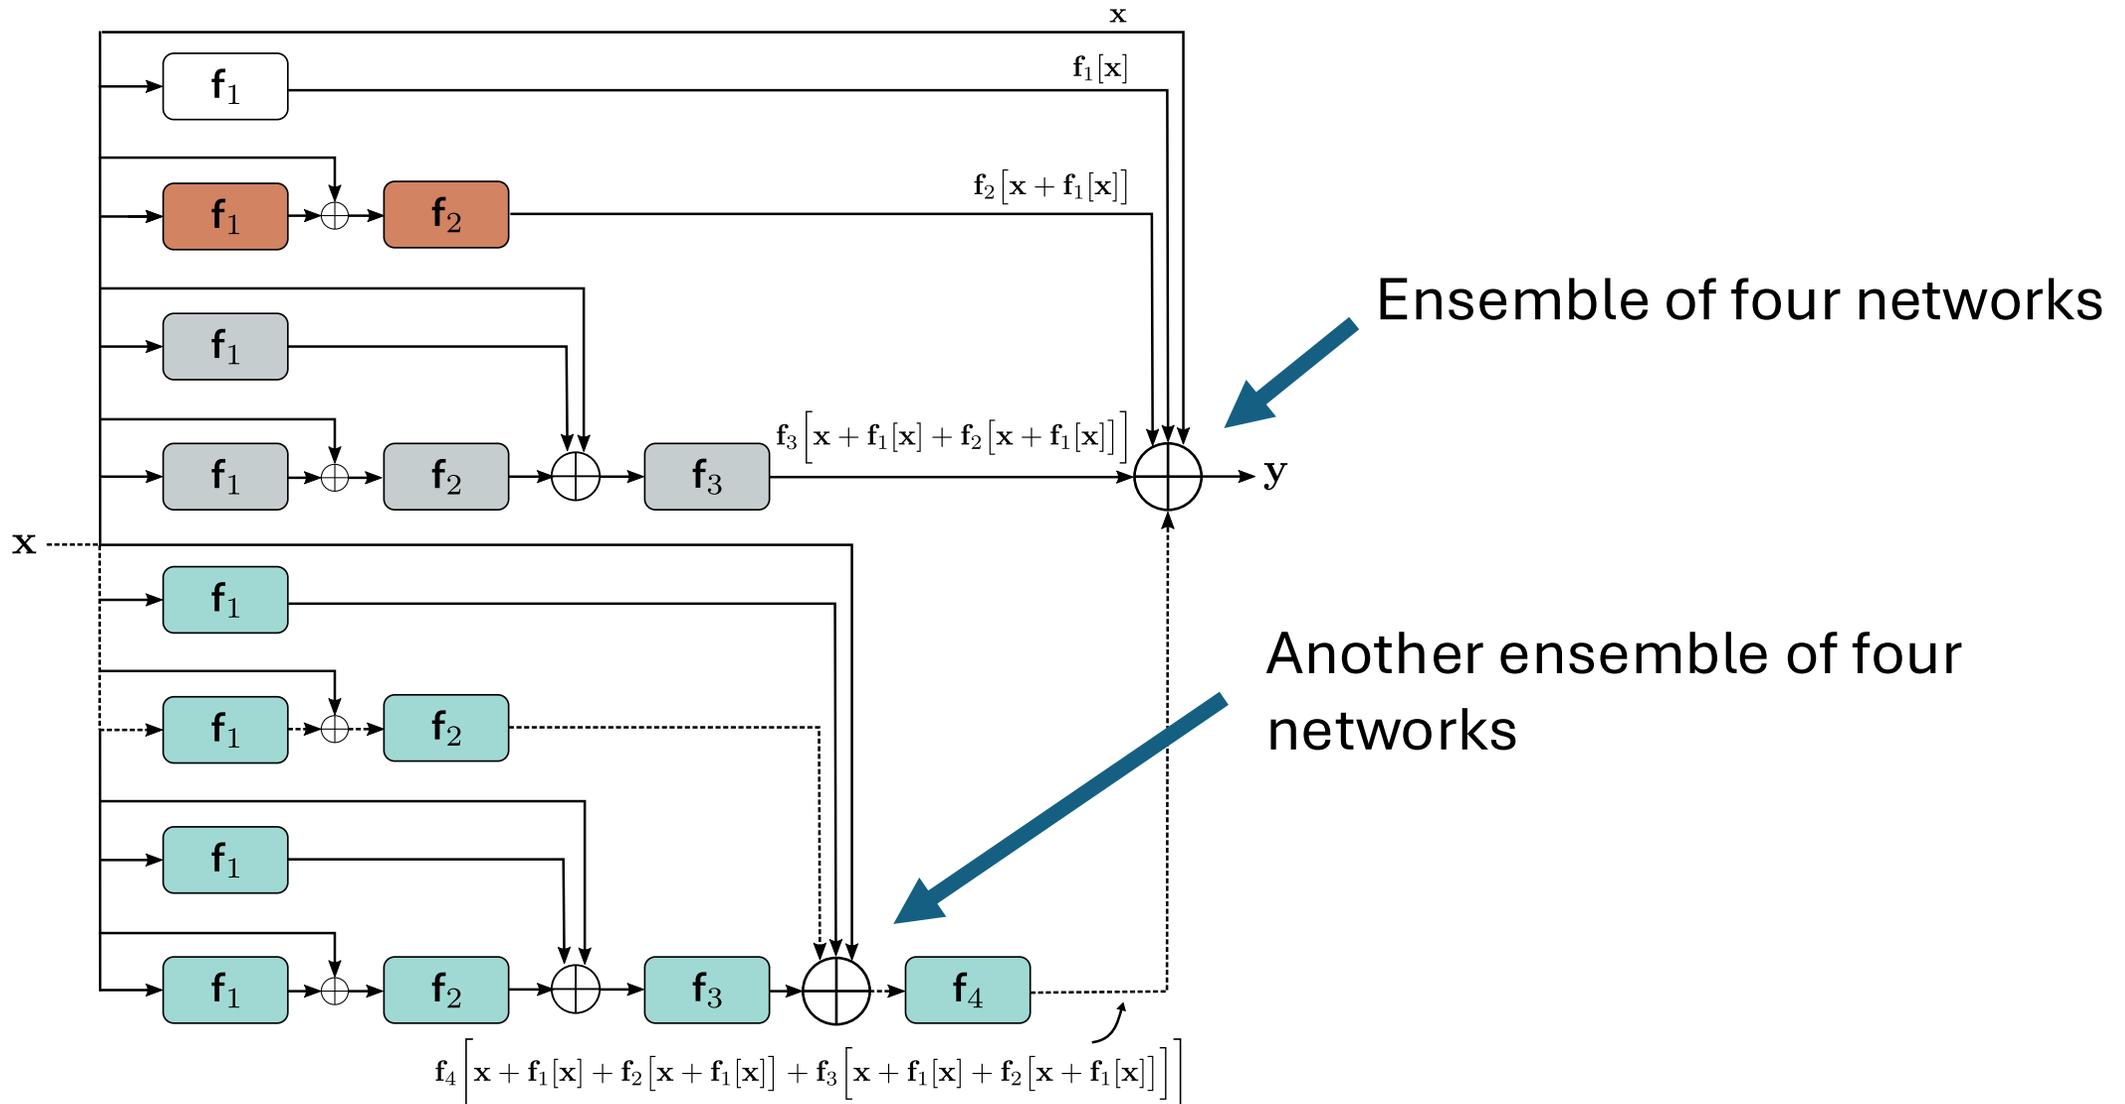
Residual Network as Ensemble of Networks
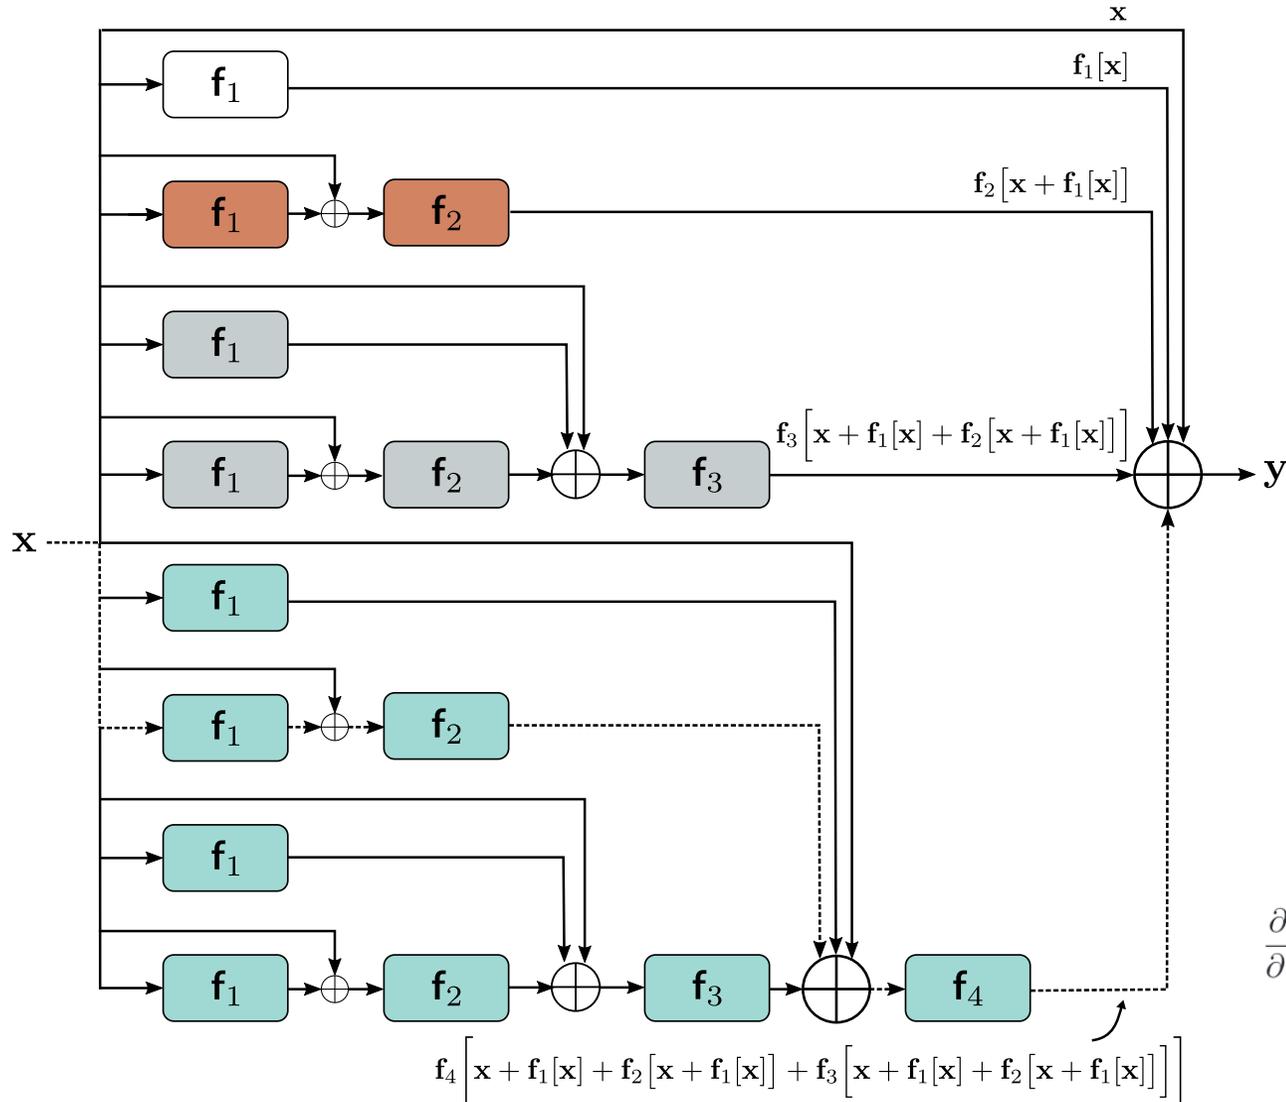
- 16 possible paths through the network!

- 8 paths include f_1

- The influence of f_1 on $\partial y / \partial f_1$ takes 8 different forms

- Gradients on shorter paths generally better behaved.

$$\frac{\partial \mathbf{y}}{\partial \mathbf{f}_1} = \mathbf{I} + \frac{\partial \mathbf{f}_2}{\partial \mathbf{f}_1} + \left(\frac{\partial \mathbf{f}_3}{\partial \mathbf{f}_1} + \frac{\partial \mathbf{f}_3}{\partial \mathbf{f}_2} \frac{\partial \mathbf{f}_2}{\partial \mathbf{f}_1} \right) + \left(\frac{\partial \mathbf{f}_4}{\partial \mathbf{f}_1} + \frac{\partial \mathbf{f}_4}{\partial \mathbf{f}_2} \frac{\partial \mathbf{f}_2}{\partial \mathbf{f}_1} + \frac{\partial \mathbf{f}_4}{\partial \mathbf{f}_3} \frac{\partial \mathbf{f}_3}{\partial \mathbf{f}_1} + \frac{\partial \mathbf{f}_4}{\partial \mathbf{f}_3} \frac{\partial \mathbf{f}_3}{\partial \mathbf{f}_2} \frac{\partial \mathbf{f}_2}{\partial \mathbf{f}_1} \right)$$

Residual Network as Ensemble of Networks
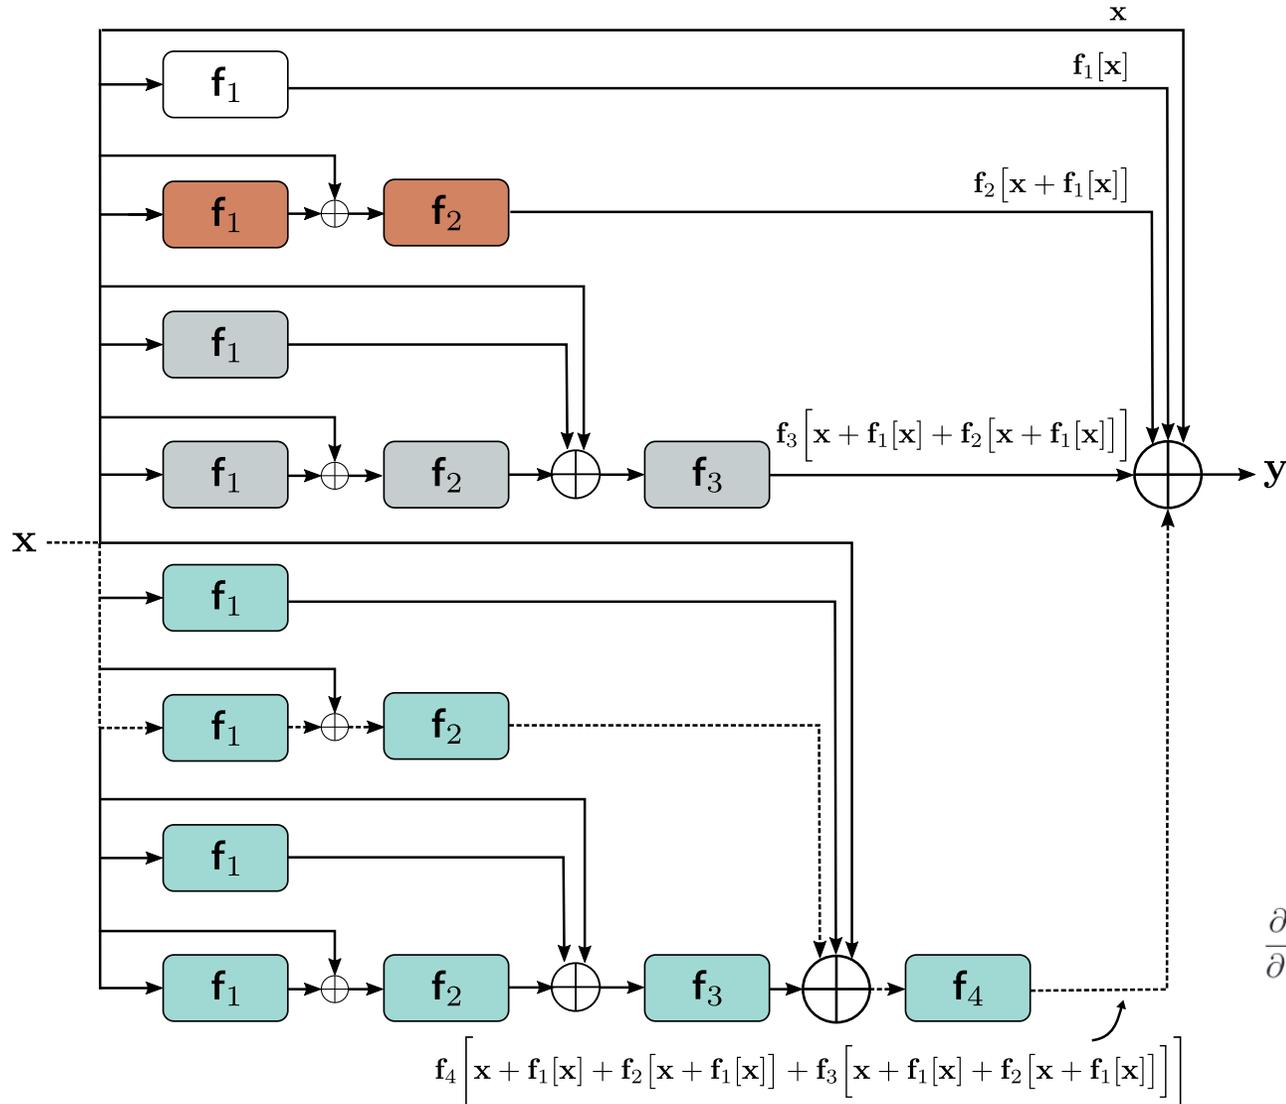
During training, the model can amplify or attenuate the different paths to achieve the best results

$$\frac{\partial \mathbf{y}}{\partial \mathbf{f}_1} = \mathbf{I} + \frac{\partial \mathbf{f}_2}{\partial \mathbf{f}_1} + \left(\frac{\partial \mathbf{f}_3}{\partial \mathbf{f}_1} + \frac{\partial \mathbf{f}_3}{\partial \mathbf{f}_2} \frac{\partial \mathbf{f}_2}{\partial \mathbf{f}_1} \right) + \left(\frac{\partial \mathbf{f}_4}{\partial \mathbf{f}_1} + \frac{\partial \mathbf{f}_4}{\partial \mathbf{f}_2} \frac{\partial \mathbf{f}_2}{\partial \mathbf{f}_1} + \frac{\partial \mathbf{f}_4}{\partial \mathbf{f}_3} \frac{\partial \mathbf{f}_3}{\partial \mathbf{f}_1} + \frac{\partial \mathbf{f}_4}{\partial \mathbf{f}_3} \frac{\partial \mathbf{f}_3}{\partial \mathbf{f}_2} \frac{\partial \mathbf{f}_2}{\partial \mathbf{f}_1} \right)$$

Order of operations is important

a)
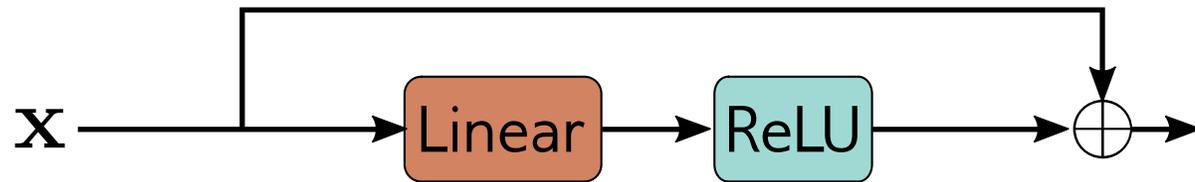
Can only add to the residual because of the ReLU

b)
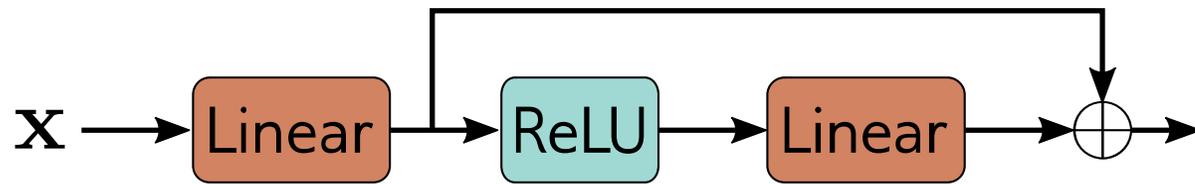
More flexible approach to end with linear block.
Starting with linear block gives us some flexibility on spatial resolution.

c)
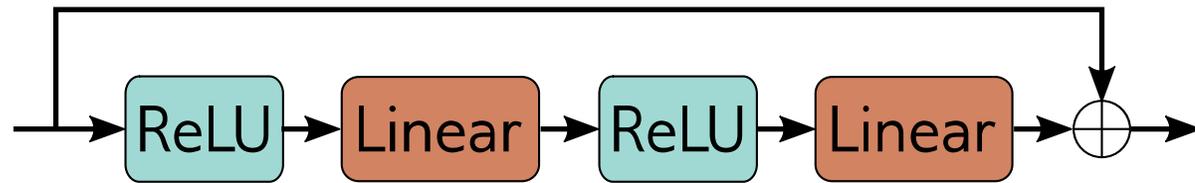
Note: if we start with a ReLU, then will clamp negative values and so do nothing

This helps increase depth up to a point...

Any Questions?

???

- Challenges of deep networks
- Residual connections and residual blocks
- Exploding gradients in residual networks
- Batch normalization
- Common residual architectures

Exploding Gradients in Residual Networks

x 1 1 $\mathbf{f}_1$ 1 1 2 2 $\mathbf{f}_2$ 2 2 4 4 $\mathbf{f}_3$ 4 4 8
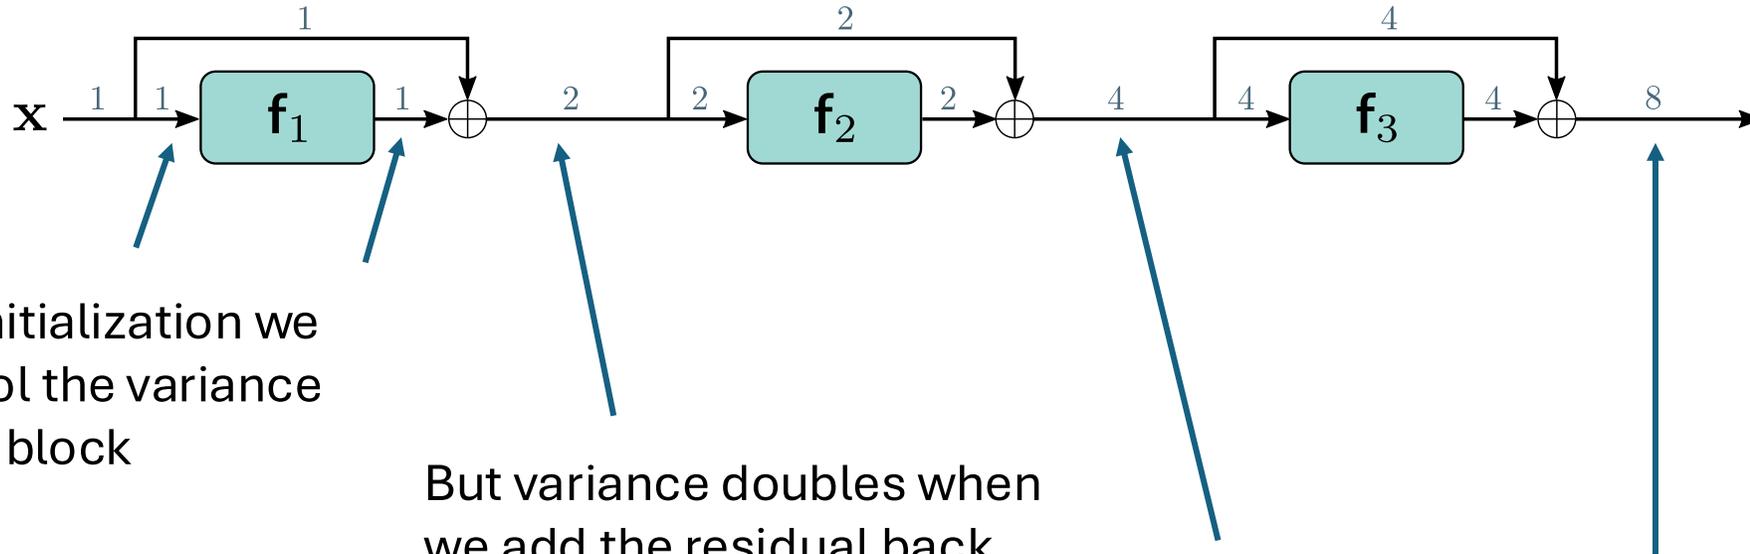
With He initialization we can control the variance inside the block

But variance doubles when we add the residual back in.

And then grows exponentially.

Exploding Gradients in Residual Networks
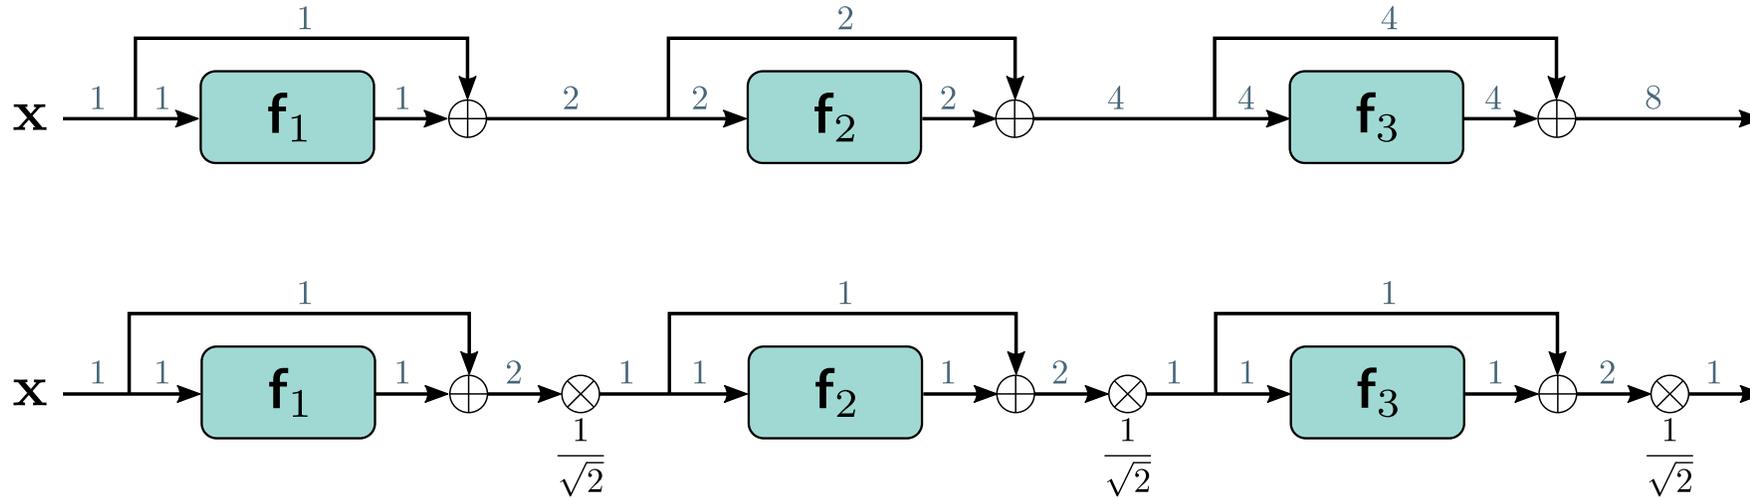
Could stabilize by renormalizing after adding each residual.

More common to apply *batch normalization*.

Plan for Today

- Finish up CNN examples (from last time)
- Residual networks

- Challenges of deep networks
- Residual connections and residual blocks
- Exploding gradients in residual networks
- Batch normalization
- Common residual architectures

Batch Normalization (a.k.a. *BatchNorm*)
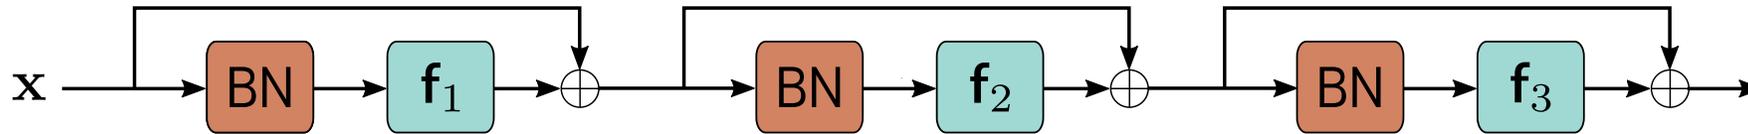
- Shifts and rescales each activation so that its mean and variance across the batch become values that are learned during training

S. Ioffe and C. Szegedy, "Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift," *arXiv:1502.03167 [cs]*, Mar. 2015, http://arxiv.org/abs/1502.03167

Batch Normalization (a.k.a. *BatchNorm*)
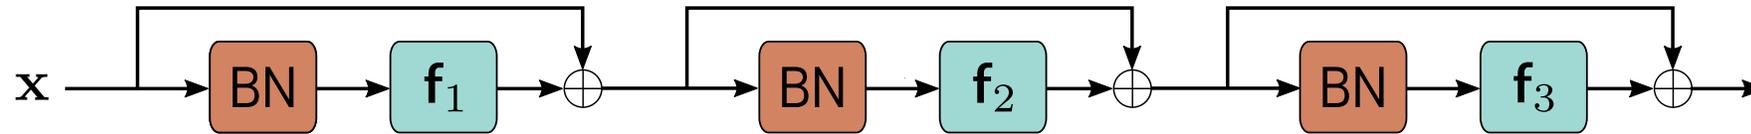
- Shifts and rescales each activation so that its mean and variance across the batch become values that are learned during training

Calculate the sample *mean* and *standard deviation* for each hidden unit across samples of the batch.

$$m_h = \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} h_i$$

$$s_h = \sqrt{\frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} (h_i - m_h)^2}$$

Batch Normalization (a.k.a. *BatchNorm*)
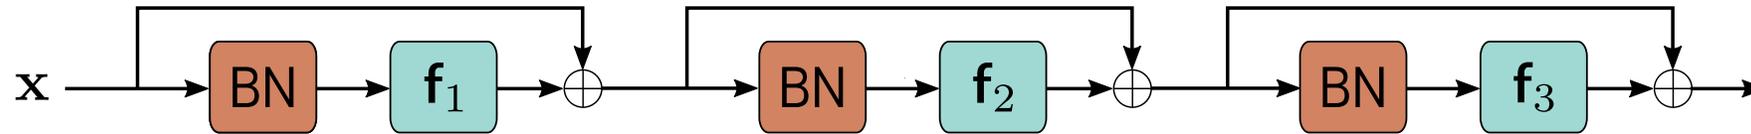
- Shifts and rescales each activation so that its mean and variance across the batch become values that are learned during training

Calculate the sample *mean* and *standard deviation* for each hidden unit across samples of the batch.

$$m_h = \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} h_i$$

$$s_h = \sqrt{\frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} (h_i - m_h)^2}$$

Standardize (*normalize*) to zero-mean and unit standard deviation.

$$\hat{h}_i \leftarrow \frac{h_i - m_h}{s_h + \epsilon} \qquad \forall i \in \mathcal{B},$$

Scale by γ and shift by δ, which are *learned* parameters.

$$h_i \leftarrow \gamma h_i + \delta \qquad \forall i \in \mathcal{B}.$$

Batch Normalization (a.k.a. *BatchNorm*)
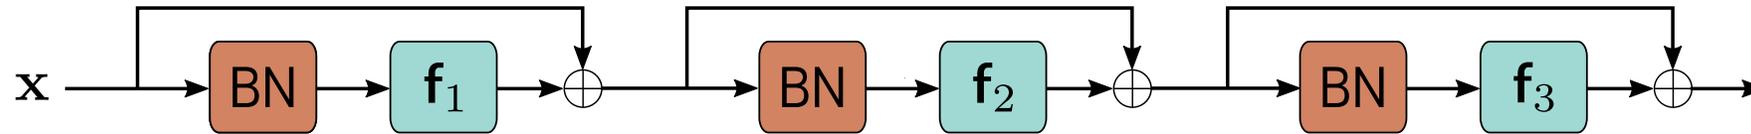
- Applied independently to each hidden unit

- Standard FC Network with K layers, each with D hidden units:
  KD learned scales, γ , and KD learned offset, δ

- Convolutional Network with K layers, each with C channels:
  KC learned scales, γ , and KC learned offset, δ

Benefits of BatchNorm
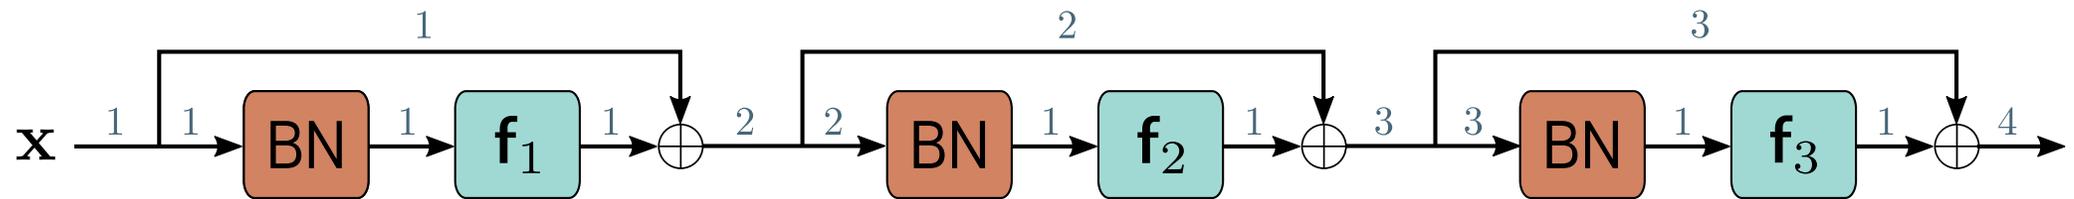
**Stable forward propagation**

- Initialize offsets δ to zero and scales γ to 1

- Variance now increases linearly

- k^{th} block adds *one unit of variance* to variance of k

- At initialization, later layers make smaller relative change to overall variation

- During training, the scales can increase in later layers if helpful
  $\rightarrow$control the effective depth

Benefits of BatchNorm

Supports higher learning rates

Makes the loss surface smoother (reduces shattered gradients)

H. Li, Z. Xu, G. Taylor, C. Studer, and T. Goldstein, "Visualizing the Loss Landscape of Neural Nets," arXiv.org,   https://arxiv.org/abs/1712.09913v3
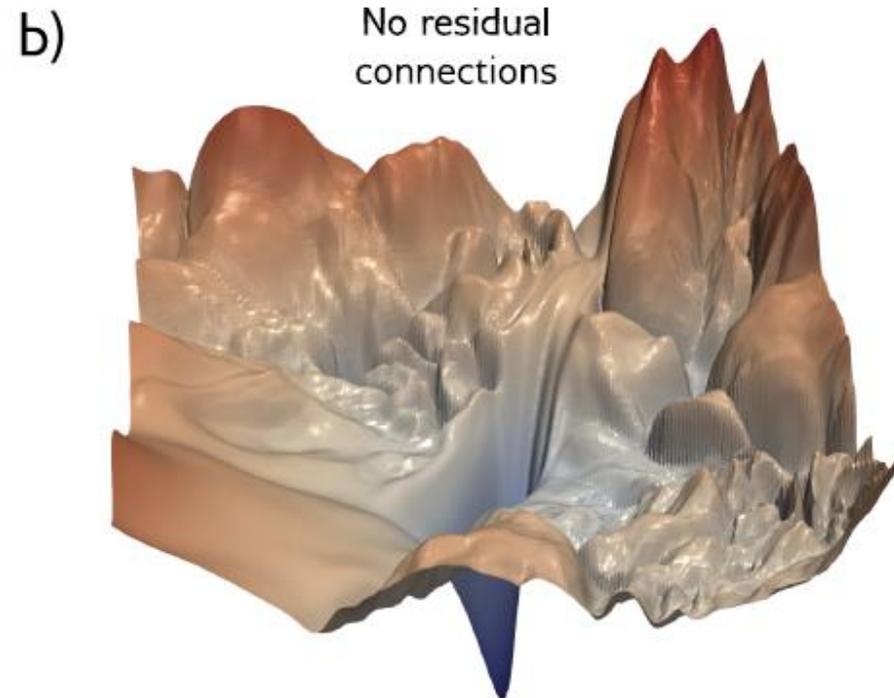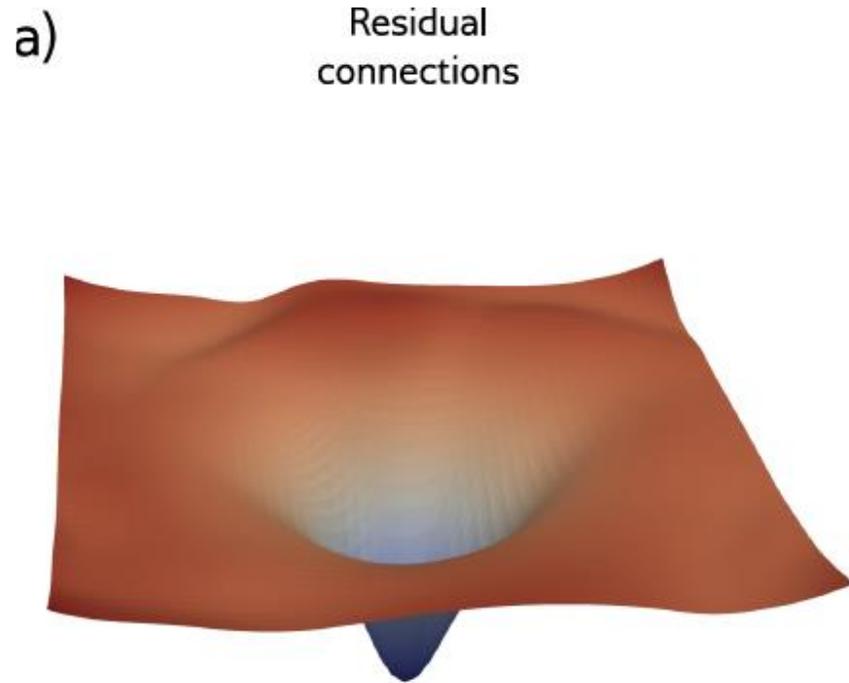
Benefits of BatchNorm

Regularization via added noise

BatchNorm injects noise since BN scale and shift depend on batch statistics.

Disadvantages of Batch Normalization

- Batch normalization makes results dependent on what else is in the same batch.
  - Much more likely if you group on target value.
- Layer normalization instead?
  - Similar spirit.
  - Normalize hidden layer activations of same input.
- Both actively used nowadays.

Any questions?

???

- Challenges of deep networks
- Residual connections and residual blocks
- Exploding gradients in residual networks
- Batch normalization
- Common residual architectures

ResNet (2015)

ResNet Block

Bottleneck Residual

K. He, X. Zhang, S. Ren, and J. Sun, "Deep Residual Learning for Image Recognition," *arXiv:1512.03385 [cs]*, Dec. 2015,
http://arxiv.org/abs/1512.03385
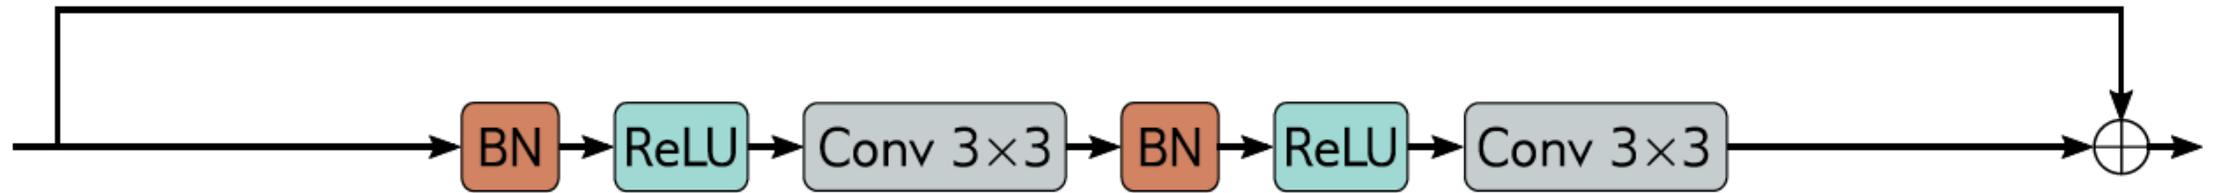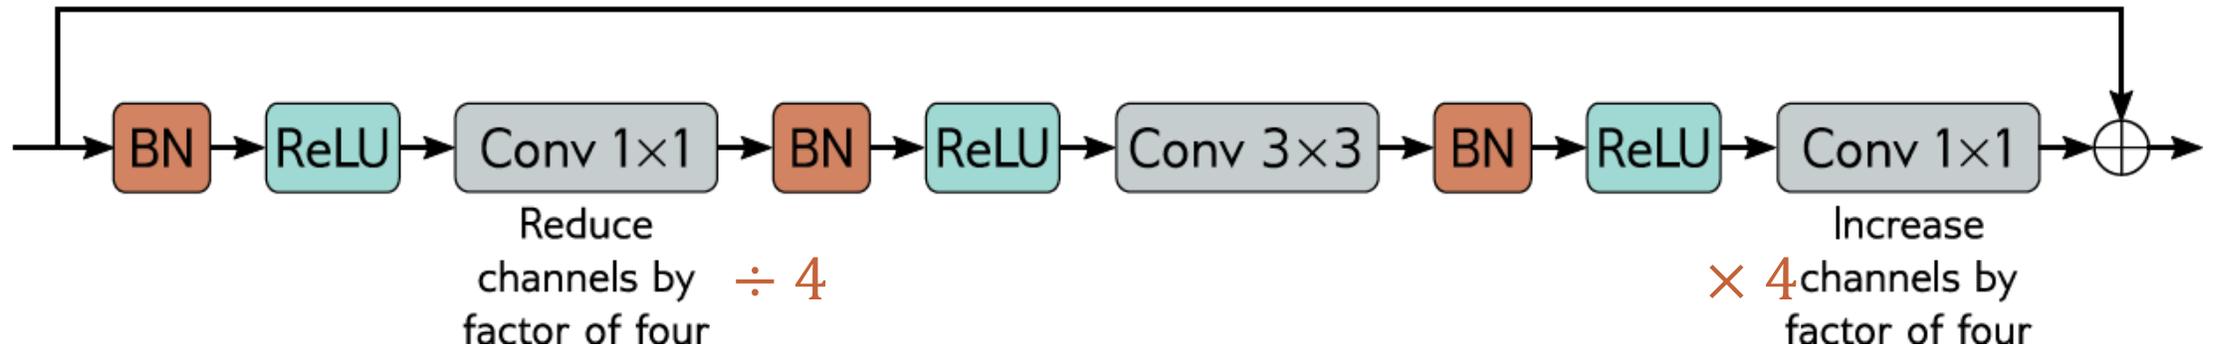
Resnet 200 (2016) for ImageNet Classification
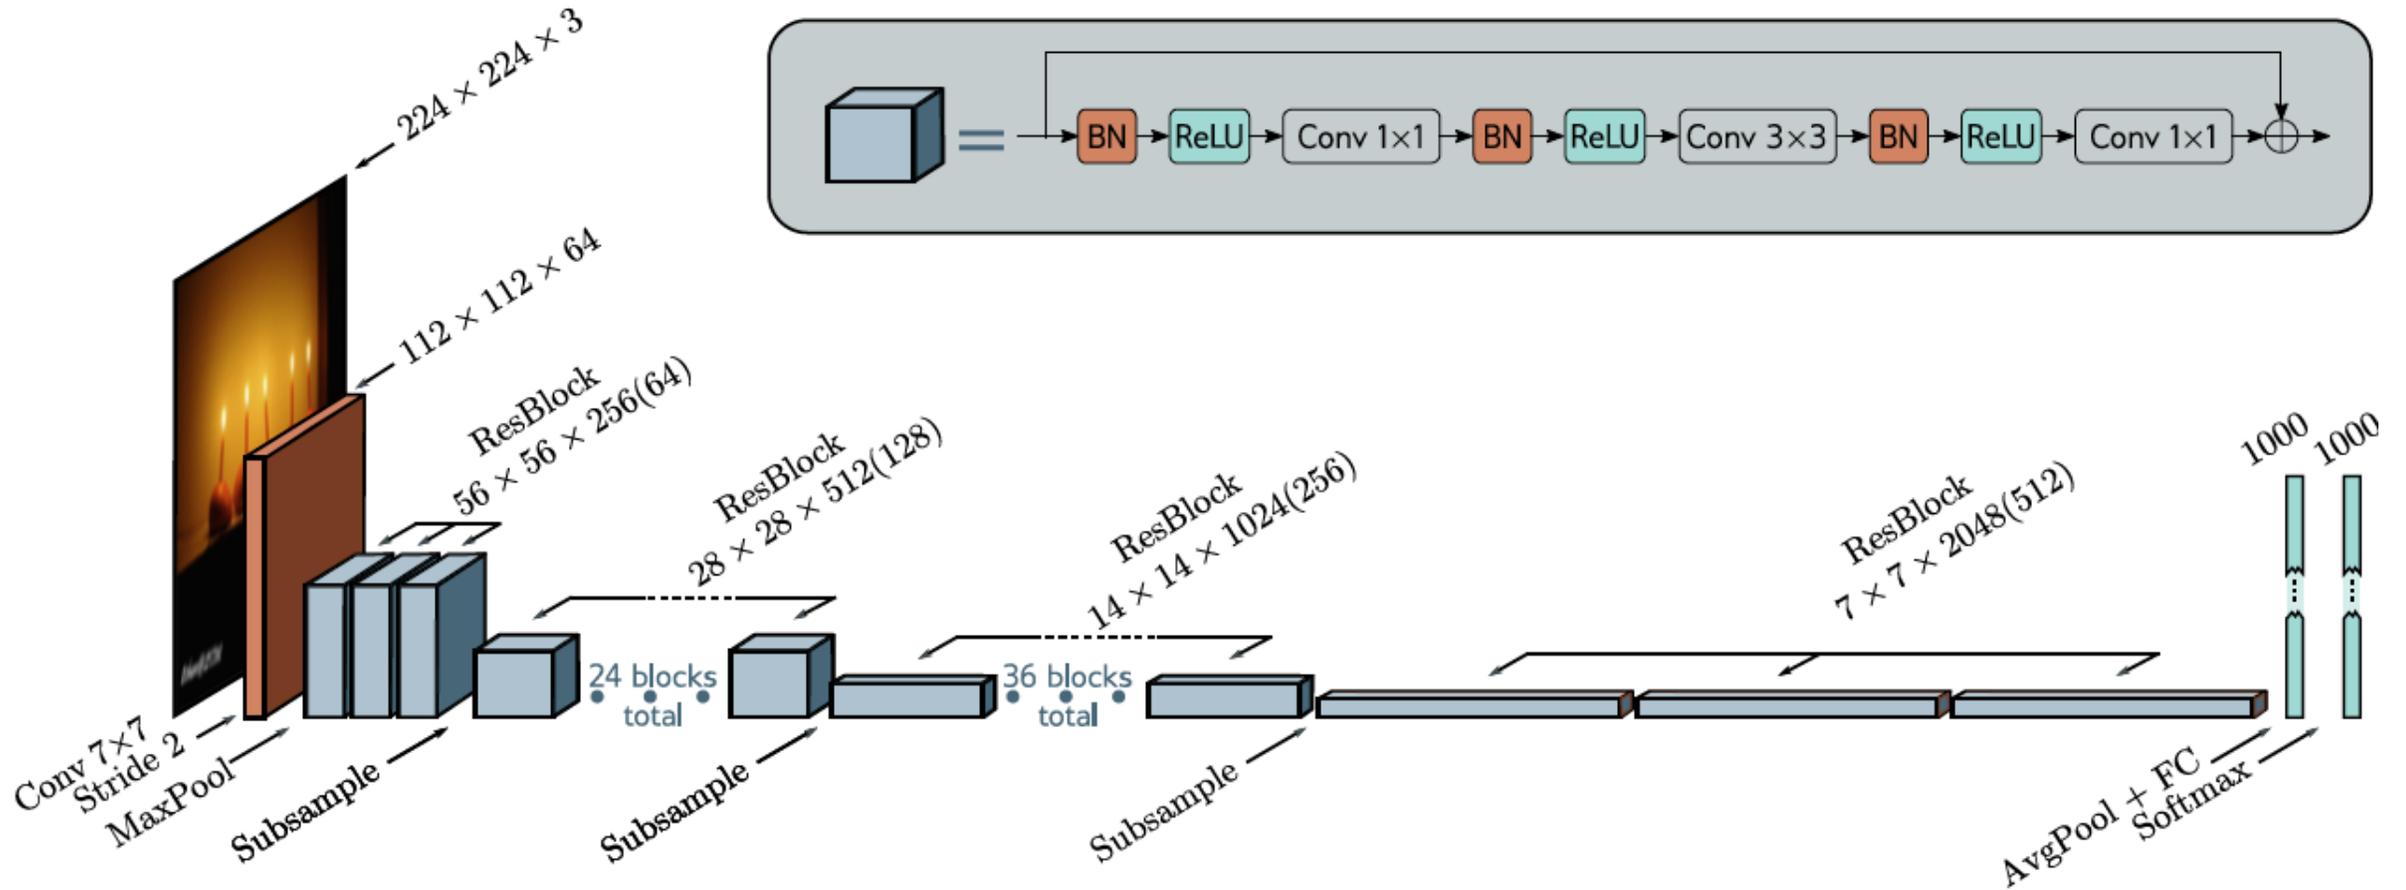
K. He, X. Zhang, S. Ren, and J. Sun, "Deep Residual Learning for Image Recognition," *arXiv:1512.03385 [cs]*, Dec. 2015, http://arxiv.org/abs/1512.03385

ImageNet History

DenseNet

Concatenate to output   Concatenate to output   Concatenate to output

BN → ReLU → Conv 3×3   BN → ReLU → Conv 3×3   BN → ReLU → Conv 3×3

3 channels   32+3=35 channels   32+35=67 channels   32+67=99 channels
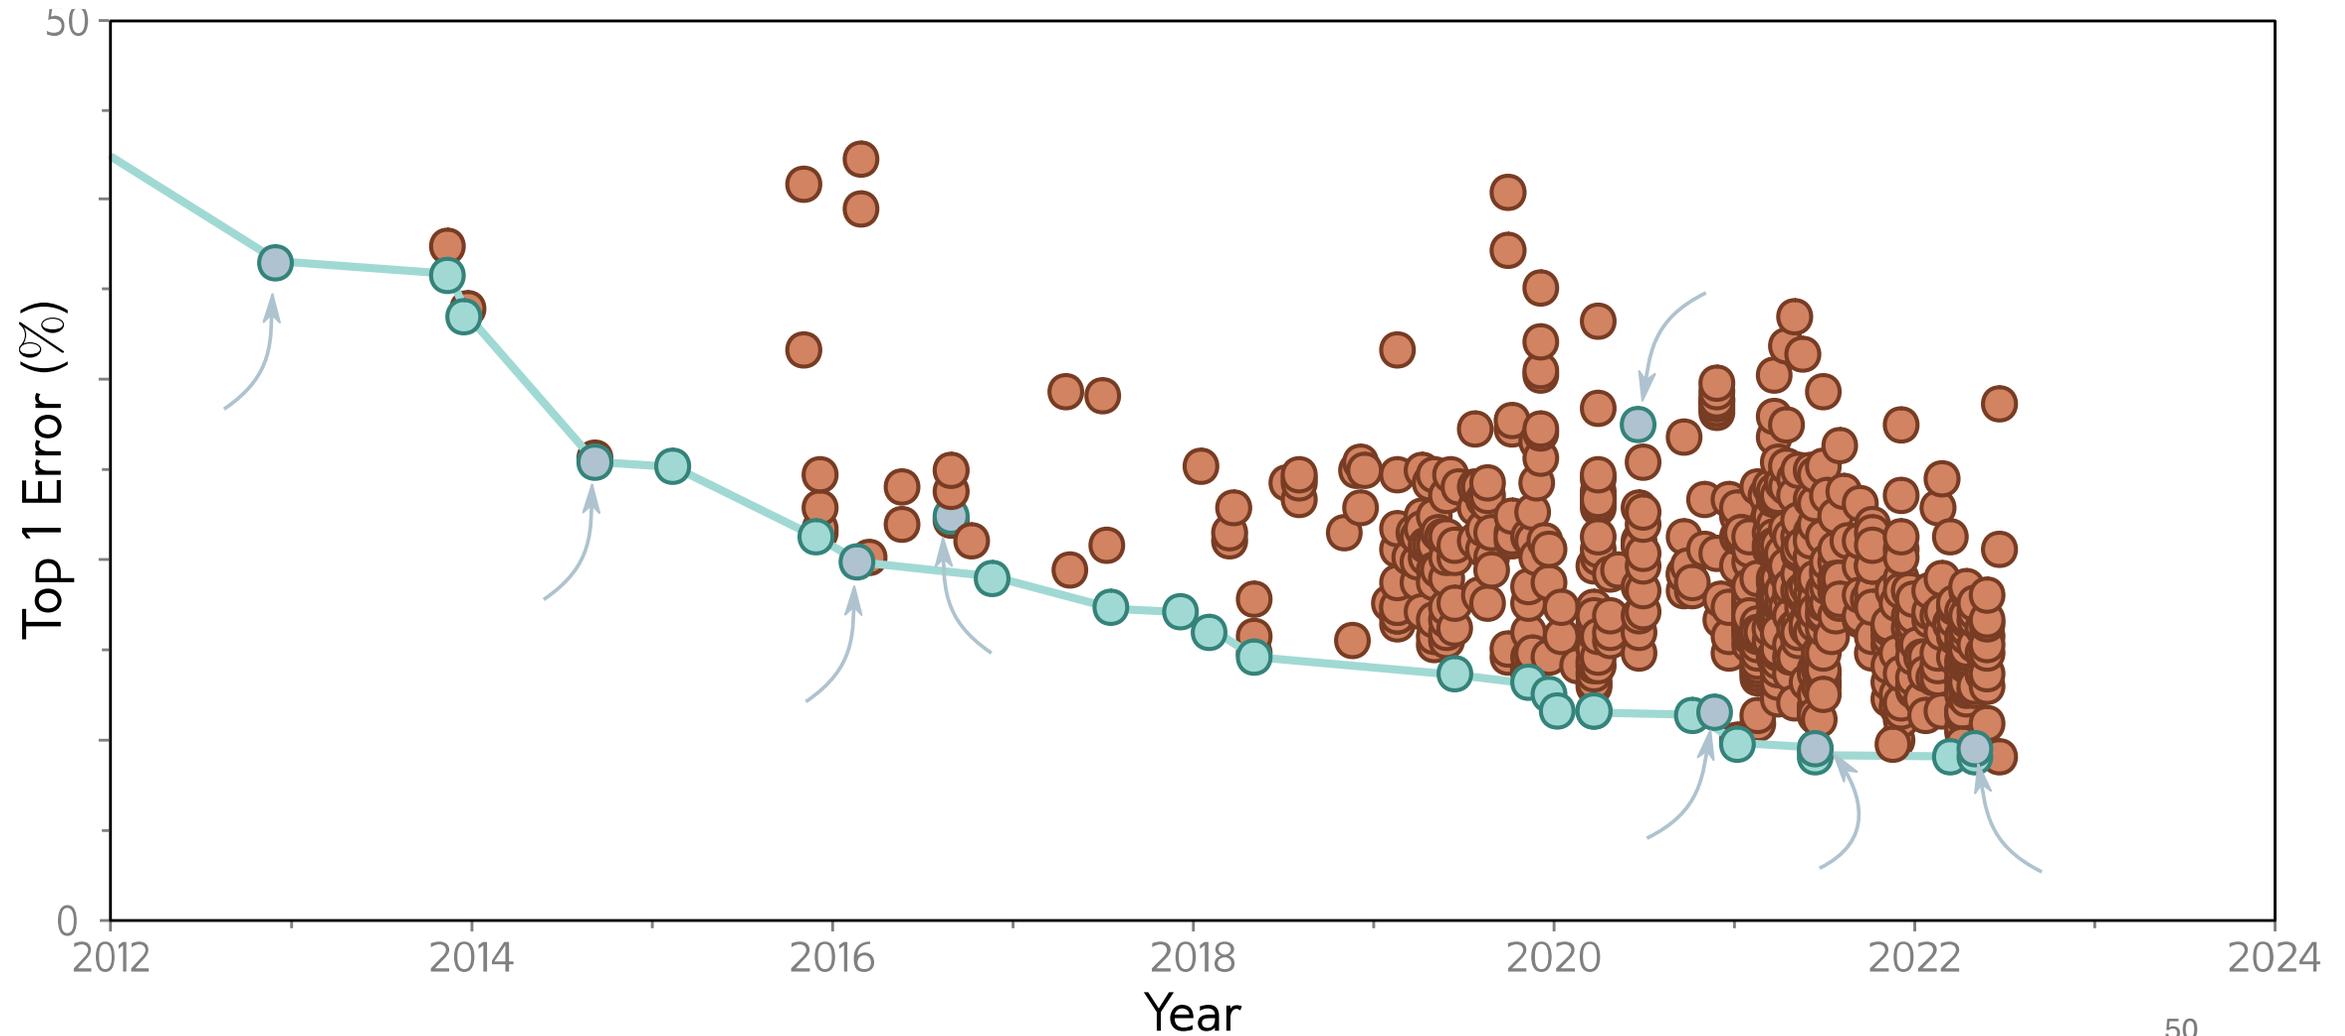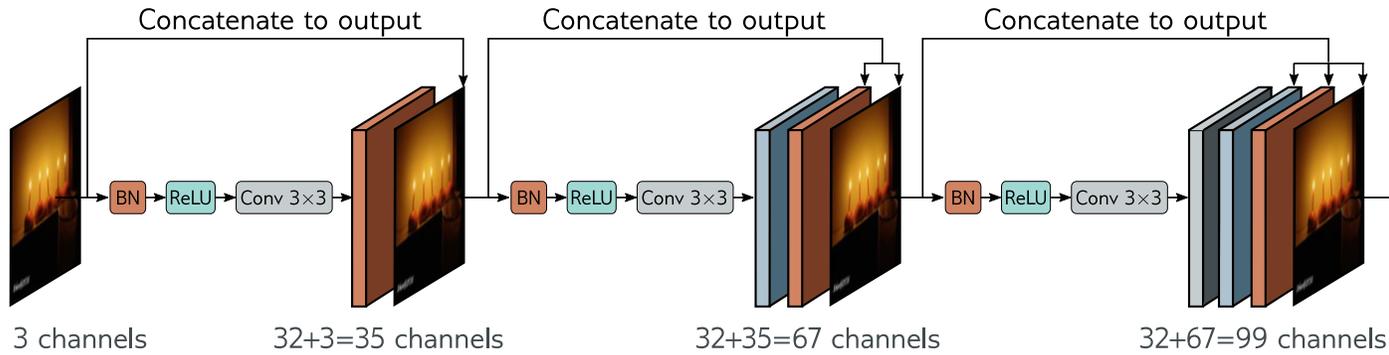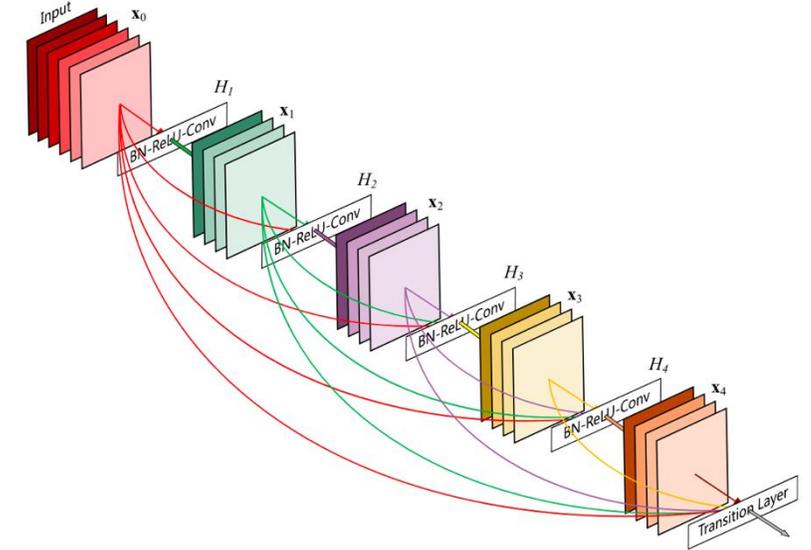
Figure from UDL

Figure 1: A 5-layer dense block with a growth rate of $k = 4$. Each layer takes all preceding feature-maps as input.

Figure from paper

Huang, G., Liu, Z., Van Der Maaten, L., & Weinberger, K. Q. (2017b). Densely connected convolutional networks. IEEE/CVF Computer Vision & Pattern Recognition, 4700–4708.

U-Net (2016)

Ronneberger, O., Fischer, P., & Brox, T. (2015). U-Net: Convolutional networks for biomedical image segmentation. International Conference on Medical Image Computing and ComputerAssisted Intervention, 234–241.
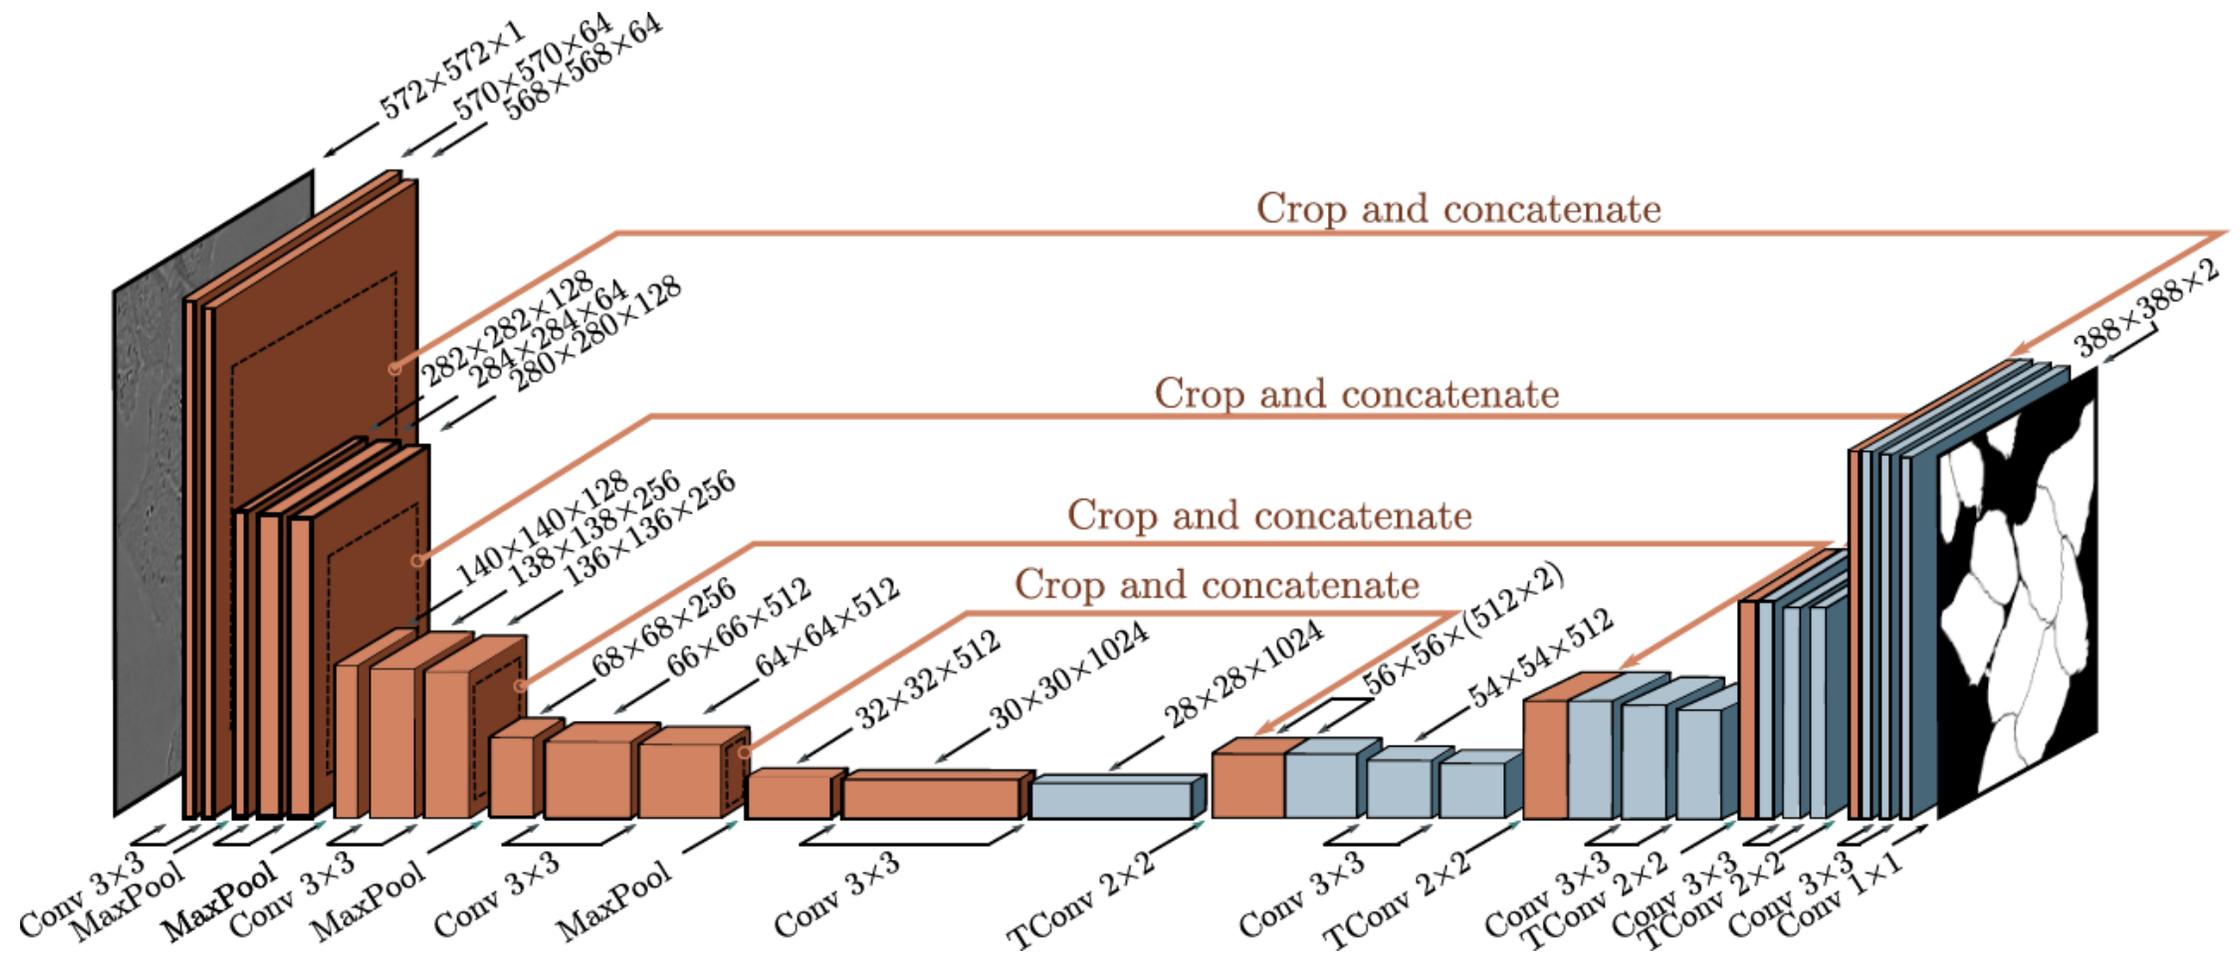
U-Net Results
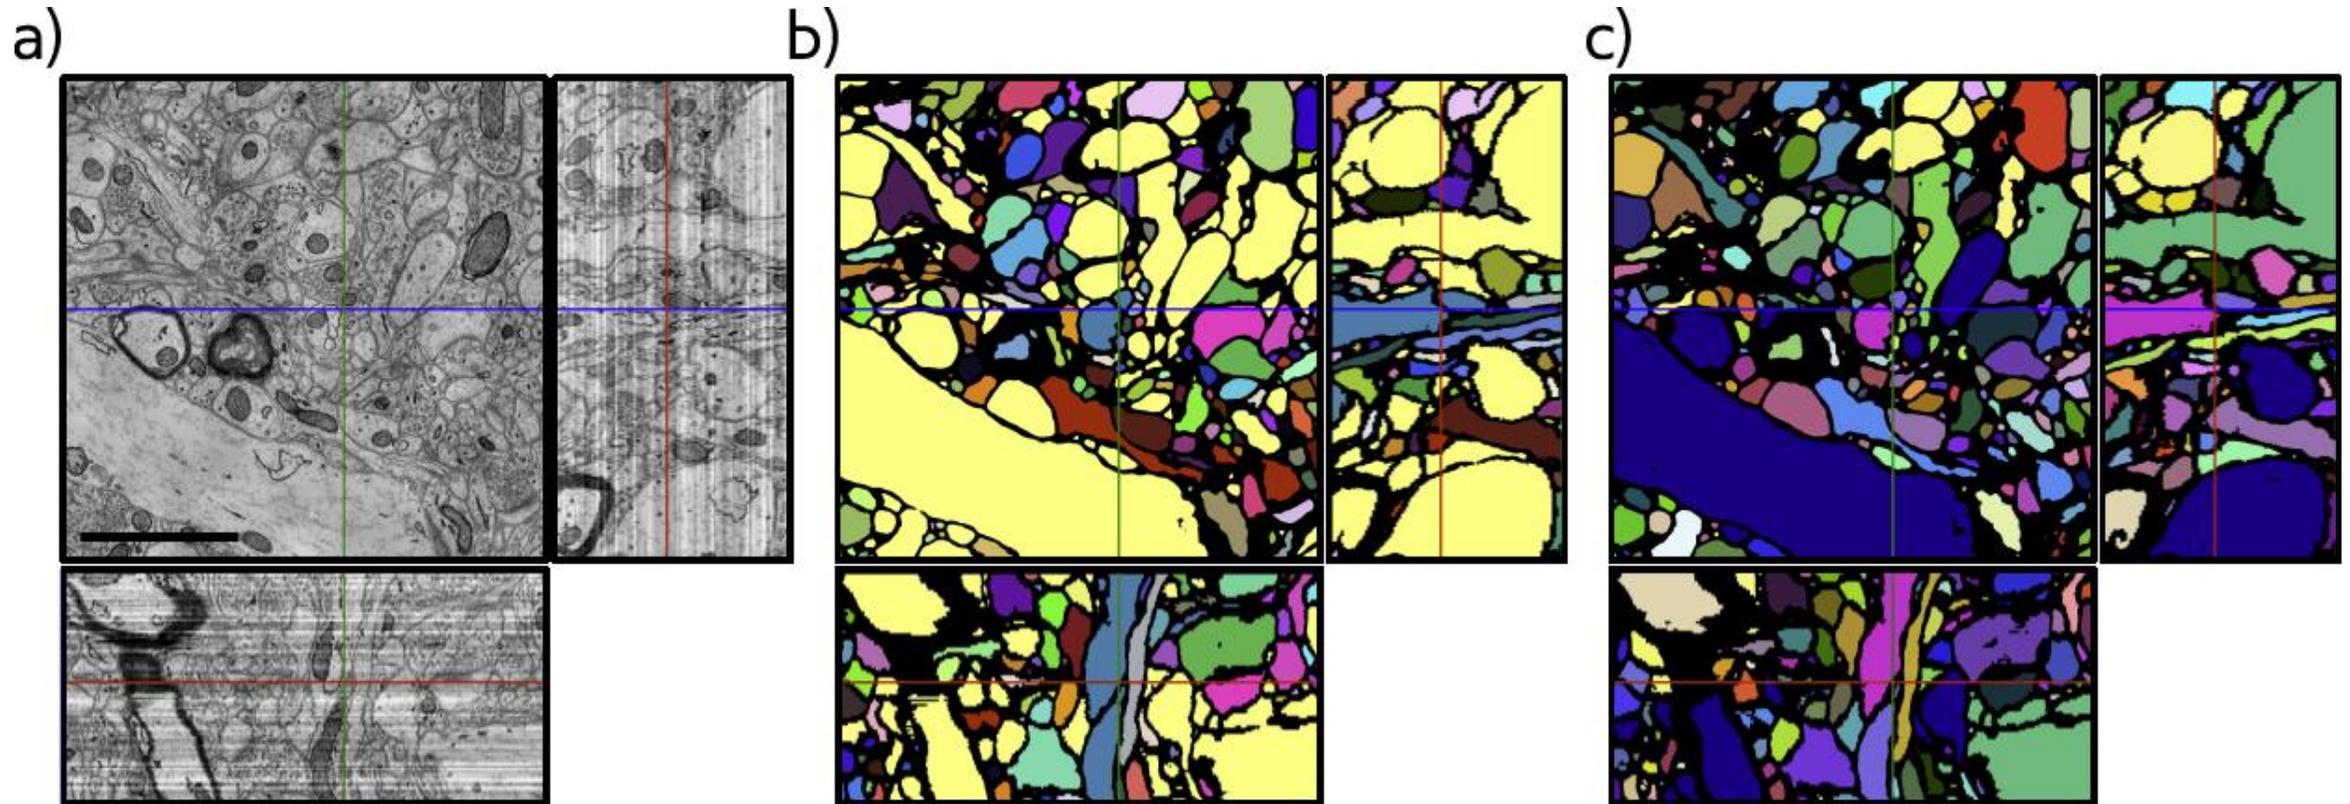
Figure 11.11 Segmentation using U-Net in 3D. a) Three slices through a 3D volume of mouse cortex taken by scanning electron microscope. b) A single U-Net is used to classify voxels as being inside or outside neurites. Connected regions are identified with different colors. c) For a better result, an ensemble of five U-Nets is trained, and a voxel is only classified as belonging to the cell if all five networks agree. Adapted from Falk et al. (2019).

Stacked hourglass networks for Pose Estimation

Newell, A., Yang, K., & Deng, J. (2016). Stacked hourglass networks for human pose estimation. European Conference on Computer Vision, 483–499.

Any Questions?

???

- Challenges of deep networks
- Residual connections and residual blocks
- Exploding gradients in residual networks
- Batch normalization
- Common residual architectures