

# Deep Learning for Data Science

## DS 542

<https://dl4ds.github.io/fa2025/>

Transformer Details

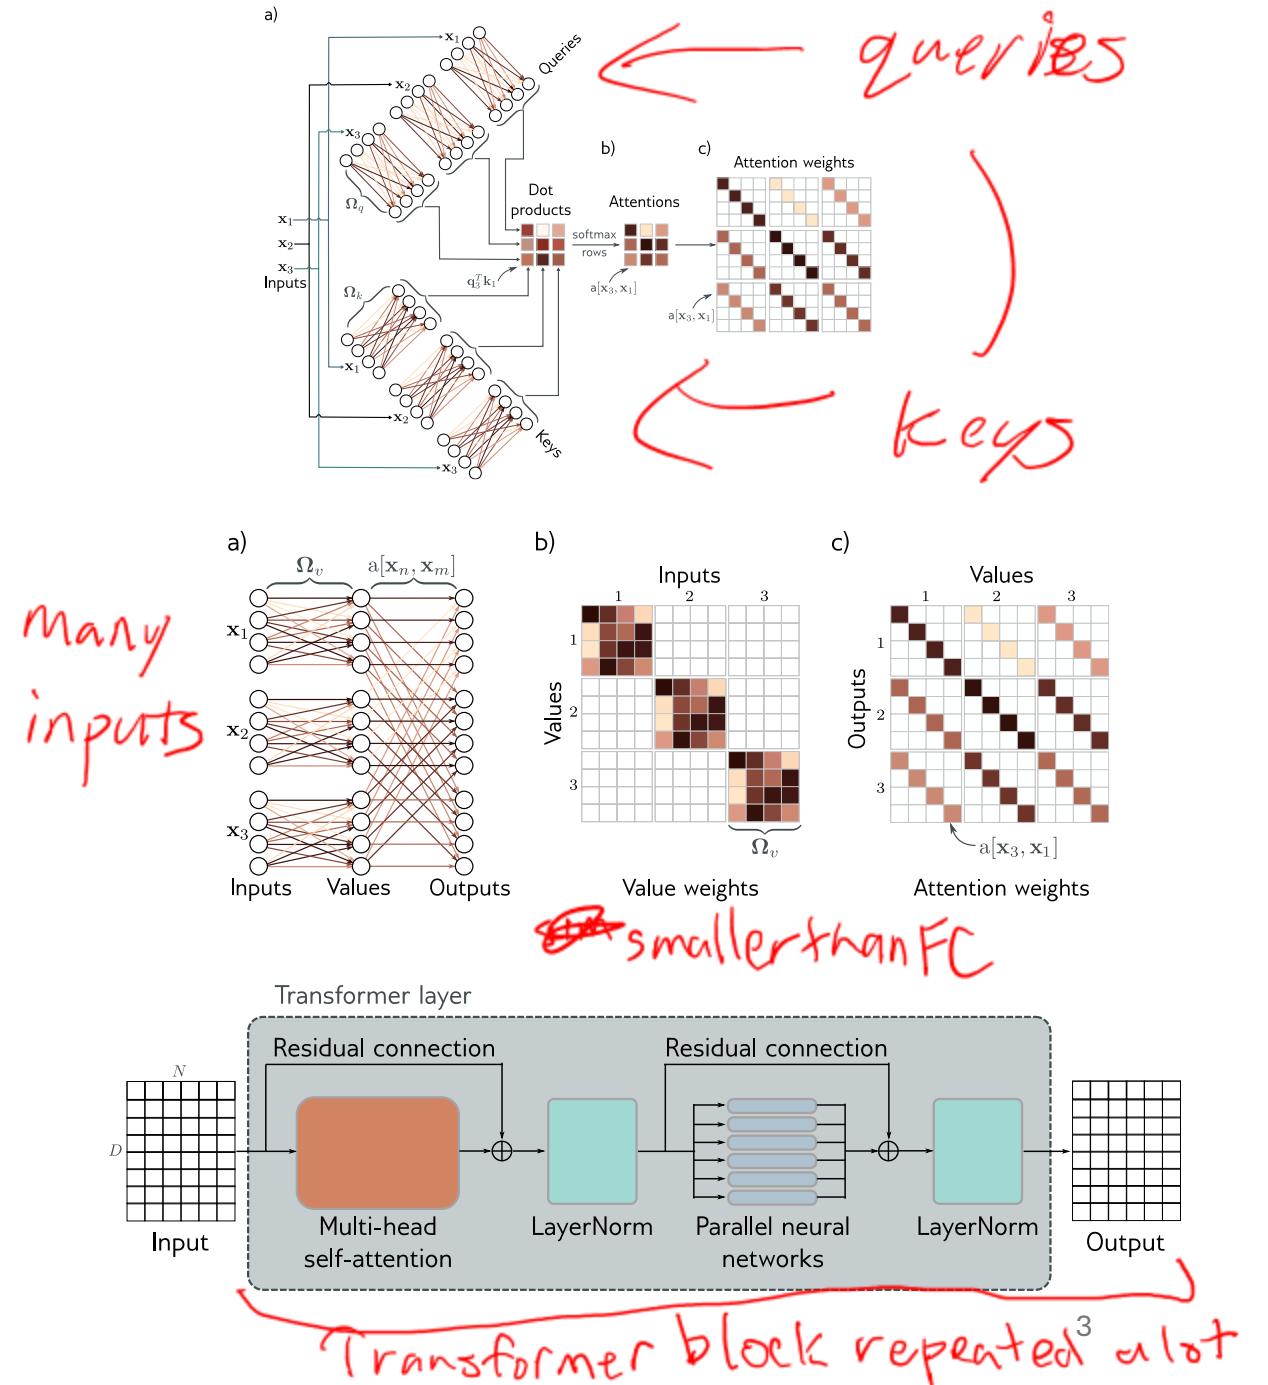


# Plan for Today

- Transformer recap
- What are tokens?
- Tokenization and word embedding
- Next token selection
- Transformers for long sequences

# Recap From Part 1

- Motivation
- Dot-product self-attention
- Applying Self-Attention
- The Transformer Architecture
- Three Types of NLP Transformer Models
  - Encoder
  - Decoder
  - Encoder-Decoder

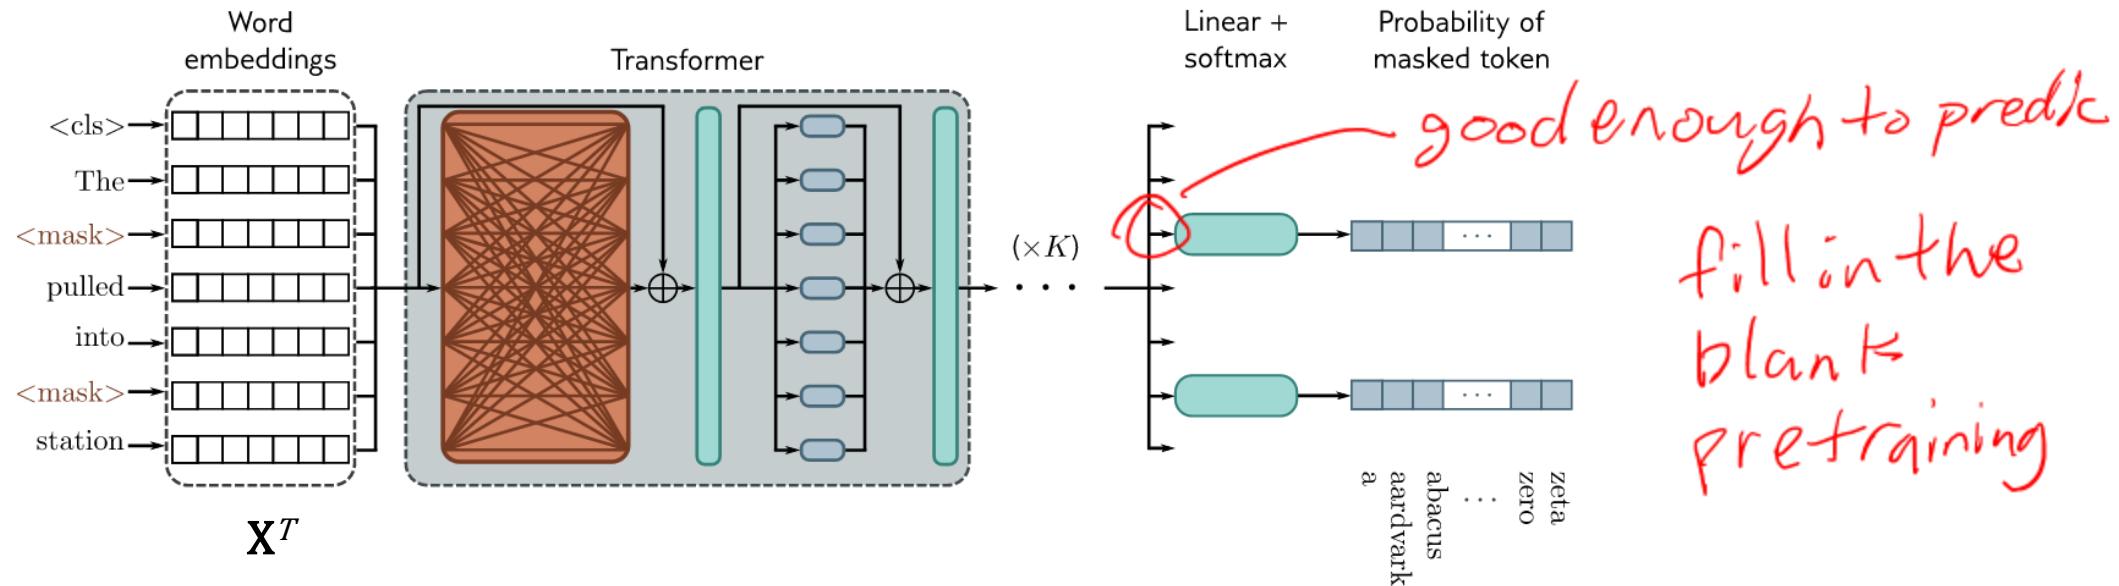


# 3 Types of Transformer Models

1. *Encoder* – transforms text embeddings into representations that support variety of tasks (e.g. sentiment analysis, classification)
  - ❖ Model Example: BERT
2. *Decoder* – predicts the next token to continue the input text (e.g. ChatGPT, AI assistants)
  - ❖ Model Example: GPT4, GPT4
3. *Encoder-Decoder* – used in sequence-to-sequence tasks, where one text string is converted to another (e.g. machine translation)

# Encoder Pre-Training

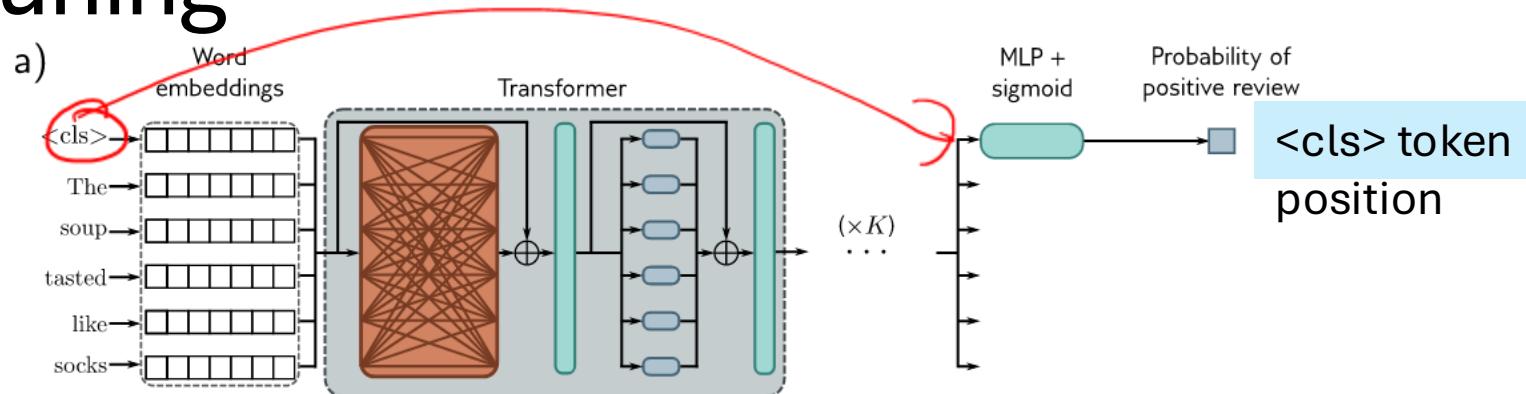
Special <cls> token used for aggregate sequence representation for classification



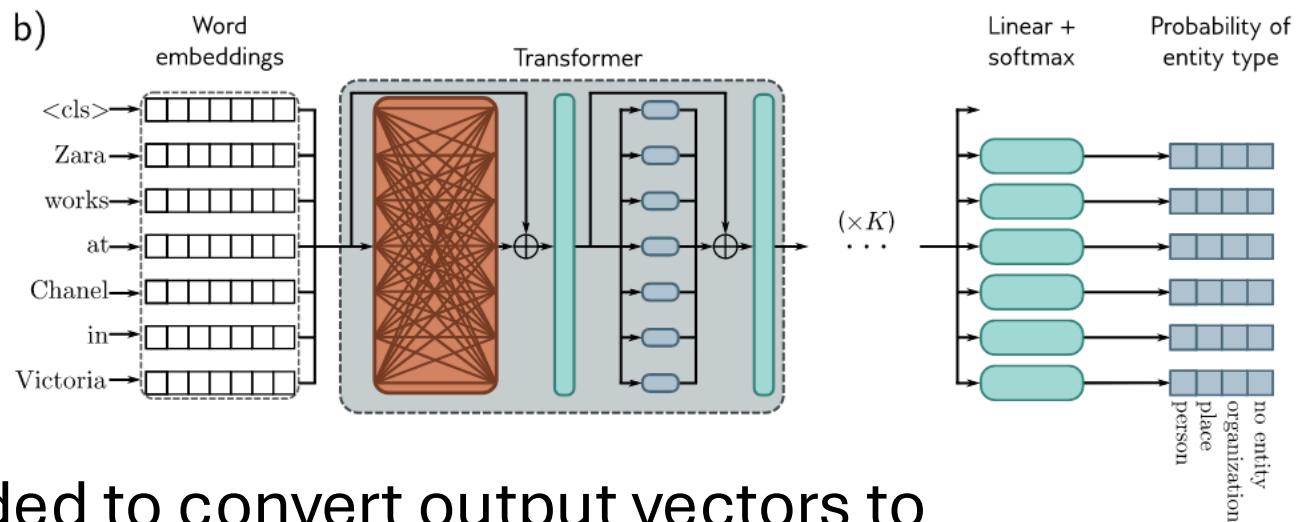
- A small percentage of input embedding replaced with a generic <mask> token
- Predict missing token from output embeddings
- Added linear layer and softmax to generate probabilities over vocabulary
- Trained on BooksCorpus (800M words) and English Wikipedia (2.5B words)

# Encoder Fine-Tuning

Sentiment Analysis

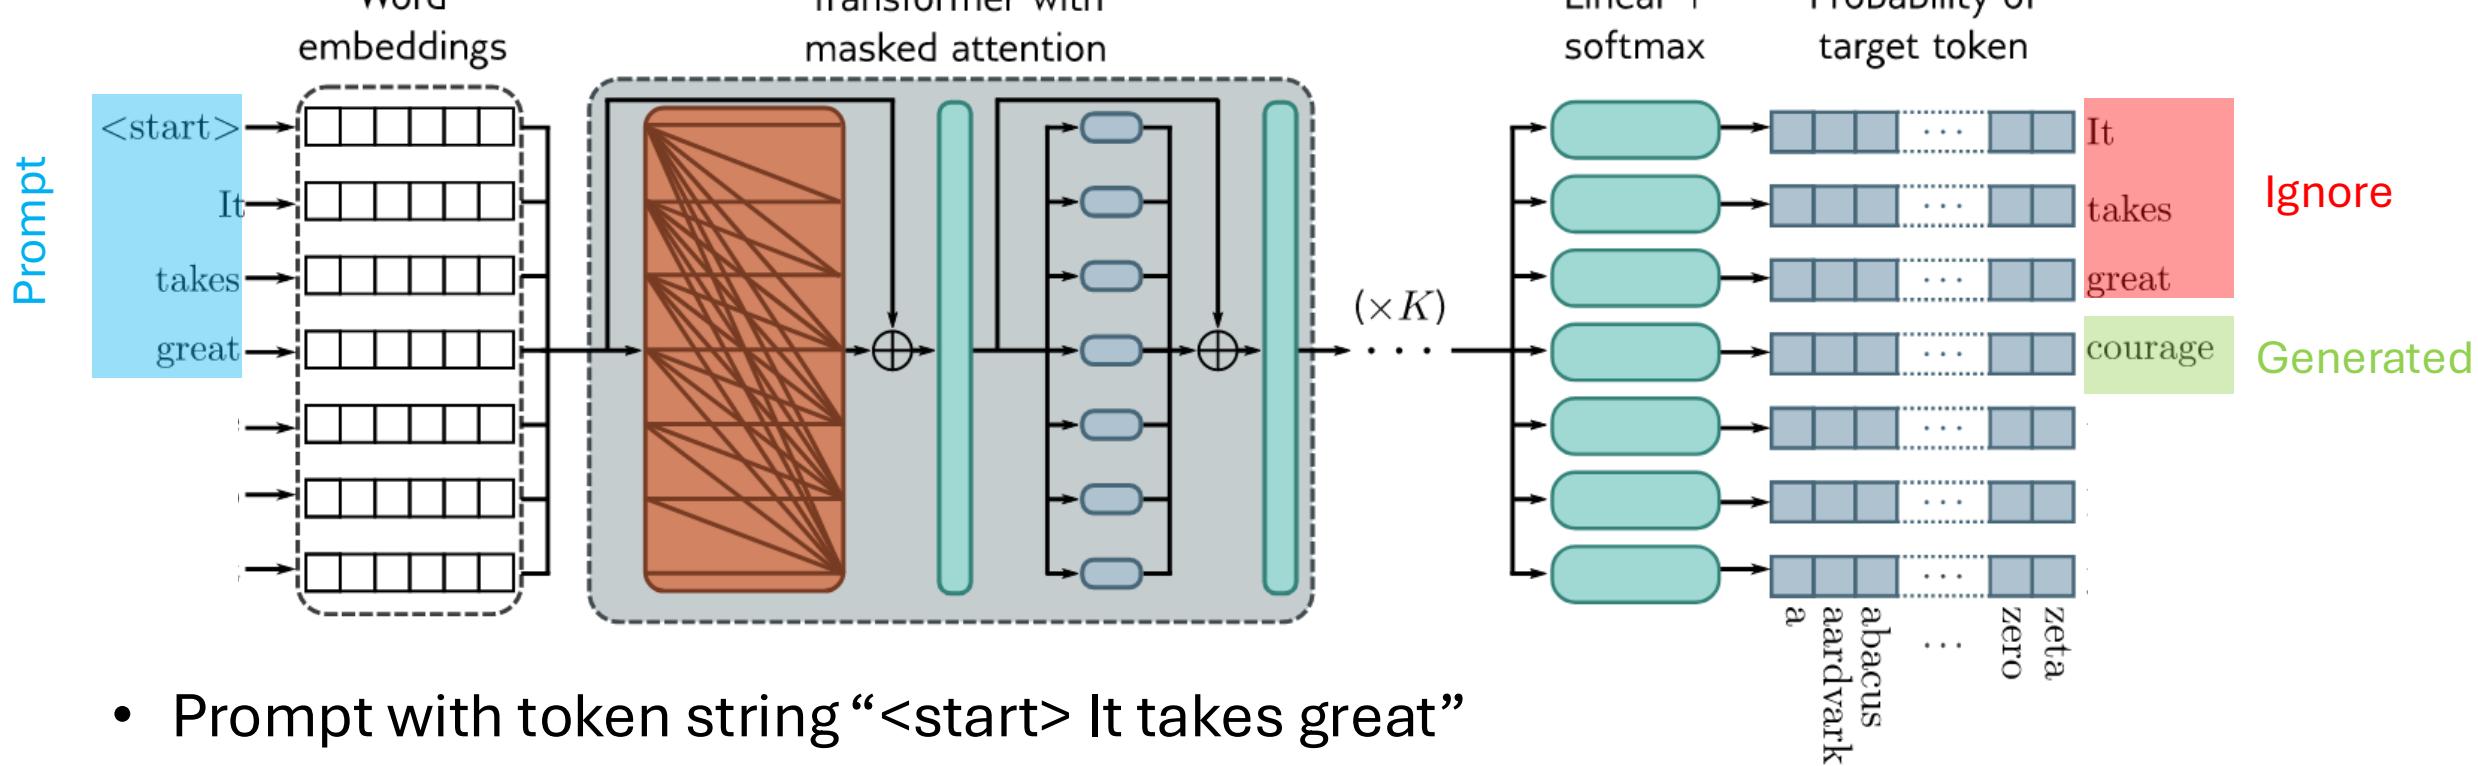


Named Entity Recognition  
(NER)



- Extra layer(s) appended to convert output vectors to desired output format
- 3<sup>rd</sup> Example: Text span prediction -- predict start and end location of answer to a question in passage of Wikipedia, see <https://rajpurkar.github.io/SQuAD-explorer/>

# Decoder: Text Generation (Generative AI)



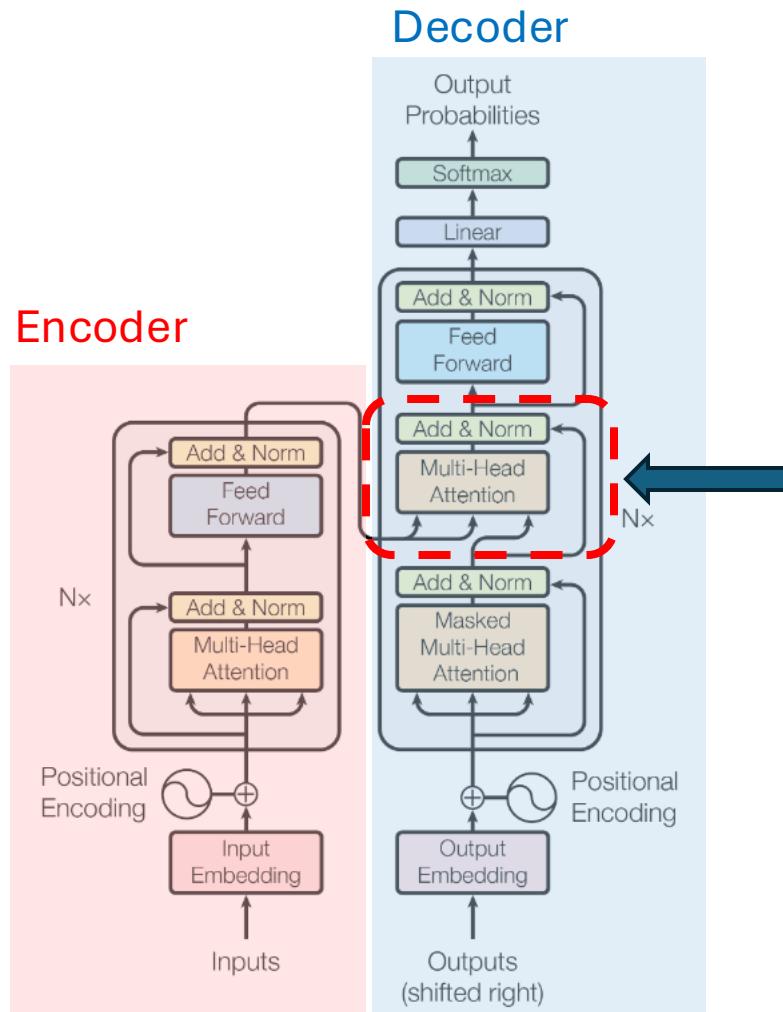
- Prompt with token string “<start> It takes great”
- Generate next token for the sequence by
  - picking most likely token
  - sample from the probability distribution
    - alternative *top-k* sampling to avoid picking from the long tail
  - beam search – select the most likely sentence rather than greedily pick

# Encoder-Decoder Model

- Used for *machine translation*, which is a *sequence-to-sequence* task



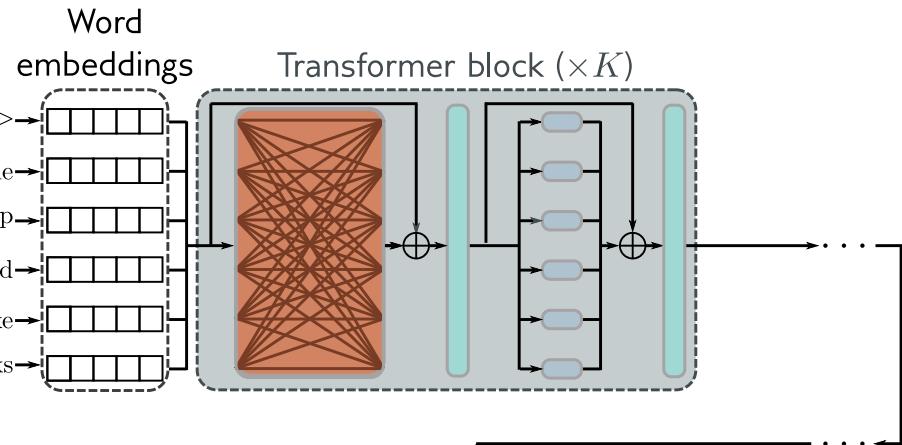
# Encoder Decoder Model



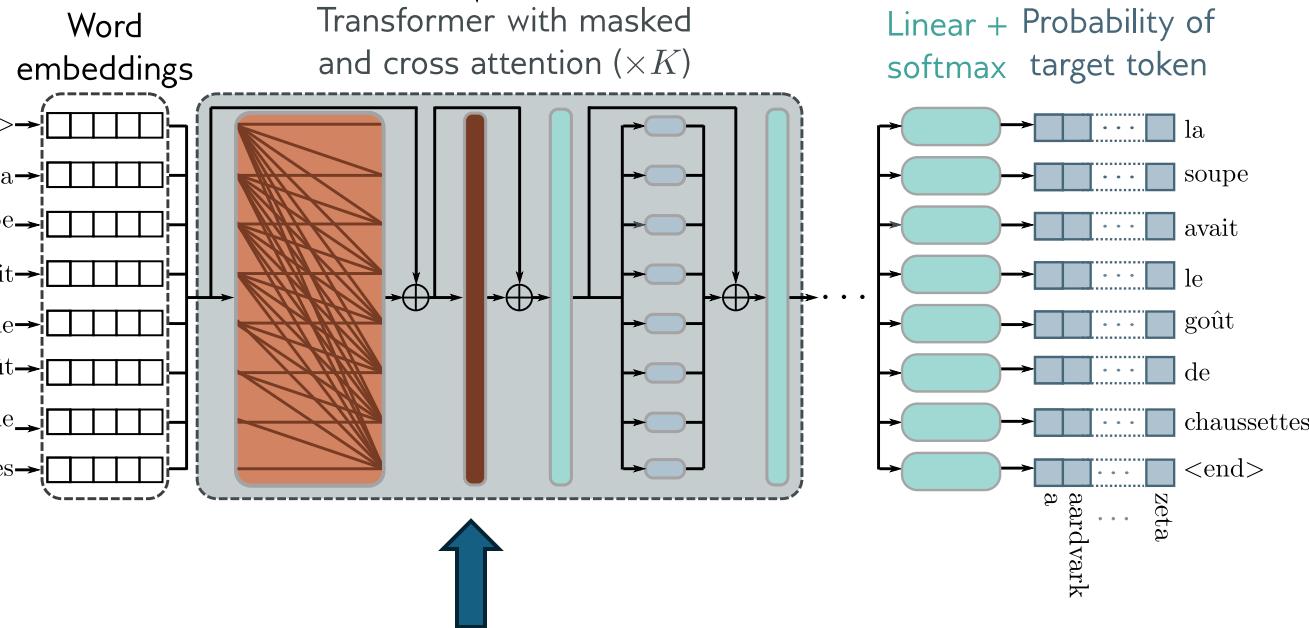
- The transformer layer in the decoder of the encoder-decoder model has an extra stage
- Attends to the input of the encoder with **cross attention** using Keys and Values from the output of the encoder
- Shown here on original diagram from “Attention is all you need” paper

# Encoder Decoder Model

a)

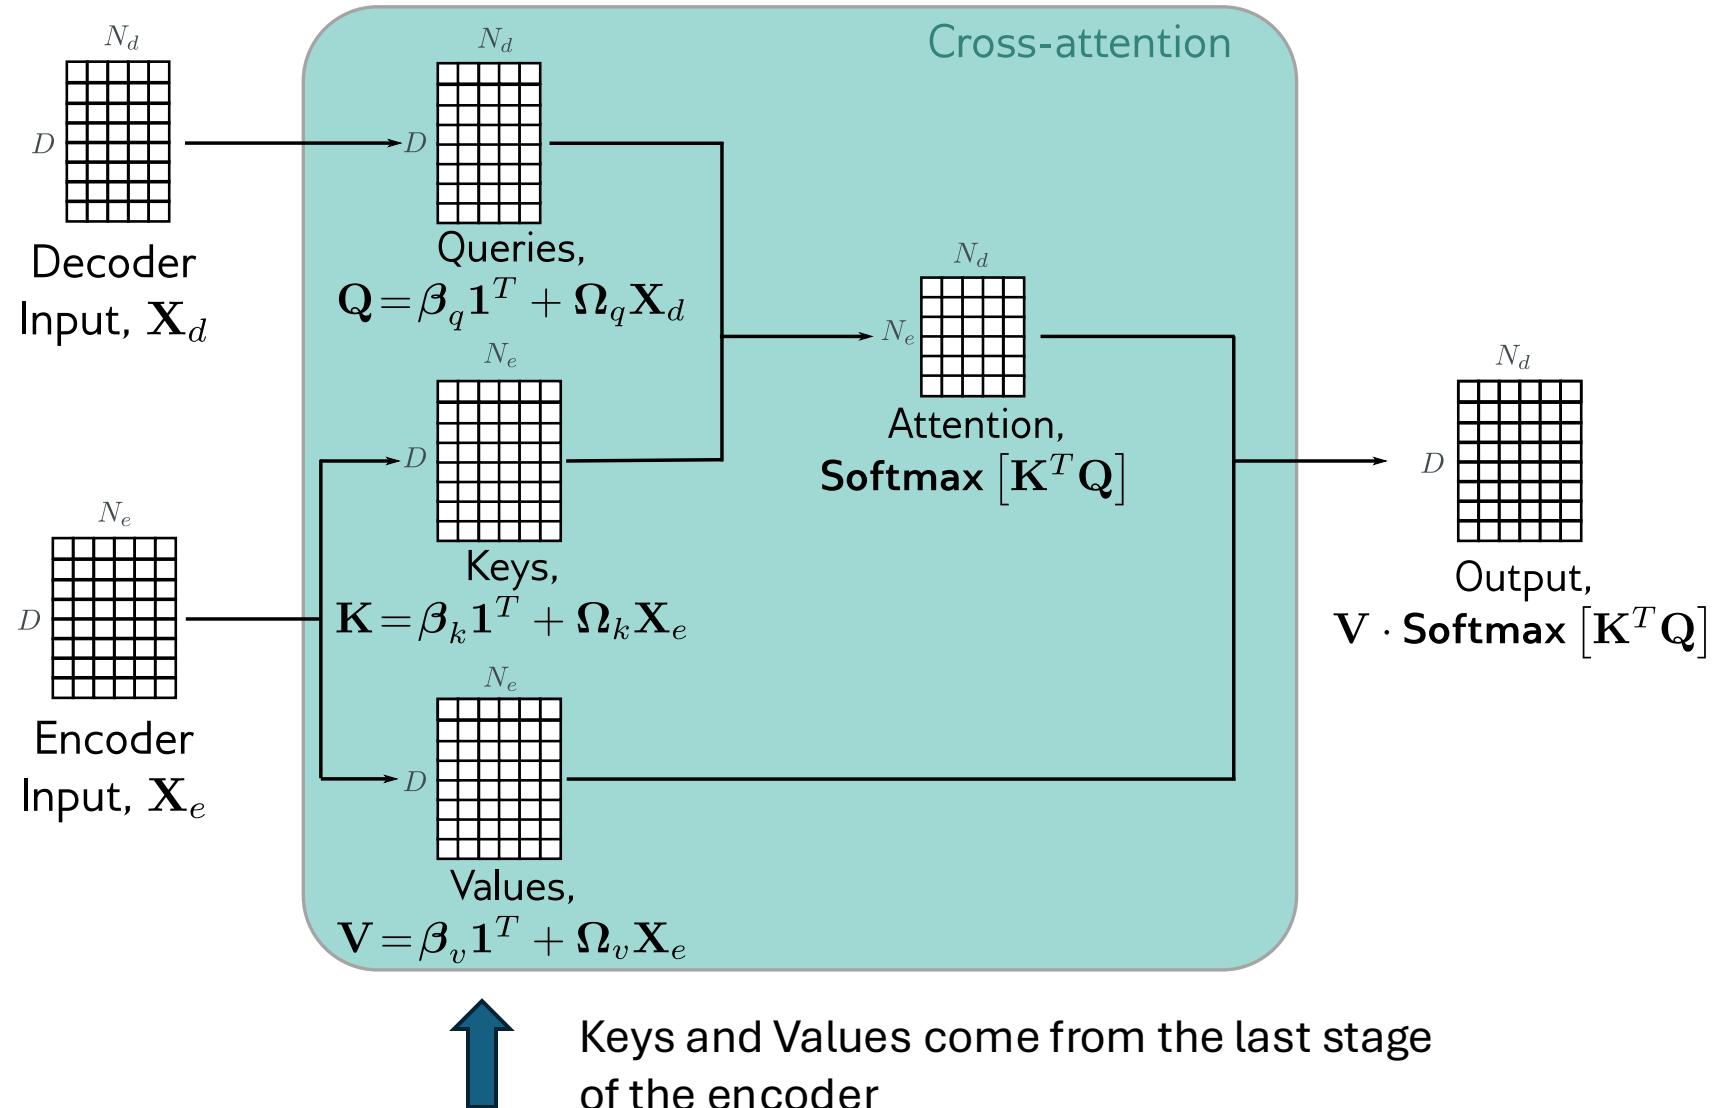


b)



- Same view per UDL book

# Cross-Attention



# Any Questions?

???

## Moving on

- Transformer recap
- **What are tokens?**
- Tokenization and word embedding
- Next token selection
- Transformers for long sequences

# What's a Token?

A small chunk of text that we use to aid language modeling.

- Represents one or more bytes
- Input texts are greedily divided into tokens.
  - Longest prefix matching a token.
- Token set also constructed greedily.
  - Start with 256 possible bytes.
  - Then greedily pick the most common pairs of adjacent tokens.

# Why Tokens?

Instead of...

- Bits - not enough semantics\* and missing intrabyte positioning
  - Bytes - not enough semantics\* for Unicode
  - Characters - too many of them if we try to support all languages
  - Words - even more words than characters
- which ever chosen*

Remember:

- One-hot/Softmax tactic means we will have at least one output per possible output value, and many more parameters in practice.

# Unicode Standard and UTF-8

- **Unicode** – *variable length* character encoding standard. currently defines 149,813 characters and 161 scripts, including emoji, symbols, etc.
- **Unicode Codepoint** – can represent up to  $17 \times 2^{16} = 1,114,112$  entries. e.g. U+0000 – U+10FFFF in hexadecimal *defined so far, not all used*
- **Unicode Transformation Standard (e.g. UTF-8)** – is a *variable length encoding* using one to four bytes
  - First 128 chars same as ASCII

Code point ↔ UTF-8 conversion						
First code point	Last code point	Byte 1	Byte 2	Byte 3	Byte 4	
U+0000	U+007F	0xxxxxxx				
U+0080	U+07FF	110xxxxx	10xxxxxx			
U+0800	U+FFFF	1110xxxx	10xxxxxx	10xxxxxx		
U+010000	[b]U+10FFFF	11110xxx	10xxxxxx	10xxxxxx	10xxxxxx	

Covers ASCII

Covers remainder of almost all Latin-script alphabets

Basic Multilingual Plane including Chinese, Japanese and Korean characters

Emoji, historic scripts, math symbols

<https://en.wikipedia.org/wiki/Unicode>

<https://en.wikipedia.org/wiki/UTF-8>

variable length

# Example Tokens

```
▶ import tiktoken
[6] enc = tiktoken.encoding_for_model("gpt-4o")
▶ for i in range(1024):
    d = enc.decode([i])
    if len(d) >= 4:
        print(i, enc.decode([i]))
```

257	530	699
269	553	are
290	561	import
309	562	able
	583	ight
	584	ublic
	387	ation
	395	for
	406	con
	408	new
	440	pro
	447	port
	452	com
	475	ction
	481	you
	483	with
	484	that
	495	this
	506	
	508	ment
	518	----
	529	turn
	673	was
	677	int
	679	have
	686	par
	694	res
		.get
		-----
		ance
		ould
		ient
		form
		get
		all
		ject
		des
		alue
		will
		();
		class
		public
		ions
		}

# Tokenizer



Tokenizer chooses input “units”, e.g. words, sub-words, characters via *tokenizer training*

*which tokens do we want?*

In tokenizer training, commonly occurring substrings are greedily merged based on their frequency, starting with character pairs

# Tokenization Issues

“A lot of the issues that may look like issues with the neural network architecture actually trace back to tokenization. Here are just a few examples” – Andrej Karpathy

- Why can't LLM spell words? Tokenization.
- Why can't LLM do super simple string processing tasks like reversing a string? Tokenization.
- Why is LLM worse at non-English languages (e.g. Japanese)? Tokenization.
- Why is LLM bad at simple arithmetic? Tokenization.
- Why did GPT-2 have more than necessary trouble coding in Python? Tokenization.
- Why did my LLM abruptly halt when it sees the string "<|endoftext|>"? Tokenization.
- What is this weird warning I get about a "trailing whitespace"? Tokenization.
- Why did the LLM break if I ask it about "SolidGoldMagikarp"? Tokenization.
- Why should I prefer to use YAML over JSON with LLMs? Tokenization.
- Why is LLM not actually end-to-end language modeling? Tokenization.
- What is the real root of suffering? Tokenization.

# Any Questions?

???

## Moving on

- Transformer recap
- What are tokens?
- Tokenization and word embedding
- Next token selection
- Transformers for long sequences

# Tokenization Matters

From the gpt-4o announcement,

“It matches GPT-4 Turbo performance on text in English and code, with significant improvement on text in non-English languages, while also being much faster and 50% cheaper in the API.”

Gains were from increasing the number of tokens in the updated tokenizer.

<https://openai.com/index/hello-gpt-4o/>

Russian 1.7x fewer tokens (from 39 to 23)

Привет, меня зовут GPT-4o. Я — новая языковая модель, приятно познакомиться!

Korean 1.7x fewer tokens (from 45 to 27)

안녕하세요, 제 이름은 GPT-4o입니다. 저는 새로운 유형의 언어 모델입니다, 만나서 반갑습니다!

Vietnamese 1.5x fewer tokens (from 46 to 30)

Xin chào, tên tôi là GPT-4o. Tôi là một loại mô hình ngôn ngữ mới, rất vui được gặp bạn!

Chinese 1.4x fewer tokens (from 34 to 24)

你好，我的名字是GPT-4o。我是一种新型的语言模型，很高兴见到你！

Japanese 1.4x fewer tokens (from 37 to 26)

こんにちは、私の名前はGPT-4oです。私は新しいタイプの言語モデルです。初めまして！

# Tokenizer

Two common tokenizers:

- • Byte Pair Encoding (BPE) – Used by OpenAI GPT2, GPT4, etc.
  - The BPE algorithm is "byte-level" because it runs on UTF-8 encoded strings.
  - This algorithm was popularized for LLMs by the [GPT-2 paper](#) and the associated [GPT-2 code release](#) from OpenAI. [Sennrich et al. 2015](#) is cited as the original reference for the use of BPE in NLP applications. Today, all modern LLMs (e.g. GPT, Llama, Mistral) use this algorithm to train their tokenizers.\*
- sentencepiece
  - (e.g. Llama, Mistral) use [sentencepiece](#) instead. Primary difference being that sentencepiece runs BPE directly on Unicode code points instead of on UTF-8 encoded bytes.

\* <https://github.com/karpathy/minbpe/tree/master>

# BPE Pseudocode

almost always bytes, not full characters



Initialize vocabulary with individual characters in the text and their frequencies

While desired vocabulary size not reached:

← this is width of token prediction output

Identify the most frequent pair of adjacent tokens/characters in the vocabulary

Merge this pair to form a new token

Update the vocabulary with this new token

Recalculate frequencies of all tokens including the new token

Return the final vocabulary

# Enforce a Token Split Pattern

```
GPT2_SPLIT_PATTERN = r"""\b(?:[sdmt]|ll|ve|re)| \p{L}+|\p{N}+|\p{N}\p{L}\p{N}]+|\s+(?!S)|\s+\b"""
```

```
GPT4_SPLIT_PATTERN = r"""\b(?:i:[sdmt]|ll|ve|re)|[\r\n]\p{L}\p{N}]?+\p{L}+|\p{N}{1,3}|  
?\p{L}\p{N}]+[\r\n]*|\s*[r\n]|\s+(?!S)|\s+\b""""
```

- Do not allow tokens to merge across certain characters or patterns
- Common contraction endings: ‘ll, ‘ve, ‘re
- Match words with a leading space
- Match numeric sequences
- carriage returns, new lines

# GPT4 Tokenizer

## Tiktokenizer

```
a sailor went to sea sea sea  
to see what he could see see see  
but all that he could see see see  
was the bottom of the deep blue sea sea sea
```

Token count  
36

```
a·sailor·went·to·sea·sea·sea\n  
to·see·what·he·could·see·see·see\n  
but·all·that·he·could·see·see·see\n  
was·the·bottom·of·the·deep·blue·sea·sea·sea
```

```
[64, 93637, 4024, 311, 9581, 9581, 9581, 198, 99  
8, 1518, 1148, 568, 1436, 1518, 1518, 1518, 198,  
8248, 682, 430, 568, 1436, 1518, 1518, 1518, 198,  
16514, 279, 5740, 315, 279, 5655, 6437, 9581, 958  
1, 9581]
```

Show whitespace

cl100k\_base is the GPT4  
tokenizer

cl100k\_base



# GPT2 Tokenizer

## Tiktokenizer

```
class Tokenizer:  
    """Base class for Tokenizers"""\n  
  
    def __init__(self):  
        # default: vocab size of 256 (all bytes), no merges,  
        no patterns  
        self.merges = {} # (int, int) -> int  
        self.pattern = "" # str  
        self.special_tokens = {} # str -> int, e.g.  
        {'<|endoftext|>': 100257}  
        self.vocab = self._build_vocab() # int -> bytes
```

gpt2

Token count

146

```
class ·Tokenizer:\n...."""Base ·class ·for ·Tokenizers"""\n\n\n....def ·__init__ (self):\n....# ·default: ·vocab ·size ·of ·256 ·(all ·bytes), ·no ·m  
erges, ·no ·patterns\n....self ·merges = ·{} ·# ·(int, ·int) ·-> ·int\n....self ·pattern = ·"" ·# ·str\n....self ·special ·tokens = ·{} ·# ·str ·-> ·int, ·e.g. ·\n{<|endoftext|> : ·100257}\n....self ·vocab = ·self ·_build ·vocab () ·# ·int ·-> ·byte  
s
```

You can see some issues with the GPT2 tokenizer with respect to python code

```
[4871, 29130, 7509, 25, 198, 220, 220, 220, 37227, 148  
81, 1398, 329, 29130, 11341, 37811, 628, 220, 220, 22  
0, 825, 11593, 15003, 834, 7, 944, 2599, 198, 220, 22  
0, 220, 220, 220, 220, 1303, 4277, 25, 12776, 39  
7, 2546, 286, 17759, 357, 439, 9881, 828, 645, 4017, 3  
212, 11, 645, 7572, 198, 220, 220, 220, 220, 220,  
220, 2116, 13, 647, 3212, 796, 23884, 1303, 357, 600,  
11, 493, 8, 4613, 493, 198, 220, 220, 220, 220, 220, 2  
20, 220, 2116, 13, 33279, 796, 13538, 1303, 965, 198,  
220, 220, 220, 220, 220, 220, 2116, 13, 20887, 6  
2, 83, 482, 641, 796, 23884, 1303, 965, 4613, 493, 11,  
304, 13, 70, 13, 1391, 6, 50256, 10354, 1802, 28676, 9  
2, 198, 220, 220, 220, 220, 220, 2116, 13, 1  
8893, 397, 796, 2116, 13557, 11249, 62, 18893, 397, 34  
19, 1303, 493, 4613, 9881]
```

<https://tiktokenizer.vercel.app/>

Show whitespace

# GPT4 Tokenizer

## Tiktokenizer

```
class Tokenizer:  
    """Base class for Tokenizers"""\n\n    def __init__(self):  
        # default: vocab size of 256 (all bytes), no merges,  
        no patterns  
        self.merges = {} # (int, int) -> int  
        self.pattern = "" # str  
        self.special_tokens = {} # str -> int, e.g.  
        {'<|endoftext|>': 100257}  
        self.vocab = self._build_vocab() # int -> bytes
```

cl100k\_base

Token count  
96

```
class Tokenizer:\n    """Base class for Tokenizers"""\n\n    def __init__(self):\n        # default: vocab size of 256 (all bytes), no merges,  
        no patterns\n        self.merges = {} # (int, int) -> int\n        self.pattern = "" # str\n        self.special_tokens = {} # str -> int, e.g.\n        {'<|endoftext|>': 100257}\n        self.vocab = self._build_vocab() # int -> bytes
```

```
[1058, 9857, 3213, 512, 262, 4304, 4066, 538, 369, 985  
7, 12509, 15425, 262, 711, 1328, 2381, 3889, 726, 997,  
286, 674, 1670, 25, 24757, 1404, 315, 220, 4146, 320,  
543, 5943, 705, 912, 82053, 11, 912, 12912, 198, 286,  
659, 749, 2431, 288, 284, 4792, 674, 320, 396, 11, 52  
8, 8, 1492, 528, 198, 286, 659, 40209, 284, 1621, 674,  
610, 198, 286, 659, 64308, 29938, 284, 4792, 674, 610,  
1492, 528, 11, 384, 1326, 13, 5473, 100257, 1232, 220,  
1041, 15574, 534, 286, 659, 78557, 284, 659, 1462, 595  
7, 53923, 368, 674, 528, 1492, 5943]
```

Show whitespace

<https://tiktokenizer.vercel.app/>

a)

a\_sailor\_went\_to\_sea\_sea\_sea\_  
to\_see\_what\_he\_could\_see\_see\_see\_  
but\_all\_that\_he\_could\_see\_see\_see\_  
was\_the\_bottom\_of\_the\_deep\_blue\_sea\_sea\_sea\_

_	e	s	a	t	o	h		u	b	d	w	c	f	i	m	n	p	r
33	28	15	12	11	8	6	6	4	3	3	3	2	1	1	1	1	1	1

## Byte Pair Encoding (BPE) Example

Minimal starting vocabulary of subset of lower case latin alphabet and space ` \_ `.

a) a\_sailor\_went\_to\_sea\_sea\_sea\_  
to\_see\_what\_he\_could\_see\_see\_see\_  
but\_all\_that\_he\_could\_see\_see\_see\_  
was\_the\_bottom\_of\_the\_deep\_blue\_sea\_sea\_sea\_

_	e	s	a	t	o	h		u	b	d	w	c	f	i	m	n	p	r
33	28	15	12	11	8	6	6	4	3	3	3	2	1	1	1	1	1	1

b) a\_sailor\_went\_to\_sea\_sea\_sea\_  
to\_see\_what\_he\_could\_see\_see\_see\_  
but\_all\_that\_he\_could\_see\_see\_see\_  
was\_the\_bottom\_of\_the\_deep\_blue\_sea\_sea\_sea\_

_	e	se	a	t	o	h		u	b	d	w	c	s	f	i	m	n	p	r
33	15	13	12	11	8	6	6	4	3	3	3	2	2	1	1	1	1	1	1

## Byte Pair Encoding (BPE)

### Example

Find the most frequent pair of adjacent tokens,  
**'se'**, in this case and form new token.

sea or see

a) a\_sailor\_went\_to\_sea\_sea\_sea\_  
to\_see\_what\_he\_could\_see\_see\_see\_  
but\_all\_that\_he\_could\_see\_see\_see\_  
was\_the\_bottom\_of\_the\_deep\_blue\_sea\_sea\_sea\_

_	e	s	a	t	o	h	l	u	b	d	w	c	f	i	m	n	p	r
33	28	15	12	11	8	6	6	4	3	3	3	2	1	1	1	1	1	1

b) a\_sailor\_went\_to\_sea\_sea\_sea\_  
to\_see\_what\_he\_could\_see\_see\_see\_  
but\_all\_that\_he\_could\_see\_see\_see\_  
was\_the\_bottom\_of\_the\_deep\_blue\_sea\_sea\_sea\_

_	e	se	a	t	o	h	l	u	b	d	w	c	s	f	i	m	n	p	r
33	15	13	12	11	8	6	6	4	3	3	3	2	2	1	1	1	1	1	1

c) a\_sailor\_went\_to\_sea\_sea\_sea\_  
to\_see\_what\_he\_could\_see\_see\_see\_  
but\_all\_that\_he\_could\_see\_see\_see\_  
was\_the\_bottom\_of\_the\_deep\_blue\_sea\_sea\_sea\_

_	se	a	e_	t	o	h	l	u	b	d	e	w	c	s	f	i	m	n	p	r
21	13	12	12	11	8	6	6	4	3	3	3	3	2	2	1	1	1	1	1	1

## Byte Pair Encoding (BPE) Example

Next most frequent pair of tokens is `e\_`

a) a\_sailor\_went\_to\_sea\_sea\_sea\_  
 to\_see\_what\_he\_could\_see\_see\_see\_  
 but\_all\_that\_he\_could\_see\_see\_see\_  
 was\_the\_bottom\_of\_the\_deep\_blue\_sea\_sea\_sea\_

_	e	s	a	t	o	h		u	b	d	w	c	f	i	m	n	p	r
33	28	15	12	11	8	6	6	4	3	3	3	2	1	1	1	1	1	1

b) a\_sailor\_went\_to\_sea\_sea\_sea\_  
 to\_see\_what\_he\_could\_see\_see\_see\_  
 but\_all\_that\_he\_could\_see\_see\_see\_  
 was\_the\_bottom\_of\_the\_deep\_blue\_sea\_sea\_sea\_

_	e	se	a	t	o	h		u	b	d	w	c	s	f	i	m	n	p	r
33	15	13	12	11	8	6	6	4	3	3	3	2	2	1	1	1	1	1	1

c) a\_sailor\_went\_to\_sea\_sea\_sea\_  
 to\_see\_what\_he\_could\_see\_see\_see\_  
 but\_all\_that\_he\_could\_see\_see\_see\_  
 was\_the\_bottom\_of\_the\_deep\_blue\_sea\_sea\_sea\_

_	se	a	e_	t	o	h		u	b	d	e	w	c	s	f	i	m	n	p	r
21	13	12	12	11	8	6	6	4	3	3	3	3	2	2	1	1	1	1	1	1

⋮              ⋮

d)

see_	sea_	e	b		w	a	could_	hat_	he_	o	t	t_	the_	to_	u	a	d	f	m	n	p	s	sailor_	to	
7	6	4	3	3	3	3	2	2	2	2	2	2	2	2	2	1	1	1	1	1	1	1	1	1	1

2x 4 char tokens  
 that are common      worth it?

## Byte Pair Encoding (BPE) Example

# Byte Pair Encoding (BPE) Example

a) a\_sailor\_went\_to\_sea\_sea\_sea\_  
 to\_see\_what\_he\_could\_see\_see\_see\_  
 but\_all\_that\_he\_could\_see\_see\_see\_  
 was\_the\_bottom\_of\_the\_deep\_blue\_sea\_sea\_sea\_

_	e	s	a	t	o	h	l	u	b	d	w	c	f	i	m	n	p	r
33	28	15	12	11	8	6	6	4	3	3	3	2	1	1	1	1	1	1

b) a\_sailor\_went\_to\_sea\_sea\_sea\_  
 to\_see\_what\_he\_could\_see\_see\_see\_  
 but\_all\_that\_he\_could\_see\_see\_see\_  
 was\_the\_bottom\_of\_the\_deep\_blue\_sea\_sea\_sea\_

_	e	se	a	t	o	h	l	u	b	d	w	c	s	f	i	m	n	p	r
33	15	13	12	11	8	6	6	4	3	3	3	2	2	1	1	1	1	1	1

c) a\_sailor\_went\_to\_sea\_sea\_sea\_  
 to\_see\_what\_he\_could\_see\_see\_see\_  
 but\_all\_that\_he\_could\_see\_see\_see\_  
 was\_the\_bottom\_of\_the\_deep\_blue\_sea\_sea\_sea\_

_	se	a	e_	t	o	h	l	u	b	d	e	w	c	s	f	i	m	n	p	r
21	13	12	12	11	8	6	6	4	3	3	3	3	2	2	1	1	1	1	1	1

⋮      ⋮

d) see\_sea\_e\_b\_l\_w\_a\_could\_hat\_he\_o\_t\_t\_the\_to\_u\_a\_d\_f\_m\_n\_p\_s\_sailor\_to

7	6	4	3	3	3	3	2	2	2	2	2	2	2	2	1	1	1	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

⋮      ⋮      ⋮

e) see\_sea\_could\_he\_the\_a\_all\_blue\_bottom\_but\_deep\_of\_sailor\_that\_to\_was\_went\_what

7	6	2	2	2	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

a) a\_sailor\_went\_to\_sea\_sea\_sea\_  
 to\_see\_what\_he\_could\_see\_see\_see\_  
 but\_all\_that\_he\_could\_see\_see\_see\_  
 was\_the\_bottom\_of\_the\_deep\_blue\_sea\_sea\_sea\_

_	e	s	a	t	o	h	l	u	b	d	w	c	f	i	m	n	p	r
33	28	15	12	11	8	6	6	4	3	3	3	2	1	1	1	1	1	1

b) a\_sailor\_went\_to\_sea\_sea\_sea\_  
 to\_see\_what\_he\_could\_see\_see\_see\_  
 but\_all\_that\_he\_could\_see\_see\_see\_  
 was\_the\_bottom\_of\_the\_deep\_blue\_sea\_sea\_sea\_

_	e	se	a	t	o	h	l	u	b	d	w	c	s	f	i	m	n	p	r
33	15	13	12	11	8	6	6	4	3	3	3	2	2	1	1	1	1	1	1

c) a\_sailor\_went\_to\_sea\_sea\_sea\_  
 to\_see\_what\_he\_could\_see\_see\_see\_  
 but\_all\_that\_he\_could\_see\_see\_see\_  
 was\_the\_bottom\_of\_the\_deep\_blue\_sea\_sea\_sea\_

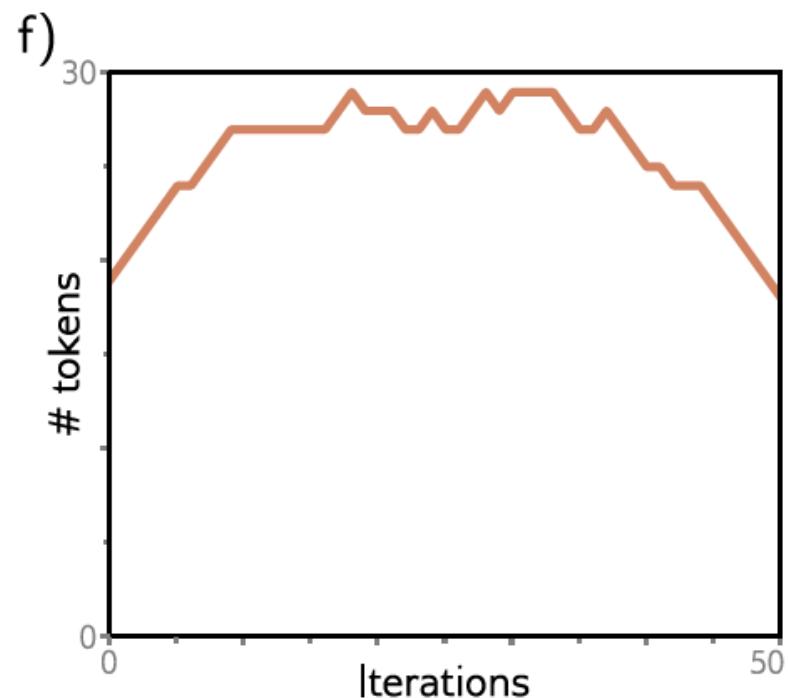
_	se	a	e_	t	o	h	l	u	b	d	e	w	c	s	f	i	m	n	p	r
21	13	12	12	11	8	6	6	4	3	3	3	3	2	2	1	1	1	1	1	1

d)

see_	sea_	e	b	l	w	a	could_	hat_	he_	o	t	t_	the_	to_	u	a	d	f	m	n	p	s	sailor_	to	
7	6	4	3	3	3	3	2	2	2	2	2	2	2	2	2	1	1	1	1	1	1	1	1	1	1

e)

see_	sea_	could_	he_	the_	a_	all_	blue_	bottom_	but_	deep_	of_	sailor_	that_	to_	was_	went_	what_
7	6	2	2	2	1	1	1	1	1	1	1	1	1	1	1	1	1

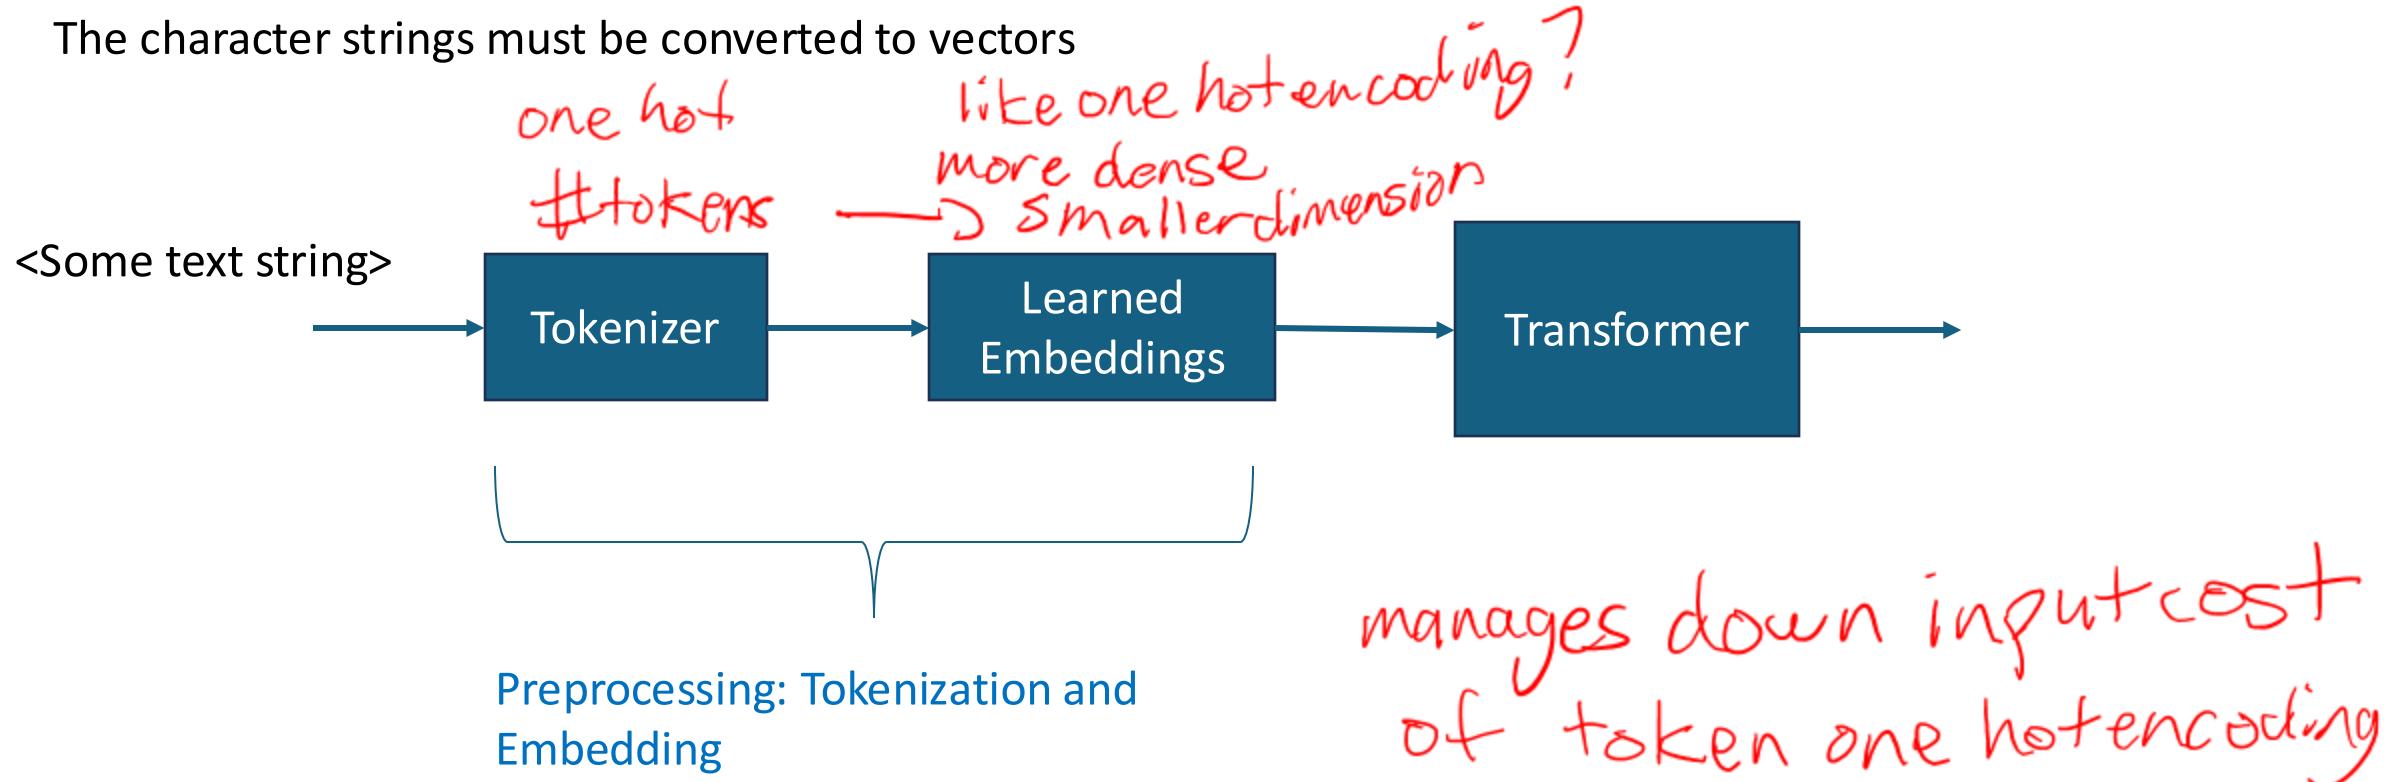


Generally # of tokens increases  
 and then starts decreasing after  
 continuing to merge tokens

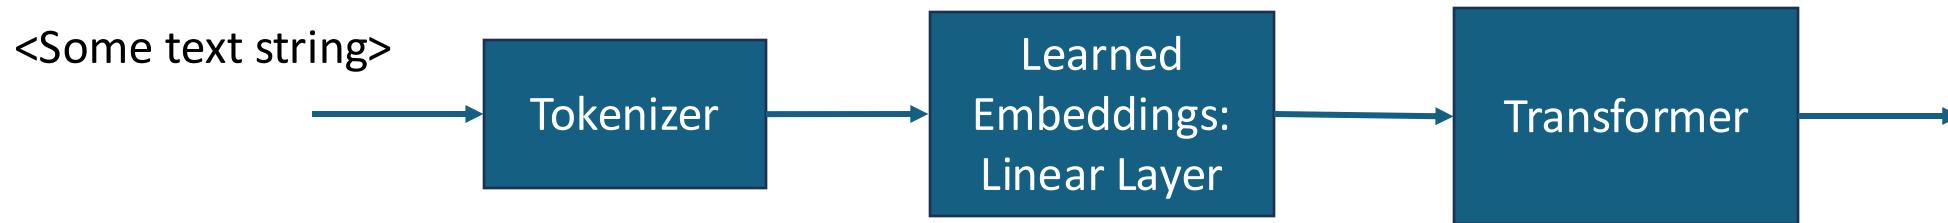
# NLP Preprocessing Pipeline

Transformers don't work on character string directly, but rather on vectors.

The character strings must be converted to vectors



# Learned Embeddings



- After the tokenizer, you have an updated “vocabulary” indexed by token ID
- Next step is to translate the token into an embedding vector
- Translation is done via a linear layer which is typically learned with the rest of the transformer model

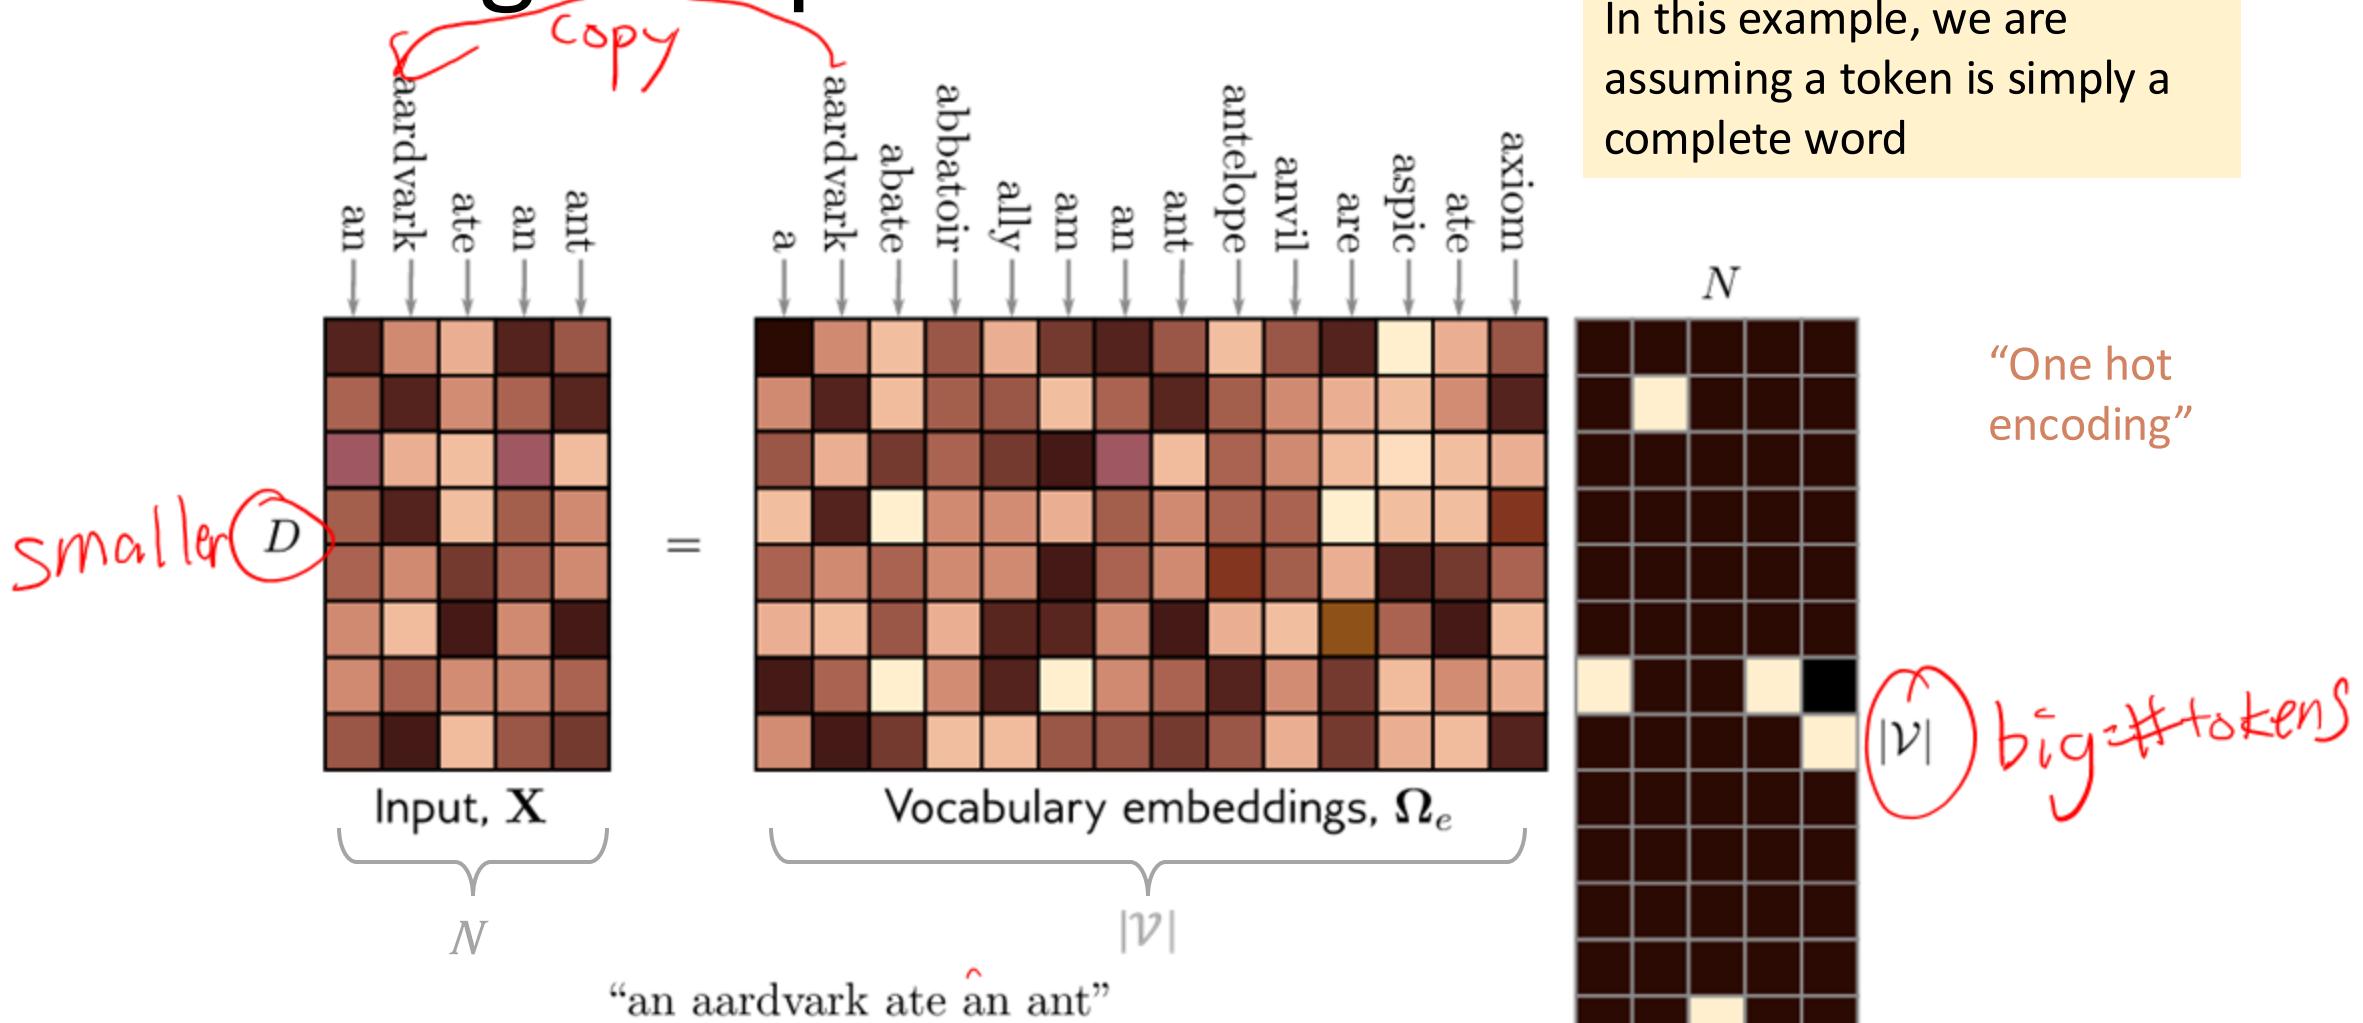
`self.embedding = nn.Embedding(vocab_size, embedding_dim)`

*#tokens*  
*big*      *small*

- Special layer definition, likely to exploit sparsity of input

*b/c onehot tokens*

# Embeddings Output



- Typical embedding size,  $D$ , is 1024
- Typical vocabulary size,  $|V|$ , is 30,000 *Actually, closer to 200K now!*
- So 30M parameters just for this matrix!

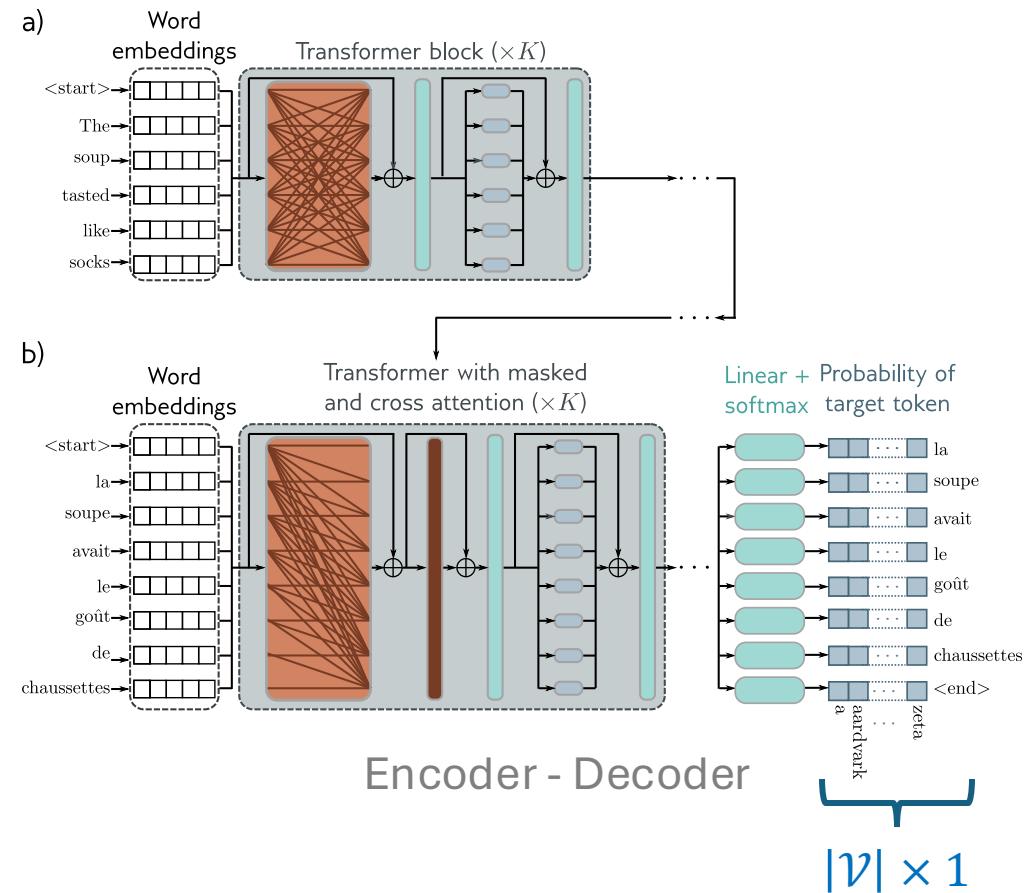
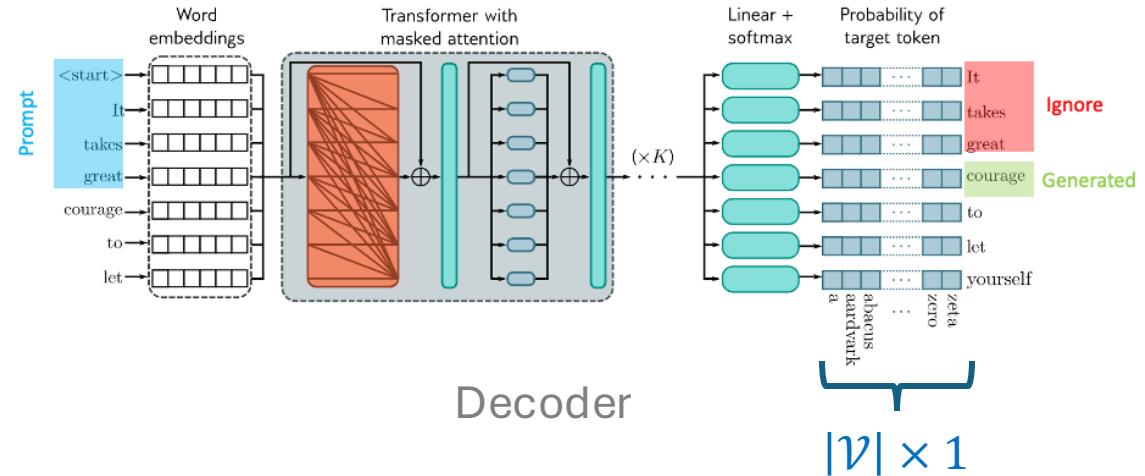
# Any Questions?

???

## Moving on

- Transformer recap
- What are tokens?
- Tokenization and word embedding
- Next token selection
- Transformers for long sequences

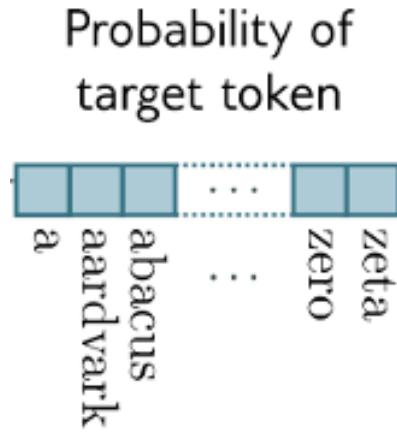
# Next Token Selection



- Recall: output is a  $|\mathcal{V}| \times 1$  vector of probabilities
- How should we pick the next token?
- Trade off between **accuracy** and **diversity**

# Next Token Selection

Recall: output is a  $|\mathcal{V}| \times 1$  vector of probabilities



Selectin methods:

- Greedy selection
- Top-K
- Nucleus
- Temperature
- Beam search

# Next Token Selection – Greedy

Pick most likely token (greedy)

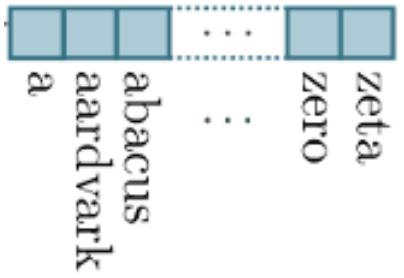
Simple to implement. Just take the `max()`.

$$\hat{y}_t = \operatorname{argmax}_{w \in \mathcal{V}} [Pr(y_t = w | \hat{\mathbf{y}}_{<t}, \mathbf{x}, \phi)]$$

```
# in PyTorch  
outputs = model(inputs)  
value, index = outputs.max(1)
```

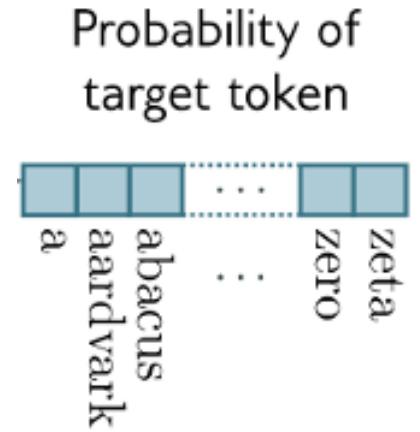
Might pick first token  $y_0$ , but then there is no  $y_1$  where  $Pr(y_1 | y_0)$  is high.

Result is generic and predictable. Same output for a given input *← repeatable*



# Next Token Selection -- Sampling

Sample from the probability distribution



Get a bit more diversity in the output

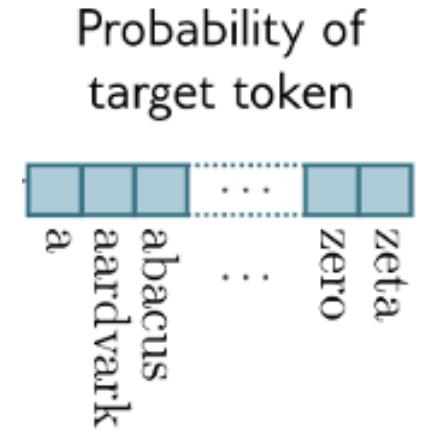
Will occasionally sample from the long tail of the distribution, producing some unlikely word combinations.

But real text does have unlikely words too.



# Next Token Selection – Top $K$ Sampling

1. Generate the probability vector as usual
2. Sort tokens by likelihood
3. Discard all but top  $k$  most probable words
4. Renormalize the probabilities to be valid probability distribution (e.g. sum to 1)
5. Sample from the new distribution



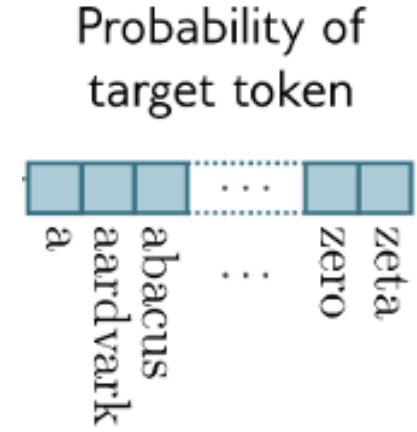
Diversifies word selection.

Depends on the distribution. Could be low variance, reducing diversity.

If only one good choice, forced  $k-1$  tokens would be funky

# Next Token Selection – Nucleus Sampling

Instead of keeping top- $k$ , keep the top  $p$  percent of the probability mass.



Choose from the smallest set from the vocabulary such that

$$\sum_{w \in V^{(p)}} P(w | \mathbf{w}_{\leq t}) \geq p.$$

Diversifies word selection with less dependence on nature of distribution.

Depends on the distribution. Could be low variance, reducing diversity.

# Next Token Selection - Temperature

- Before applying softmax to calculate probabilities, divide the logit outputs by a temperature  $T$ . *new param*

*normal*

$$\text{softmax}_i(\mathbf{z}) = \frac{e^{z_i}}{\sum_j e^{z_j}}$$

*temperature*

$$\text{softmax}_i(\mathbf{z}, T) = \frac{e^{z_i/T}}{\sum_j e^{z_j/T}}$$

- What happens as  $T \rightarrow \infty$ ?

*exponents  $\rightarrow 0$*

*$e^{z_i/T} \rightarrow 1$*

$$P_i = \frac{1}{\# \text{tokens}}$$

- What happens as  $T \rightarrow 0$ ?

*$\rightarrow$  greedy sampling,*

*$\emptyset$  works in serious implementation*

- Offered in most LLM interfaces.

# Next Token Selection – Beam Search

Commonly used in *machine translation*

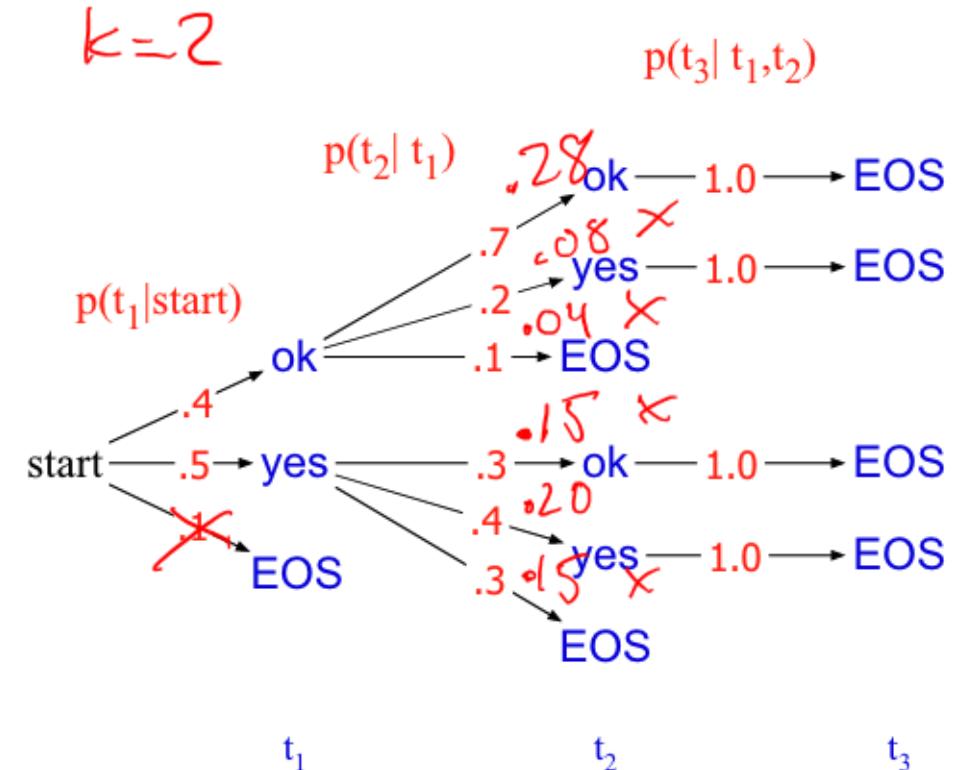
Maintain multiple output choices  
and then choose best  
combinations later via tree search

$$V = \{\text{yes}, \text{ok}, \text{<eos>}\}$$

We want to maximize  $p(t_1, t_2, t_3)$ .

Greedy:  $0.5 \times 0.4 \times 1.0 = 0.20$

Optimal:  $0.4 \times 0.7 \times 1.0 = 0.28$



**TLDR: keep best k paths at each level of tree.  
Less popular as models have gotten bigger.**

# Dummy's Guide to LLM Sampling

- <https://rentry.co/samplers>

Dummy's Guide to Modern LLM Sampling	
1.	Intro Knowledge
1.	Short Glossary
2.	Why tokens?
1.	Why not letters?
2.	Why not whole words?
3.	How are the sub-words chosen?
4.	How does the model generate text?
2.	Sampling
1.	Notes on Algorithm Presentations
1.	Notation Guide
2.	Implementation Considerations
2.	Temperature
3.	Presence Penalty
4.	Frequency Penalty
5.	Repetition Penalty
6.	DRY (Don't Repeat Yourself)
7.	Top-K
8.	Top-P
9.	Min-P
10.	Top-A
11.	XTC (eXclude Top Choices)
12.	Top-N-Sigma
13.	Tall-Free Sampling
14.	Eta Cutoff
15.	Epsilon Cutoff
16.	Locally Typical Sampling
17.	Quadratic Sampling
18.	Mirostat Sampling
19.	Dynamic Temperature Sampling
20.	Beam Search
21.	Contrastive Search
3.	Sampler Order
1.	The Typical Sampling Pipeline
2.	Effects and Interactions of Samplers with Each Other
1.	How Samplers Transform Distributions
2.	Critical Order-Dependent Interactions
1.	Temperature Before vs. After Filtering
2.	Penalties Before vs. After Other Samplers
3.	DRY's Position Matters
3.	Synergies and Conflicts
1.	Synergistic Combos
2.	Conflicting Combos
4.	Advanced: More on Tokenizers
1.	Building the Vocabulary
1.	Byte Pair Encoding (BPE)
2.	SentencePiece
3.	Vocabulary Size and Granularity
4.	Token Boundaries and Phrasing
5.	Handling Rare or New Words

# Any Questions?

???

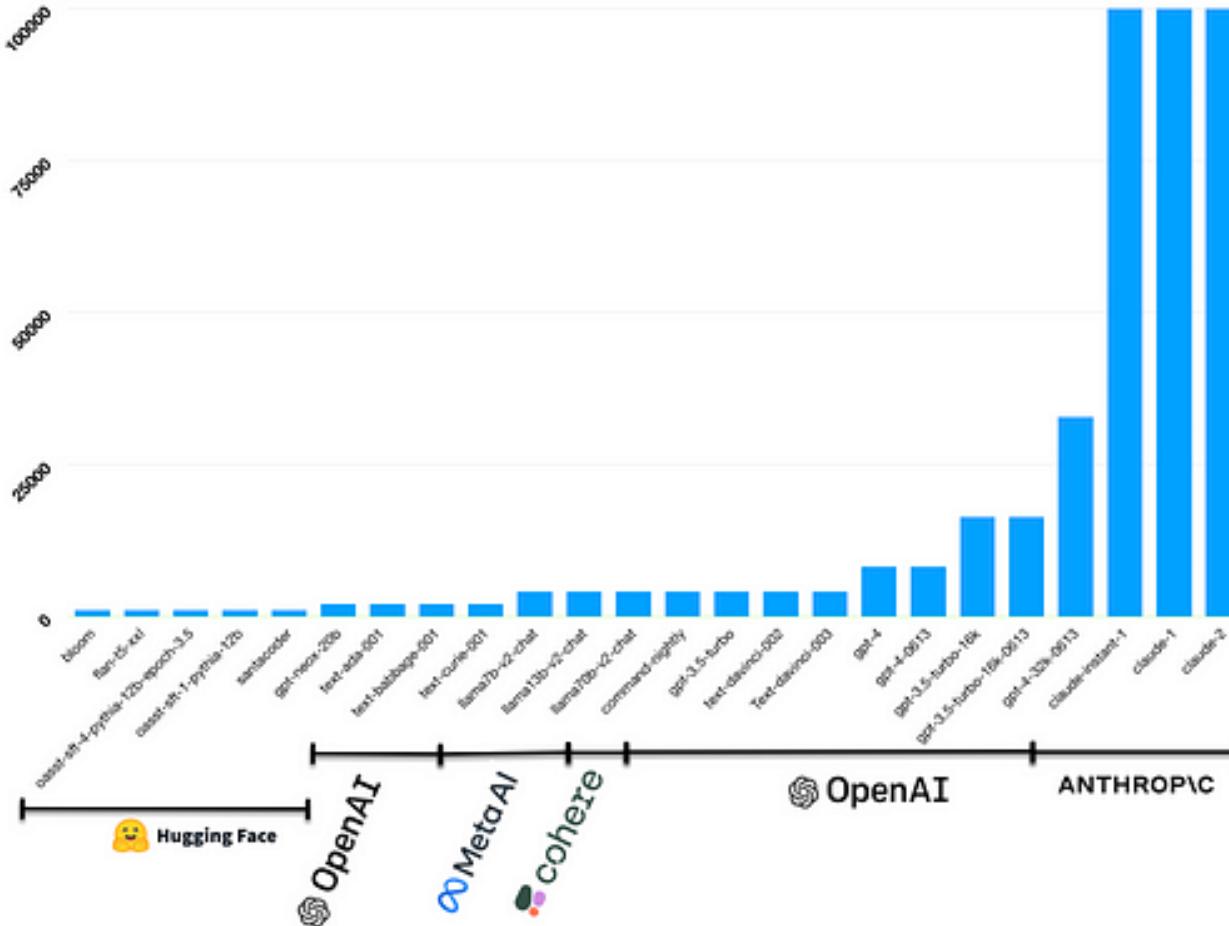
## Moving on

- Transformer recap
- What are tokens?
- Tokenization and word embedding
- Next token selection
- Transformers for long sequences

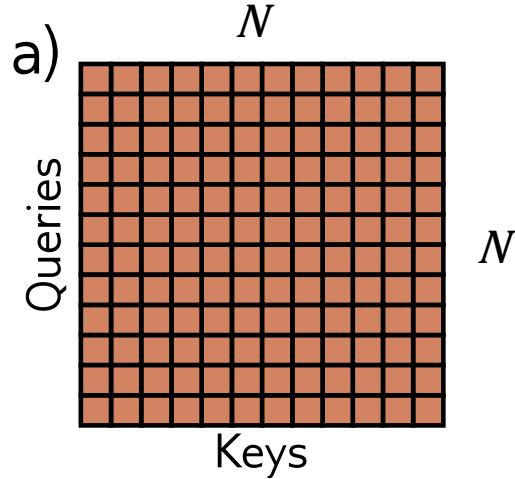
# Context Length of LLMs

## Large Language Model Context Size

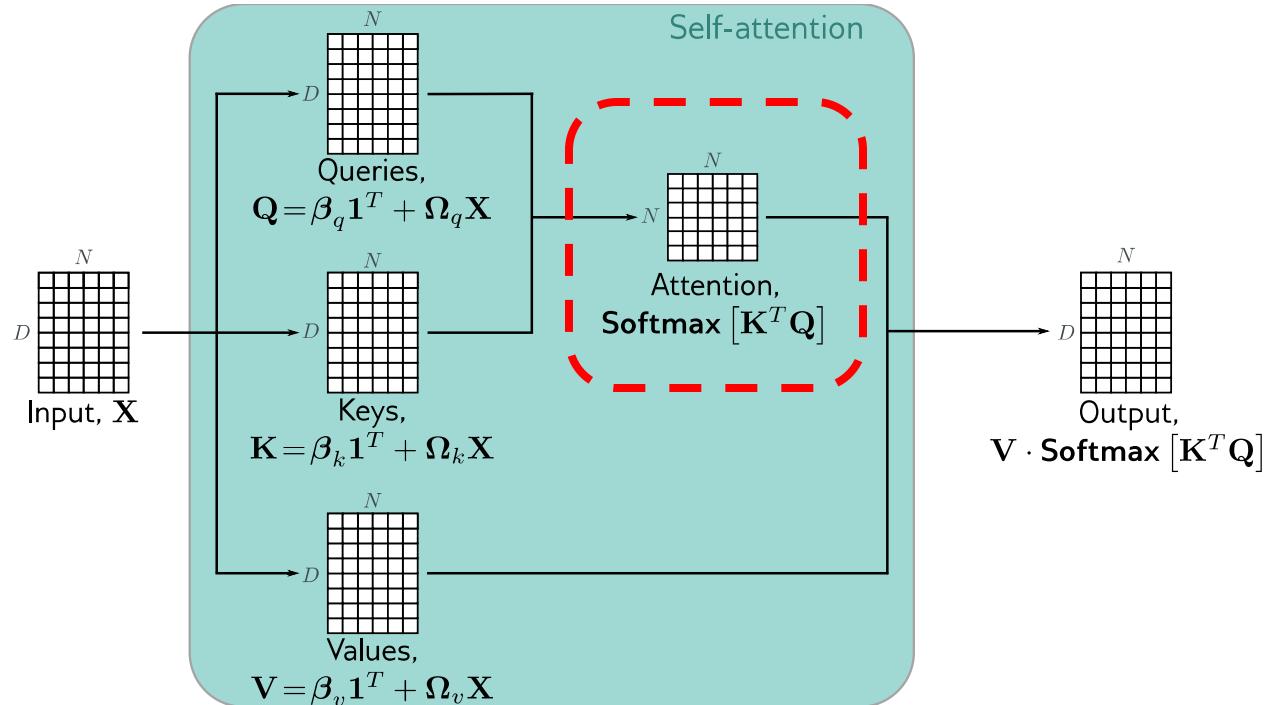
Model	Context Length
Llama 2	32K
GPT4	32K
GPT-4 Turbo, Llama 3.1	128K
Claude 3.5	200K
Sonnet	
Google Gemini 1.5 Pro	Millions



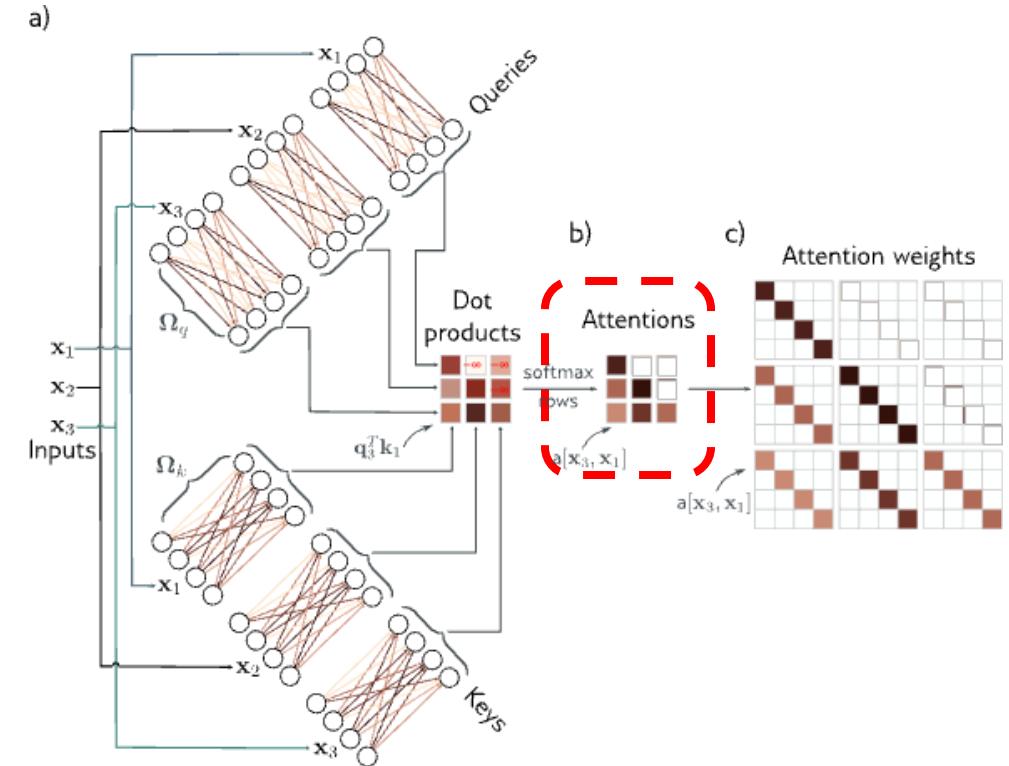
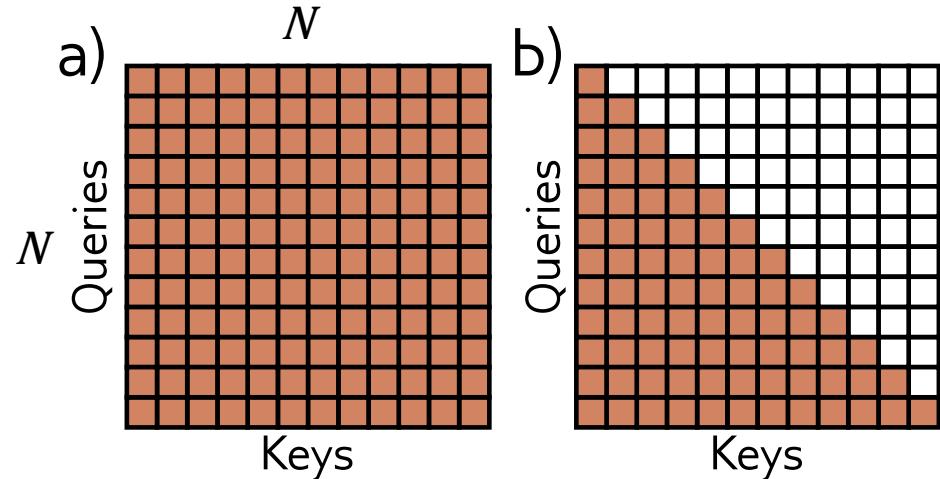
# Attention Matrix



Scales quadratically with sequence length  $N$ , e.g.  $N^2$ .

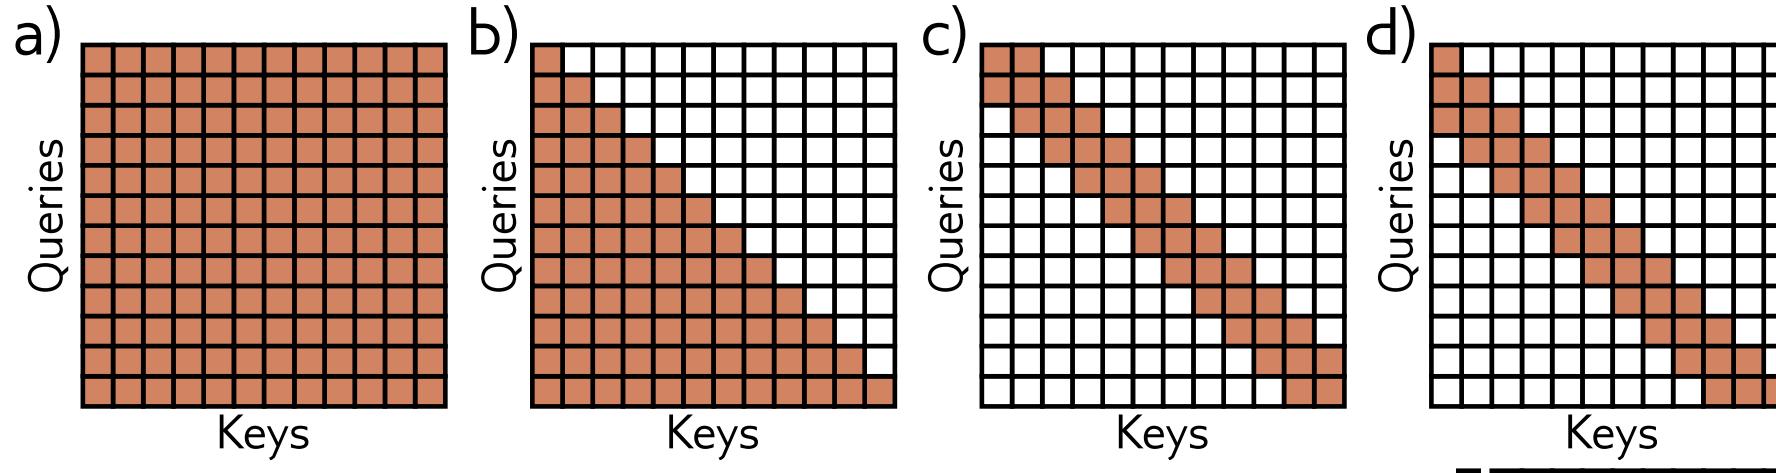


# Masked Attention



~1/2 the interactions but  
still scales quadratically

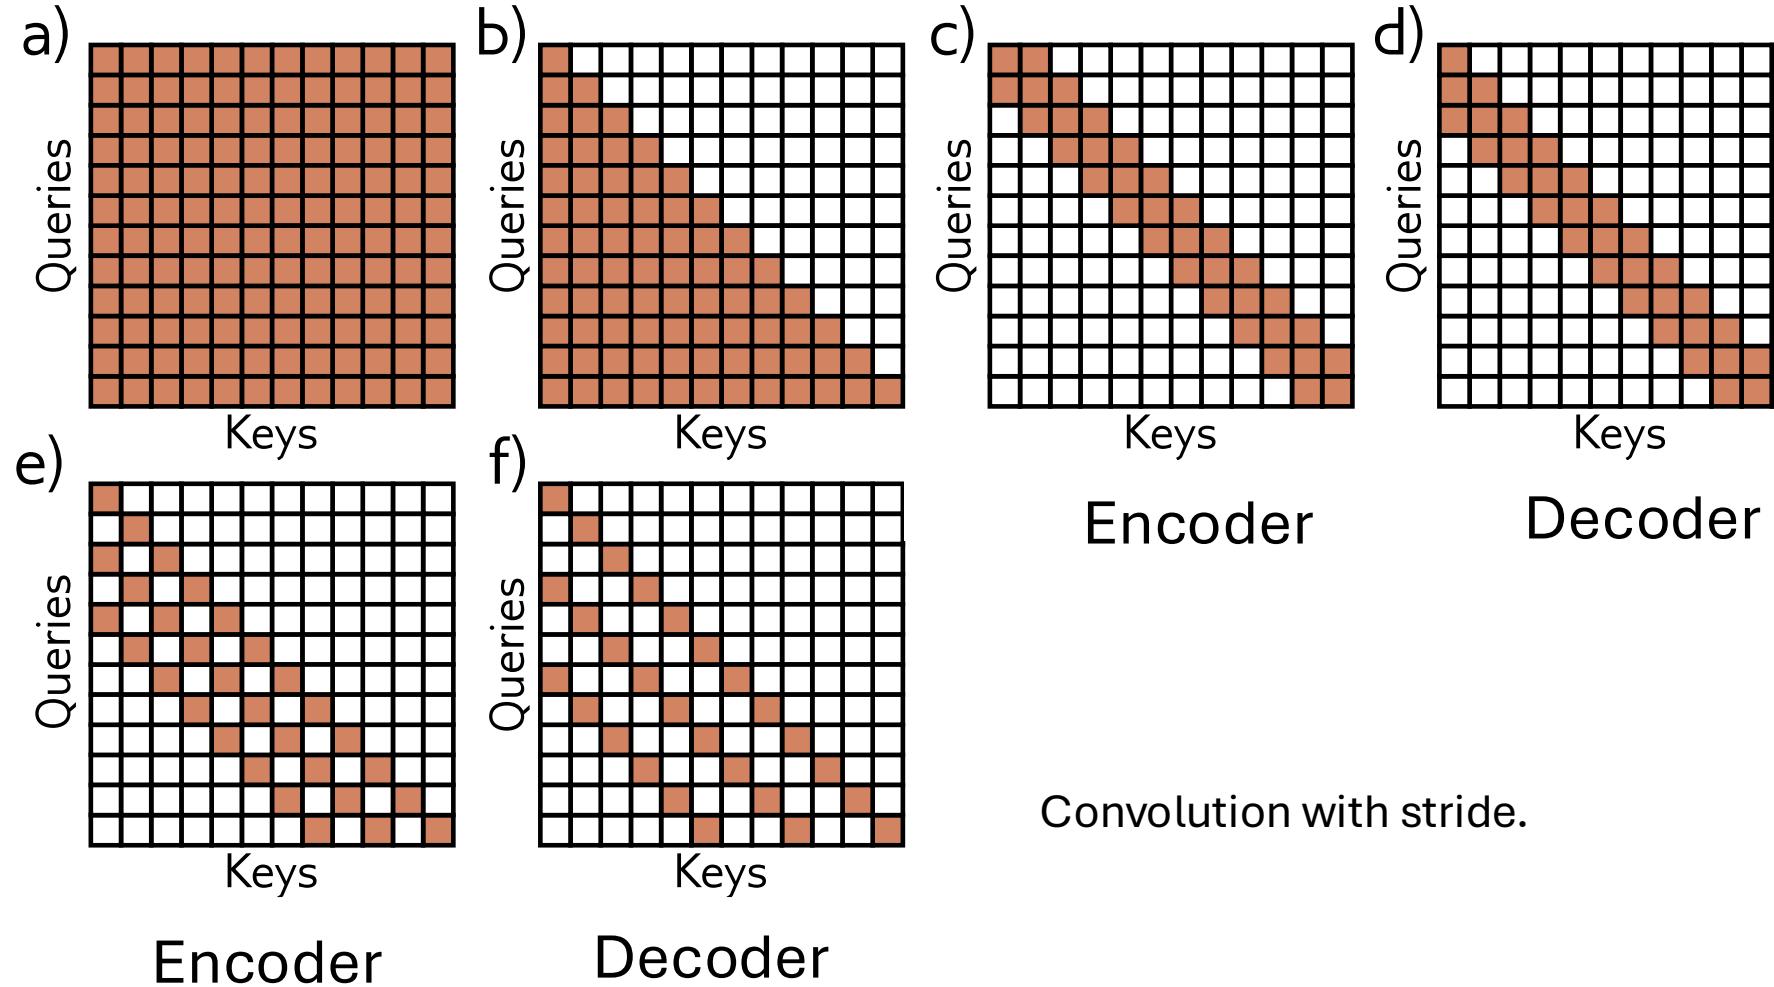
# Use Convolutional Structure in Attention



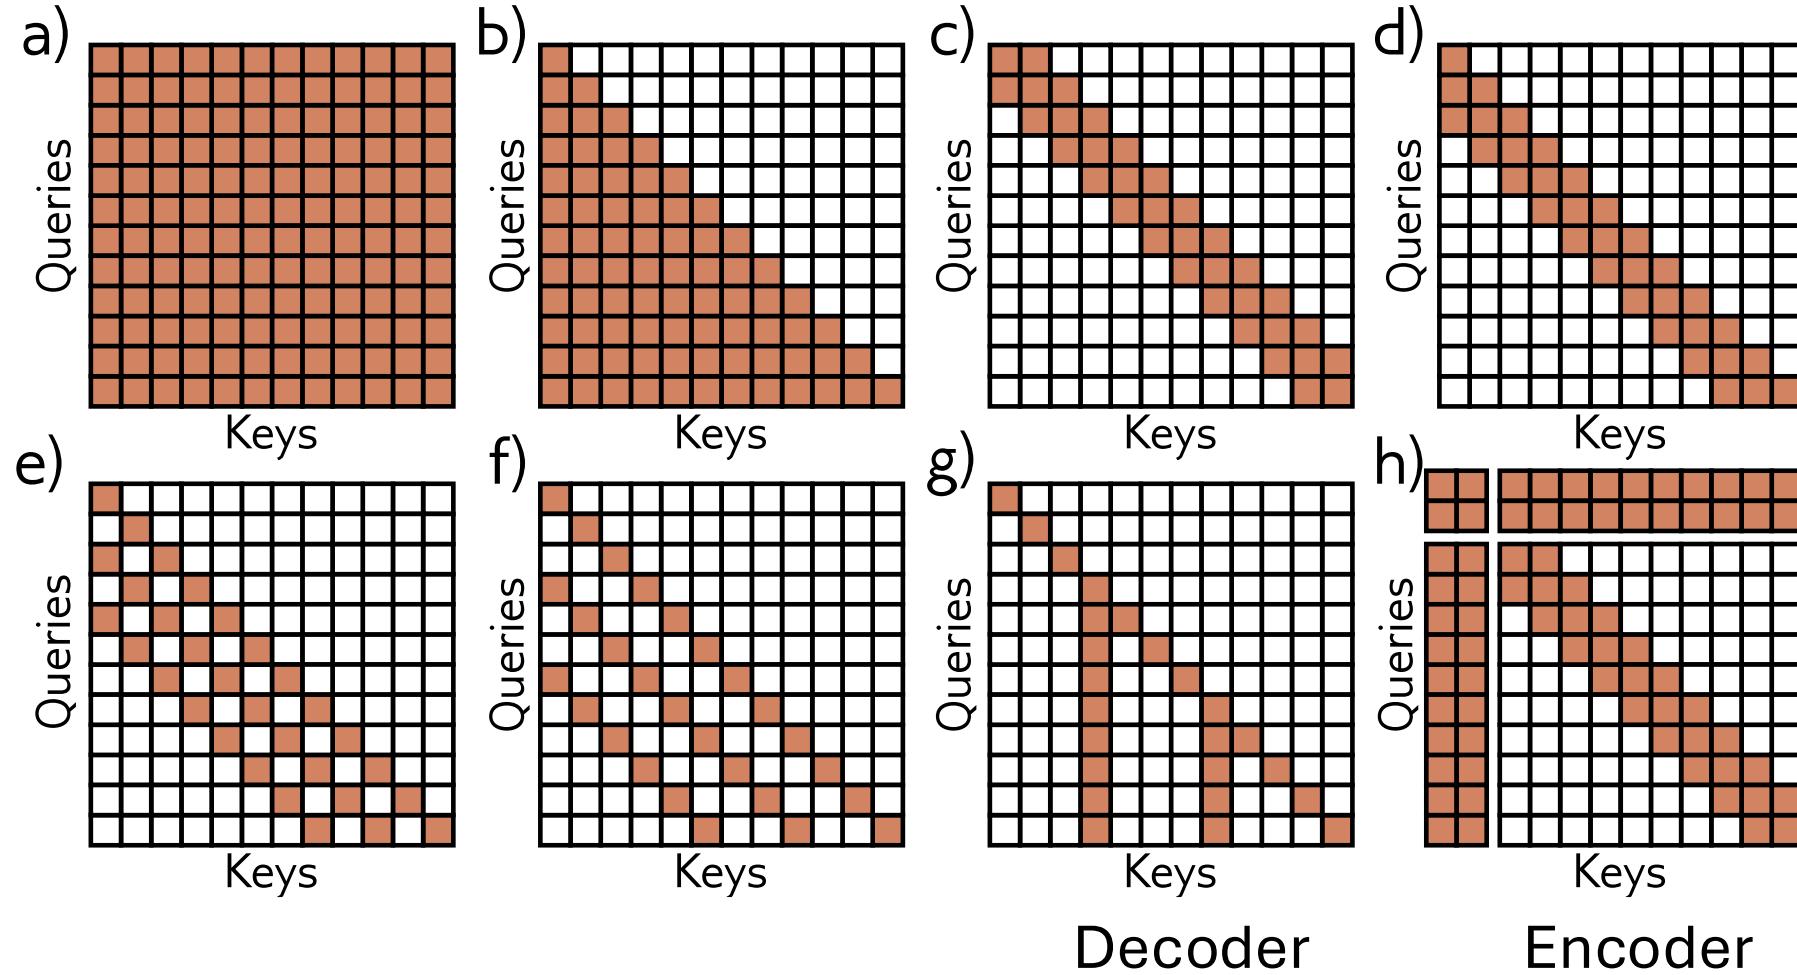
Encoder  
(non-causal)

Decoder  
(causal)

# Dilated Convolutional Structures



# Have some tokens interact globally



# Many of Attempts at Sub-Quadratic Attention

- AFAIK none in state-the-art-models
  - Many published papers claiming linear or  $n \log n$  scaling with comparable performance, but none demonstrated at the same scale.
    - 10-20B parameters vs 500B parameters.
- Many practical speedups for leaner quadratic attention.
  - FlashAttention is popular. Same calculations but optimized to be IO-aware.

# Any Questions?

???

## Moving on

- Transformer recap
- What are tokens?
- Tokenization and word embedding
- Next token selection
- Transformers for long sequences