

# Deep Learning for Data Science

## DS 542

<https://dl4ds.github.io/fa2025/>

Latent Diffusion Models



# Plan for Today

- Diffusion model math
- Training process (again)
- Motivation for latent diffusion
- Latent diffusion
- Conditional generation



# Mathematics of Diffusion Models

Assume the forward and reverse process operate in  $T$  steps.

Both forward and reverse process are discrete so becomes a *Markov chain with Gaussian transition probability*.

$$\mathcal{N}(\mu, \sigma^2)$$

# Mathematics of Diffusion Models

$$x_0 \rightarrow x_1 \rightarrow \dots \rightarrow x_T$$

Denote  $x_0$  as a sample from a distribution  $q(x_0)$ .

original distribution  
complicated!

Forward process: gaussian transition probability

$$q(x_t|x_{t-1}) = \mathcal{N}(x_t; \underbrace{\sqrt{(1 - \beta_t)} x_{t-1}}_{\text{shrink factor}}, \underbrace{\beta_t I}_{\text{noise scale}}) \quad \text{where } t \in \mathbb{N}$$

and where  $\beta_t$  indicates trade-off between info to be kept from previous step and new noise added.

$\beta_t$  chosen  
noise scale  
at time  $t$

$$\beta_1 \sim 1e^{-4}$$

mostly signal,  
small noise

$$\beta_T \sim 0.02$$

paper recommendations,

$T=1000$   
often for training  
sometimes smaller  
T used for speed

linearly interpolate  
in between

# Mathematics of Diffusion Models

Denote  $x_0$  as a sample from a distribution  $q(x_0)$ .

Forward process: gaussian transition probability

$$q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{(1 - \beta_t)} x_{t-1}, \beta_t I) \quad \text{where } t \in \mathbb{N}$$

and where  $\beta_t$  indicates trade-off between info to be kept from previous step and new noise added.

We can equivalently write

$$x_t = \sqrt{(1 - \beta_t)} x_{t-1} + \sqrt{\beta_t} \epsilon_t \quad \epsilon_t \sim \mathcal{N}(0, I)$$

Discretized diffusion process

# Mathematics of Diffusion Models

Through recurrence, we can represent any step in the chain as directly represented from  $x_0$ :

$$q(x_t|x_0) = \mathcal{N}(x_t; \underbrace{\sqrt{\alpha_t} x_0}_{\text{cumulative variance}}, \underbrace{(1 - \bar{\alpha}_t)I}_{\text{cumulative scaling from steps}})$$

where

$$\alpha_t = (1 - \beta_t) \quad \text{and} \quad \bar{\alpha}_t = \prod_{i=1}^t \alpha_i = \prod_{i=1}^t (1 - \beta_i)$$

and from the Markov property, the entire forward trajectory is

$$q(x_{0:T}) = q(x_0) \prod_{t=1}^T q(x_t|x_{t-1})$$

# The reverse process

With the assumption on the drift and diffusion coefficients, the reverse of the diffusion process takes the same form.

Reverse gaussian transition probability

$q(x_{t-1} | x_t)$   
can then be approximated by

$p_\theta(x_{t-1} | x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$

where  $\mu_\theta$  and  $\Sigma_\theta$  are two functions parameterized by  $\theta$  and learned.

*said this was  
a diffusion process  
too last time  
but need to  
learn func*

# The reverse process

*conditional  
give independence so we can just multiply*

Using the Markov property, the probability of a given backward trajectory can be approximated by

$$p_\theta(x_{0:T}) = \underbrace{p(x_T)}_{\text{"start" of reverse process}} \prod_{t=1}^T p_\theta(x_{t-1}|x_t)$$

where  $p(x_T)$  is an isotropic gaussian distribution that does not depend on  $\theta$

$$p(x_T) = \mathcal{N}(x_T; 0, I)$$

b/c we targeted this w/ forward process design

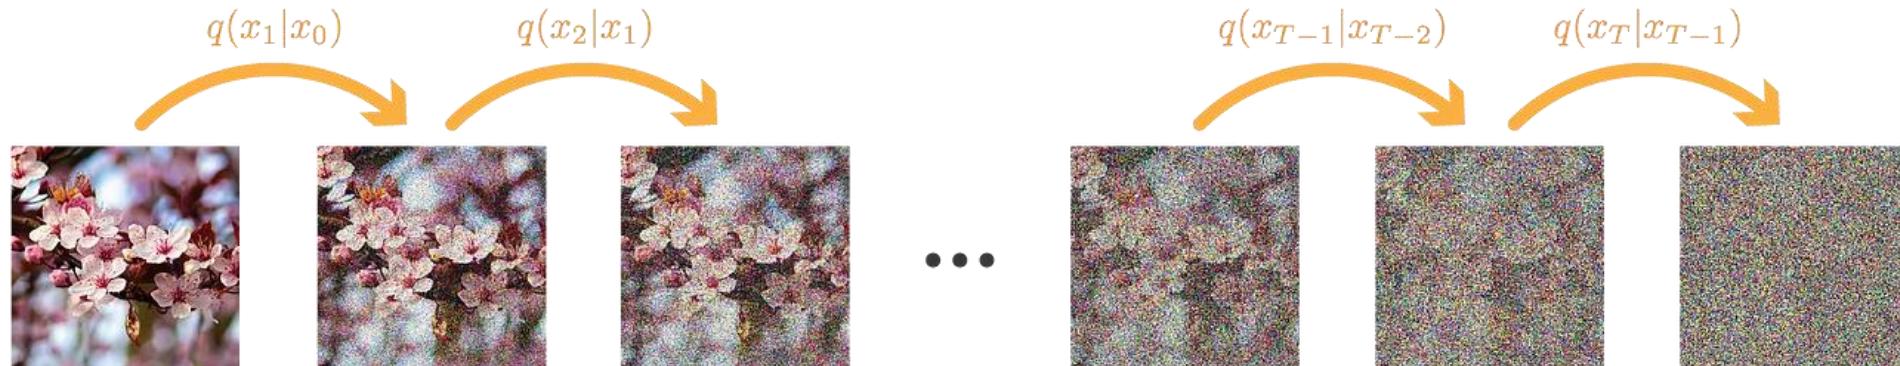
## FIXED FORWARD PROCESS

Initial distribution

$$q(x_0)$$

Gaussian transition kernel

$$q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I)$$



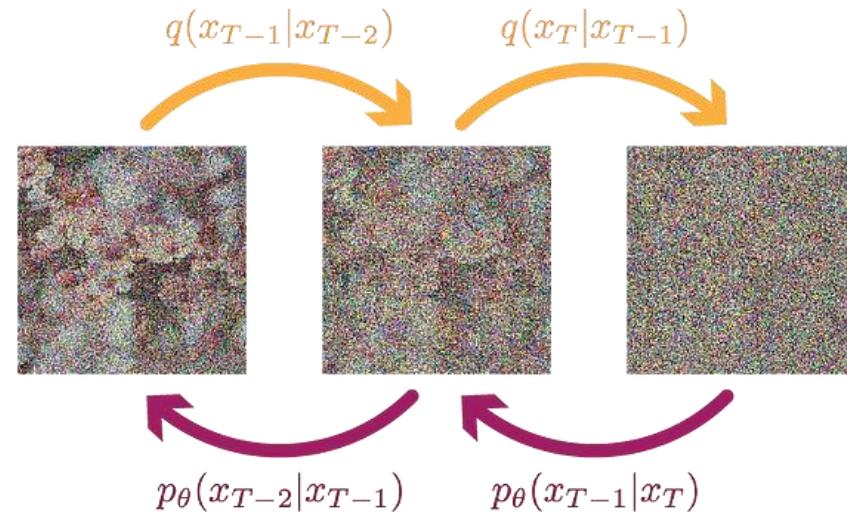
$p_\theta(x_0)$  should match  $q(x_0)$

Approximation of

$$q(x_{t-1}|x_t)$$

Gaussian transition kernel with parameters to be learned

$$p_\theta(x_{t-1}|x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$$



$\uparrow$   
learned func

Initial distribution

$$p(x_T) = \mathcal{N}(x_T; 0, I)$$

## LEARNED BACKWARD PROCESS

# Questions

How do we learn the parameters  $\theta$  for  $\mu_\theta$  and  $\Sigma_\theta$ ?

What is the loss to be optimized?

- We hope that  $p_\theta(x_0)$ , the distribution of the last step of the reverse process, will be close to  $q(x_0)$

make these match ...

probability distributions comparison

★ KL divergence

Earth Mover's / Wasserstein

# Optimization Objective

*truth*      ↓      *distribution being optimized*      ↓

$$\mu_\theta^*, \Sigma_\theta^* = \arg \min_{\mu_\theta, \Sigma_\theta} (D_{KL}(q(\underline{x}_0) || p_\theta(\underline{x}_0)))$$

*pick best  
params  
to make  
distributions  
close*

$$= \arg \min_{\mu_\theta, \Sigma_\theta} \left( - \int q(x_0) \log \left( \frac{p_\theta(x_0)}{q(x_0)} \right) dx_0 \right)$$

*flipped from  
normal  
formula*

$$= \arg \min_{\mu_\theta, \Sigma_\theta} \left( - \int q(x_0) \log(p_\theta(x_0)) dx_0 \right)$$

*b/c  $q(x_0) \log(q(x_0))$*

*does not have*

~~0~~ ⊖

# Skipping a lot more math

- Expand  $p$ -theta as marginalization integral
- Use Jensen's inequality to define a slightly simpler upper bound to the loss
- Some manipulations with Bayes' Theorem
- Properties of KL divergence of two gaussian distributions
- An additional simplification suggested by [Ho et al 2020]

# Diffusion models in practice

Simple formulas  
best if iterative.

We have the forward process

$$q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I)$$

$$q(x_t|x_0) = \mathcal{N}(x_t; \sqrt{\bar{\alpha}_t}x_0, (1 - \bar{\alpha}_t)I)$$

and our reverse process

$$p_\theta(x_{t-1}|x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$$

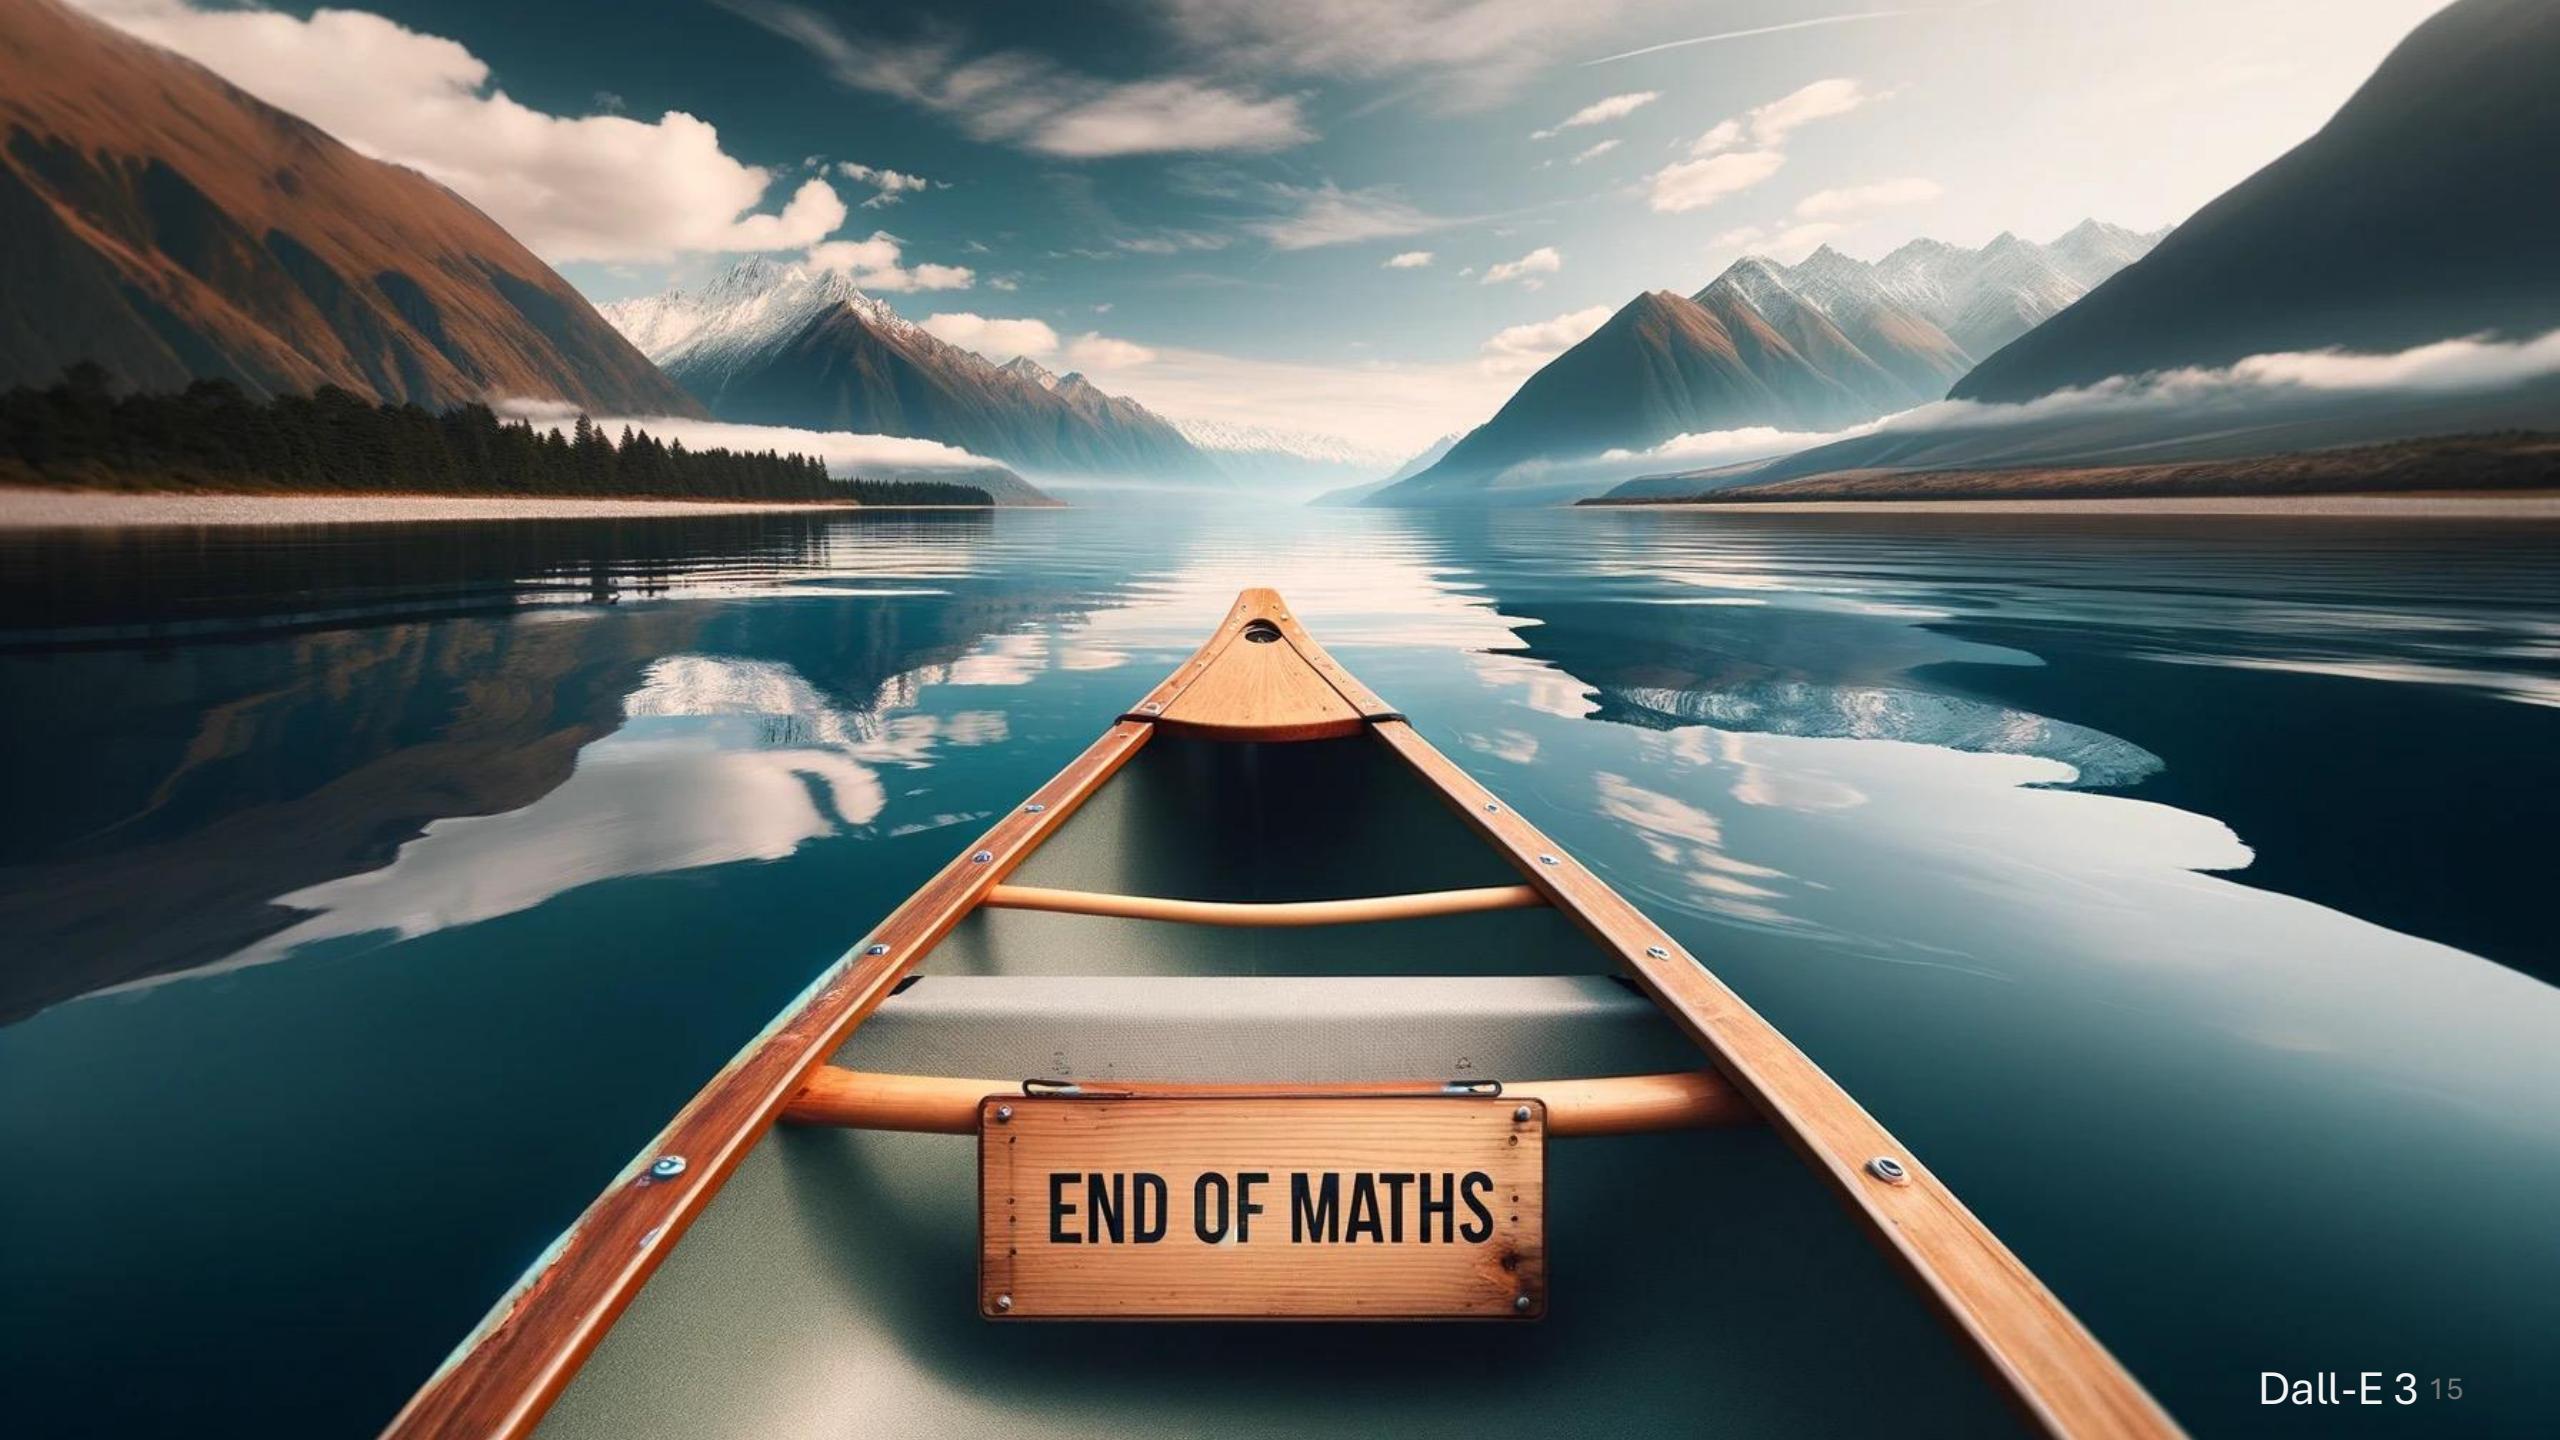
and we want to train to minimize this simplified upper bound ~~\*~~

$$\mathbb{E}_{x_0, t, \epsilon} (\|\epsilon - \underbrace{\epsilon_\theta(x_t, t)}_{\substack{\text{actual noise} \\ \text{added} \\ \text{to sample}}} \|^2) = \mathbb{E}_{x_0, t, \epsilon} (\|\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t)\|^2)$$

*easier formula from paper  
or simpler variational bound"*

★ Disclaimer.  
Even though total noise was estimated, only remove about 1 step of noise at a time.  
One step output is poor.

Alternative:  
Estimate  $x_0$  or  $\mu(\cdot)$  directly...



# Any Questions?

???

## Moving on

- Diffusion model math
- Training process (again)
- Motivation for latent diffusion
- Latent diffusion
- Conditional generation

# Training Process

---

## Algorithm 1 Training

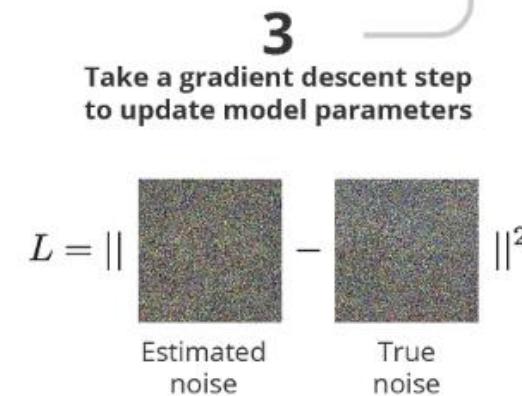
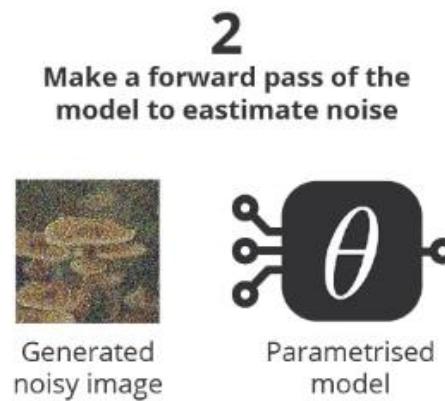
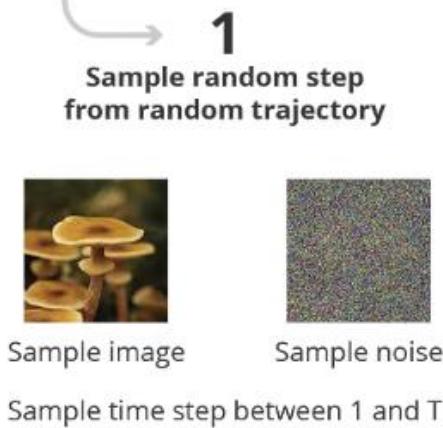
---

```
1: repeat
2:    $x_0 \sim q(x_0)$ 
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$ 
4:    $\epsilon \sim \mathcal{N}(0, I)$ 
5:    $x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon$ 
6:   Take gradient descent step on  $\nabla_{\theta} \|\epsilon - \epsilon_{\theta}(x_t, t)\|^2$ 
7: until converged
```

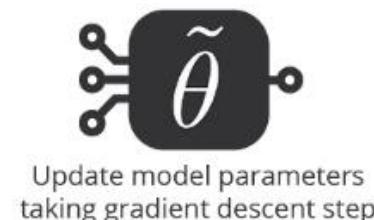
---

▷ Sample random initial data  
▷ Sample random step  
▷ Sample random noise  
▷ Rand. step of rand. trajectory  
▷ Optimisation

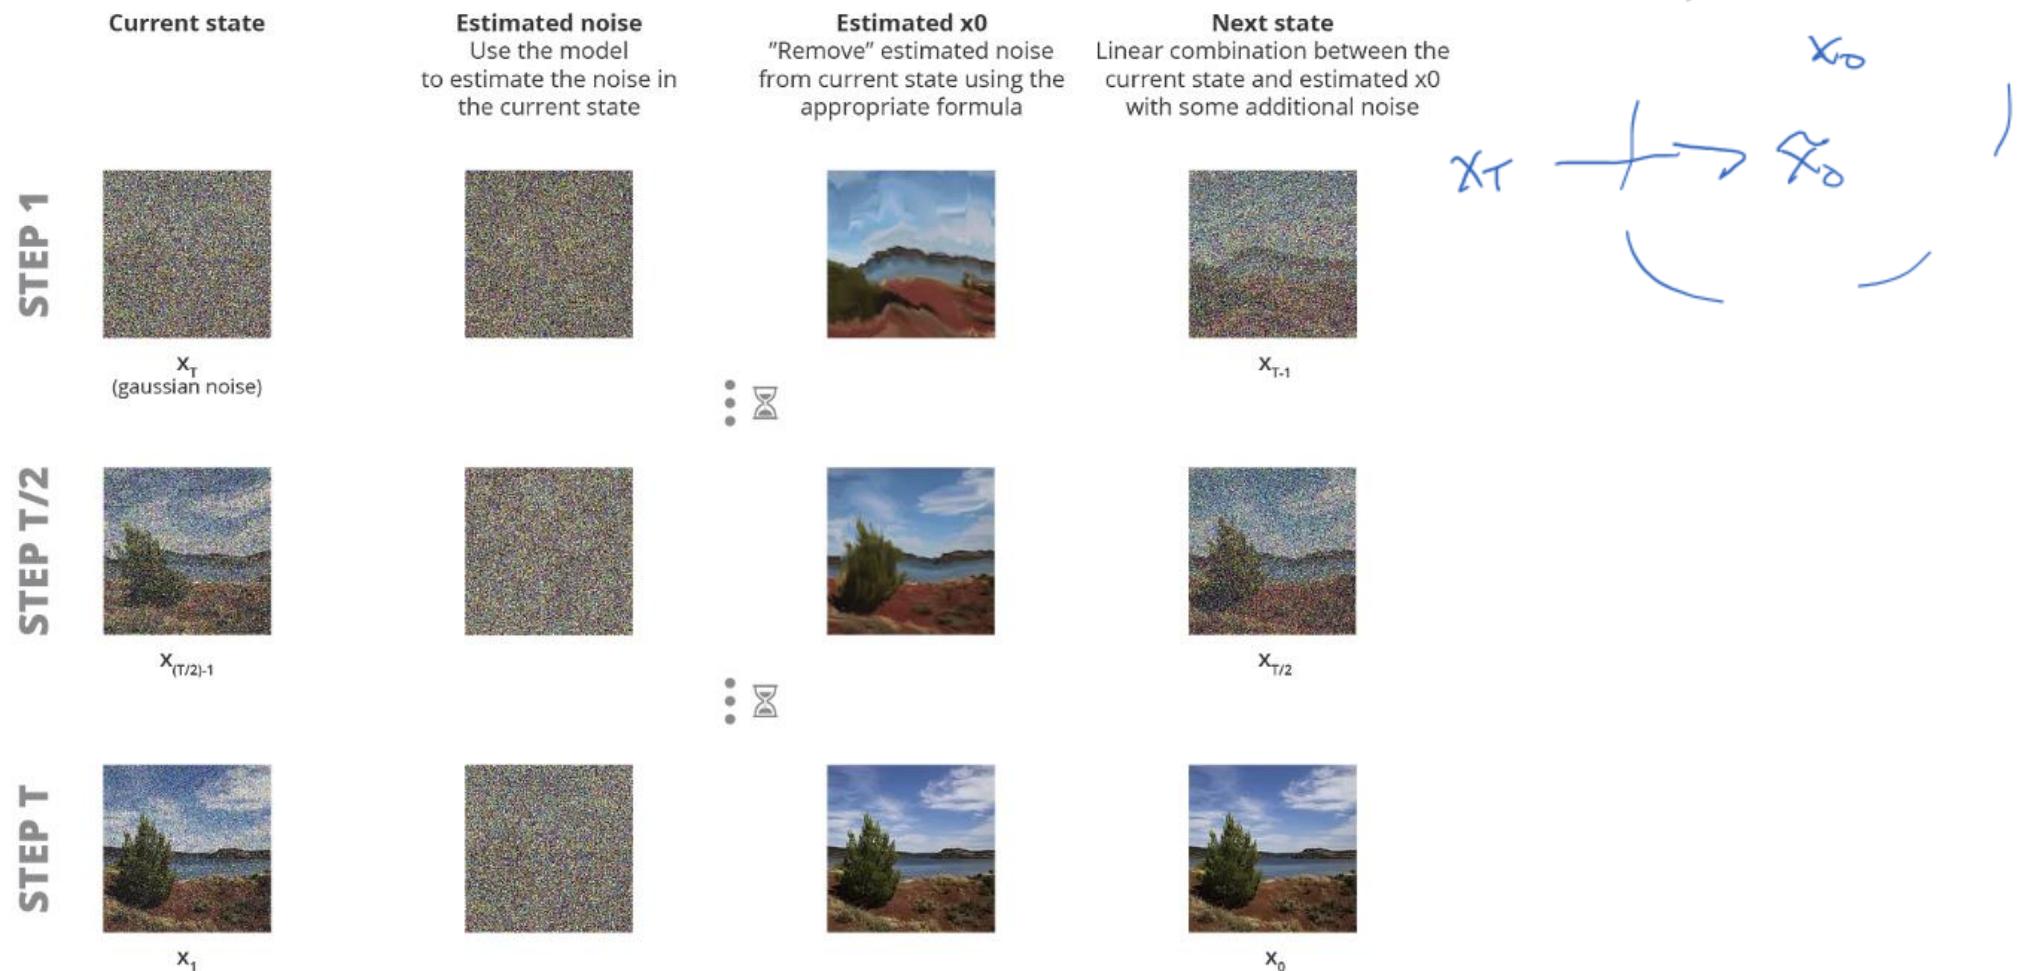
Iterate until convergence



Update model parameters taking gradient descent step



# To sample/generate



# To sample/generate

---

## Algorithm 2 Sampling

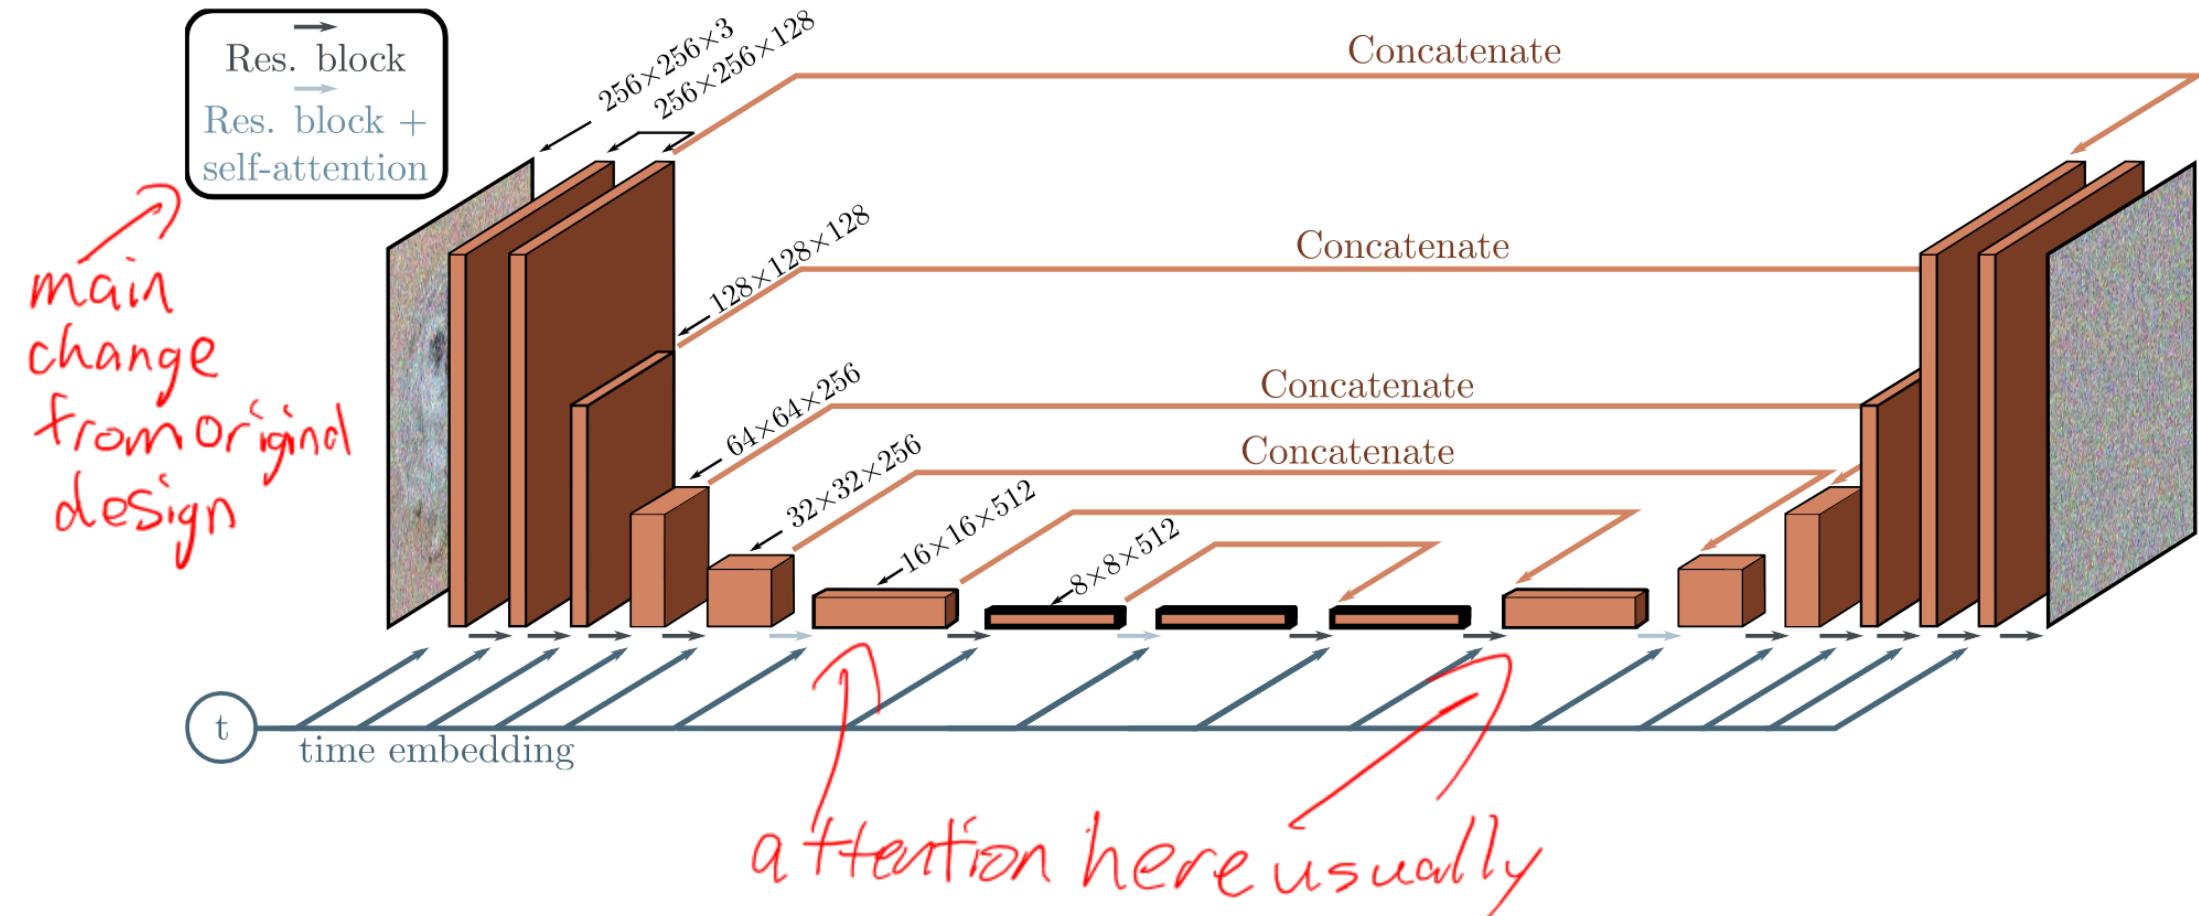
---

```
1:  $x_T \sim \mathcal{N}(0, I)$                                      ▷Initial isotropic gaussian noise sampling  
2: for  $t = T, \dots, 1$  do  
3:    $z \sim \mathcal{N}(0, I)$  if  $t > 1$  else  $z = 0$            ▷Sample random noise (if not last step)  
4:    $\tilde{\epsilon} = \epsilon_\theta(x_t, t)$                       ▷Estimated noise in current noisy data  
5:    $\tilde{x}_0 = \frac{1}{\sqrt{\bar{\alpha}_t}}(x_t - \sqrt{1 - \bar{\alpha}_t}\tilde{\epsilon})$     ▷Estimated  $x_0$  from estimated noise  
6:    $\tilde{\mu} = \mu_t(x_t, \tilde{x}_0) \left(= \frac{1}{\sqrt{\alpha_t}} \left( x_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right) \right)$     ▷Mean for previous step sampling  
7:    $x_{t-1} = \tilde{\mu} + \sigma_t z$                          ▷Previous step sampling  
8: end for  
9: return  $x_0$ 
```

---

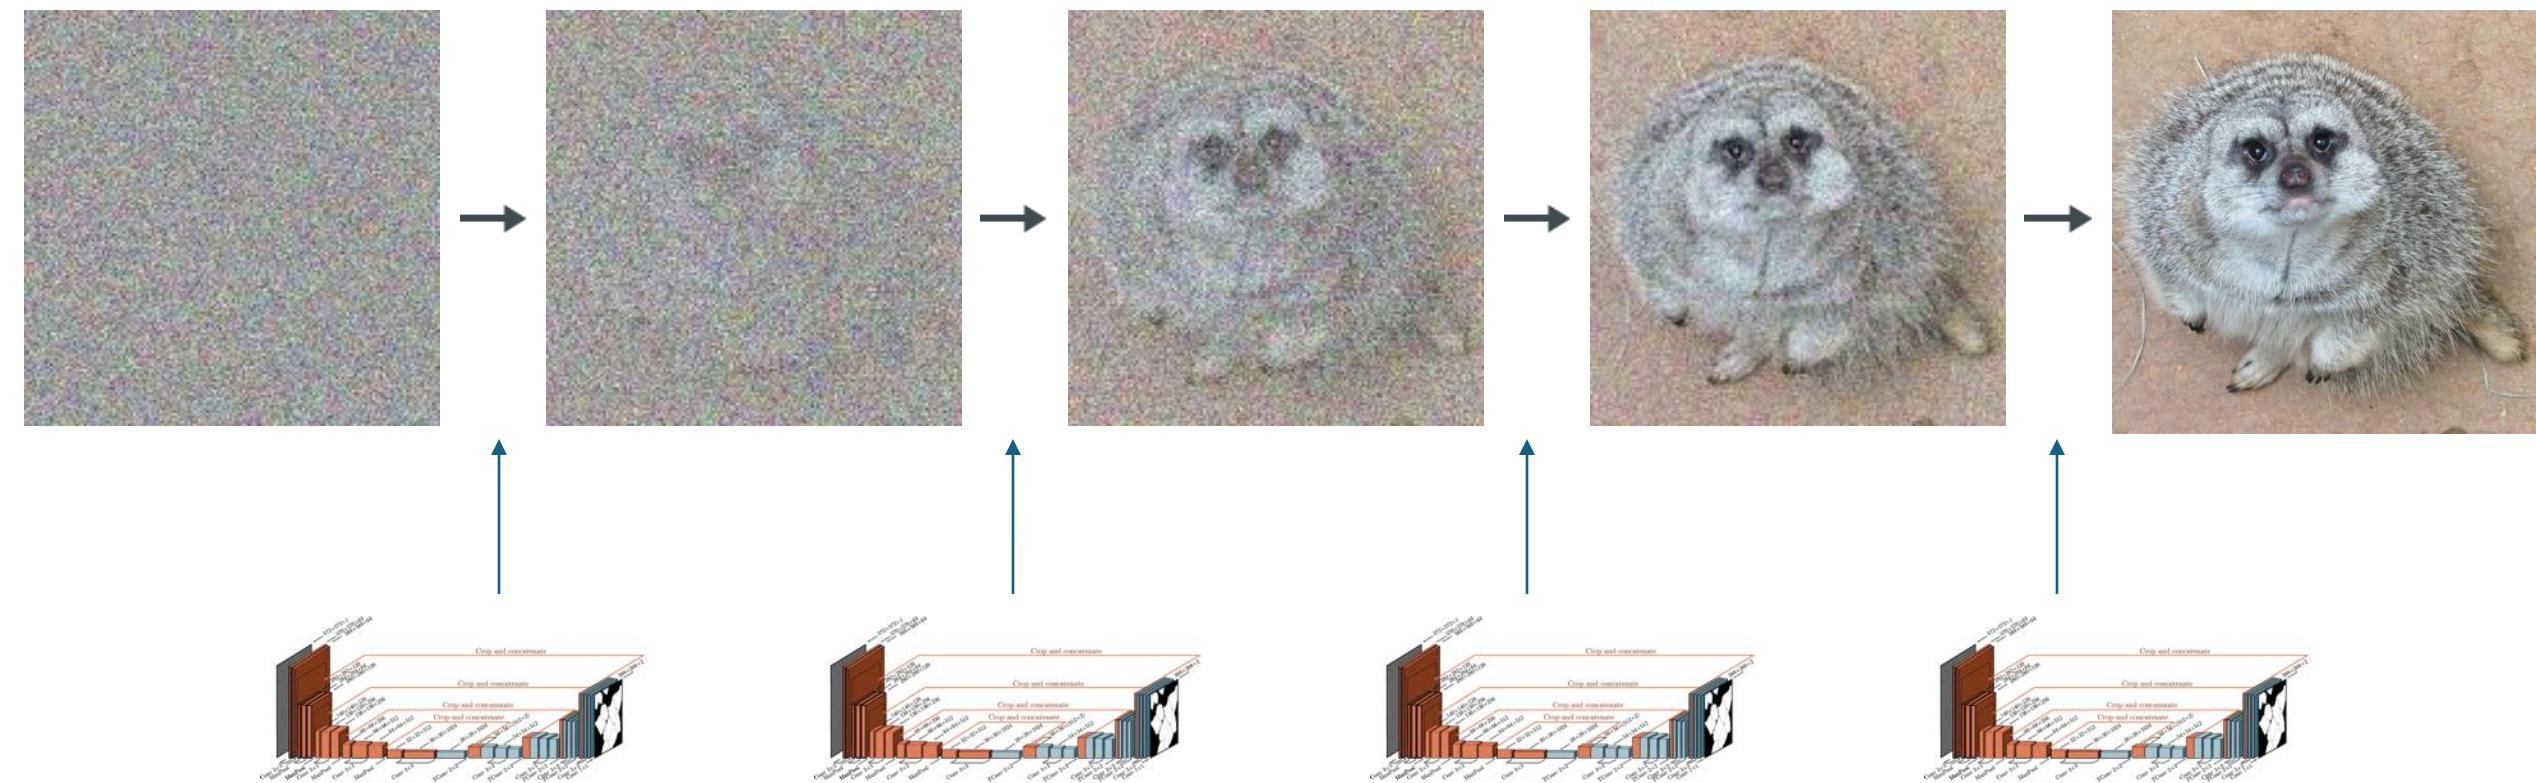
# U-Net (2016) as the Model

Output is same size as the input.



# U-Net for reverse diffusion

dog or weasel?



# Any Questions?

???

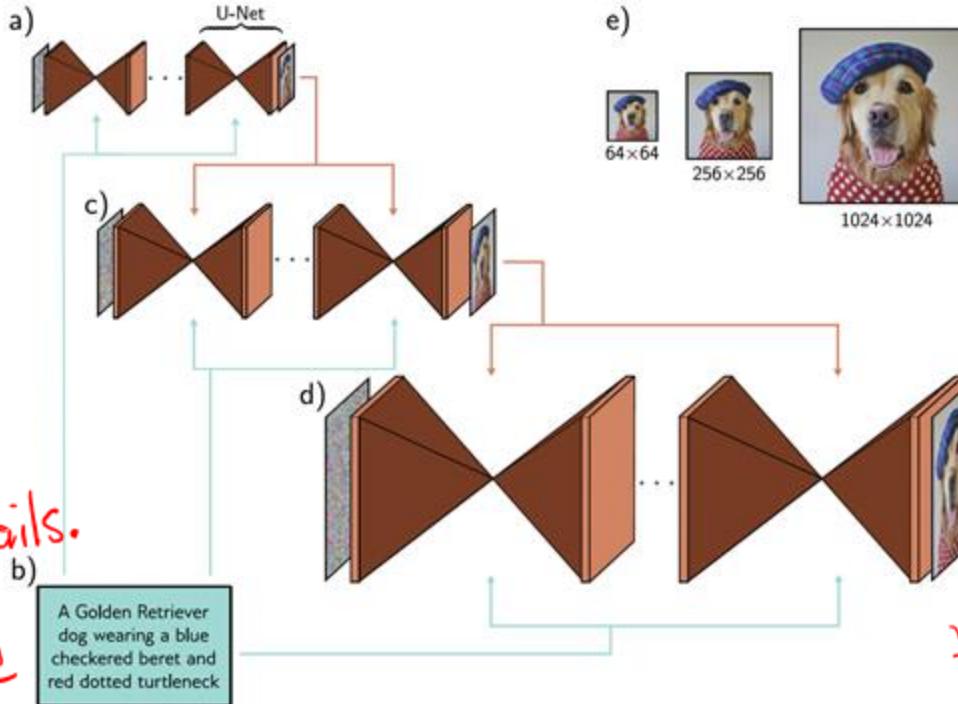
## Moving on

- Diffusion model math
- Training process
- Motivation for latent diffusion
- Latent diffusion
- Conditional generation

# Hi-Res Conditional Generation via Multi-Resolution Generation

Positive interpretation  
lowres output specifies overall layout.  
highres output fills in consistent details.

Negative interpretation  
high resolution version was too hard to do directly



**Figure 18.11** Cascaded conditional generation based on a text prompt. a) A diffusion model consisting of a series of U-Nets is used to generate a 64×64 image. b) This generation is conditioned on a sentence embedding computed by a language model. c) A higher resolution 256×256 image is generated and conditioned on the smaller image and the text encoding. d) This is repeated to create a 1024×1024 image. e) Final image sequence. Adapted from Saharia et al. (2022b).

leverage small working model / generation to bootstrap bigger generation.

Pyramid complexity often gives net speedup. Almost guarantees consistency.

# Diffusion Models

Eventually got quality comparable to GANs...

- Better than VAEs and normalizing flows, but very very slow...
- Every step of the reverse process is working with full size images.
  - Full size input
  - Full size output
- Multi-resolution architectures are often a sign cost is an issue.
  - But also useful for global consistency, so do not disregard.

# Do we have good models?

	GANs	VAEs	Flows	Diffusion
Efficient sampling	✓	✓	✓	✗
High quality	✓	✗	✗	✓
Coverage	✗	?	?	?
Well-behaved latent space	✓	✓	✓	✗
Interpretable latent space	?	?	?	✗
Efficient likelihood	n/a	✗	✓	✗

Diffusion on pixels is slow.

Diffusion on latents will be faster.

# Any Questions?

???

## Moving on

- Diffusion model math
- Training process (again)
- Motivation for latent diffusion
- **Latent diffusion**
- Conditional generation

# Different Models for Different Tradeoffs

Focus on lower dimension latent model that gets the semantic details right...

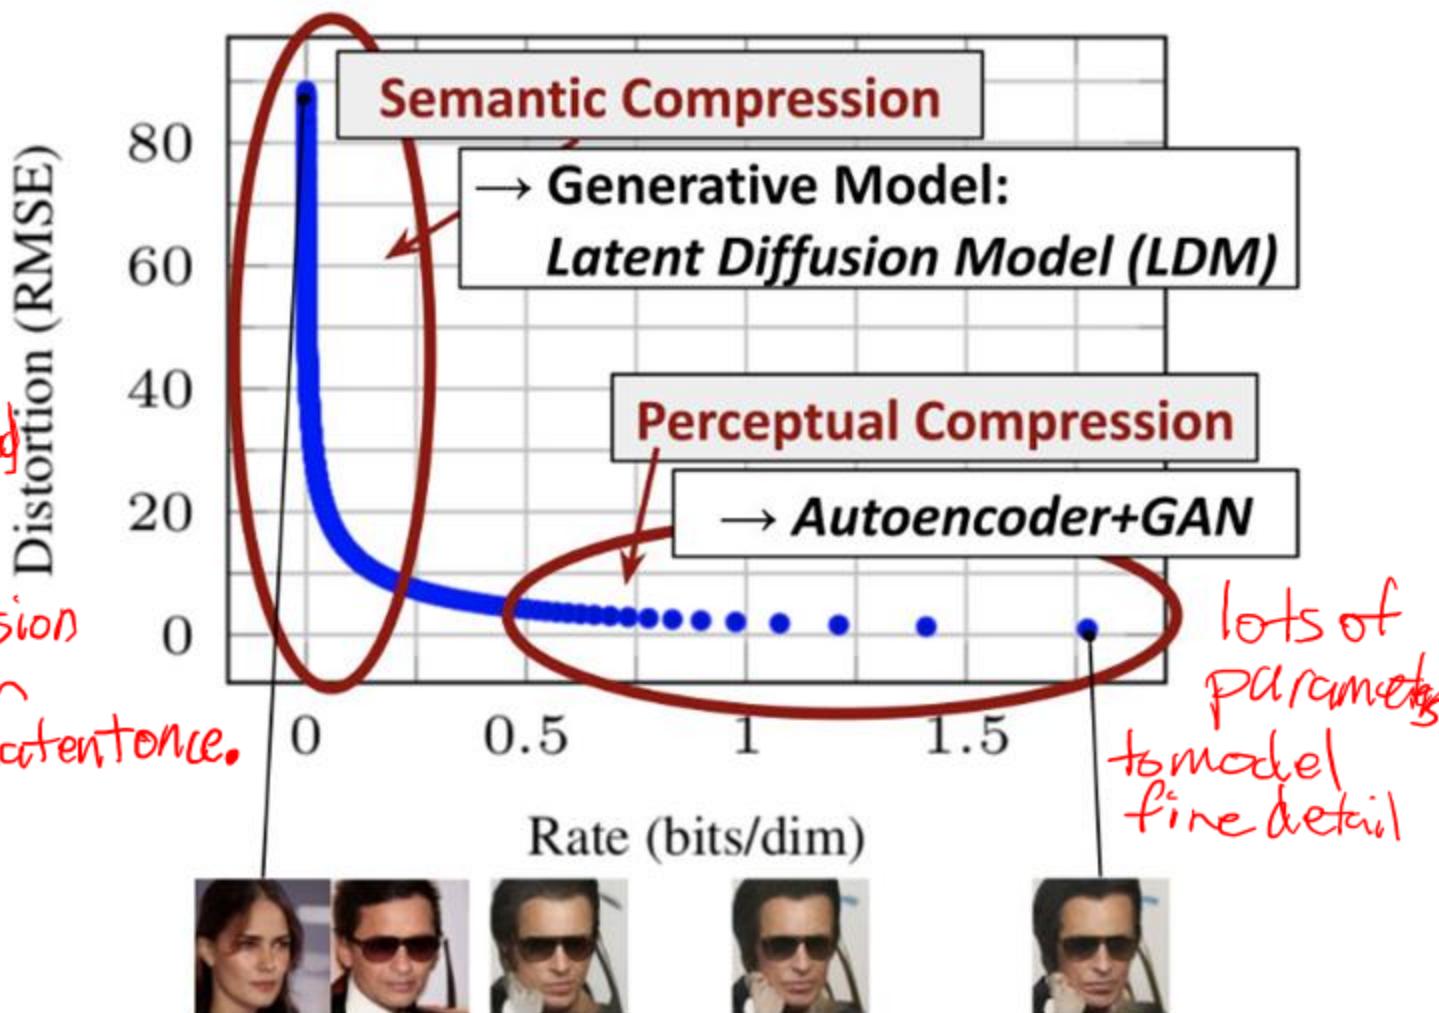
- Diffusion in latent space
- Use another high quality model for image generation.

borrow a pretrained  
model w/small  
latents, do diffusion

“High-Resolution Image  
Synthesis with Latent Diffusion  
Models”

By Rombach, Blattman, Lorenz,  
Esser and Ommer (2021)

you can see objects but blurry  
all the people/objects



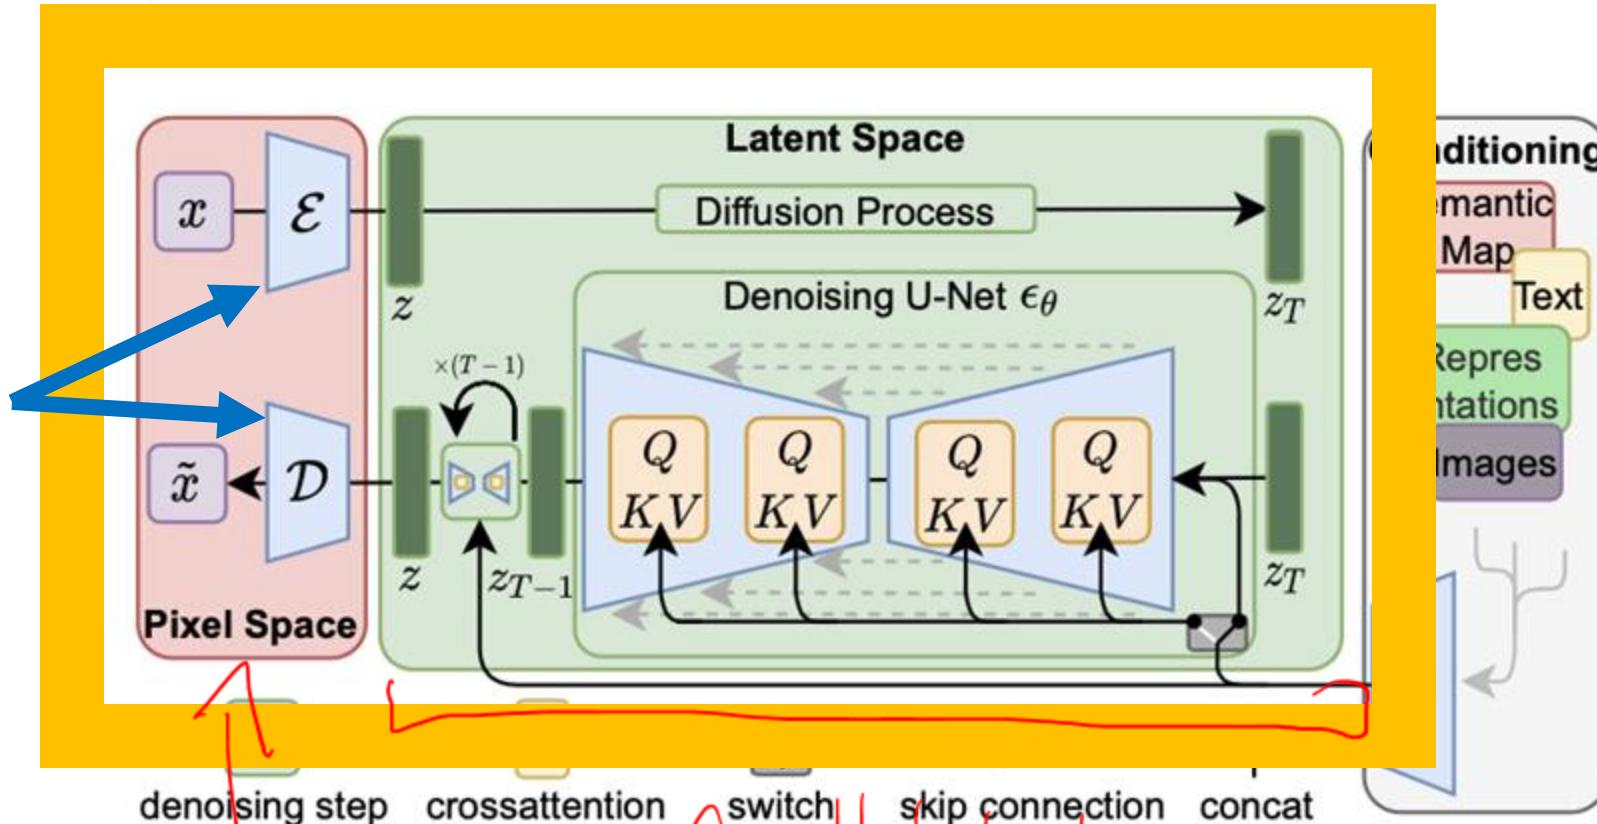
# Combining Auto-Encoder Latents with Diffusion

Pick your favorite autoencoder with a small latent space.



someone  
already  
pre-trained

big images



# Latent Diffusion

Key idea:

- Pixel space is big.
- Run diffusion process in smaller latent space.

“High-Resolution Image Synthesis with Latent Diffusion Models”

By Rombach, Blattman,  
Lorenz, Esser and Ommer  
(2021)



# Latent Diffusion Components

Latent diffusion models have two main components.

- Function **mapping latent codes to high quality images**.
  - This function **does not need to have all the qualities we want** from generative models.
  - In particular, much of the latent space may not make sense. *bad output* ↗ *ok.*
- Function **mapping noise to latent codes of high quality images**.
  - This is the latent diffusion part.
  - Just this part gets retrained for different applications.  
*reverse diffusion ↗ on latents*

# A Random Latent from Stable Diffusion

```
latent = torch.randn(1, 32, 32, 4)  
latent_image = decode_latent(latent)
```

Image shape is 1x256x256x3.

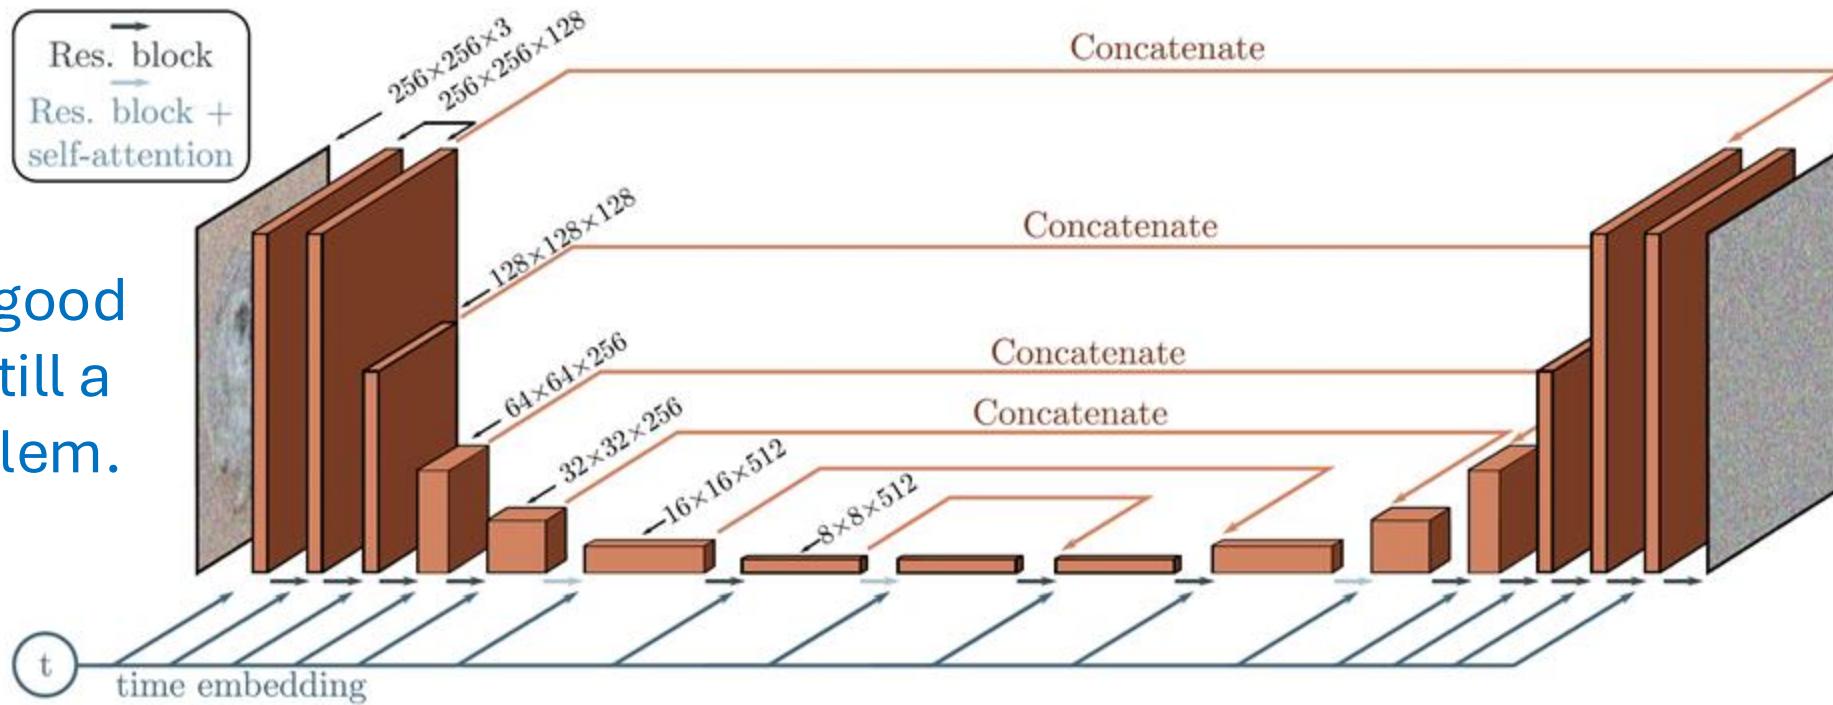
This is that high quality image model?

<https://stability.ai/>

~~Stable Diffusion~~  
~~X3~~



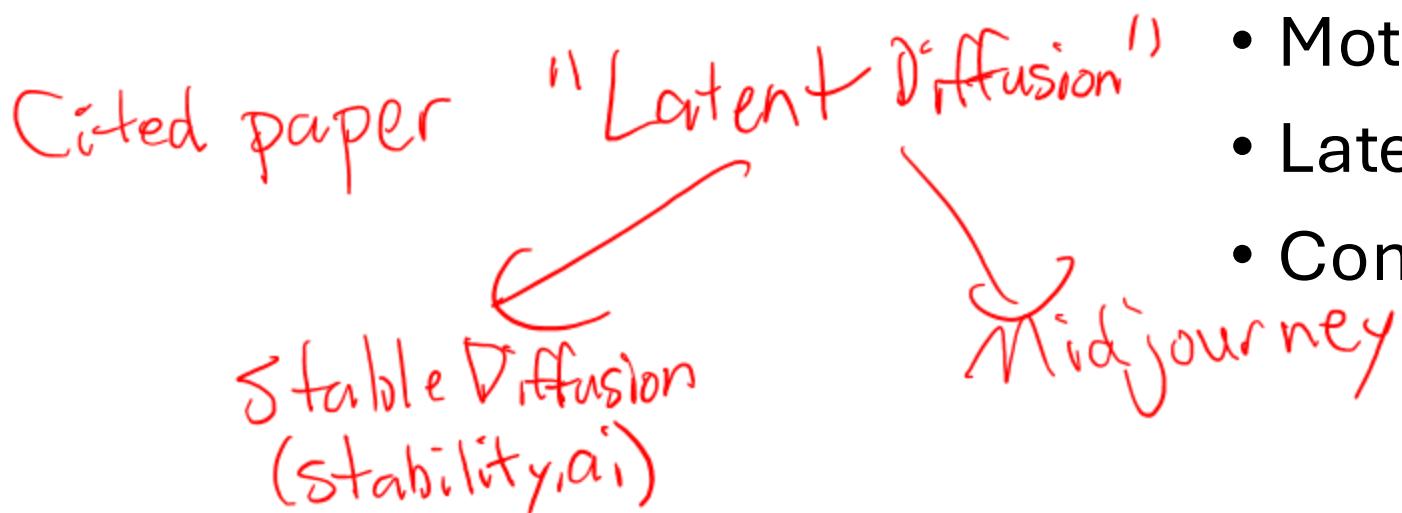
Picking a good latent is still a hard problem.



**Figure 18.9** U-Net as used in diffusion models for images. The network aims to predict the noise that was added to the image. It consists of an encoder which reduces the scale and increases the number of channels and a decoder which increases the scale and reduces the number of channels. The encoder representations are concatenated to their partner in the decoder. Connections between adjacent representations consist of residual blocks, and periodic global self-attention in which every spatial position interacts with every other spatial position. A single network is used for all time steps, by passing a sinusoidal time embedding (figure 12.5) through a shallow neural network and adding the result to the channels at every spatial position at every stage of the U-Net.

# Any Questions?

???



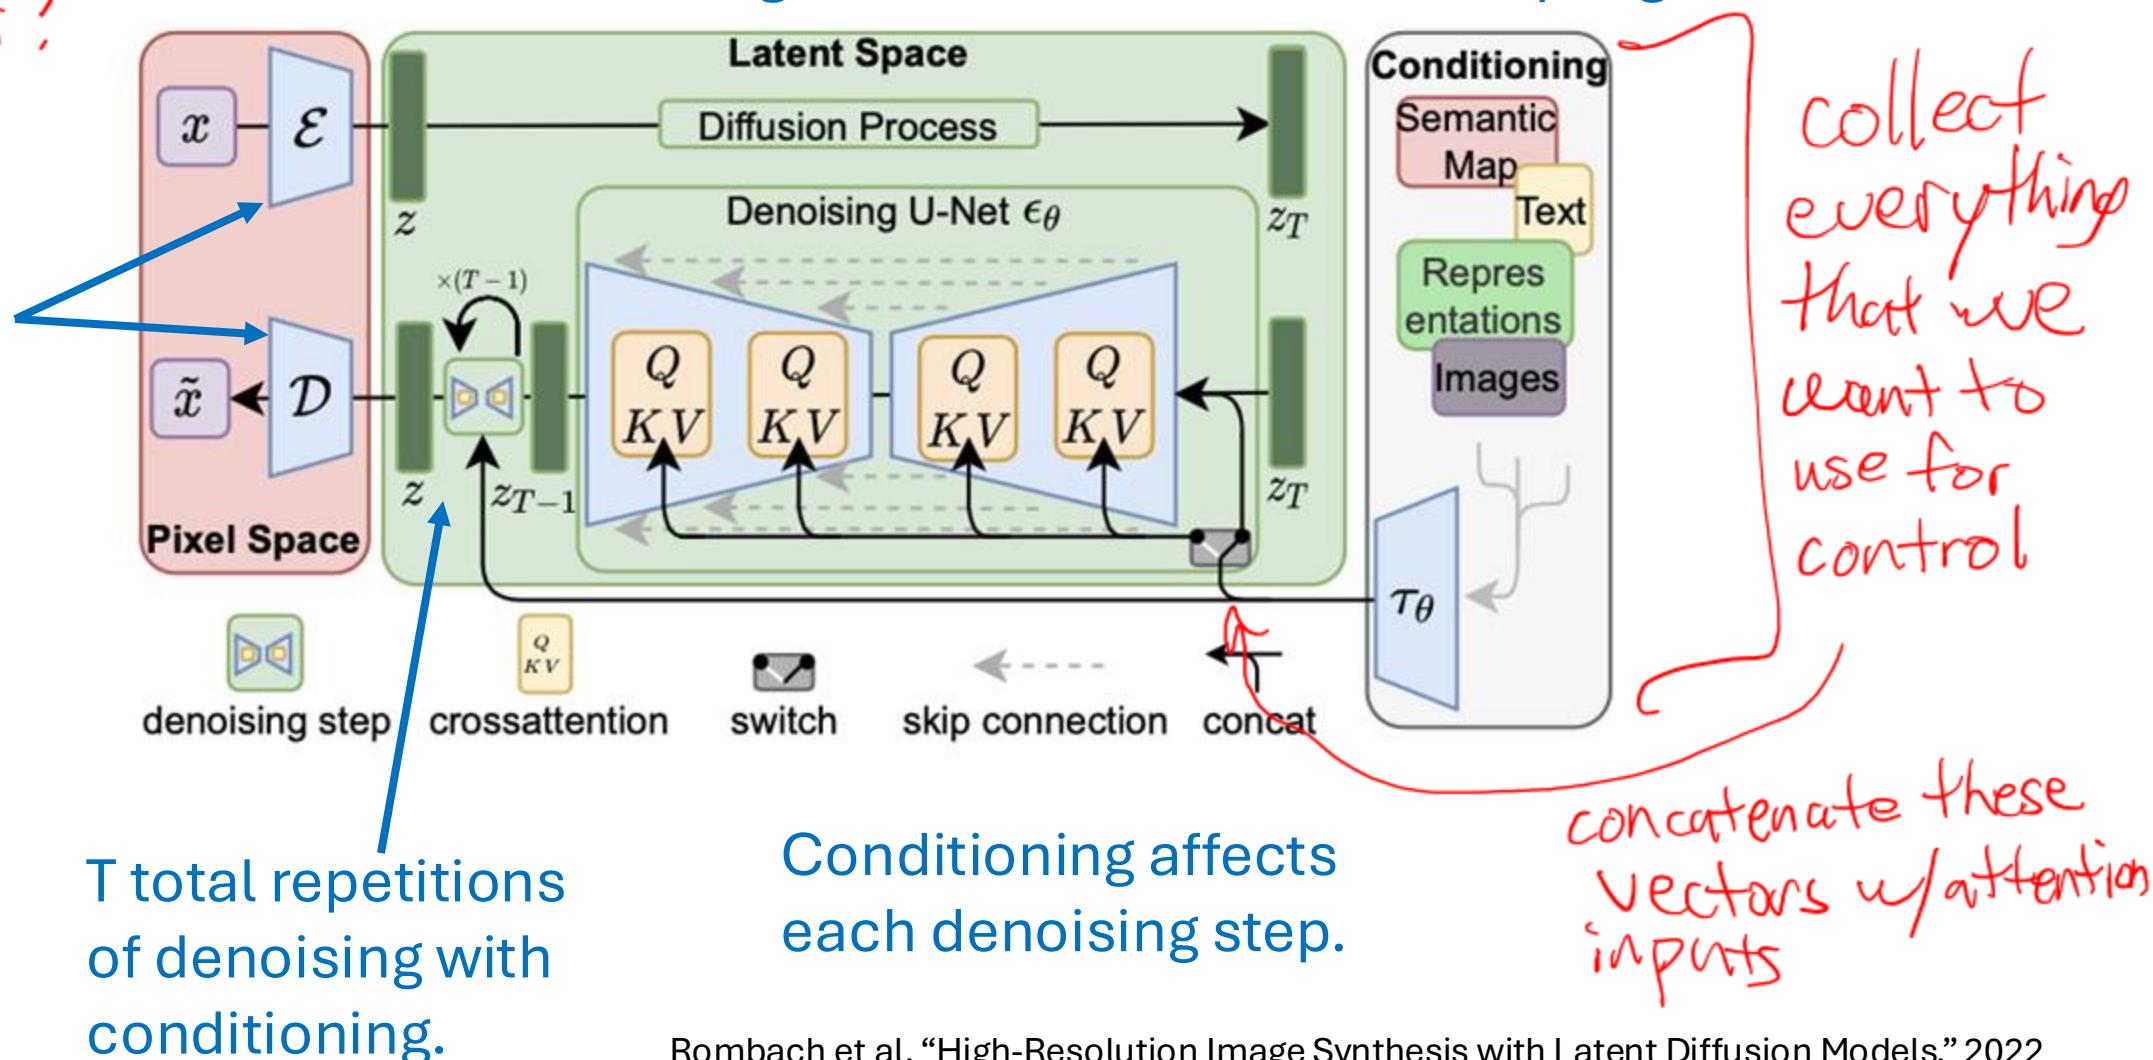
## Moving on

- Diffusion model math
- Training process
- Motivation for latent diffusion
- Latent diffusion
- Conditional generation

# Conditioning in Latent Space

How do we control  
the output?

Pick your favorite autoencoder with a small latent space.



# Conditional generation using classifier guidance



**Figure 18.12** Conditional generation using classifier guidance. Image samples conditioned on different ImageNet classes. The same model produces high quality samples of highly varied image classes. Adapted from Dhariwal & Nichol (2021).

Build a classifier  
for images,  
and adjust  
diffusion steps  
to also move  
towards target  
class.

Also works w/  
a classifier  
for latents

Limited to cases  
where classifier  
~~exists~~ or  
can be trained

# Conditional generation using text prompts

Stable Diffusion  
Conditionning  
prompt  $\Rightarrow$  latent vector.

add noise to  
that latent.  
then diffuse  
like normal.

language model  
embedding + latent  
codes need to be  
compatible. CLIP is a good choice.



**Figure 18.13** Conditional generation using text prompts. Synthesized images from a cascaded generation framework, conditioned on a text prompt encoded by a large language model. The stochastic model can produce many different images compatible with the prompt. The model can count objects and incorporate text into images. Adapted from Saharia et al. (2022b).

General Version

Text prompt  
 $\Rightarrow$  language model  
embedding

$\Rightarrow$  includes  
input into  
diffusion process

(mostly used  
by attention  
parts?)

Must be included  
in training process.

# Super Resolution

- Downsample training images 4x
- Upsample with bicubic interpolation.
- Train latent diffusion model to recover the finer-grained details.

“High-Resolution Image Synthesis with Latent Diffusion Models”  
By Rombach, Blattman, Lorenz, Esser and Ommer (2021)



Figure 9. ImageNet 64→256 super-resolution on ImageNet-Val. *LDM-SR* has advantages at rendering realistic textures but SR3 can synthesize more coherent fine structures. See appendix for additional samples and dropouts. SR3 results from [70].

# In Painting

Mask some of the image and reconstruct rest of the image to be consistent.

- Don't change the unmasked part!
- Keeping unmasked part mostly easy because latent code has spatial structure.

“High-Resolution Image Synthesis with Latent Diffusion Models”  
By Rombach, Blattman, Lorenz, Esser and Ommer (2021)

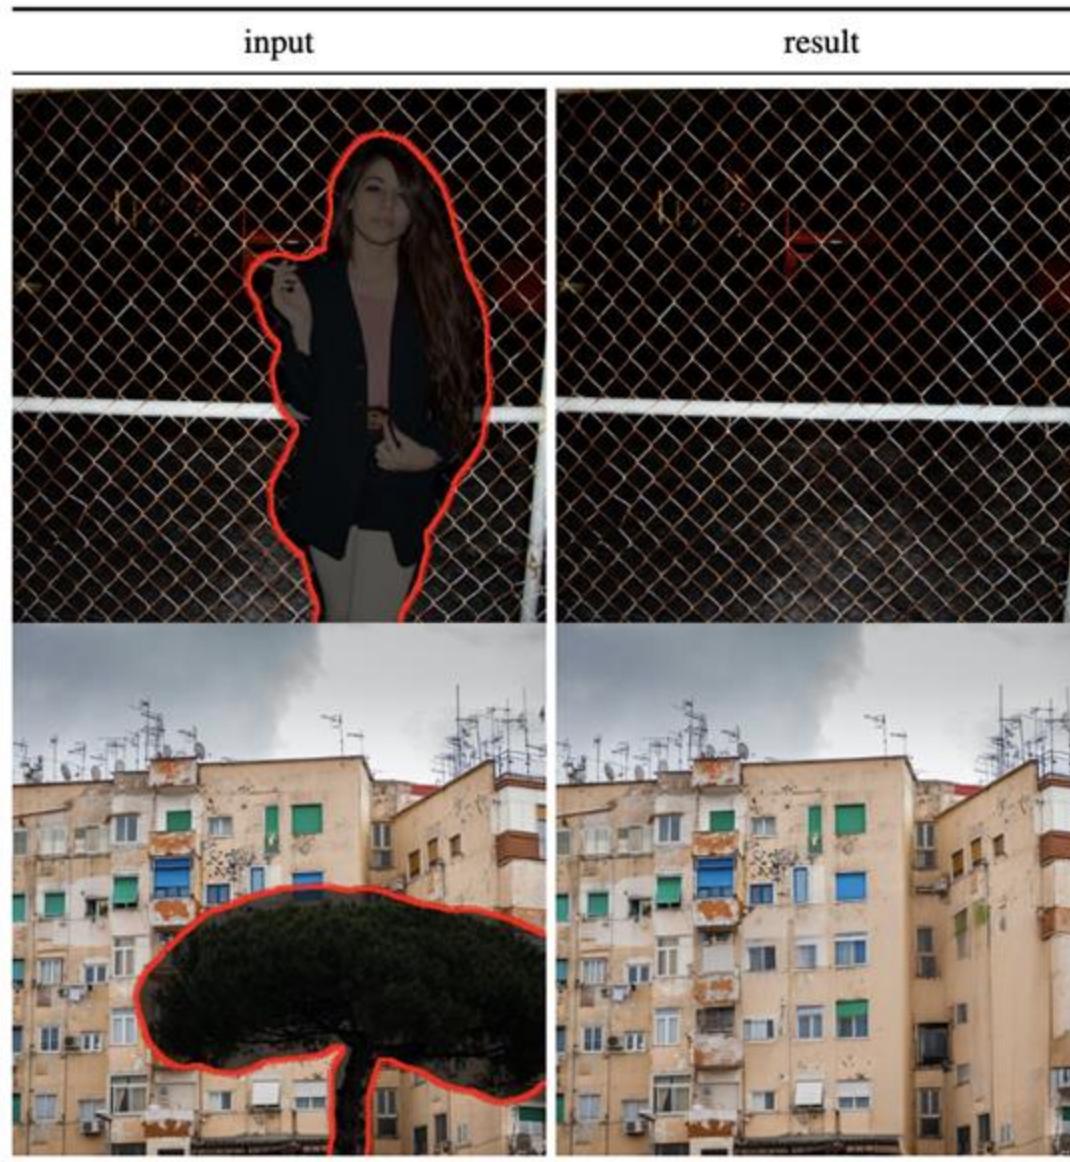


Figure 10. Qualitative results on object removal with our *big, w/ ft* inpainting model. For more results, see Fig. 21.

# Example Notebooks from the Book

- Notebook 18.1 – Diffusion Encoder in 1D   [Open in Colab](#)
- Notebook 18.2 – Training Decoder   [Open in Colab](#)
- Notebook 18.3 – 1D Reparameterized Model (more robust)   [Open in Colab](#)
- Notebook 18.4 – Families of Diffusion Models (DDIM)   [Open in Colab](#)